

Spartan-6 FPGA

Integrated

Endpoint Block for

PCI Express

User Guide

UG672 January 18, 2012



Notice of Disclaimer

The information disclosed to you hereunder (the "Materials") is provided solely for the selection and use of Xilinx products. To the maximum extent permitted by applicable law: (1) Materials are made available "AS IS" and with all faults, Xilinx hereby DISCLAIMS ALL WARRANTIES AND CONDITIONS, EXPRESS, IMPLIED, OR STATUTORY, INCLUDING BUT NOT LIMITED TO WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT, OR FITNESS FOR ANY PARTICULAR PURPOSE; and (2) Xilinx shall not be liable (whether in contract or tort, including negligence, or under any other theory of liability) for any loss or damage of any kind or nature related to, arising under, or in connection with, the Materials (including your use of the Materials), including for any direct, indirect, special, incidental, or consequential loss or damage (including loss of data, profits, goodwill, or any type of loss or damage suffered as a result of any action brought by a third party) even if such damage or loss was reasonably foreseeable or Xilinx had been advised of the possibility of the same. Xilinx assumes no obligation to correct any errors contained in the Materials or to notify you of updates to the Materials or to product specifications. You may not reproduce, modify, distribute, or publicly display the Materials without prior written consent. Certain products are subject to the terms and conditions of the Limited Warranties which can be viewed at <http://www.xilinx.com/warranty.htm>; IP cores may be subject to warranty and support terms contained in a license issued to you by Xilinx. Xilinx products are not designed or intended to be fail-safe or for use in any application requiring fail-safe performance; you assume sole risk and liability for use of Xilinx products in Critical Applications: <http://www.xilinx.com/warranty.htm#critapps>.

© Copyright 2010–2012 Xilinx, Inc. Xilinx, the Xilinx logo, Artix, ISE, Kintex, Spartan, Virtex, Zynq, and other designated brands included herein are trademarks of Xilinx in the United States and other countries. PCI, PCI Express, PCIe, and PCI-X are trademarks of PCI-SIG. All other trademarks are the property of their respective owners.

Revision History

The following table shows the revision history for this document.

Date	Version	Revision
10/05/10	1.0	Initial Xilinx release.
06/22/11	1.1	Updated core to v2.3 and ISE tool to v13.2. Removed Chapter 3, Licensing the Core, because a license is not required. Changed tstrb to tkeep in signal names. In Table 2-3 , revised the System Reset signal description and added table note about 100 MHz reference clock. Added s_axis_tx_tuser[0] and s_axis_tx_tkeep[3:0] to Table 2-6 . Added m_axis_rx_tkeep[3:0] to Table 2-7 . Added note about the interrupt status bit to Legacy Interrupt Mode, page 101 . Added Note 2 on page 102 about power-on-reset modeling. Added sentence about the 100 MHz reference clock to Clocking, page 102 . In Table J-1 , updated Difference column for trn_teof_n/s_axis_tx_tlast to indicate polarity change.
01/18/12	1.2	Updated core to v2.4 and ISE tool to v13.4. <ul style="list-style-type: none">• Removed Preface and added Appendix A, Additional Resources.• Removed Active Low indication from descriptions of cfg_rd_en and cfg_pm_wake in Table 2-8.• Indicated cfg_command[6] is hardwired to 0 in Table 5-6.• Replaced Figure 5-25 and Figure 5-26.• In Appendix D, Debugging Designs, replaced the "FPGA Configuration Time Debug" section with Clock Debug, page 171, and removed ACS violation from list of uncorrectable non-fatal errors on page 178.

Table of Contents

Revision History	2
Chapter 1: Introduction	
About the Core	7
Supported Tools and System Requirements	7
Recommended Design Experience	7
Additional Core Resources	8
Chapter 2: Core Overview	
Overview	9
Protocol Layers	10
PCI Configuration Space	12
Core Interfaces	14
Transaction Interface	15
Configuration Interface	19
Error Reporting Signals	23
Chapter 3: Getting Started Example Design	
Overview	27
Generating the Core	30
Simulating the Example Design	33
Implementing the Example Design	34
Directory Structure and File Contents	41
Chapter 4: Generating and Customizing the Core	
Customizing the Core through the CORE Generator Software	47
Chapter 5: Designing with the Core	
Design with Configuration Space Registers and Configuration Interface	82
Additional Packet Handling Requirements	88
Reporting User Error Conditions	89
Flow Control Credit Information	93
Power Management	96
Generating Interrupt Requests	98
Clocking and Reset of the Integrated Endpoint Block Core	102
Chapter 6: Core Constraints	
Contents of the User Constraints File	107

Required Modifications	108
Device Selection	108
Core I/O Assignments	109
Core Physical Constraints	109
Core Timing Constraints	109
Relocating the Integrated Endpoint Block	110
Supported Core Pinouts	110
Chapter 7: FPGA Configuration	
Configuration Terminology	113
Configuration Access Time	114
Board Power in Real-World Systems	116
Recommendations	117
Chapter 8: Known Restrictions	
Master Data Parity Error Bit Set Incorrectly	123
Non-Posted UpdateFC During PPM Transition	123
Appendix A: Additional Resources	
Xilinx Resources	125
Solution Centers	125
Additional Core Resources	125
Appendix B: Programmed Input/Output Example Design	
System Overview	127
PIO Hardware	128
PIO Operation	139
Summary	140
Root Port Model Test Bench for Endpoint	141
Appendix C: Migration Considerations	
Integrated PHY	157
System Clocking and Reset	157
Interface Changes	157
Block RAM Settings	159
Signal Change Summary	159
Appendix D: Debugging Designs	
Finding Help on Xilinx.com	161
Contacting Xilinx Technical Support	163
Debug Tools	163
Debug Ports	167

Hardware Debug	169
Simulation Debug	178

Appendix E: Managing Receive-Buffer Space for Inbound Completions

General Considerations and Concepts	183
Methods of Managing Completion Space	185

Appendix F: Board Design Guidelines

Overview	189
Example PCB Reference	189
Board Stackup	190
Data Routing Guidelines	191
Reference Clock Considerations	195

Appendix G: PCIE_A1 Port Descriptions

Clock and Reset Interface	199
Transaction Layer Interface	200
Block RAM Interface	203
GTP Transceiver Interface	203
Configuration Management Interface	206
Debug Interface Ports	216

Appendix H: PCIE_A1 Attribute Descriptions

Appendix I: PCIE_A1 Timing Parameter Descriptions

Appendix J: TRN to AXI Interface Migration Considerations

High-Level Summary	243
Step-by-Step Migration Guide	243
Signal Changes	244
Data Path DWORD Ordering	246
Start-Of-Frame Signaling	247
Remainder/Keep Signaling	247
Packet Transfer Discontinue on Receive	249
Packet Reordering on Receive	250
System Reset	250

Introduction

This user guide describes the function and operation of the Spartan®-6 FPGA Integrated Endpoint Block for PCI Express® core, including how to design, customize, and implement the core. This chapter introduces the core and provides related information including system requirements, recommended design experience, additional core resources, technical support, and submitting feedback to Xilinx.

About the Core

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express core is a reliable, high-bandwidth, scalable serial interconnect building block for use with the Spartan-6 FPGA family. The core instantiates the Spartan-6 FPGA Integrated Endpoint Block for PCI Express found in the Spartan-6 family, and supports both Verilog-HDL and VHDL.

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express core is a CORE Generator™ IP core, included in the latest IP Update on the Xilinx IP Center. For detailed information about the core, see the Spartan-6 FPGA Integrated Endpoint Block for PCI Express [product page](#).

A license is not required to use the Spartan-6 FPGA Integrated Endpoint Block for PCI Express core. The core is provided under the terms of the [Xilinx End User License Agreement](#).

Supported Tools and System Requirements

Operating System Requirements

For a list of system requirements, see the [ISE Design Suite 13: Release Notes Guide](#).

Tools

- ISE® v13.4 software

Recommended Design Experience

Although the Spartan-6 FPGA Integrated Endpoint Block for PCI Express core is a fully verified solution, the challenge associated with implementing a complete design varies depending on the configuration and functionality of the application. For best results, previous experience building high-performance, pipelined FPGA designs using Xilinx implementation software and User Constraints Files (UCF) is recommended.

Additional Core Resources

For detailed information and updates about the integrated Endpoint block core, see these documents listed on the Spartan-6 FPGA Integrated Endpoint Block for PCI Express [product page](#):

- *LogiCORE™ IP Spartan-6 FPGA Integrated Endpoint Block for PCI Express Data Sheet*
- *XTP025, LogiCORE IP Spartan-6 FPGA Integrated Endpoint Block for PCI Express Release Notes*

Additional information and resources related to the PCI Express technology are available from the following websites:

- [PCI Express at PCI-SIG](#)
- [PCI Express Developer's Forum](#)

Core Overview

This chapter describes the main components of the Spartan®-6 FPGA Integrated Endpoint Block for PCI Express® core architecture.

Overview

[Table 2-1](#) defines the Spartan-6 FPGA Integrated Endpoint Block for PCI Express solution.

Table 2-1: Product Overview

Product Name	FPGA Architecture	User Interface Width	Lane Widths Supported	Link Speeds Supported	PCI Express Base Specification Compliance
1-lane Integrated Endpoint Block	Spartan-6	32	x1	2.5 Gb/s	v1.1

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express core internally instantiates the Spartan-6 FPGA Integrated Endpoint Block for PCI Express. See [Appendix G, PCIE_A1 Port Descriptions](#), and [Appendix H, PCIE_A1 Attribute Descriptions](#), for information about the software primitive, PCIE_A1, which represents the hardened-IP integrated Endpoint block. The integrated Endpoint block follows the *PCI Express Base Specification* layering model, which consists of the Physical, Data Link, and Transaction Layers.

[Figure 2-1](#) illustrates the interfaces to the core, as defined below:

- System (SYS) interface
- PCI Express (PCI_EXP) interface
- Configuration (CFG) interface
- Transaction (AXI4-Stream) interface

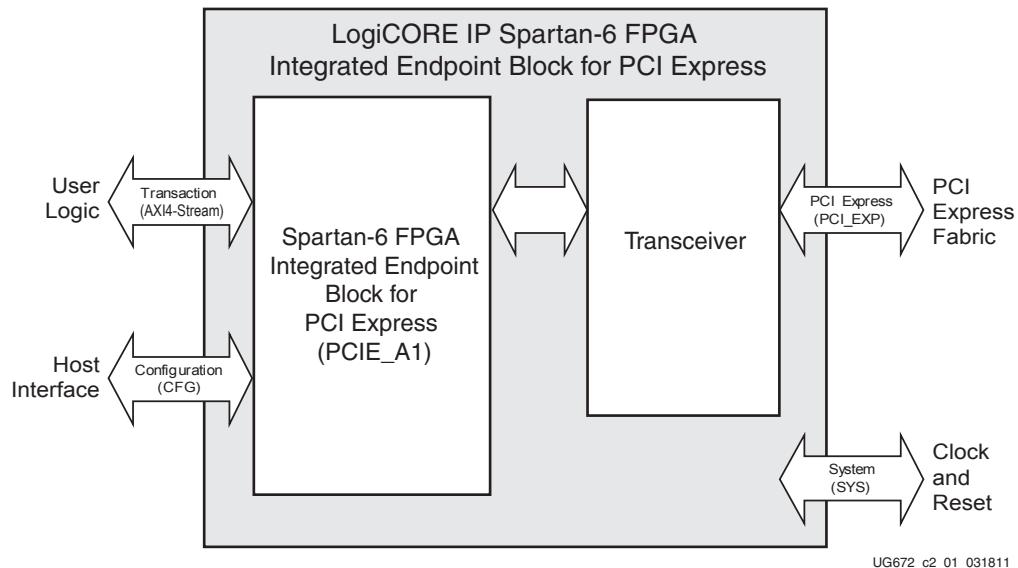


Figure 2-1: Top-Level Functional Blocks and Interfaces

The core uses packets to exchange information between the various modules. Packets are formed in the Transaction and Data Link Layers to carry information from the transmitting component to the receiving component. Necessary information is added to the packet being transmitted, which is required to handle the packet at those layers. At the receiving end, each layer of the receiving element processes the incoming packet, strips the relevant information and forwards the packet to the next layer.

As a result, the received packets are transformed from their Physical Layer representation to their Data Link Layer representation and the Transaction Layer representation.

Protocol Layers

The functions of the protocol layers, as defined by the *PCI Express Base Specification*, include generation and processing of Transaction Layer Packets (TLPs), flow control management, initialization, power management, data protection, error checking and retry, physical link interface initialization, maintenance and status tracking, serialization, deserialization and other circuitry for interface operation. Each layer is defined in the remainder of this section.

Transaction Layer

The Transaction Layer is the upper layer of the PCI Express architecture, and its primary function is to accept, buffer, and disseminate Transaction Layer packets or TLPs. TLPs communicate information through the use of memory, I/O, configuration, and message transactions. To maximize the efficiency of communication between devices, the Transaction Layer enforces PCI compliant Transaction ordering rules and manages TLP buffer space via credit-based flow control.

Data Link Layer

The Data Link Layer acts as an intermediate stage between the Transaction Layer and the Physical Layer. Its primary responsibility is to provide a reliable mechanism for the exchange of TLPs between two components on a link.

Services provided by the Data Link Layer include data exchange (TLPs), error detection and recovery, initialization services and the generation and consumption of Data Link Layer Packets (DLLPs). DLLPs are used to transfer information between Data Link Layers of two directly connected components on the link. DLLPs convey information such as Power Management, Flow Control, and TLP acknowledgments.

Physical Layer

The Physical Layer interfaces the Data Link Layer with signalling technology for link data interchange, and is subdivided into the Logical sub-block and the Electrical sub-block.

- The Logical sub-block is responsible for framing and deframing of TLPs and DLLPs. It also implements the Link Training and Status State machine (LTSSM) which handles link initialization, training, and maintenance. Scrambling, descrambling and 8B/10B encoding and decoding of data is also performed in this sub-block.
- The Electrical sub-block defines the input and output buffer characteristics that interfaces the device to the PCIe® link.

The Physical Layer also supports Lane Polarity Inversion, as indicated in the *PCI Express Base Specification rev 1.1* requirement.

Configuration Management

The Configuration Management layer maintains the PCI Type0 Endpoint configuration space and supports these features:

- Implements PCI Configuration Space
- Supports Configuration Space accesses
- Power Management functions
- Implements error reporting and status functionality
- Implements packet processing functions
 - Receive
 - Configuration Reads and Writes
 - Transmit
 - Completions with or without data
 - TLM Error Messaging
 - User Error Messaging
 - Power Management Messaging/Handshake
- Implements MSI and INTx interrupt emulation
- Implements the Device Serial Number Capability in the PCIe Extended Capability space

PCI Configuration Space

The configuration space consists of three primary parts, illustrated in [Table 2-4](#). These include:

- Legacy PCI v3.0 Type 0 Configuration Header
- Legacy Extended Capability Items
 - PCIe Capability Item
 - Power Management Capability Item
 - Message Signaled Interrupt (MSI) Capability Item
- PCIe Extended Capabilities
- Device Serial Number Extended Capability Structure (optional)

The core implements three legacy extended capability items. The remaining legacy extended capability space from address 0x6C to 0xFF is reserved. The core returns 0x00000000 when this address range is read.

The core also optionally implements one PCIe Extended Capability. The remaining PCIe Extended Capability space is reserved. If the Device Serial Number Capability is implemented, addresses from 0x10C to 0xFFFF are reserved; otherwise addresses from 0x104 to 0xFFFF are reserved. The core returns a Completion with Data of 0x00000000 if there is a configuration read to addresses in the reserved space range; writes are ignored.

Table 2-2: PCI Configuration Space Header

	31	16	15	0
	Device ID		Vendor ID	000h
	Status		Command	004h
		Class Code	Rev ID	008h
	BIST	Header	Lat Timer	Cache Ln
				00Ch
				010h
				014h
				018h
				01Ch
				020h
				024h
				028h
				02Ch
				030h
				034h
				038h
	Max Lat	Min Gnt	Intr Pin	Intr Line
				03Ch
		PM Capability	NxtCap	PM Cap
				040h
	Data	BSE		PMCSR
				044h
		MSI Control	NxtCap	MSI Cap
				048h
			Message Address (Lower)	04Ch
			Message Address (Upper)	050h
		Reserved	Message Data	054h
		PE Capability	NxtCap	PE Cap
				058h
			PCI Express Device Capabilities	05Ch
		Device Status	Device Control	060h
			PCI Express Link Capabilities	064h
		Link Status	Link Control	068h
			Reserved Legacy Configuration Space (Returns 0x00000000)	06Ch-0FFh
Optional Returns 0 if not implemented	Next Cap	Capability Version	PCI Express Extended Capability - DSN	100h
			PCI Express Device Serial Number (1st)	104h
			PCI Express Device Serial Number (2nd)	108h
			Reserved Extended Configuration Space (Returns Completion with 0x00000000)	10Ch-FFFh

Core Interfaces

The Integrated Endpoint Block for PCI Express core includes top-level signal interfaces that have sub-groups for the receive direction, transmit direction, and signals common to both directions.

System Interface

The System (SYS) interface consists of the system reset signal, sys_reset, the system clock signal, sys_clk, and a hot reset indicator, received_hot_reset, as described in [Table 2-3](#).

Table 2-3: System Interface Signals

Function	Signal Name	Direction	Description
System Reset	sys_reset	Input	Asynchronous signal. This signal is used to reset the PCI Express block, the GTP transceiver, the block RAM, and the PLL. The minimum pulse width for this input signal is 1500 ns.
System Clock	sys_clk	Input	Reference clock: 100 MHz ⁽¹⁾ or 125 MHz.
Hot Reset	received_hot_reset	Output	The core received a hot reset.

Notes:

1. Standard PCI Express slots provide a compliant 100 MHz reference clock that can be routed directly to the FPGA without the need for any external devices.

The system reset signal is an asynchronous input. The assertion of sys_reset causes a hard reset of the entire core. The system input clock must be either 100 MHz or 125 MHz, as selected in the CORE Generator software GUI. (see Note 1 in [Table 2-3](#)).

The reset provided by the PCI Express system is typically active Low (for example, PERST#) and needs to be inverted before connecting to the sys_reset signal.

PCI Express Interface

The PCI Express (PCI_EXP) interface consists of differential transmit and receive pairs. A PCI Express lane consists of a pair of transmit differential signals {pci_exp_txp, pci_exp_txn} and a pair of receive differential signals {pci_exp_rxp, pci_exp_rxn}. The 1-lane core supports only Lane 0. Transmit and receive signals of the PCI_EXP interface are defined in [Table 2-4](#).

Table 2-4: PCI Express Interface Signals for the 1-Lane Endpoint Core

Lane Number	Name	Direction	Description
0	pci_exp_txp0	Output	PCI Express Transmit Positive: Serial Differential Output 0 (+)
0	pci_exp_txn0	Output	PCI Express Transmit Negative: Serial Differential Output 0 (-)

Table 2-4: PCI Express Interface Signals for the 1-Lane Endpoint Core (Cont'd)

Lane Number	Name	Direction	Description
0	pci_exp_rxp0	Input	PCI Express Receive Positive: Serial Differential Input 0 (+)
0	pci_exp_rxn0	Input	PCI Express Receive Negative: Serial Differential Input 0 (-)

Transaction Interface

The Transaction interface provides a mechanism for the user design to generate and consume TLPs. The signal descriptions for this interface are provided in [Table 2-5](#), [Table 2-6](#), and [Table 2-7](#).

Common Interface

[Table 2-5](#) defines the common interface signals.

Table 2-5: Common Transaction Interface Signals

Name	Direction	Description
user_clk_out	Output	<p>Transaction Clock: Transaction and Configuration interface operations are referenced to and synchronous with the rising edge of this clock. user_clk_out is unavailable when the core sys_reset is held asserted. user_clk_out is guaranteed to be stable at the nominal operating frequency only after user_reset_out is deasserted. The user_clk_out clock output is a fixed frequency clock output. user_clk_out does not change frequencies in case of link recovery.</p> <ul style="list-style-type: none"> • 1-lane Integrated Endpoint Block Frequency: 62.5 MHz
user_reset_out	Output	<p>Transaction Reset: User logic interacting with the Transaction and Configuration interfaces must use user_reset_out to return to its quiescent state. user_reset_out is deasserted synchronously with respect to user_clk_out, user_reset_out is asserted asynchronously with sys_reset assertion. The user_reset_out signal is asserted for core in-band reset events like Hot Reset or Link Disable.</p>
user_lnk_up	Output	<p>Transaction Link Up: Transaction link-up is asserted when the core and the connected upstream link partner port are ready and able to exchange data packets. Transaction link-up is deasserted when the core and link partner are attempting to establish communication, or when communication with the link partner is lost due to errors on the transmission channel. user_lnk_up is also deasserted when the core is driven to Hot Reset or Link Disable states by the link partner, and all TLPs stored in the core are lost.</p>

Table 2-5: Common Transaction Interface Signals (Cont'd)

Name	Direction	Description
fc_sel[2:0]	Input	Flow Control Informational Select: Selects the type of flow control information presented on the fc_* signals. Possible values: 000 == receive buffer available space 001 == receive credits granted to the link partner 010 == receive credits consumed 100 == transmit user credits available 101 == transmit credit limit 110 == transmit credits consumed
fc_ph[7:0]	Output	Posted Header Flow Control Credits: The number of Posted Header FC credits for the selected flow control type
fc_pd[11:0]	Output	Posted Data Flow Control Credits: The number of Posted Data FC credits for the selected flow control type.
fc_nph[7:0]	Output	Non-Posted Header Flow Control Credits: The number of Non-Posted Header FC credits for the selected flow control type.
fc_npdi[11:0]	Output	Non-Posted Data Flow Control Credits: The number of Non-Posted Data FC credits for the selected flow control type.
fc_cplh[7:0]	Output	Completion Header Flow Control Credits: The number of Completion Header FC credits for the selected flow control type.
fc_cpld[11:0]	Output	Completion Data Flow Control Credits: The number of Completion Data FC credits for the selected flow control type.

Transmit Interface

Table 2-6 defines the transmit (TX) interface signals. Note that the bus s_axis_tx_tuser consists of unrelated signals. Mnemonics for these signals are used throughout this document in place of the TUSER signal names.

Table 2-6: Transmit Interface Signals

Name	Mnemonic	Direction	Description
s_axis_tx_tlast		Input	Transmit End-of-Frame (EOF): Signals the end of a packet. Valid only along with assertion of s_axis_tx_tvalid.
s_axis_tx_tdata[31:0]		Input	Transmit Data: Packet data to be transmitted.
s_axis_tx_tvalid		Input	Transmit Source Ready: Indicates that the user application is presenting valid data on s_axis_tx_tdata[31:0].
s_axis_tx_tready		Output	Transmit Destination Ready: Indicates that the core is ready to accept data on s_axis_tx_tdata[31:0]. The simultaneous assertion of s_axis_tx_tvalid and s_axis_tx_tready marks the successful transfer of one data beat on s_axis_tx_tdata[31:0].

Table 2-6: Transmit Interface Signals (Cont'd)

Name	Mnemonic	Direction	Description
s_axis_tx_tuser[3]	src_dsc	Input	Transmit Source Discontinue: Can be asserted any time starting on the first cycle after SOF to EOF, inclusive.
tx_buf_av[5:0]		Output	Transmit Buffers Available: Indicates the number of transmit buffers available in the core. Each free transmit buffer can accommodate one TLP up to the supported Maximum Payload Size. Maximum number of Transmit buffers is determined by the Supported Maximum Payload Size and block RAM configuration selected.
tx_terr_drop		Output	Transmit Error Drop: Indicates that the core discarded a packet because of a length violation or, when streaming, data was not presented on consecutive clock cycles. Length violations include packets longer than supported.
s_axis_tx_tuser[2]	str	Input	Transmit Streamed: Indicates a packet is presented on consecutive clock cycles and transmission on the link can begin before the entire packet has been written to the core. Commonly referred to as transmit cut-through mode.
tx_cfg_req		Output	Transmit Configuration Request: Asserted when the core is ready to transmit a Configuration Completion or other internally generated TLP.
tx_cfg_gnt		Input	Transmit Configuration Grant: Asserted by the user application in response to tx_cfg_req to allow the core to transmit an internally generated TLP. Holding tx_cfg_gnt deasserted after tx_cfg_req allows user-initiated TLPs to be given higher priority of transmission over core generated TLPs. tx_cfg_req is asserted once for each internally generated packet. It cannot be deasserted immediately following tx_cfg_gnt if there are no transmit buffers available. If the user does not wish to alter the prioritization of the transmission of internally generated TLPs, this signal can be continuously asserted.
s_axis_tx_tuser[1]	terr_fwd	Input	Transmit Error Forward: This input marks the current packet in progress as error-poisoned. It can be asserted any time between SOF and EOF, inclusive.

Table 2-6: Transmit Interface Signals (Cont'd)

Name	Mnemonic	Direction	Description
s_axis_tx_tuser[0]		Input	Not used for Spartan-6 devices. Drive this signal Low.
s_axis_tx_tkeep[3:0]		Input	Not used for Spartan-6 devices. Drive to 0xF.

Receive Interface

Table 2-7 defines the receive (RX) interface signals. The bus m_axis_rx_tuser consists of unrelated signals. Mnemonics for these signals are used throughout this document in place of the TUSER signal names.

Table 2-7: Receive Transaction Interface Signals

Name	Mnemonic	Direction	Description
m_axis_rx_tlast		Output	Receive End-of-Frame (EOF): Signals the end of a packet. Valid only if m_axis_rx_tvalid is also asserted.
m_axis_rx_tdata[31:0]		Output	Receive Data: Packet data being received. Valid only if m_axis_rx_tvalid is also asserted.
m_axis_rx_tuser[1]	rerr_fwd	Output	Receive Error Forward: Marks the packet in progress as error poisoned. Asserted by the core for the entire length of the packet.
m_axis_rx_tkeep[3:0]		Output	Typically not needed for Spartan-6 devices because all bytes are always valid; however, needed for AXI compliance.
m_axis_rx_tvalid		Output	Receive Source Ready: Indicates the core is presenting valid data on m_axis_rx_tdata[31:0]
m_axis_rx_tready		Input	Receive Destination Ready: Indicates the user application is ready to accept data on m_axis_rx_tdata[31:0]. The simultaneous assertion of m_axis_rx_tvalid and m_axis_rx_tready marks the successful transfer of one data beat on s_axis_tx_tdata[31:0].

Table 2-7: Receive Transaction Interface Signals (*Cont'd*)

Name	Mnemonic	Direction	Description
rx_np_ok		Input	<p>Receive Non-Posted OK: The user application asserts this signal when it is ready to accept a Non-Posted Request TLP. rx_np_ok must be deasserted when the user application cannot process received Non-Posted TLPs, so that these can be buffered within the core's receive queue. In this case, Posted and Completion TLPs received after the Non-Posted TLPs will bypass the blocked Non-Posted TLPs.</p> <p>When the user application approaches a state where it is unable to service Non-Posted Requests, it must deassert rx_np_ok two clock cycles before the core presents m_axis_rx_tlast of the second-to-last Non-Posted TLP the user application can accept.</p>
m_axis_rx_tuser[9:2]	bar_hit[6:0]	Output	<p>Receive BAR Hit: Indicates BAR(s) targeted by the current receive transaction. Asserted throughout the packet, from beginning of the packet to m_axis_rx_tlast.</p> <ul style="list-style-type: none"> • m_axis_rx_tuser[2] => BAR0 • m_axis_rx_tuser[3] => BAR1 • m_axis_rx_tuser[4] => BAR2 • m_axis_rx_tuser[5] => BAR3 • m_axis_rx_tuser[6] => BAR4 • m_axis_rx_tuser[7] => BAR5 • m_axis_rx_tuser[8] => Expansion ROM Address. <p>If two BARs are configured into a single 64-bit address, both corresponding m_axis_rx_tuser bits are asserted.</p> <p>m_axis_rx_tuser[9] is reserved for future use.</p>
m_axis_rx_tuser[0], m_axis_rx_tuser[16:15]			Reserved.

Configuration Interface

The Configuration (CFG) interface enables the user design to inspect the state of the Endpoint for PCIe configuration space. The user provides a 10-bit configuration address, which selects one of the 1024 configuration space double word (DWORD) registers. The endpoint returns the state of the selected register over the 32-bit data output port. [Table 2-8](#)

defines the Configuration interface signals. See [Design with Configuration Space Registers and Configuration Interface, page 82](#) for usage.

Table 2-8: Configuration Interface Signals

Name	Direction	Description				
cfg_do[31:0]	Output	Configuration Data Out: A 32-bit data output port used to obtain read data from the configuration space inside the core.				
cfg_rd_wr_done	Output	Configuration Read Write Done: Read-write done signal. Indicates a successful completion of the user configuration register access operation. <ul style="list-style-type: none"> • For a user configuration register read operation, the signal validates the cfg_do[31:0] data-bus value. • Writes to the configuration space are not supported. 				
cfg_dwaddr[9:0]	Input	Configuration DWORD Address: A 10-bit address input port used to provide a configuration register DWORD address during configuration register accesses.				
cfg_rd_en	Input	Configuration Read Enable: Read enable for configuration register access. Note: cfg_rd_en must be asserted for no more than 1 user_clk_out cycle for each access.				
cfg_interrupt	Input	Configuration Interrupt: Interrupt request signal. The user application can assert this to cause the selected interrupt message type to be transmitted by the core. The signal should be held Low until cfg_interrupt_rdy is asserted.				
cfg_interrupt_rdy	Output	Configuration Interrupt Ready: Interrupt grant signal. The simultaneous assertion of cfg_interrupt_rdy and cfg_interrupt indicates that the core has successfully transmitted the requested interrupt message.				
cfg_interrupt_assert	Input	Configuration Legacy Interrupt Assert/Deassert Select: Selects between Assert and Deassert messages for Legacy interrupts when cfg_interrupt is asserted. Not used for MSI interrupts. Value Message Type <table style="margin-left: 20px;"> <tr> <td>0</td> <td>Deassert</td> </tr> <tr> <td>1</td> <td>Assert</td> </tr> </table>	0	Deassert	1	Assert
0	Deassert					
1	Assert					

Table 2-8: Configuration Interface Signals (Cont'd)

Name	Direction	Description										
cfg_interrupt_di[7:0]	Input	<p>Configuration Interrupt Data In: For Message Signaling Interrupts (MSI), the portion of the Message Data that the endpoint must drive to indicate MSI vector number, if Multi-Vector Interrupts are enabled. The value indicated by cfg_interrupt_mmenable[2:0] determines the number of lower-order bits of Message Data that the endpoint provides; the remaining upper bits of cfg_interrupt_di[7:0] are not used.</p> <p>For Single-Vector Interrupts, cfg_interrupt_di[7:0] is not used. For Legacy interrupt messages (Assert_INTx, Deassert_INTx), this list defines the type of message to be sent:</p> <table> <thead> <tr> <th>Value</th><th>Legacy Interrupt</th></tr> </thead> <tbody> <tr> <td>00h</td><td>INTA</td></tr> <tr> <td>01h</td><td>INTB</td></tr> <tr> <td>02h</td><td>INTC</td></tr> <tr> <td>03h</td><td>INTD</td></tr> </tbody> </table>	Value	Legacy Interrupt	00h	INTA	01h	INTB	02h	INTC	03h	INTD
Value	Legacy Interrupt											
00h	INTA											
01h	INTB											
02h	INTC											
03h	INTD											
cfg_interrupt_do[7:0]	Output	Configuration Interrupt Data Out: The value of the lowest eight bits of the Message Data field in the endpoint's MSI capability structure. This value is not used and is provided for informational purposes and backwards compatibility.										
cfg_interrupt_mmenable[2:0]	Output	<p>Configuration Interrupt Multiple Message Enable: This is the value of the Multiple Message Enable field and defines the number of vectors the system allows for multi-vector MSI. Values range from 000b to 101b. A value of 000b indicates that single vector MSI is enabled, while other values indicate the number of lower-order bits that can be used for cfg_interrupt_di[7:0].</p> <p>cfg_interrupt_mmenable[2:0] values:</p> <ul style="list-style-type: none"> • 000b, 0 bits • 001b, 1 bit • 010b, 2 bits • 011b, 3 bits • 100b, 4 bits • 101b, 5 bits 										
cfg_interrupt_msienable	Output	Configuration Interrupt MSI Enabled: Indicates that the Message Signaling Interrupt (MSI) messaging is enabled. If 0, then only Legacy (INTx) interrupts can be sent. If 1, only MSI interrupts can be sent.										

Table 2-8: Configuration Interface Signals (*Cont'd*)

Name	Direction	Description
cfg_bus_number[7:0]	Output	Configuration Bus Number: Provides the assigned bus number for the device. The user application must use this information in the Bus Number field of outgoing TLP requests. Default value after reset is 00h. Refreshed whenever a Type 0 Configuration Write packet is received.
cfg_device_number[4:0]	Output	Configuration Device Number: Provides the assigned device number for the device. The user application must use this information in the Device Number field of outgoing TLP requests. Default value after reset is 00000b. Refreshed whenever a Type 0 Configuration Write packet is received.
cfg_function_number[2:0]	Output	Configuration Function Number: Provides the function number for the device. The user application must use this information in the Function Number field of outgoing TLP request. Function number is hardwired to 000b.
cfg_status[15:0]	Output	Configuration Status: Status register from the Configuration Space Header.
cfg_command[15:0]	Output	Configuration Command: Command register from the Configuration Space Header.
cfg_dstatus[15:0]	Output	Configuration Device Status: Device Status register from the PCI Express Extended Capability Structure.
cfg_dcommand[15:0]	Output	Configuration Device Command: Device Control register from the PCI Express Extended Capability Structure.
cfg_lstatus[15:0]	Output	Configuration Link Status: Link Status register from the PCI Express Extended Capability Structure.
cfg_lcommand[15:0]	Output	Configuration Link Command: Link Control register from the PCI Express Extended Capability Structure.
cfg_to_turnoff	Output	Configuration To Turnoff: This output notifies the user that a PME_TURN_Off message has been received and the CMM will start polling the cfg_turnoff_ok input coming in from the user. After cfg_turnoff_ok is asserted, CMM sends a PME_To_Ack message to the upstream device.
cfg_turnoff_ok	Input	Configuration Turnoff OK: The user application can assert this to notify the integrated Endpoint block core that it is safe to turn the power off.

Table 2-8: Configuration Interface Signals (Cont'd)

Name	Direction	Description
cfg_pm_wake	Input	Configuration Power Management Wake: A one-clock cycle assertion signals the core to generate and send a Power Management Wake Event (PM_PME) Message TLP to the upstream link partner. Note: You must assert this input only under stable link conditions as reported on the cfg_PCIE_link_state[2:0] bus. Assertion of this signal when the PCI Express Link is in transition results in incorrect behavior on the PCI Express Link.
cfg_PCIE_link_state[2:0]	Output	PCI Express Link State: This encoded bus reports the PCIe Link State Information to the user. <ul style="list-style-type: none"> • 110b - PCI Express Link State is "L0" • 101b - PCI Express Link State is "L0s" • 011b - PCI Express Link State is "L1" • 111b - PCI Express Link State is "in transition"
cfg_trn_pending	Input	User Transaction Pending: If asserted, sets the Transactions Pending bit in the Device Status register. Note: The user is required to assert this input if the user application has not received a completion to an upstream request.
cfg_dsn[63:0]	Input	Configuration Device Serial Number: Serial Number register fields of the Device Serial Number extended capability. Not used if DSN capability is disabled.
cfg_ltssm_state[4:0]	Output	LTSSM State: Indicates the current state of the Link Training and Status State Machine. For state encodings, see CFGLTSSMSTATE in Table G-10, page 214 .

Error Reporting Signals

Table 2-9 defines the user application error-reporting signals.

Table 2-9: User Application Error-Reporting Signals

Port Name	Direction	Description
cfg_err_ecrc	Input	ECRC Error Report: The user can assert this signal to report an ECRC error (end-to-end CRC).
cfg_err_ur	Input	Configuration Error Unsupported Request: The user can assert this signal to report that an unsupported request was received. This signal is ignored if cfg_err_cpl_rdy is deasserted.

Table 2-9: User Application Error-Reporting Signals (Cont'd)

Port Name	Direction	Description
cfg_err_cpl_timeout	Input	<p>Configuration Error Completion Timeout: The user can assert this signal to report a completion timed out.</p> <p>Note: The user should assert this signal only if the device power state is D0. Asserting this signal in non-D0 device power states might result in an incorrect operation on the PCIe link. For additional information, see the <i>PCI Express Base Specification</i>, Rev.1.1, Section 5.3.1.2.</p>
cfg_err_cpl_abort	Input	Configuration Error Completion Aborted: The user can assert this signal to report that a completion was aborted. This signal is ignored if cfg_err_cpl_rdy is deasserted.
cfg_err_posted	Input	Configuration Error Posted: This signal is used to further qualify any of the cfg_err_* input signals. When this input is asserted concurrently with one of the other signals, it indicates that the transaction which caused the error was a posted transaction.
cfg_err_cor	Input	Configuration Error Correctable Error: The user can assert this signal to report that a correctable error was detected.
cfg_err_tlp_cpl_header[47:0]	Input	<p>Configuration Error TLP Completion Header: Accepts the header information from the user when an error is signaled. This information is required so that the core can issue a correct completion, if required.</p> <p>This information should be extracted from the received error TLP and presented in the indicated format:</p> <ul style="list-style-type: none"> [47:41] Lower Address [40:29] Byte Count [28:26] TC [25:24] Attr [23:8] Requester ID [7:0] Tag

Table 2-9: User Application Error-Reporting Signals (Cont'd)

Port Name	Direction	Description
cfg_err_cpl_rdy	Output	Configuration Error Completion Ready: When asserted, this signal indicates that the core can accept assertions on cfg_err_ur and cfg_err_cpl_abort for Non-Posted transactions. Assertions on cfg_err_ur and cfg_err_cpl_abort are ignored when cfg_err_cpl_rdy is deasserted.
cfg_err_locked	Input	Configuration Error Locked: This signal is used to further qualify any of the cfg_err_* input signals. When this input is asserted concurrently with one of the other signals, it indicates that the transaction that caused the error was a locked transaction. This signal is intended to be used in Legacy mode. If the user needs to signal an unsupported request or an aborted completion for a locked transaction, this signal can be used to return a Completion Locked with UR or CA status. Note: When not in Legacy mode, the core automatically returns a Completion Locked, if appropriate.

Getting Started Example Design

This chapter provides an overview of the Spartan®-6 FPGA Integrated Endpoint Block for PCI Express® example design and instructions for generating the core. It also includes information about simulating and implementing the example design using the provided demonstration test bench.

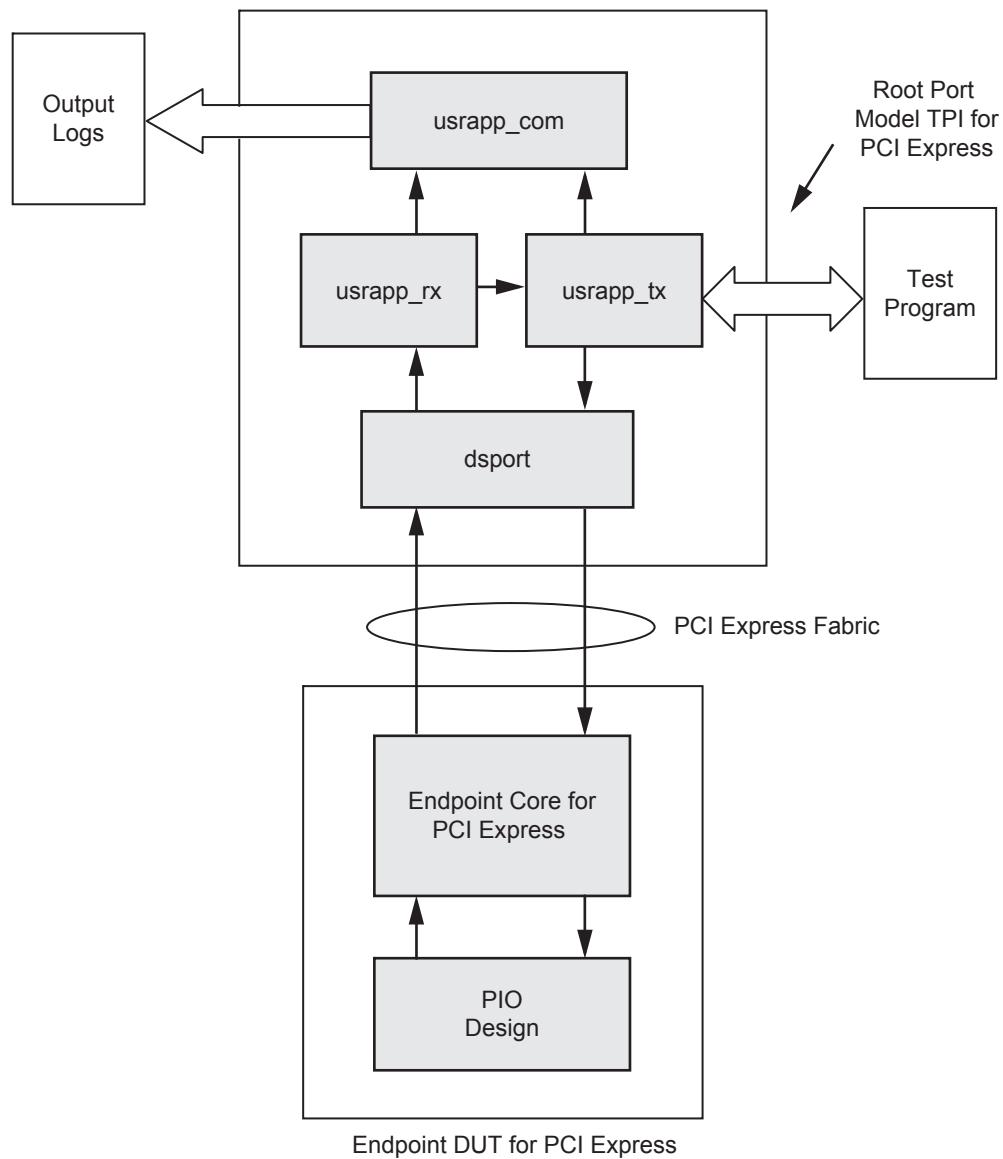
Overview

The example simulation design consists of two discrete parts:

- The Root Port Model, a test bench that generates, consumes, and checks PCI Express bus traffic.
- The Programmed Input/Output (PIO) example design, a complete application for PCI Express. The PIO example design responds to Read and Write requests to its memory space and can be synthesized for testing in hardware.

Simulation Design Overview

For the simulation design, transactions are sent from the Root Port Model to the integrated Endpoint block core and processed by the PIO example design. [Figure 3-1](#) illustrates the simulation design provided with the integrated Endpoint block core. For more information about the Root Port Model, see [Root Port Model Test Bench for Endpoint, page 141](#).



UG672_c4_01_083110

Figure 3-1: Simulation Example Design Block Diagram

Implementation Design Overview

The implementation design consists of a simple PIO example that can accept read and write transactions and respond to requests, as illustrated in [Figure 3-2](#). Source code for the example is provided with the core. For more information about the PIO example design, see [Appendix B, Programmed Input/Output Example Design](#).

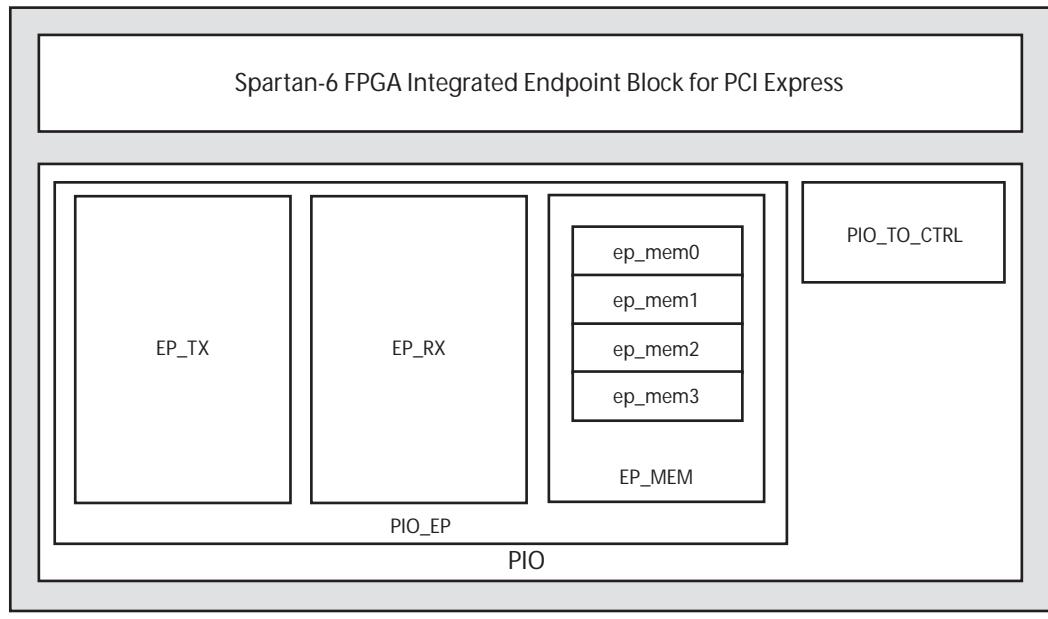


Figure 3-2: Implementation Example Design Block Diagram

Example Design Elements

The PIO example design elements include:

- Core wrapper
- An example Verilog HDL or VHDL wrapper (instantiates the cores and example design)
- A customizable demonstration test bench to simulate the example design

The example design has been tested and verified with Xilinx ISE® v13.4 software and these simulators:

- Mentor Graphics ModelSim
- Cadence Incisive Enterprise Simulator (IES)
- Synopsys VCS and VCS MX
- ISE Simulator (ISim)

Note: The VHDL example design supports only ModelSim.

Generating the Core

To generate a core using the default values in the CORE Generator software GUI, follow these steps:

1. Start the CORE Generator tool.

For help starting and using the CORE Generator tool, see the *Xilinx CORE Generator Guide*, available from the [ISE Design Suite](#) web page.

2. Choose **File > New Project**.

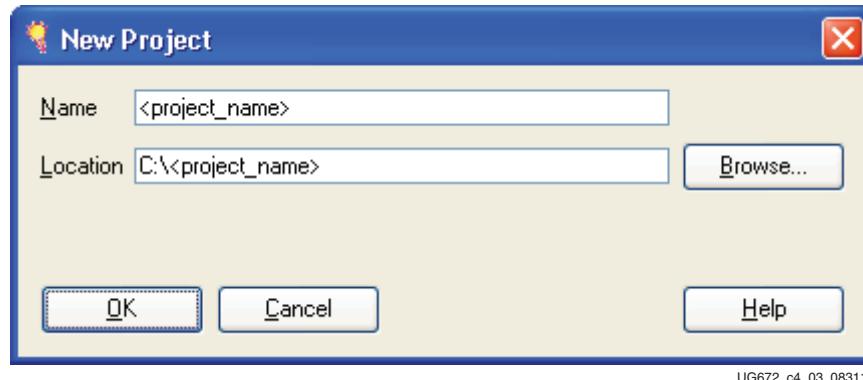


Figure 3-3: New Project Dialog Box

3. Enter a project name and location, then click **OK**. `<project_dir>` is used in this example. The Project Options dialog box appears.

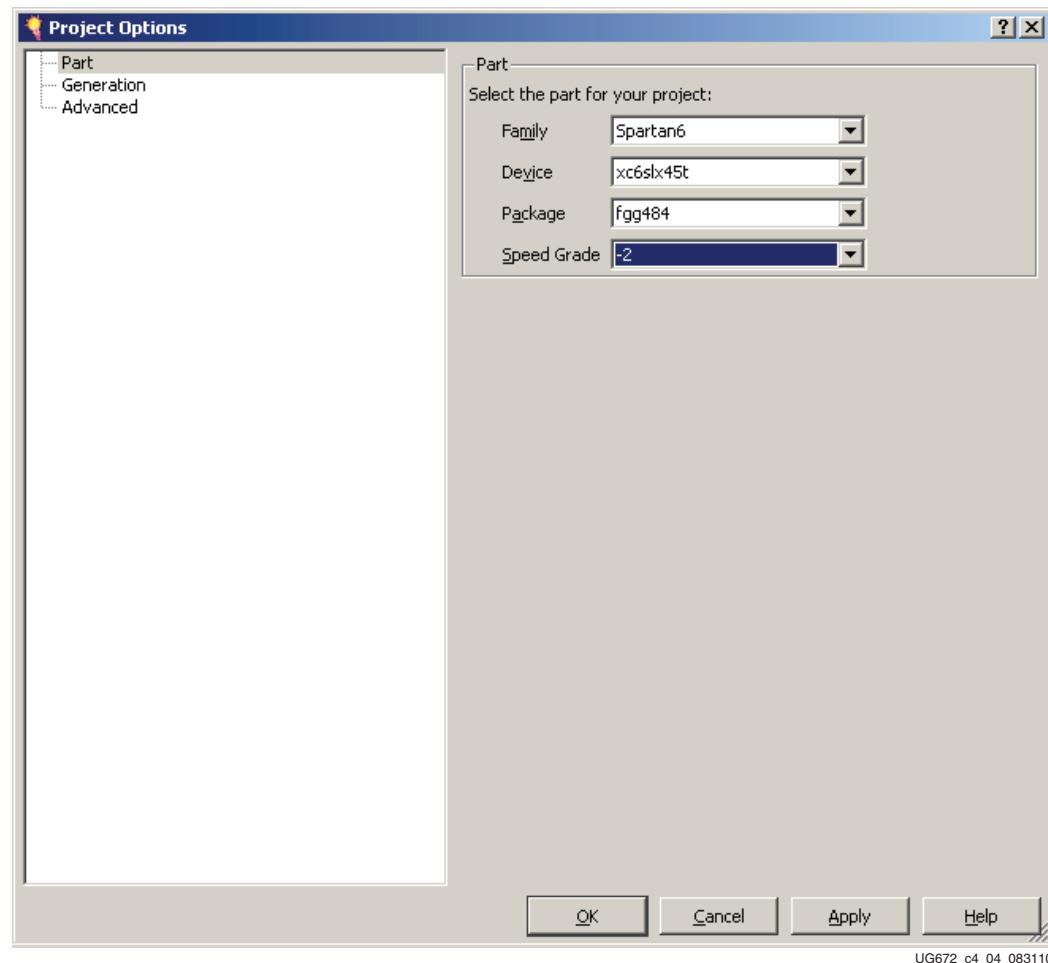


Figure 3-4: Project Options

4. Set the project options:

From the Part tab, select these options:

- **Family:** Spartan6
- **Device:** xc6slx45t
- **Package:** fgg484
- **Speed Grade:** -2

Note: If an unsupported silicon device is selected, the core is dimmed (unavailable) in the list of cores.

From the Generation tab, select these parameters, and then click **OK**.

- **Design Entry.** Select **Verilog** or **VHDL**.
- **Vendor.** Select **Synplicity** or **ISE** (for XST).

5. Locate the core in the selection tree under Standard Bus Interfaces/PCI Express; then double-click the core name to display the integrated Endpoint block main screen.

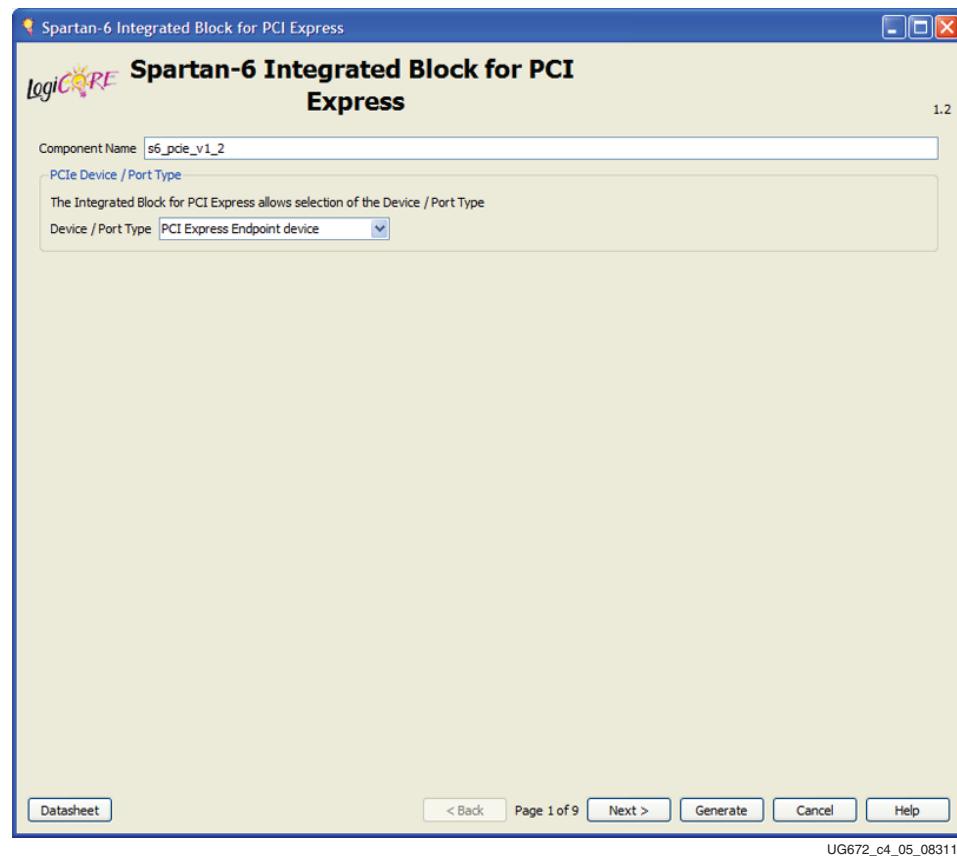


Figure 3-5: Integrated Endpoint Block Main Screen

6. In the Component Name field, enter a name for the core. <component_name> is used in this example.
7. Click **Finish** to generate the core using the default parameters. The core and its supporting files, including the PIO example design and Root Port Model test bench, are generated in the project directory.

For detailed information about the example design files and directories see [Directory Structure and File Contents, page 41](#). See the README file.

Simulating the Example Design

The example design provides a quick way to simulate and observe the behavior of the core. The simulation environment provided with the integrated Endpoint block core performs simple memory access tests on the PIO example design. Transactions are generated by the Root Port Model and responded to by the PIO example design.

- PCI Express Transaction Layer Packets (TLPs) are generated by the test bench transmit user application (`pci_exp_usrapp_tx`). As it transmits TLPs, it also generates a log file, `tx.dat`.
- PCI Express TLPs are received by the test bench receive user application (`pci_exp_usrapp_rx`). As the user application receives the TLPs, it generates a log file, `rx.dat`.

For more information about the test bench, see [Root Port Model Test Bench for Endpoint, page 141](#).

Setting up for Simulation

To run the functional simulation the Xilinx Simulation Libraries must be compiled for the user system. See the Compiling Xilinx Simulation Libraries (COMPXLIB) in the *Xilinx ISE Synthesis and Verification Design Guide*, and the *Xilinx ISE Software Manuals and Help*. Documents can be downloaded from www.xilinx.com/support/software_manuals.htm.

Simulator Requirements

Spartan-6 device designs require a Verilog LRM-IEEE 1364-2005 encryption-compliant simulator.

Note for Cadence IUS users: The work construct must be manually inserted into the CDS.LIB file as shown below.

```
DEFINE WORK WORK
```

Running the Simulation

The simulation scripts provided with the example design support pre-implementation (RTL) simulation. The existing test bench can be used to simulate with a post-implementation version of the example design.

The pre-implementation simulation consists of these components:

- Verilog or VHDL model of the test bench
 - Verilog or VHDL RTL example design
 - The Verilog or VHDL model of the Integrated Endpoint Block for PCI Express
1. To run the simulation, go to this directory:
`<project_dir>/<component_name>/simulation/functional`
 2. Run the script that corresponds to the user simulation tool using one of these:
 - **ModelSim:** `vsim -do simulate_mti.do`
 - **VCS:** `>./simulate_vcs.sh`

- **IES:** > ./simulate_ncsim.sh
- **ISIM (UNIX):** > ./simulate_isim.sh
- **ISIM (Windows):** > simulate_isim.bat

Implementing the Example Design

After generating the core, the netlists and the example design can be processed using the Xilinx implementation tools. The generated output files include scripts to assist in running the Xilinx software.

To implement the example design:

Open a command prompt or terminal window and type:

Windows

```
ms-dos> cd <project_dir>\<component_name>\implement  
ms-dos> implement.bat
```

Linux

```
% cd <project_dir>/<component_name>/implement  
% ./implement.sh
```

These commands execute a script that synthesizes, builds, maps, and place-and-routes the example design, and then generates a post-par simulation model for use in timing simulation. The resulting files are placed in the `results` directory and execute these processes:

1. Removes data files from the previous runs.
2. Synthesizes the example design using XST or Synplify.
3. `ngdbuild`. Builds a Xilinx design database for the example design.
 - Inputs:
Part-Package-Speed Grade selection:
For example, XC6SLX45T-FFG484-1
 - Example design UCF:**
`xilinx_pcie_1_lane_ep_<device>.ucf`
4. `map`: Maps design to the selected FPGA using the constraints provided.
5. `par`: Places cells onto FPGA resources and routes connectivity.
6. `trce`: Performs static timing analysis on design using constraints specified.
7. `netgen`: Generates a logical Verilog HDL or VHDL representation of the design and an SDF file for post-layout verification.
8. `bitgen`: Generates a bitstream file for programming the FPGA.

These FPGA implementation related files are generated in the `results` directory:

- `routed.bit`
FPGA configuration information.
- `routed.v[hd]`
Verilog or VHDL functional Model.
- `routed.sdf`
Timing model Standard Delay File.

- mapped.mrp
Xilinx map report.
- routed.par
Xilinx place and route report.
- routed.twr
Xilinx timing analysis report.

The script file starts from Verilog or VHDL source files and results in a bitstream file.

Users can also use the ISE Project Navigator GUI tool to implement designs. An example ISE software project file is provided when the core is generated.

Using the ISE Project Navigator GUI Tool

To build a core and PIO example design with the ISE Project Navigator GUI tool:

1. Start the ISE Project Navigator GUI tool.

For help starting and using the ISE Project Navigator tool, see the *ISE Project Navigator Guide*, available from the ISE tool documentation web page (http://www.xilinx.com/support/documentation/dt_ise.htm).

2. Choose **File → New Project**.
3. Enter a project name and location, then click **Next >** (see [Figure 3-6](#)).

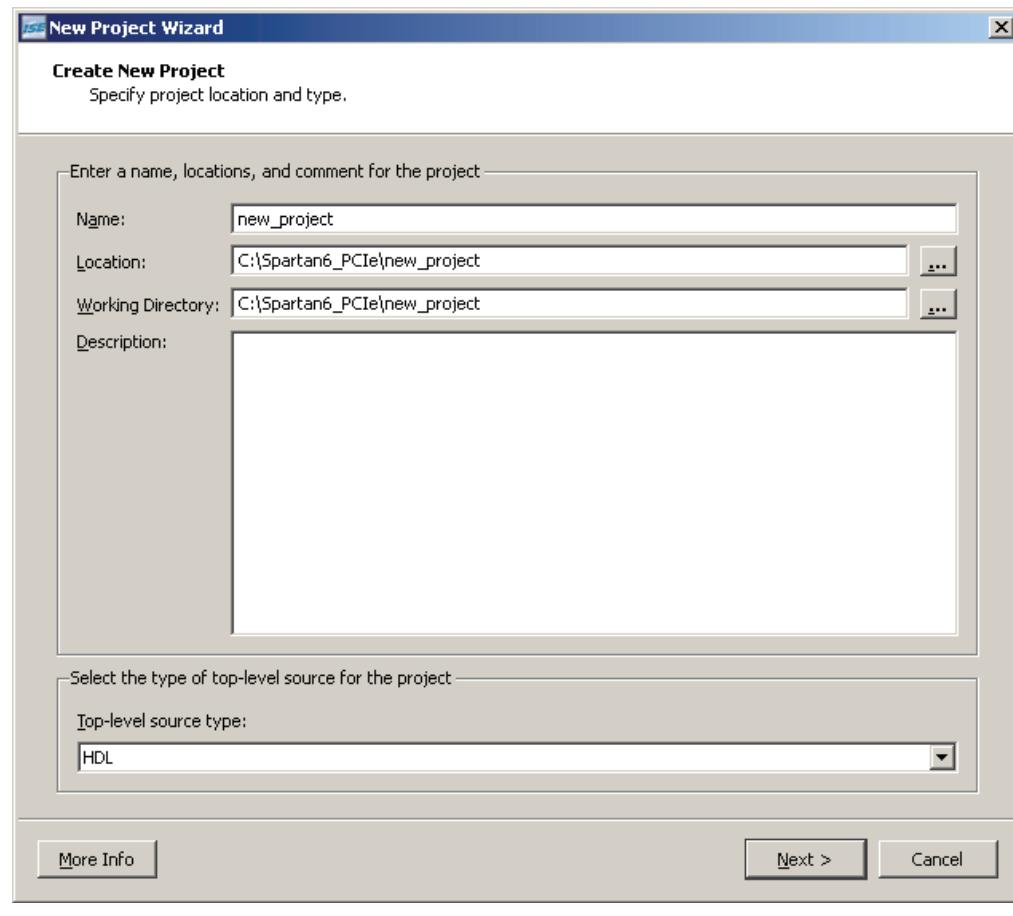
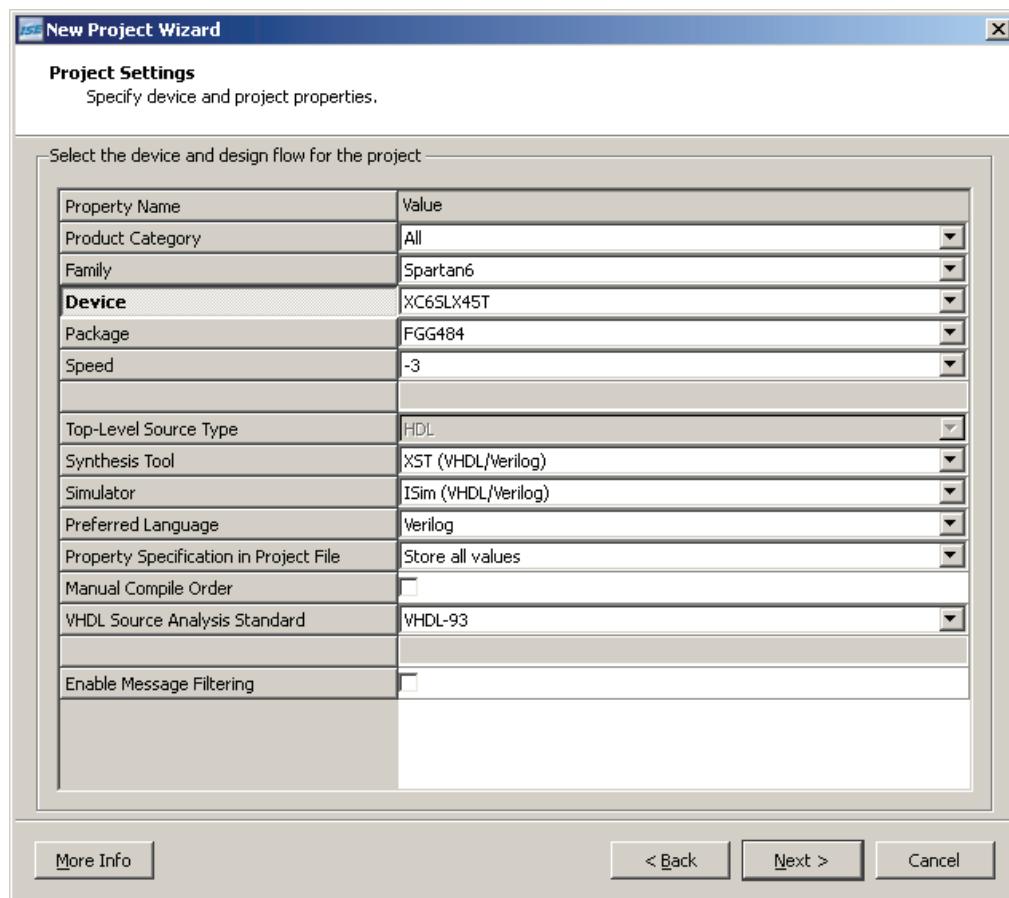


Figure 3-6: Create New Project

4. Set the project options (see [Figure 3-7](#)):

Family: **Spartan6**

Device: Any LXT device



UG762_c4_07_083110

Figure 3-7: Project Settings

5. Click **Next >** and then **Finish** to create the project.
6. Choose **Project → New Source**.
7. Select **IP (Core Generator & Architecture Wizard)**.

8. Enter a file name and ensure the “Add to project” checkbox is checked (see Figure 3-8).

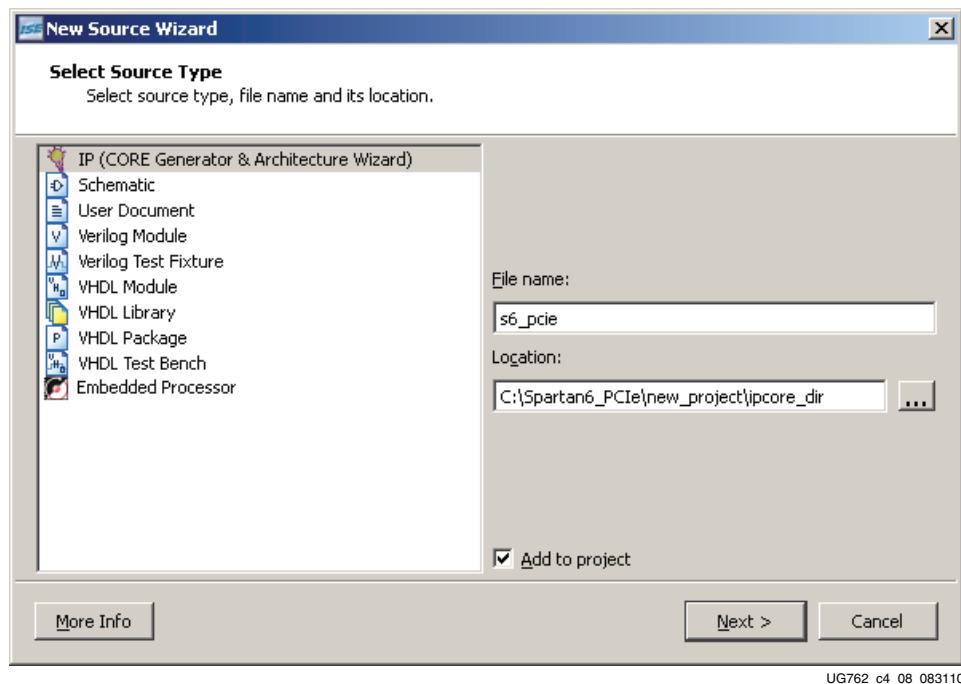


Figure 3-8: Select Source Type

9. Select **Spartan-6 Integrated Block for PCI Express**. Click **Next >** and then **Finish** (see Figure 3-9).

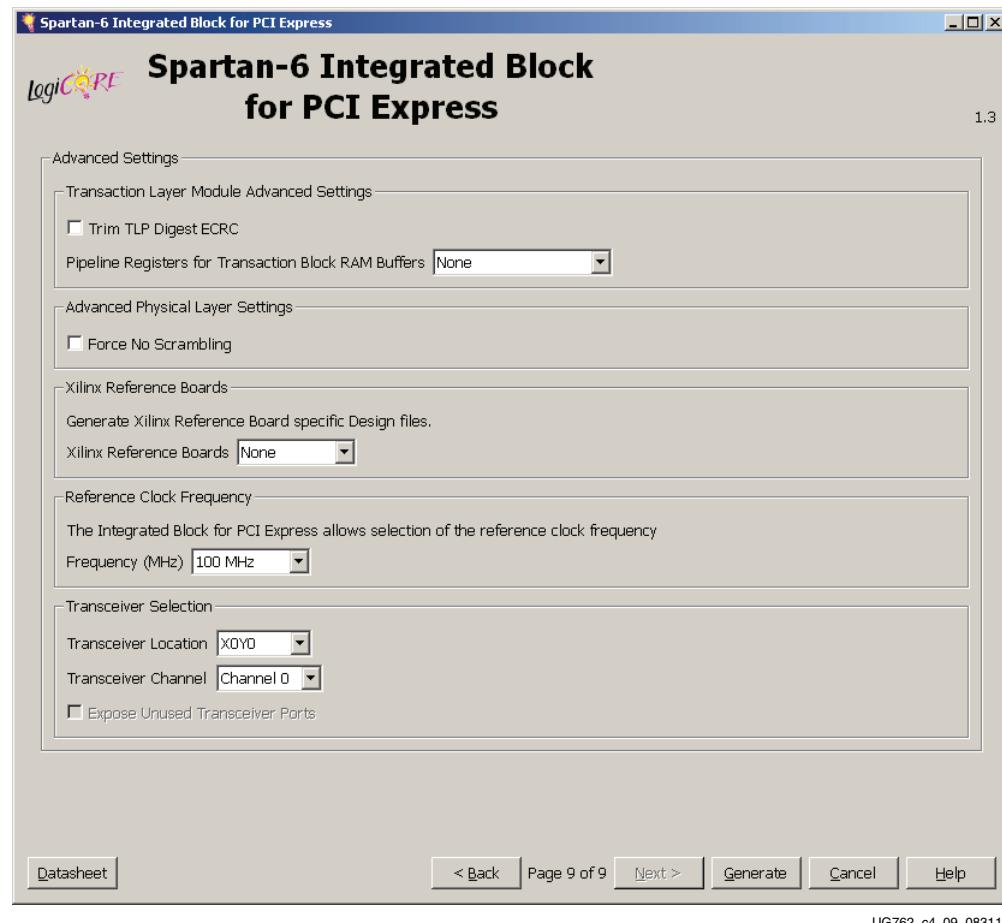


Figure 3-9: Select IP

10. Configure the core as described in [Chapter 4, Generating and Customizing the Core](#).
11. Choose **File → Open Project**.
12. Enter the `ipcore_dir` directory (see [Figure 3-10](#)).

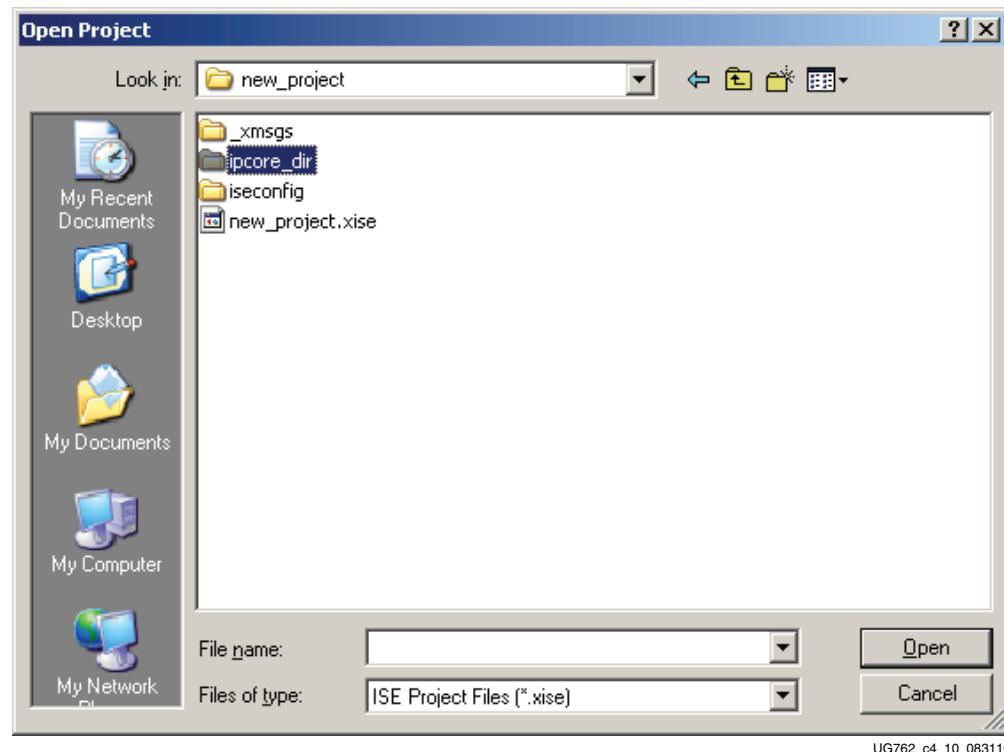


Figure 3-10: Directory ipcore_dir

13. Select the <component_name>_<lang>_example_project.xise file to load an example project with the PIO example design along with the core built in step 1 through step 9 (see Figure 3-11).

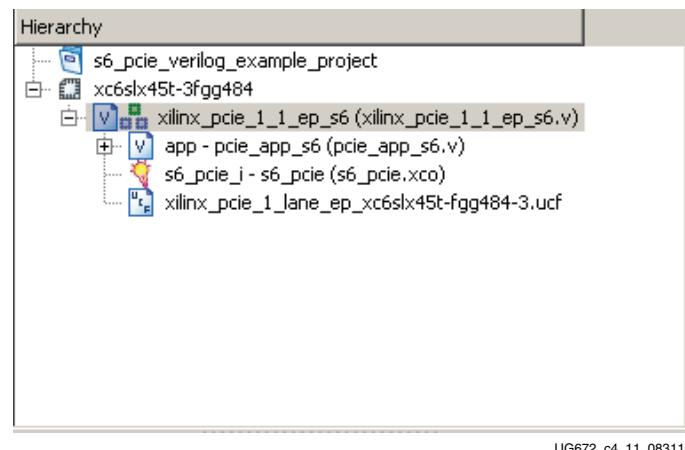


Figure 3-11: Load Example Project

Directory Structure and File Contents

The integrated Endpoint block example design directories and their associated files are defined in the sections that follow. Click a directory name to go to the desired directory and its associated files.

Example Design

 <code><project directory></code>	Top-level project directory; name is user-defined
 <code><project directory>/<component name></code>	Core release notes readme file
 <code><component name>/doc</code>	Product documentation
 <code><component name>/example_design</code>	Verilog or VHDL design files
 <code><component name>/implement</code>	Implementation script files
 <code>implement/results</code>	Results directory, created after implementation scripts are run, and contains implement script results
 <code><component name>/simulation</code>	
 <code>simulation/dsport</code>	Root Port Bus Functional Model
 <code>simulation/functional</code>	Functional simulation files
 <code>simulation/tests</code>	Test command files
 <code><component name>/source</code>	Core source files

`<project directory>`

The project directory contains all the CORE Generator tool project files.

Table 3-1: Project Directory

Name	Description
<code><project_dir></code>	
<code><component_name>.xco</code>	CORE Generator tool project-specific option file; can be used as an input to the CORE Generator software.
<code><component_name>_flist.txt</code>	List of files delivered with core.
<code><component_name>_<lang>_example_project.xise</code>	ISE software Project Navigator project file for the PIO example design.
<code><component_name>.v[eo ho]</code>	Verilog or VHDL instantiation template.

[Back to Top](#)

<project directory>/<component name>

The component name directory contains the release notes readme file provided with the core, which can include tool requirements, last-minute changes, updates, and issue resolution.

Table 3-2: Component Name Directory

Name	Description
<project_dir>/<component_name>	
s6_pcie_v2_3_readme.txt	Readme file.

[Back to Top](#)

<component name>/doc

The doc directory contains the PDF documentation provided with the core.

Table 3-3: Doc Directory

Name	Description
<project_dir>/<component_name>/doc	
ug672_S6_IntEndptBlk_PCIe.pdf	<i>Spartan-6 FPGA Integrated Endpoint Block for PCI Express User Guide.</i>
ds801_s6_pcie.pdf	<i>LogiCORE™ IP Spartan-6 FPGA Integrated Endpoint Block for PCI Express Data Sheet.</i>

[Back to Top](#)

<component name>/example_design

The example design directory contains the example design files provided with the core.

Table 3-4: Example Design Directory

Name	Description
<project_dir>/<component_name>/example_design	
xilinx_pcie_1_lane_ep_<device>.ucf	Example design UCF. Filename varies by part, package, and speed grade.
xilinx_pcie_1_1_ep_s6.v[hd]	Top-level PIO example design files for 1-lane cores.
pcie_app_s6.v[hd] PIO_EP_MEM.v[hd] PIO.v[hd] PIO_EP.v[hd] PIO_EP_MEM_ACCESS.v[hd] PIO_TO_CTRL.v[hd] PIO_32.v[hd] PIO_32_RX_ENGINE.v[hd] PIO_32_TX_ENGINE.v[hd]	PIO example design files.

[Back to Top](#)

<component name>/implement

The `implement` directory contains the core implementation script files.

Table 3-5: Implement Directory

Name	Description
<project_dir>/<component_name>/implement	
xst.scr	XST synthesis script.
implement.bat implement.sh	DOS and Linux implementation scripts.
synplify.prj	Synplify synthesis script.
xst.prj	XST project file for the example design.

[Back to Top](#)

implement/results

The `results` directory is created by the `implement` script, after which the `implement` script results are placed in the `results` directory.

Table 3-6: Results Directory

Name	Description
<project_dir>/<component_name>/implement/results	
Implement script result files.	

[Back to Top](#)

<component name>/simulation simulation/dsport

The dsport directory contains the Root Port Bus Functional model files provided with the core.

Table 3-7: dsport Directory

Name	Description
<project_dir>/<component_name>/simulation/dsport gtx_drp_chanalign_fix_3752_v6.v[hd] gtx_rx_valid_filter_v6.v[hd] gtx_tx_sync_rate_v6.v[hd] gtx_wrapper_v6.v[hd] pci_exp_usrapp_cfg.v[hd] pci_exp_usrapp_com.v pci_exp_usrapp_pl.v[hd] pci_exp_usrapp_rx.v[hd] pci_exp_usrapp_tx.v[hd] pcie_2_0_rport_v6.v[hd] pcie_2_0_v6_rp.v[hd] pcie_bram_top_v6.v[hd] pcie_bram_v6.v[hd] pcie_brams_v6.v[hd] pcie_clocking_v6.v[hd] pcie_gtx_v6.v[hd] pcie_pipe_lane_v6.v[hd] pcie_pipe_misc_v6.v[hd] pcie_pipe_v6.v[hd] pcie_reset_delay_v6.v[hd] pcie_upconfig_fix_3451_v6.v[hd] test_interface.vhd xilinx_pcie_2_0_rport_v6.v[hd]	Root port model files.

[Back to Top](#)

simulation/functional

The functional directory contains functional simulation scripts provided with the core.

Table 3-8: Functional Directory

Name	Description
<project_dir>/<component_name>/simulation/functional	
board_common.v	Contains test bench definitions.
board.f	List of files for RTL simulations.
board.v[hd]	Top-level simulation module.
isim_cmd.tcl	Simulation helper script for ISim.
simulate_isim.bat/simulate_isim.sh	Simulation scripts for ISIM DOS/UNIX.
simulate_mti.do	Simulation script for ModelSim.

Table 3-8: Functional Directory (Cont'd)

Name	Description
simulate_ncsim.sh	Simulation script for Cadence IUS.
simulate_vcs.sh	Simulation script for VCS.
sys_clk_gen_ds.v[hd]	System differential clock source.
sys_clk_gen.v[hd]	System clock source.
wave.{do, sv, tcl, wcfg}	Waveform setup scripts.

[Back to Top](#)

simulation/tests

The tests directory contains test definitions for the example test bench.

Table 3-9: Tests Directory

Name	Description
<project_dir>/<component_name>/simulation/tests	
tests.v[hd]	Test definitions for example test bench.

[Back to Top](#)

<component name>/source

This directory contains the source files for the core.

Table 3-10: Source Directory

Name	Description
<project_dir>/<component_name>/source	
<component name>.v[hd]	Verilog or VHDL top-level wrapper, which instantiates the Endpoint block, block RAMs, GTP transceiver, and clocking resources.
gtpa1_dual_wrapper_tile.v[hd] gtpa1_dual_wrapper.v[hd]	Wrapper for the GTPA1, which configures the transceiver and presents the interfaces required for use with the integrated Endpoint block.
pcie_bram_top_s6.v[hd] pcie_brams_s6.v[hd] pcie_bram_s6.v[hd]	Configures and instantiates block RAMs for use with the integrated Endpoint block.

[Back to Top](#)

Generating and Customizing the Core

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express core is a fully configurable and highly customizable solution. The integrated Endpoint block is customized using the CORE Generator software GUI.

Note: The screen captures in this chapter are conceptual representatives of their subjects and provide general information only. For the latest information, see the CORE Generator tool.

Customizing the Core through the CORE Generator Software

The CORE Generator software GUI for the Spartan-6 FPGA Integrated Endpoint Block for PCI Express consists of nine screens:

- Screen 1: [Basic Parameter Settings](#)
- Screen 2: [Base Address Registers](#)
- Screen 3: [PCI Registers](#)
- Screens 4 and 5: [Configuration Register Settings](#)
- Screen 6: [Interrupt Capabilities](#)
- Screen 7: [Power Management Registers](#)
- Screen 8: [PCI Express Extended Capabilities](#)
- Screen 9: [Advanced Settings](#)

Basic Parameter Settings

The initial customization screen shown in [Figure 4-1](#) is used to define the basic parameters for the core, including the component name, lane width and link speed.

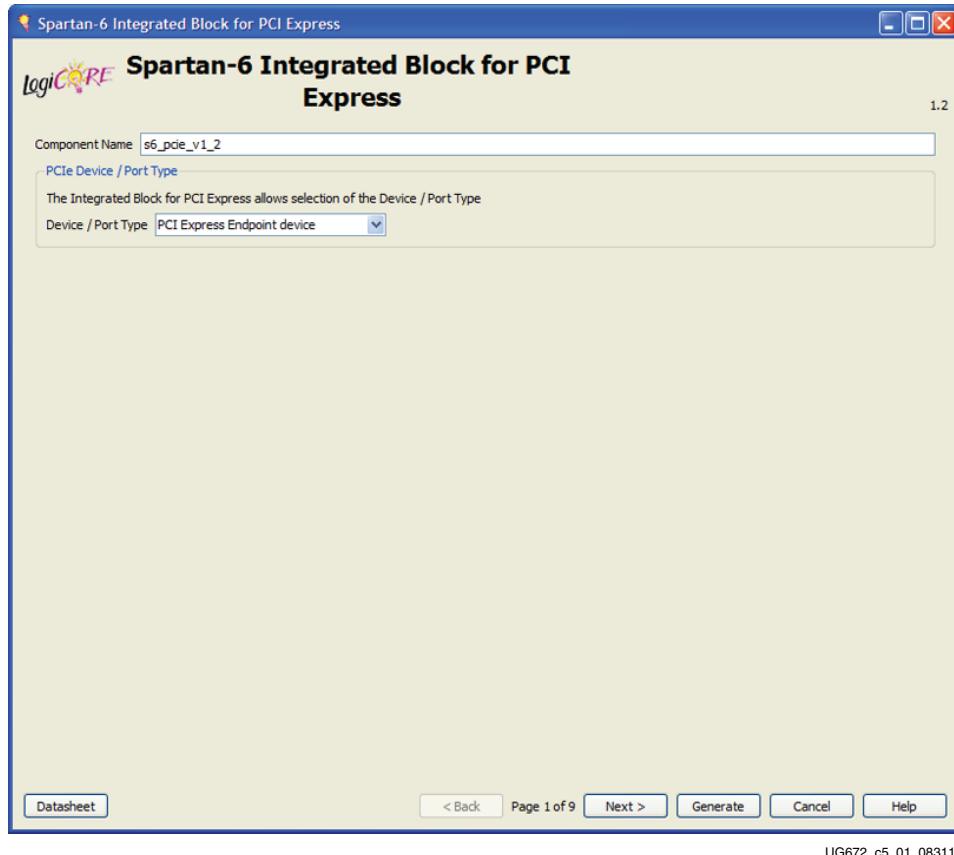


Figure 4-1: Screen 1: Integrated Endpoint Block for PCI Express Parameters

Component Name

Base name of the output files generated for the core. The name must begin with a letter and can be composed of these characters: a to z, 0 to 9, and “_.”

PCIe Device / Port Type

- **Device Port Type:** Indicates the PCI Express logical device type.

Base Address Registers

The Base Address Register (BAR) screen shown in [Figure 4-2](#) lets the user set the base address register space. Each Bar (0 through 5) represents a 32-bit parameter.

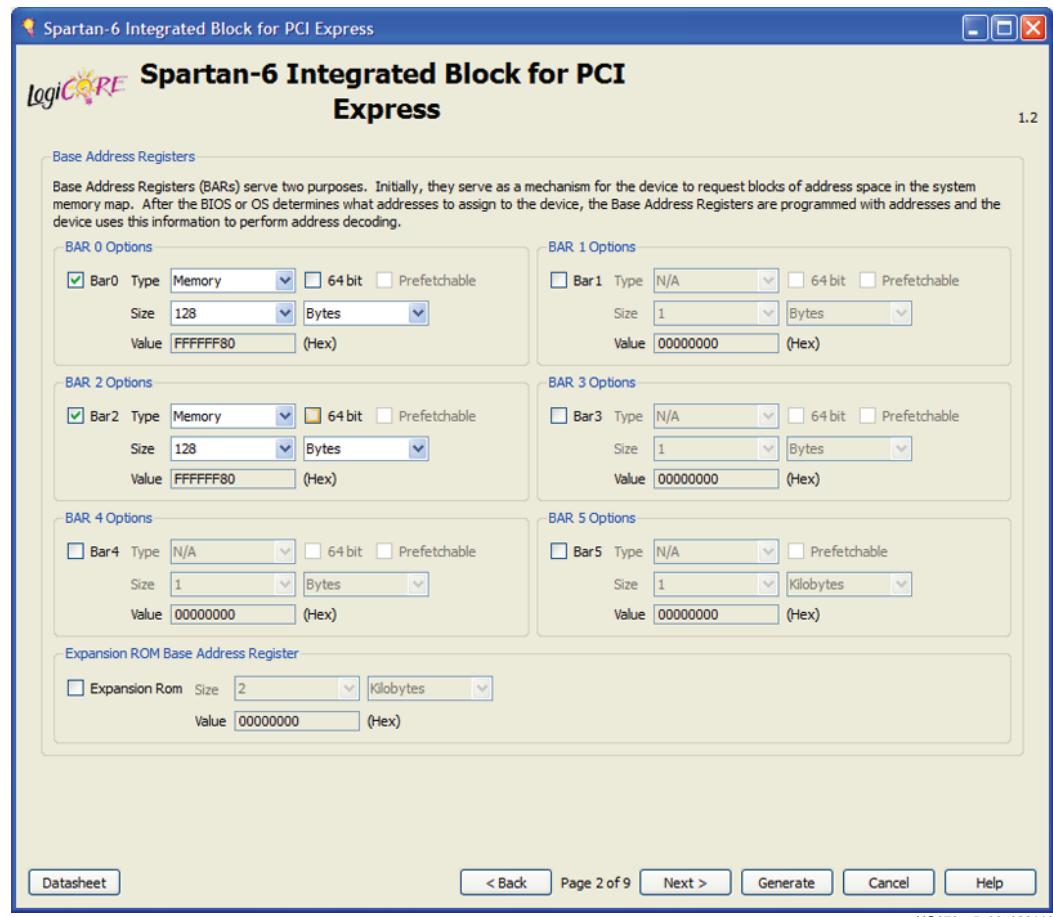


Figure 4-2: Screen 2: BAR Options

Base Address Register Overview

The Endpoint for PCIe supports up to six 32-bit Base Address Registers (BARs) or three 64-bit BARs, and the Expansion ROM BAR. BARs can be one of two sizes:

- **32-bit BARs:** The address space can be as small as 128 bytes for Memory or 16 bytes for I/O, or as large as 2 gigabytes. Used for Memory to I/O.
- **64-bit BARs:** The address space can be as small as 128 bytes or as large as 8 exabytes. Used for Memory only.

All BAR registers share these options:

- **Checkbox:** Click the checkbox to enable the BAR; deselect the checkbox to disable the BAR.
- **Type:** BARs can either be I/O or Memory.
 - *I/O:* I/O BARs can only be 32-bit; the Prefetchable option does not apply to I/O BARs.

- *Memory*: Memory BARs can be either 64-bit or 32-bit and can be prefetchable. When a BAR is set as 64 bits, it uses the next BAR for the extended address space and makes the next BAR inaccessible to the user
- *Size*
 - *Memory*: When Memory and 64-bit are not selected, the size can range from 128 bytes to 2 gigabytes. When Memory and 64-bit are selected, the size can range between 128 bytes and 8 exabytes.
 - *I/O*: When selected, the size can range from 16 bytes to 2 gigabytes.
- **Prefetchable**: Identifies the ability of the memory space to be prefetched.
- **Value**: The value assigned to the BAR based on the current selections.

For more information about managing the Base Address Register settings, see [Managing Base Address Register Settings](#).

Expansion ROM Base Address Register

If selected, the Expansion ROM is activated and can be a value from 2 KB to 4 GB.

Managing Base Address Register Settings

Memory, I/O, Type, and Prefetchable settings are handled by setting the appropriate GUI settings for the desired base address register.

Memory or I/O settings indicate whether the address space is defined as memory or I/O. The base address register only responds to commands that access the specified address space. Generally, memory spaces less than 4Kbytes in size should be avoided. The minimum I/O space allowed is 16 bytes; use of I/O space should be avoided in all new designs.

Prefetchability is the ability of memory space to be prefetched. A memory space is prefetchable if there are no side effects on reads (that is, data is not destroyed by reading, as from a RAM). Byte write operations can be merged into a single doubleword write, when applicable.

When configuring the core as an Endpoint for PCIe (non-Legacy), 64-bit addressing must be supported for all BARs (except BAR5) that have the prefetchable bit set. 32-bit addressing is permitted for all BARs that do not have the prefetchable bit set. The prefetchable bit related requirement does not apply to a Legacy Endpoint. In either of the above cases (Endpoint for PCI Express or Legacy Endpoint), the minimum memory address range supported by a BAR is 128 bytes.

Disabling Unused Resources

For best results, disable unused base address registers to conserve system resources. A base address register is disabled by deselecting unused BARs in the GUI.

PCI Registers

The PCI Registers Screen shown in [Figure 4-3](#) is used to customize the IP initial values, class code and Cardbus CIS pointer information.

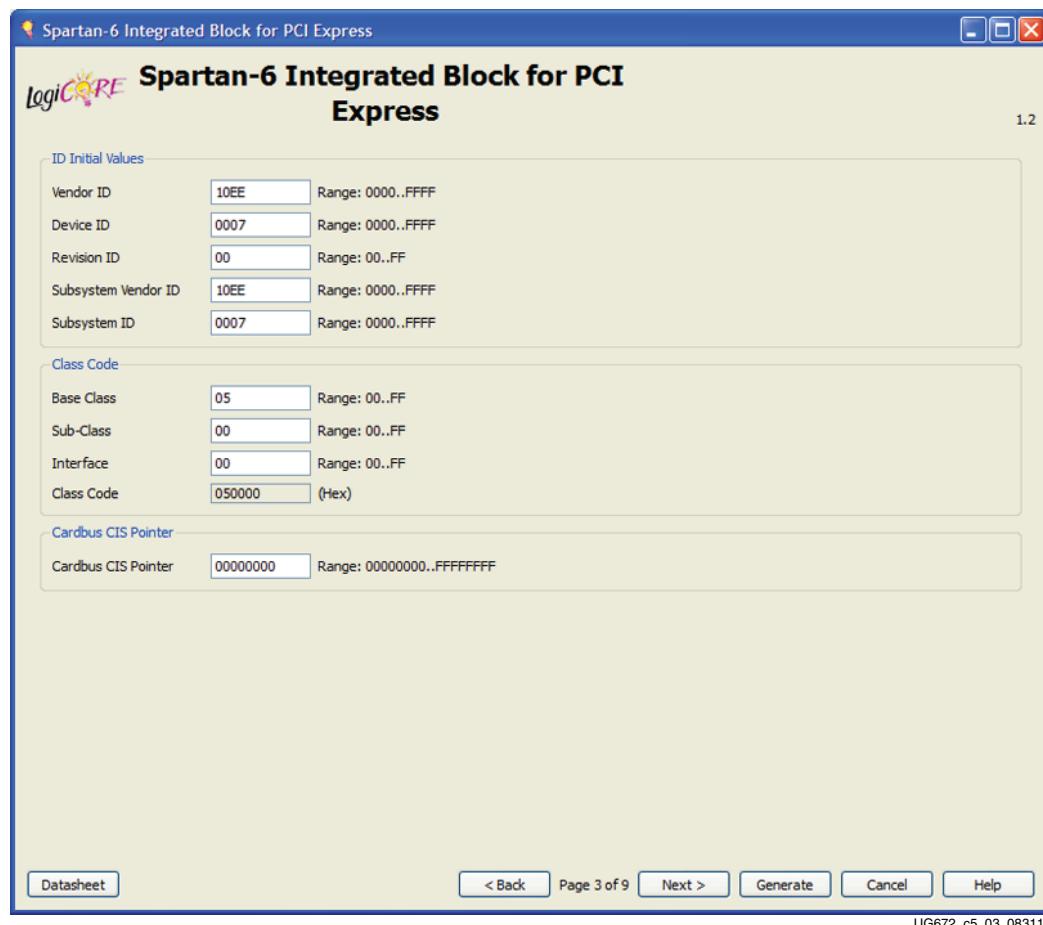


Figure 4-3: PCI Registers: Screen 3

ID Initial Values

- **Vendor ID:** Identifies the manufacturer of the device or application. Valid identifiers are assigned by the PCI Special Interest Group to guarantee that each identifier is unique. The default value, 10EEh, is the Vendor ID for Xilinx. Enter a vendor identification number here. FFFFh is reserved.
- **Device ID:** A unique identifier for the application; the default value is 0007h. This field can be any value; change this value for the application.
- **Revision ID:** Indicates the revision of the device or application; an extension of the Device ID. The default value is 00h; enter values appropriate for the application.
- **Subsystem Vendor ID:** Further qualifies the manufacturer of the device or application. Enter a Subsystem Vendor ID here; the default value is 10EE. Typically, this value is the same as Vendor ID. Setting the value to 0000h can cause compliance testing issues.

- **Subsystem ID:** Further qualifies the manufacturer of the device or application. This value is typically the same as the Device ID; default value is 0007h. Setting the value to 0000h can cause compliance testing issues.

Class Code

The Class Code identifies the general function of a device, and is divided into three byte-size fields:

- **Base Class:** Broadly identifies the type of function performed by the device.
- **Sub-Class:** More specifically identifies the device function.
- **Interface:** Defines a specific register-level programming interface, if any, allowing device-independent software to interface with the device.

Class code encoding can be found at www.pcisig.com.

Cardbus CIS Pointer

Used in cardbus systems and points to the Card Information Structure for the cardbus card. If this field is non-zero, an appropriate Card Information Structure must exist in the correct location. The default value is 0000_0000h; value range is 0000_0000h through FFFF_FFFFh.

Configuration Register Settings

The Configuration Registers screens shown in [Figure 4-4](#) and [Figure 4-5](#) show the options for the Device Capabilities and Registers, the Block RAM Configuration Options, the Link Capabilities Register, and the Link Status Register.

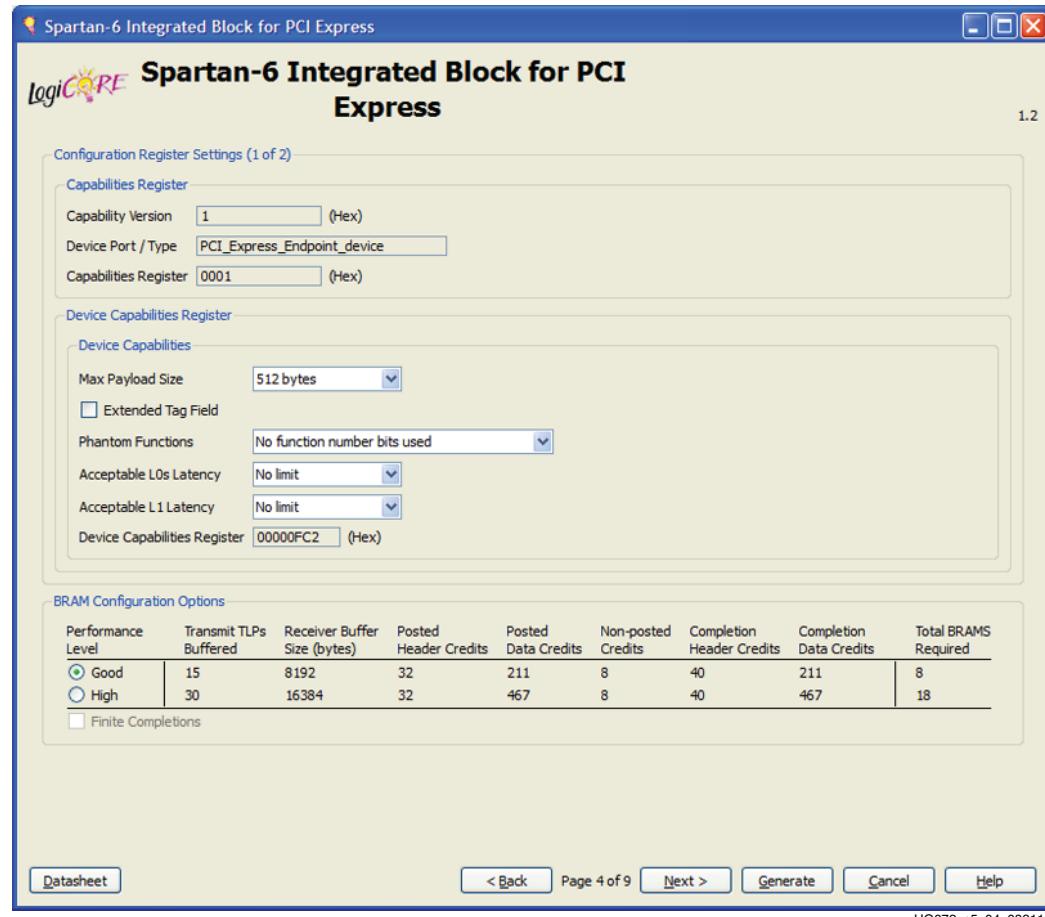


Figure 4-4: Screen 4: Configuration Settings

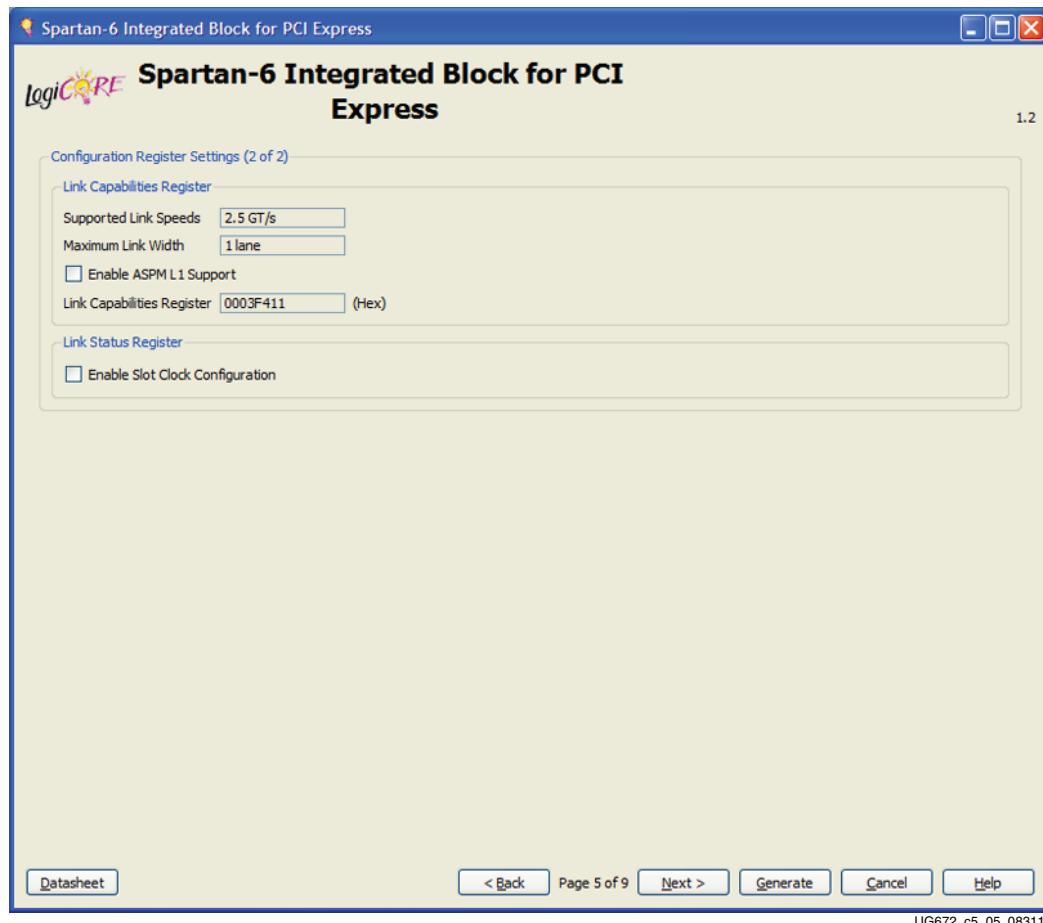


Figure 4-5: Screen 5: Configuration Settings

Capabilities Register

- **Capability Version:** Indicates PCI-SIG defined PCI Express capability structure version number; this value cannot be changed.
- **Device Port Type:** Indicates the PCI Express logical device type.
- **Capabilities Register:** Displays the value of the Capabilities register presented by the integrated Endpoint block, and is not editable.

Device Capabilities Register

- **Max Payload Size:** Indicates the maximum payload size that the device/function can support for TLPs.
- **Extended Tag Field:** Indicates the maximum supported size of the Tag field as a Requester. When selected, indicates 8-bit Tag field support. When deselected, indicates 5-bit Tag field support.
- **Phantom Functions:** Indicates the support for use of unclaimed function numbers to extend the number of outstanding transactions allowed by logically combining unclaimed function numbers (called Phantom Functions) with the Tag identifier. See Section 2.2.6.2 of the PCI Express Base Specification version 1.1 for a description of

Tag Extensions. This field indicates the number of most significant bits of the function number portion of Requester ID that are logically combined with the Tag identifier.

- **Acceptable L0s Latency:** Indicates the acceptable total latency that an Endpoint can withstand due to the transition from L0s state to the L0 state.
- **Acceptable L1 Latency:** Indicates the acceptable latency that an Endpoint can withstand due to the transition from L1 state to the L0 state.
- **Device Capabilities Register:** Displays the value of the Device Capabilities register presented by the integrated Endpoint block and is not editable.

Block RAM Configuration Options

- **Performance Level:** Selects the Performance Level settings, which determines the Receiver and Transmitter Sizes. The table displayed specifies the Receiver and Transmitter settings - number of TLPs buffered in the Transmitter, the Receiver Size, the credits advertised by the core to the Link Partner and the number of block RAMs required for the configuration, corresponding to the Max Payload Size selected, for each of the Performance Level options.
- **Finite Completions:** If selected, causes the device to advertise to the Link Partner the actual amount of space available for completions in the receiver. For an Endpoint, this is not compliant to the *PCI Express Base Specification version 1.1*, as endpoints are required to advertise an infinite amount of completion space. Finite completions are not supported in this release of the core.

Link Capabilities Register

This section is used to set the Link Capabilities register.

- **Maximum Link Speed:** Indicates the maximum link speed of the given PCI Express Link. This value is set to 2.5 Gb/s and is not editable.
- **Maximum Link Width:** This value is set to 1 lane.
- **Enable ASPM L1 Support:** Indicates the level of ASPM supported on the given PCI Express Link. L0s is always supported by the integrated Endpoint block core; L1 support is optional and is enabled if this box is checked.
- **Link Capabilities Register:** Displays the value of the Link Capabilities register presented by the Endpoint and is not editable.

Link Status Register

- **Enable Slot Clock Configuration:** Indicates that the Endpoint uses the platform-provided physical reference clock available on the connector. Must be cleared if the Endpoint uses an independent reference clock.

Interrupt Capabilities

The Interrupt Settings screen shown in Figure 4-6 sets the Legacy Interrupt Settings and MSI Capabilities.

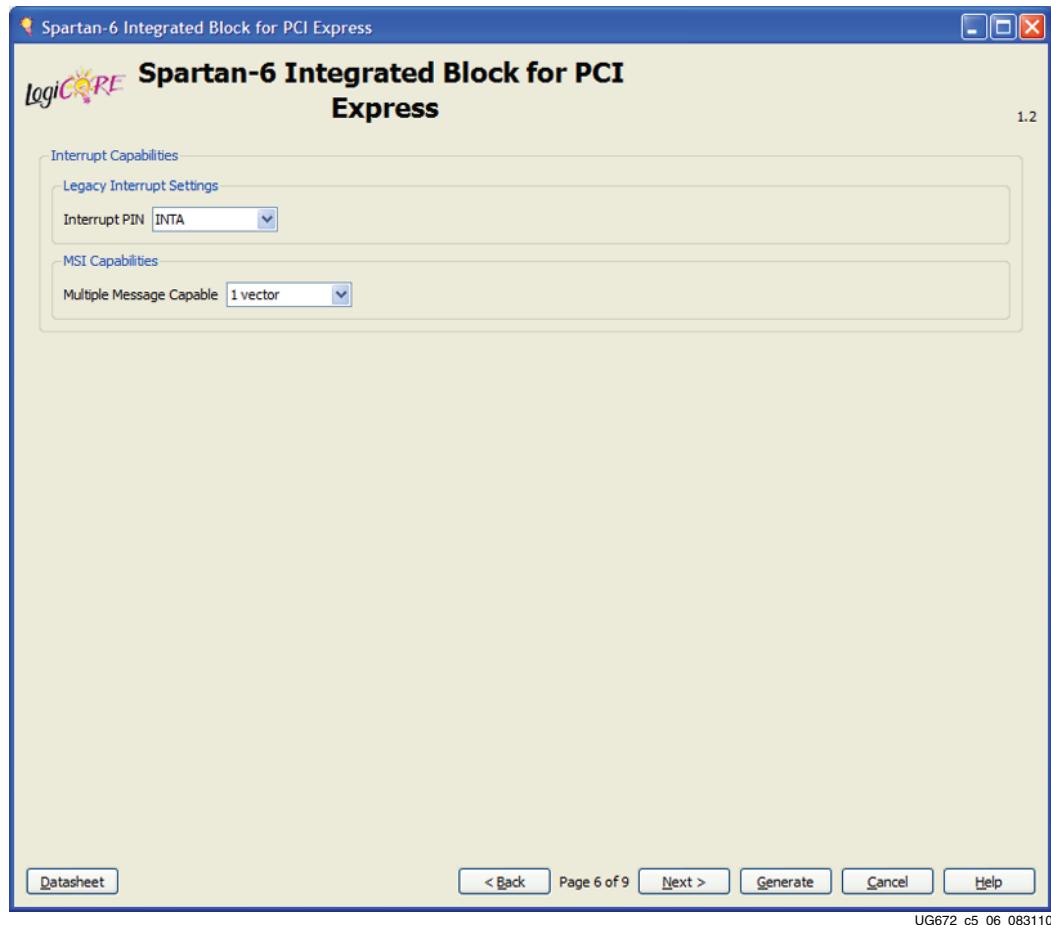


Figure 4-6: Interrupt Capabilities: Screen 6

Legacy Interrupt Settings

- **Interrupt PIN:** Indicates the mapping for Legacy Interrupt messages. A setting of "None" indicates no Legacy Interrupts are used.

MSI Capabilities

- **Multiple Message Capable:** Selects the number of MSI vectors to request from the Root Complex.

Power Management Registers

The Power Management Registers screen shown in [Figure 4-7](#) includes settings for the Power Management Registers, power consumption and power dissipation options.

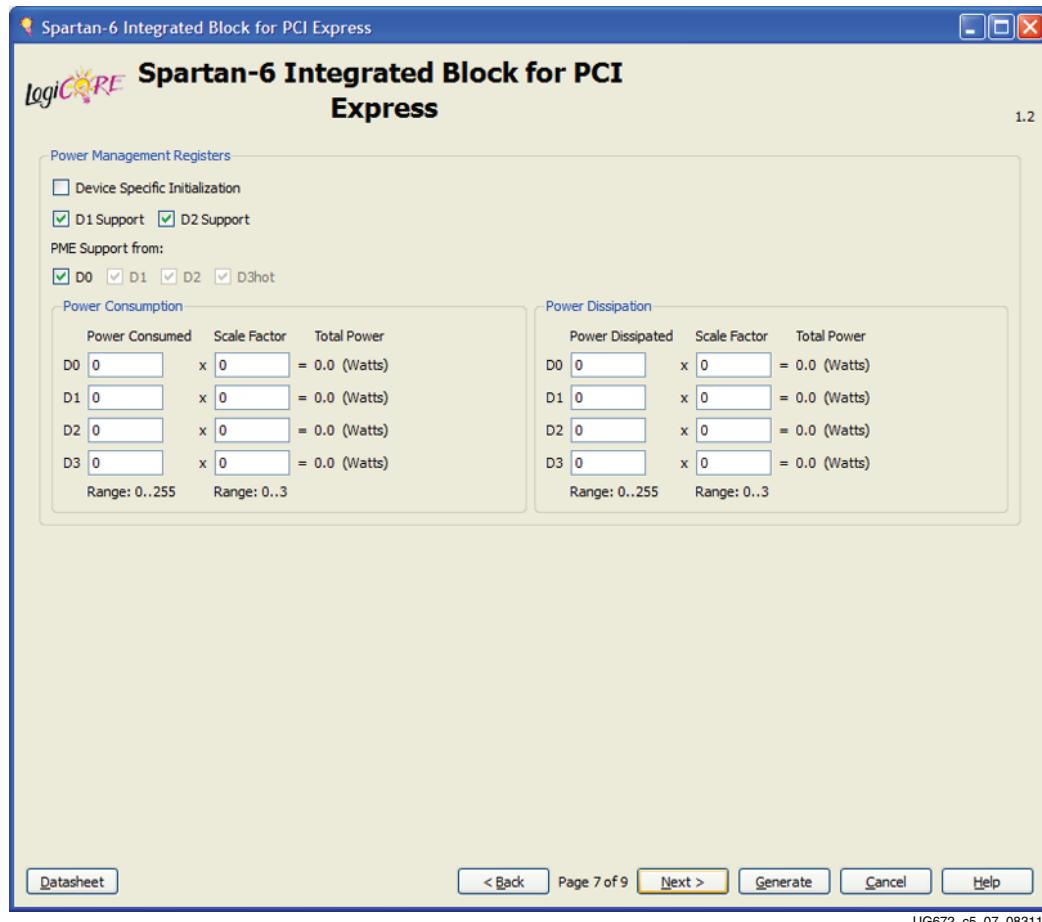


Figure 4-7: Power Management Registers: Screen 7

Power Management Registers

- Device Specific Initialization:** This bit indicates whether special initialization of this function is required (beyond the standard PCI configuration header) before the generic class device driver is able to use it. When selected, this option indicates that the function requires a device specific initialization sequence following transition to the D0 uninitialized state. See section 3.2.3 of the *PCI Bus Power Management Interface Specification Revision 1.2*.
- D1 Support:** When selected, this option indicates that the function supports the D1 Power Management State. See section 3.2.3 of the *PCI Bus Power Management Interface Specification Revision 1.2*.
- D2 Support:** When selected, this option indicates that the function supports the D2 Power Management State. See section 3.2.3 of the *PCI Bus Power Management Interface Specification Revision 1.2*.

- **PME Support From:** When this option is selected, indicates the power states in which the function can assert PME#. See section 3.2.3 of the *PCI Bus Power Management Interface Specification Revision 1.2*.

Power Consumption

For information about power consumption, see section 3.2.6 of the PCI Bus Power Management Interface Specification Revision 1.2.

Power Dissipated

For information about power dissipation, see section 3.2.6 of the PCI Bus Power Management Interface Specification Revision 1.2.

PCI Express Extended Capabilities

The PCIe Extended Capabilities screen shown in [Figure 4-8](#) includes settings for Device Serial Number Capability and optional user-defined Configuration capabilities.

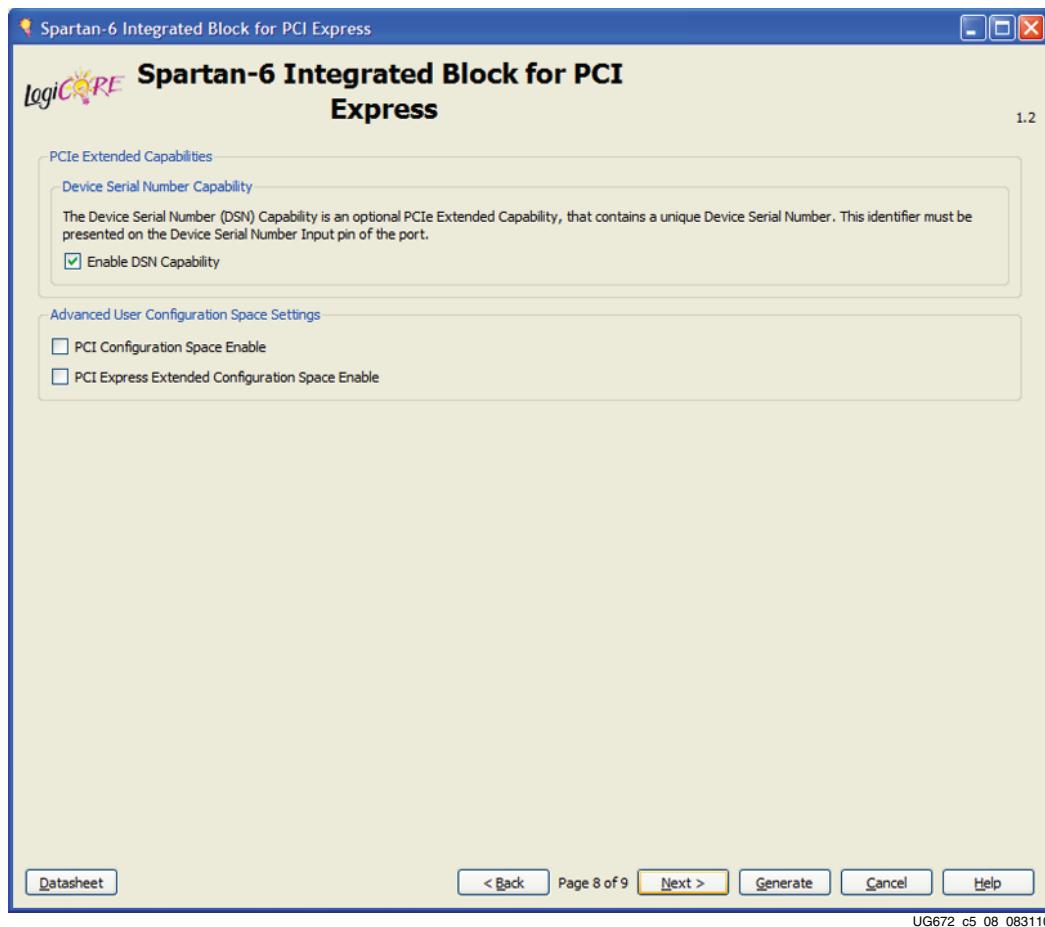


Figure 4-8: Screen 8: PCIe Extended Capabilities

Device Serial Number Capability

- **Device Serial Number Capability:** An optional PCIe Extended Capability containing a unique Device Serial Number. If enabled, the core presents the Device Serial

Number Capability using the value presented on the Device Serial Number input pin of the port. If disabled, no Device Serial Number Extended Capability is presented.

User Defined Configuration Capabilities

- **PCI Configuration Space Enable:** Allows the user application to add/implement PCI Legacy capability registers. This option should be selected if the user application implements a legacy capability configuration space. This option enables the routing of Configuration Requests to addresses outside the built-in PCI-Compatible Configuration Space address range to the AXI4-Stream interface.
- **PCI Express Extended Configuration Space Enable:** Allows the user application to add/implement PCI Express Extended capability registers. This option should be selected if the user application implements such an extended capability configuration space. This enables the routing of Configuration Requests to addresses outside the built-in PCI Express Extended Configuration Space address range to the user application.

Advanced Settings

The Advanced Settings screen shown in [Figure 4-9](#) includes settings for Transaction Layer, Physical Layer, Reference Clock Frequency and Xilinx Reference Boards options.

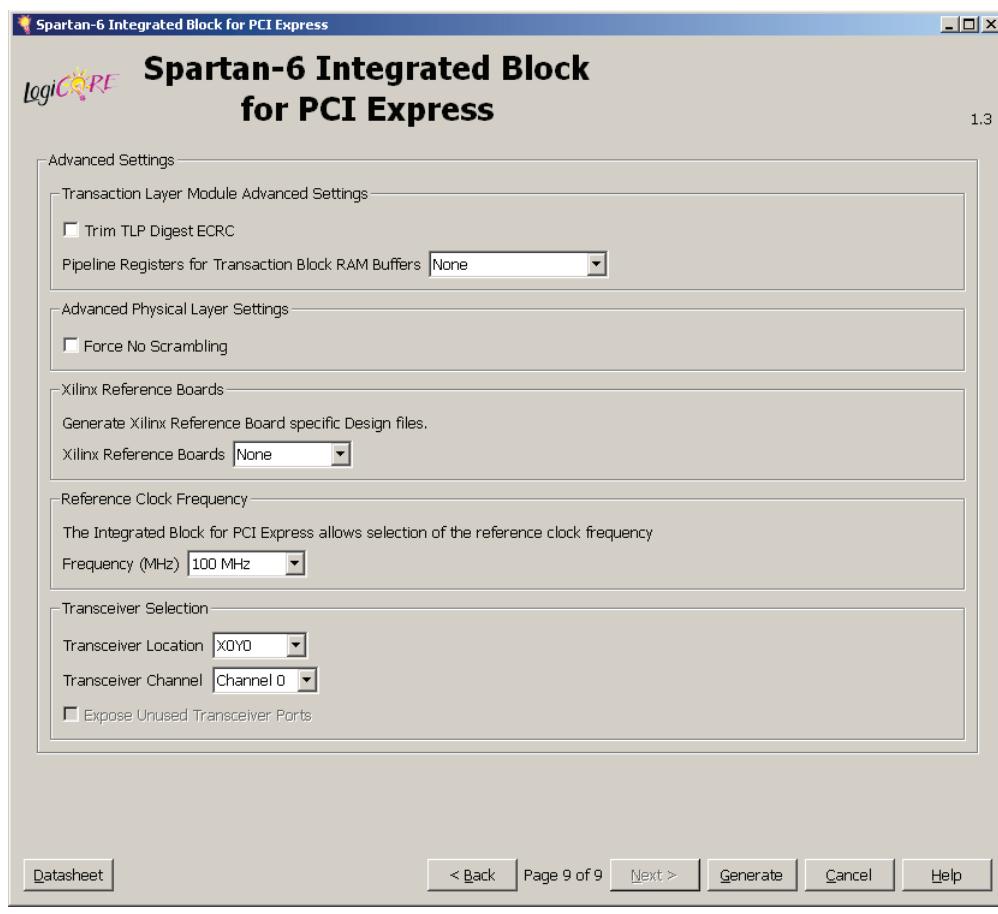


Figure 4-9: Screen 9: Advanced Settings 1

Transaction Layer Module

- **Trim TLP Digest ECRC:** Causes the core to trim any TLP digest from an inbound packet and clear the TLP Digest bit in the TLP header, before presenting it to the user.
- **Pipeline Registers for Transaction Block RAM Buffers:** Selects the Pipeline registers enabled for the Transaction Buffers. Pipeline registers can be enabled on either the Write path or both the Read and Write paths of the Transaction Block RAM buffers.

Advanced Physical Layer

- **Force No Scrambling:** Used for diagnostic purposes only and should never be enabled in a working design. Setting this bit results in the data scramblers being turned off so that the serial data stream can be analyzed.

Xilinx Reference Boards

Selecting this option enables the generation of Xilinx Reference Board specific design files. Selecting the SP605 board configures the Reference Clock Frequency, Transceiver Location and Transceiver Channel corresponding with the PCI Express edge connector on the reference board. It also sets the corresponding pin locations in the UCF file. The user must select the correct part/package combination when setting up the project to generate Xilinx reference board specific design files.

Reference Clock Frequency

Selects the frequency of the reference clock provided on sys_clk. For important information about clocking the Spartan-6 FPGA Integrated Endpoint Block for PCI Express, see [Clocking and Reset of the Integrated Endpoint Block Core, page 102](#).

Transceiver Selection

- **Transceiver Location:** Selects the GTPA1_DUAL location for the PCI Express link.
- **Transceiver Channel:** Selects the channel within the GTPA1_DUAL.

Designing with the Core

This chapter provides design instructions for the Spartan®-6 FPGA Integrated Endpoint Block for PCI Express® user interface and assumes knowledge of the PCI Express Transaction Layer Packet (TLP) header fields. Header fields are defined in *PCI Express Base Specification v1.1*, Chapter 2, Transaction Layer Specification.

This chapter includes the following design guidelines:

- [Transmitting Outbound Packets](#)
- [Receiving Inbound Packets](#)
- [Design with Configuration Space Registers and Configuration Interface](#)
- [Additional Packet Handling Requirements](#)
- [Power Management](#)
- [Generating Interrupt Requests](#)
- [Clocking and Reset of the Integrated Endpoint Block Core](#)

TLP Format on the AXI4-Stream Interface

Data is transmitted and received in Big-Endian order as required by the *PCI Express Base Specification*. See Chapter 2 of the *PCI Express Base Specification* for detailed information about TLP packet ordering. [Figure 5-1](#) represents a typical 32-bit addressable Memory Write Request TLP (as illustrated in Chapter 2 of the specification).

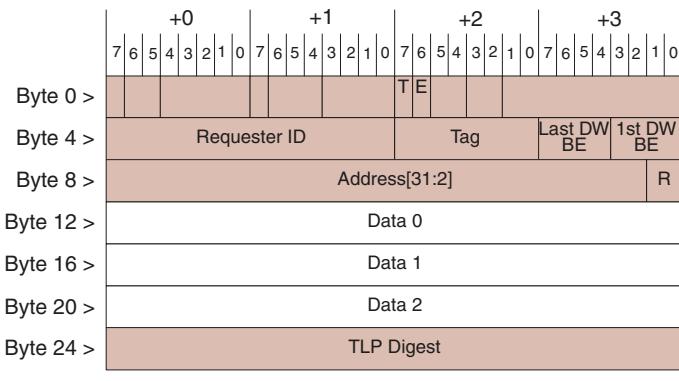


Figure 5-1: PCI Express Base Specification Byte Order

When using the 32-bit AXI4-Stream interface, packets are arranged on the 32-bit datapath in the same order as shown in [Figure 5-1](#). PCIe Byte 0 of the packet appears on s_axis_tx_tdata[31:24] (outbound) or m_axis_rx_tdata[31:24] (inbound) of the first DWORD, byte 1 on s_axis_tx_tdata[23:16] or m_axis_rx_tdata[23:16], and so forth. Byte 4 of the packet then appears on s_axis_tx_tdata[31:24] or m_axis_rx_tdata[31:24] of the second DWORD. The Header section of the packet consists of either three or four DWORDs, determined by the TLP format and type as described in section 2.2 of the *PCI Express Base Specification*.

Packets sent to the core for transmission must follow the formatting rules for Transaction Layer Packets (TLPs) as specified in Chapter 2 of the *PCI Express Base Specification*. The user application is responsible for ensuring its packets' validity. The core does not check that a packet is correctly formed and this can result in transferring a malformed TLP. The exact fields of a given TLP vary depending on the type of packet being transmitted.

The core allows the user application to add an extra level of error checking by using the optional TLP Digest field in the TLP header. The presence of a TLP Digest or ECRC is indicated by the value of TD field in the TLP Header section. When TD=1, a correctly computed CRC32 remainder DWORD is expected to be presented as the last DWORD of the packet. The CRC32 remainder DWORD is not included in the length field of the TLP header. The user application must calculate and present the TLP Digest as part of the packet when transmitting packets. Upon receiving packets with a TLP Digest present, the user application must check the validity of the CRC32 based on the contents of the packet. The core does not check the TLP Digest for incoming packets.

The *PCI Express Base Specification* requires Advanced Error Reporting (AER) capability when implementing ECRC. Although the integrated Endpoint block does not support AER, users can still implement ECRC for custom solutions that do not require *PCI Express Base Specification* Compliance.

Transmitting Outbound Packets

Basic TLP Transmit Operation

The Endpoint for PCIe automatically transmits the following types of packets:

- Completions to a remote device in response to Configuration Space requests.
- Error-message responses to inbound requests malformed or unrecognized by the core.

Note: Certain unrecognized requests, for example, unexpected completions, can only be detected by the user application, which is responsible for generating the appropriate response.

The user application is responsible for constructing these types of outbound packets:

- Memory and I/O Requests to remote devices.
- Completions in response to requests to the user application, for example, a Memory Read Request.
- Completions in response to user-implemented Configuration Space requests when enabled. These requests include PCI Legacy capability registers beyond address BF^h and PCI Express extended capability registers beyond address FF^h.

Note: For important information about accessing user-implemented Configuration Space while in a low-power state, see [Power Management, page 96](#).

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express core notifies the user application of pending internally generated TLPs that will arbitrate for the transmit datapath by asserting tx_cfg_req (1b). The user application can choose to give priority to core-generated TLPs by driving tx_cfg_gnt asserted (1b) permanently, without regard to tx_cfg_req. Doing so prevents user-application-generated TLPs from being transmitted when a core-generated TLP is pending. Alternatively, the user application can reserve priority for a user-application-generated TLP over core-generated TLPs, by holding tx_cfg_gnt deasserted (0b) until the user transaction is complete. When the user transaction is complete, the user application can assert tx_cfg_gnt (1b) for at least one clock cycle to allow the pending core-generated TLP to be transmitted. Users must not delay asserting tx_cfg_gnt indefinitely, as this might cause a completion time-out in the Requester. See the *PCI Express Base Specification* for more information on the Completion Timeout Mechanism.

[Table 2-9, page 23](#) defines the transmit-direction AXI4-Stream interface signals. To transmit a TLP, the user application must perform the following sequence of events on the transmit AXI4-Stream interface:

1. The user application logic asserts s_axis_tx_tvalid and presents the first TLP DWORD on s_axis_tx_tdata[31:0].
2. The user application asserts s_axis_tx_tvalid and presents the remainder of the TLP DWORDs on s_axis_tx_tdata[31:0] for subsequent clock cycles (for which the core asserts s_axis_tx_tready).
3. The user application asserts s_axis_tx_tvalid and s_axis_tx_tlast together with the last DWORD of data.
4. At the next clock cycle, the user application deasserts s_axis_tx_tvalid to signal the end of valid transfers on s_axis_tx_tdata[31:0].

[Figure 5-2](#) illustrates a 3-DW TLP header without a data payload; an example is a 32-bit addressable Memory Read request.

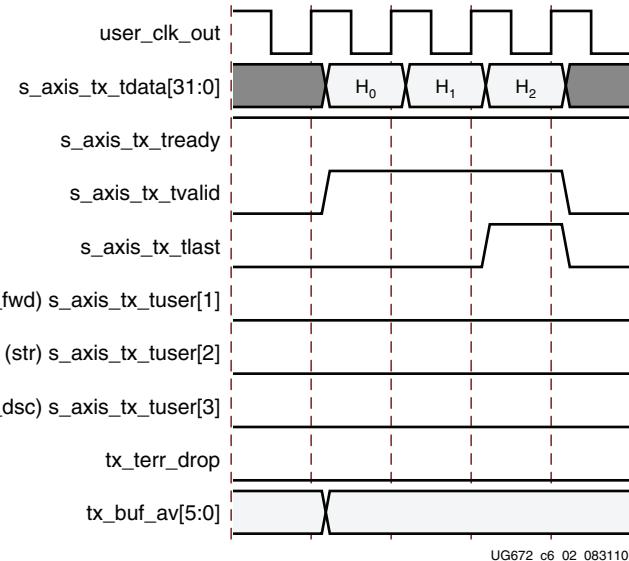


Figure 5-2: TLP 3-DW Header without Payload

[Figure 5-3](#) illustrates a 4-DW TLP header without a data payload; an example is a 64-bit addressable Memory Read request.

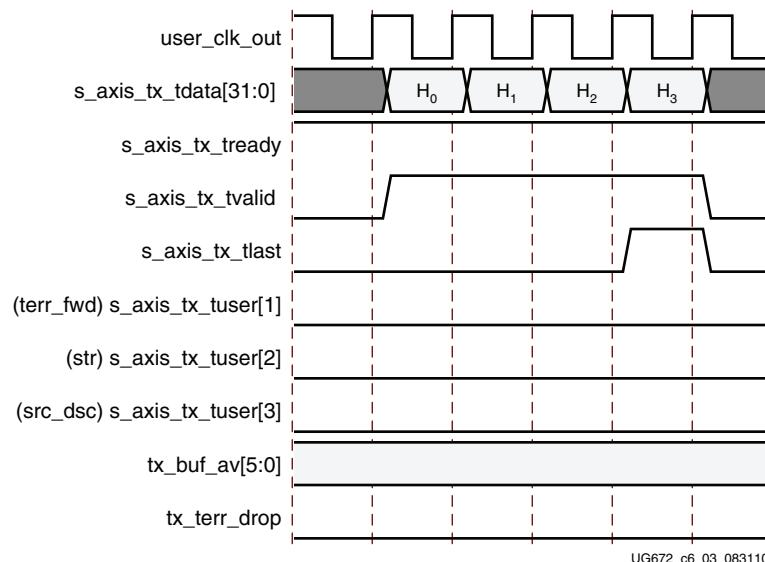


Figure 5-3: TLP 4-DW Header without Payload

[Figure 5-4](#) illustrates a 3-DW TLP header with a data payload; an example is a 32-bit addressable Memory Write request.

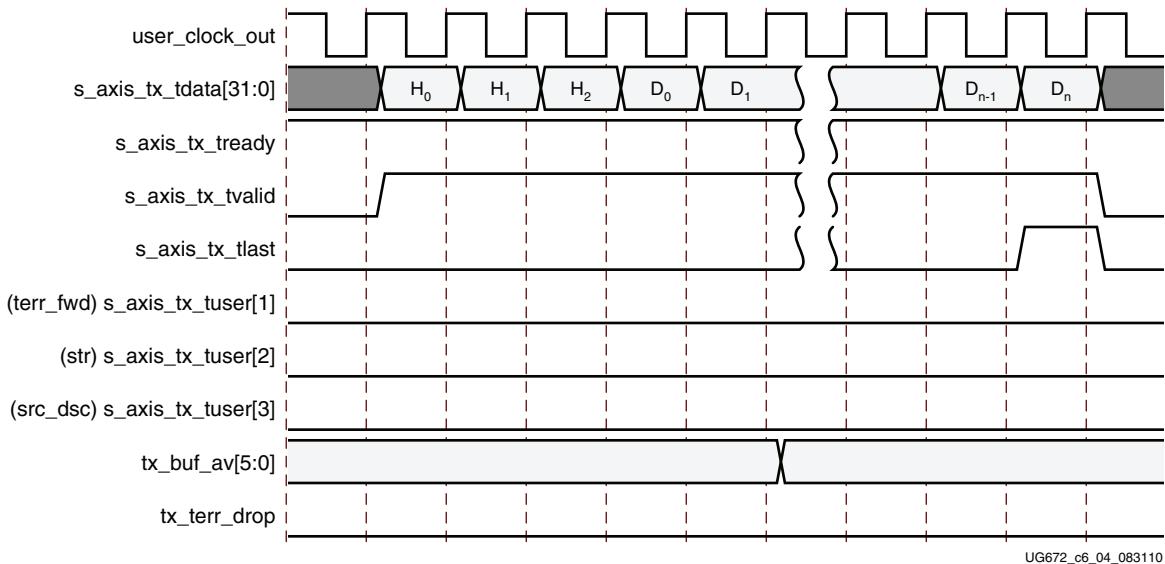


Figure 5-4: TLP with 3-DW Header with Payload

[Figure 5-5](#) illustrates a 4-DW TLP header with a data payload; an example is a 32-bit addressable Memory Write request.

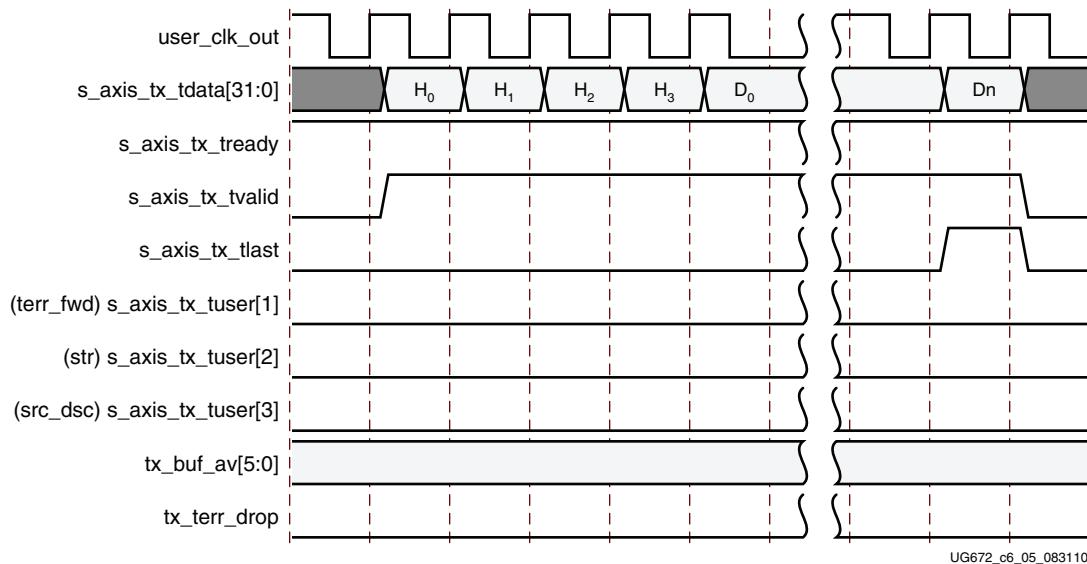


Figure 5-5: TLP with 4-DW Header with Payload

Presenting Back-to-Back Transactions on the Transmit Interface

The user application can present back-to-back TLPs on the transmit AXI4-Stream interface to maximize bandwidth utilization. [Figure 5-6](#) illustrates back-to-back TLPs presented on the transmit interface. The user application keeps s_axis_tx_tvalid asserted and presents a new TLP on the next clock cycle after asserting s_axis_tx_tlast for the previous TLP.

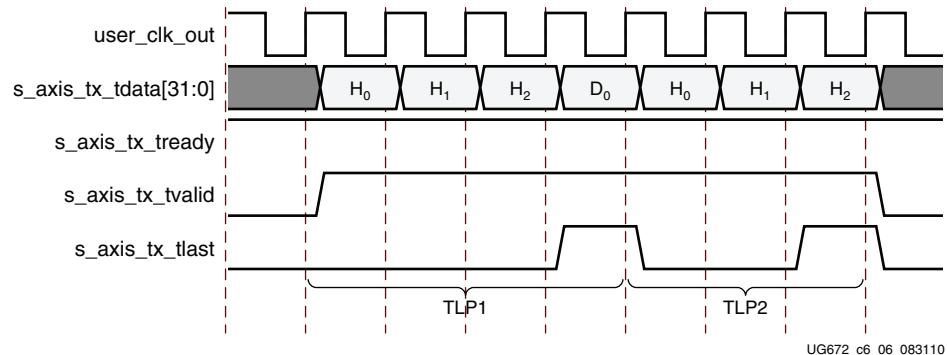


Figure 5-6: Back-to-Back Transaction on Transmit Interface

Source Throttling on the Transmit Datapath

The AXI4-Stream interface lets the user application throttle back if it has no data present on s_axis_tx_tdata[31:0]. When this condition occurs, the user application deasserts s_axis_tx_tvalid, which instructs the core AXI4-Stream interface to disregard data presented on s_axis_tx_tdata[31:0]. [Figure 5-7](#) illustrates the source throttling mechanism, where the user application does not have data to present every clock cycle, and for this reason must deassert s_axis_tx_tvalid during these cycles. The user application should not deassert s_axis_tx_tvalid during the middle of a transfer if str is asserted.

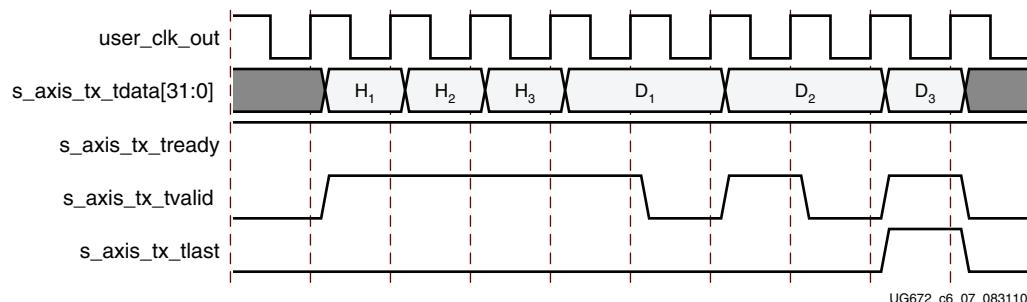


Figure 5-7: Source Throttling on the Transmit Datapath

Destination Throttling of the Transmit Datapath

The core AXI4-Stream interface throttles the user application if there is no space left for a new TLP in its transmit buffer pool. This can occur if the link partner is not processing incoming packets at a rate equal to or greater than the rate at which the user application is presenting TLPs. Figure 5-8 illustrates the deassertion of `s_axis_tx_tready` to throttle the user application when the core's internal transmit buffers are full. If the core needs to throttle the User Application, it does so after the current packet has completed. If another packet starts immediately after the current packet, the throttle occurs immediately after `tlast`.

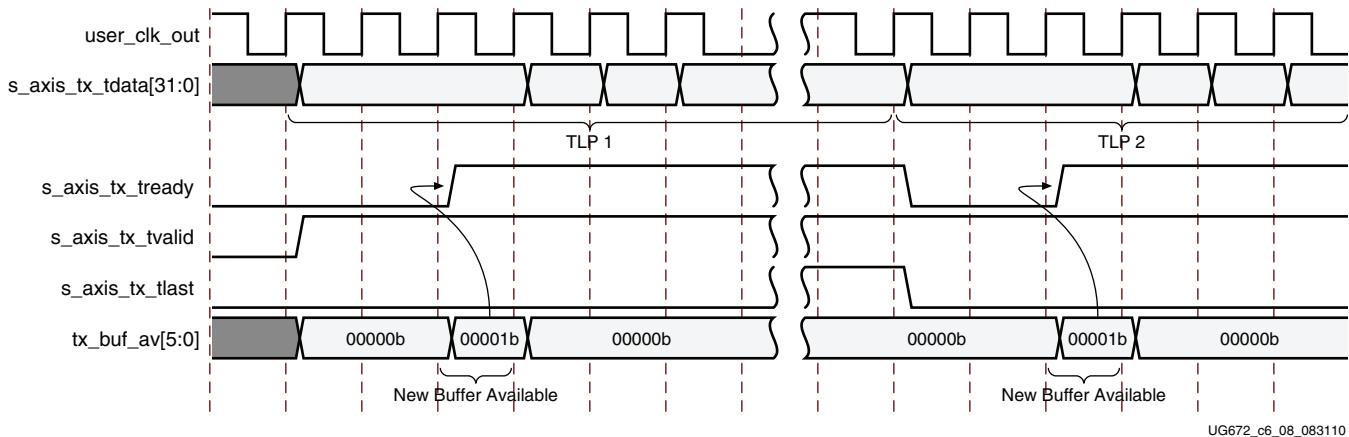


Figure 5-8: Destination Throttling of the Endpoint Transmit AXI4-Stream Interface

If the core transmit AXI4-Stream interface accepts the start of a TLP by asserting `s_axis_tx_tready`, it is guaranteed to accept the complete TLP with a size up to the value contained in the `Max_Payload_Size` field of the PCI Express Device Capability Register (offset `04H`). To stay compliant to the *PCI Express Base Specification*, users should not violate the `Max_Payload_Size` field of the PCI Express Device Control Register (offset `08H`). The core transmit AXI4-Stream interface deasserts `s_axis_tx_tready` only under these conditions:

- After it has accepted the TLP completely and has no buffer space available for a new TLP.
- When the core is transmitting an internally generated TLP (Completion TLP due to a Configuration Read or Write, error Message TLP or error response as requested by the user application on the `cfg_err` interface), after it has been granted use of the transmit datapath by the user application, by assertion of `tx_cfg_gnt`. The core subsequently asserts `s_axis_tx_tready` after transmitting the internally generated TLP.
- When the Power State field in Power Management Control/Status Register (offset `0x4`) of the PCI Power Management Capability Structure is changed to a non-D0 state. When this occurs, any ongoing TLP is accepted completely and `s_axis_tx_tready` is subsequently deasserted, disallowing the User Application from initiating any new transactions for the duration that the core is in the non-D0 power state.

On deassertion of `s_axis_tx_tready` by the core, the user application needs to hold all control and data signals until the core asserts `s_axis_tx_tready`.

Discontinuing Transmission of Transaction by Source

The core AXI4-Stream interface lets the user application terminate transmission of a TLP by asserting `src_dsc`. Both `s_axis_tx_tvalid` and `s_axis_tx_tready` must be asserted together with `src_dsc` for the TLP to be discontinued. The signal `src_dsc` must not be asserted at the beginning of a new packet. It can be asserted on any cycle after the first beat of a new packet has been accepted by the core up to and including the assertion of `s_axis_tx_tlast`. Asserting `src_dsc` has no effect if no TLP transaction is in progress on the transmit interface. [Figure 5-9](#) illustrates the user application discontinuing a packet using `src_dsc`. Asserting `src_dsc` with `s_axis_tx_tlast` is optional.

If streaming mode is not used (`str = 0b`) and the packet is discontinued, the packet is discarded before being transmitted on the serial link. If streaming mode is used (`str = 1b`), the packet is terminated with the EDB symbol on the serial link.

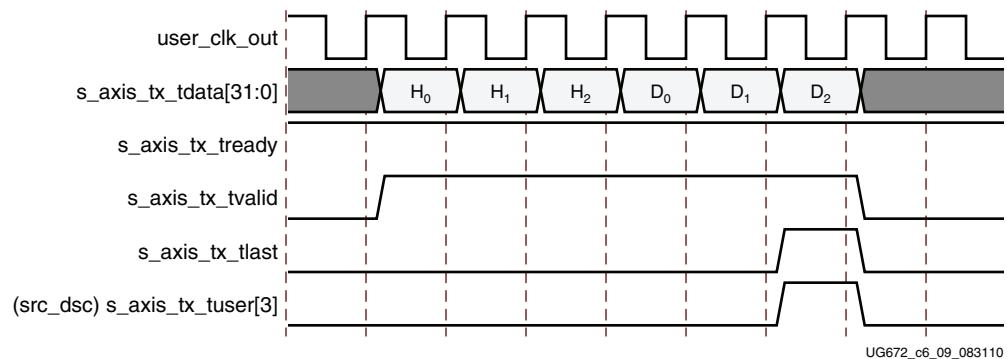


Figure 5-9: Source Driven Transaction Discontinue on Transmit Interface

Discarding of Transaction by Destination

The core transmit AXI4-Stream interface discards a TLP for three reasons:

- The PCI Express link goes down.
- Presented TLP violates the Max_Payload_Size field of the PCI Express Device Capability Register (offset 04H) (it is left to the user to not violate the Max_Payload_Size field of the Device Control Register (offset 08H)).
- str is asserted and data is not presented on consecutive clock cycles; that is, s_axis_tx_tvalid is deasserted in the middle of a TLP transfer.

When any of these occurs, the transmit AXI4-Stream interface continues to accept the remainder of the presented TLP and asserts tx_terr_drop no later than the third clock cycle following the EOF of the discarded TLP. [Figure 5-10](#) illustrates the core signaling that a packet was discarded using tx_terr_drop due to a length violation.

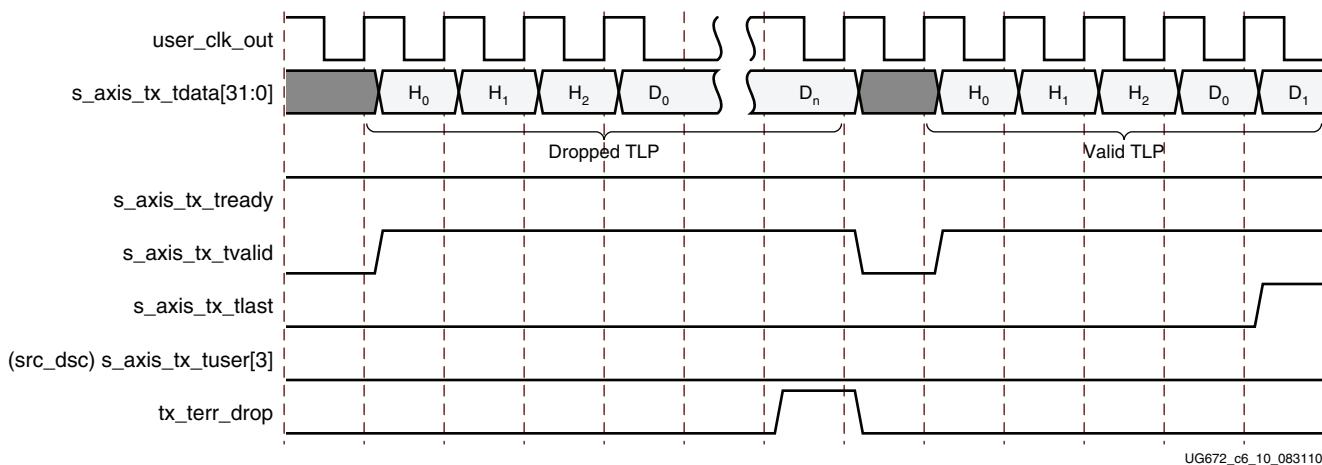


Figure 5-10: Destination Driven Transaction Discontinue on Transmit Interface

Packet Data Poisoning on the Transmit AXI4-Stream Interface

The user application uses either of these two mechanisms to mark the data payload of a transmitted TLP as poisoned:

- Set EP = 1 in the TLP header. This mechanism can be used if the payload is known to be poisoned when the first DWORD of the header is presented to the core on the AXI4-Stream interface.
- Assert terr_fwd for at least 1 valid data transfer cycle any time during the packet transmission, as shown in [Figure 5-11](#). This causes the core to set EP = 1 in the TLP header when it transmits the packet onto the PCI Express fabric. This mechanism can be used if the user application does not know whether a packet can be poisoned at the start of packet transmission. Use of terr_fwd is not supported for packets when str is asserted (streamed transmit packets).

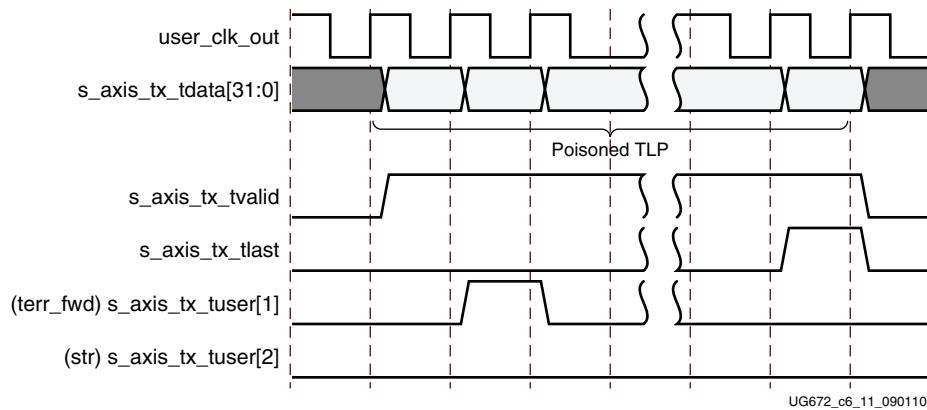
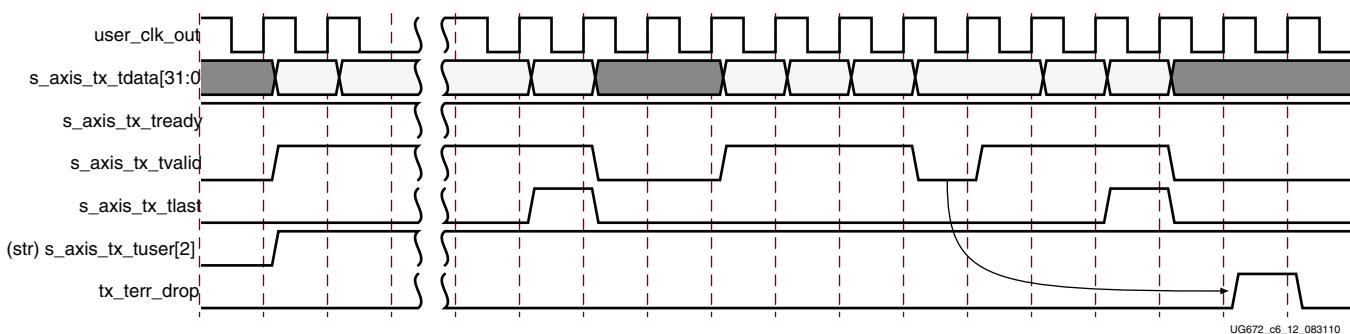


Figure 5-11: Packet Data Poisoning on the Transmit AXI4-Stream Interface

Streaming Mode for Transactions on the Transmit Interface

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express core allows the user application to enable Streaming (cut-through) mode for transmission of a TLP, when possible, to reduce latency of operation. To enable this feature, the user application must hold str asserted for the entire duration of the transmitted TLP. In addition, the user application must present valid frames on every clock cycle until the final cycle of the TLP. In other words, the user application must not deassert s_axis_tx_tvalid for the duration of the presented TLP. Source throttling of the transaction while in streaming mode of operation causes the transaction to be dropped (tx_terr_drop is asserted) and a nullified TLP to be signaled on the PCI Express link. [Figure 5-12](#) illustrates the streaming mode of operation, where the first TLP is streamed and the second TLP is dropped due to source throttling.



[Figure 5-12: Streaming Mode on the Transmit Interface](#)

Appending ECRC to Protect TLPs

If the user application needs to send a TLP Digest associated with a TLP, it must construct the TLP header such that the TD bit is set and the user application must properly compute and append the 1-DWORD TLP Digest after the last valid TLP payload section (if applicable). TLPs originating within the core, for example Completions, Error Messages, and Interrupts, do not have a TLP Digest appended.

Maximum Payload Size

TLP size is restricted by the capabilities of both link partners. After the link is trained, the root complex sets the MAX_PAYLOAD_SIZE value in the Device Control register. This value is equal to or less than the value advertised by the core's Device Capability register. The advertised value in the Device Capability register of the integrated Endpoint block core is either 128, 256, or 512 bytes, depending on the setting in the CORE Generator software GUI. For more information about these registers, see section 7.8 of the *PCI Express Base Specification*. The value of the core's Device Control register is provided to the user application on the cfg_dcommand[15:0] output. See [Design with Configuration Space Registers and Configuration Interface, page 82](#) for information about this output.

Transmit Buffers

The Endpoint for PCIe transmit AXI4-Stream interface provides tx_buf_av, an instantaneous indication of the number of Max_Payload_Size buffers available for use in the transmit buffer pool. [Table 5-1](#) defines the number of transmit buffers available and maximum supported payload size for a specific core.

Table 5-1: Transmit Buffers Available

Capability Max Payload Size (Bytes)	Performance Level ⁽¹⁾	
	Good (Minimize Block RAM Usage)	High (Maximize Performance)
128	13	27
256	14	29
512	15	30

Notes:

1. Performance level is set through a CORE Generator software GUI selection.

Each buffer can hold one maximum sized TLP. A maximum sized TLP is a TLP with a 4-DWORD header plus a data payload equal to the MAX_PAYLOAD_SIZE of the core (as defined in the Device Capability register) plus a TLP Digest. After the link is trained, the root complex sets the MAX_PAYLOAD_SIZE value in the Device Control register. This value is equal to or less than the value advertised by the core's Device Capability register. For more information about these registers, see section 7.8 of the PCI Express Base Specification. A TLP is held in the core's transmit buffer until the link partner acknowledges receipt of the packet, at which time the buffer is released and a new TLP can be loaded into it by the user application.

For example, if the Capability Max Payload Size selected for the Endpoint core is 256 bytes, and the performance level selected is high, there are 29 total transmit buffers. Each of these buffers can hold at a maximum one 64-bit Memory Write Request (4 DWORD header) plus 256 bytes of data (64 DWORDs) plus TLP Digest (1 DWORD) for a total of 69 DWORDs. This example assumes the root complex set the MAX_PAYLOAD_SIZE register of the Device Control register to 256 bytes, which is the maximum capability advertised by this core. For this reason, at any given time, this core could have 29 of these 69 DWORD TLPs awaiting transmittal. There is no sharing of buffers among multiple TLPs, so even if user is sending smaller TLPs such as 32-bit Memory Read request with no TLP Digest totaling 3 DWORDs only per TLP, each transmit buffer still holds only one TLP at any time.

The internal transmit buffers are shared between the user application and the core's configuration management module (CMM). Because of this, the tx_buf_av bus can fluctuate even if the user application is not transmitting packets. The CMM generates completion TLPs in response to configuration reads or writes, interrupt TLPs at the request of the user application, and message TLPs when needed.

The Transmit Buffers Available indication enables the user application to completely utilize the PCI transaction ordering feature of the core transmitter. The transaction ordering rules allow for Posted and Completion TLPs to bypass Non-Posted TLPs. See section 2.4 of the *PCI Express Base Specification* for more information about ordering rules.

The core supports the transaction ordering rules and promotes Posted and Completion packets ahead of blocked Non-Posted TLPs. Non-Posted TLPs can become blocked if the link partner is in a state where it momentarily has no Non-Posted receive buffers available, which it advertises through Flow Control updates. In this case, the core promotes Completion and Posted TLPs ahead of these blocked Non-Posted TLPs. However, this can only occur if the Completion or Posted TLP has been loaded into the core by the user application. By monitoring the tx_buf_av bus, the user application can ensure there is at least one free buffer available for any Completion or Posted TLP. Promotion of Completion and Posted TLPs only occurs when Non-Posted TLPs are blocked; otherwise packets are sent on the link in the order they are received from the user application.

Receiving Inbound Packets

Basic TLP Receive Operation

Table 2-7, page 18 defines the receive AXI4-Stream interface signals. This sequence of events must occur on the receive Transaction interface for the core to present a TLP to the user application logic:

1. When the user application is ready to receive data, it asserts `m_axis_rx_tready`.
2. When the core is ready to transfer data, the core asserts `m_axis_rx_tvalid` and presents the first complete TLP DWORD on `m_axis_rx_tdata[31:0]`.
3. The core then keeps `m_axis_rx_tvalid` asserted, and presents TLP DWORDs on `m_axis_rx_tdata[31:0]` on subsequent clock cycles, for which the user application logic asserts `m_axis_rx_tready`.
4. The core then asserts `m_axis_rx_tlast` and presents either the last DWORD on `s_axis_tx_tdata[31:0]`.
5. If no further TLPs are available, at the next clock cycle, the core deasserts `m_axis_rx_tvalid` to signal the end of valid transfers on `m_axis_rx_tdata[31:0]`.

Figure 5-13 illustrates a 3-DW TLP header without a data payload; an example is a 32-bit addressable Memory Read request.

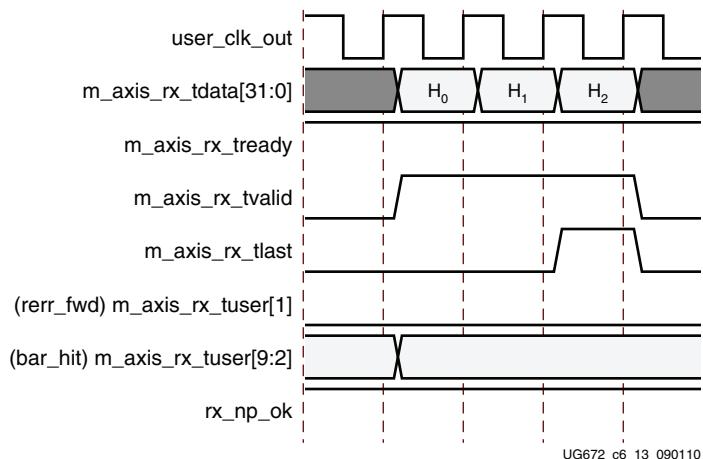


Figure 5-13: TLP 3-DW Header without Payload

[Figure 5-14](#) illustrates a 4-DW TLP header without a data payload; an example is a 64-bit addressable Memory Read request.

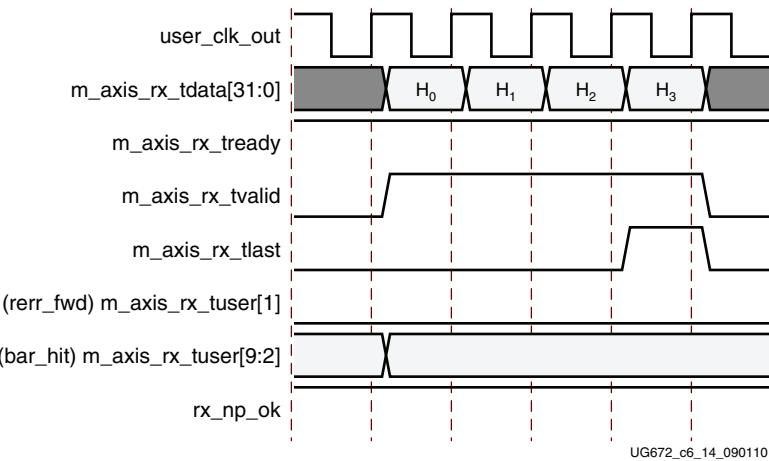


Figure 5-14: TLP 4-DW Header without Payload

[Figure 5-15](#) illustrates a 3-DW TLP header with a data payload; an example is a 32-bit addressable Memory Write request.

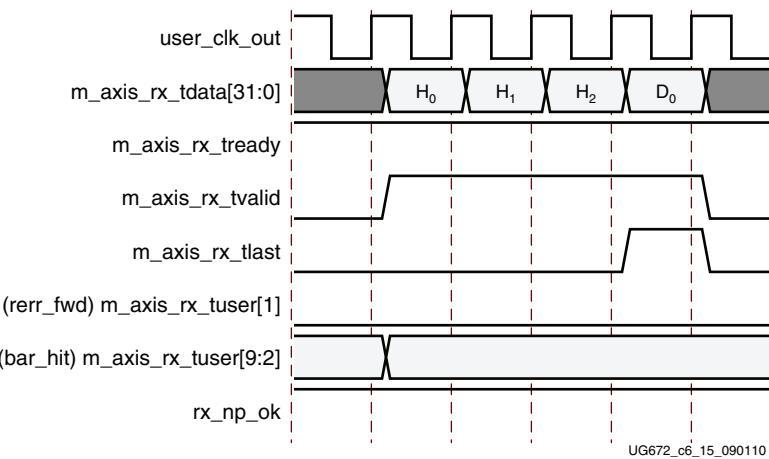
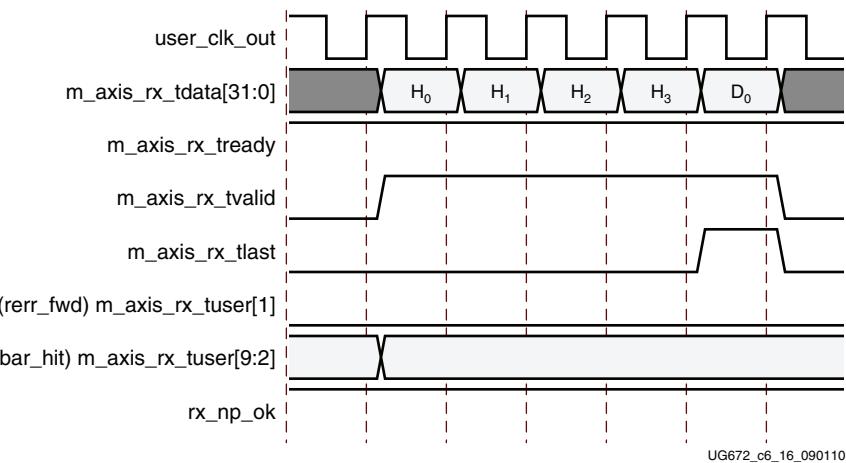


Figure 5-15: TLP 3-DW Header with Payload

[Figure 5-16](#) illustrates a 4-DW TLP header with a data payload; an example is a 64-bit addressable Memory Write request.

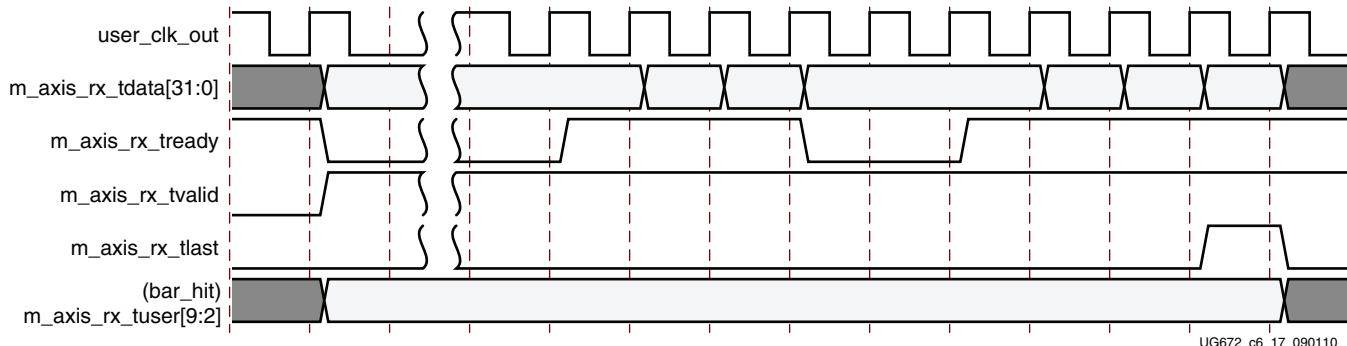


[Figure 5-16: TLP 4-DW Header with Payload](#)

Throttling the Datapath on the Receive AXI4-Stream Interface

The user application can stall the transfer of data from the core at any time by deasserting `m_axis_rx_tready`. If the user deasserts `m_axis_rx_tready` while no transfer is in progress and if a TLP becomes available, the core asserts `m_axis_rx_tvalid` and presents the first TLP DWORD on `m_axis_rx_tdata[31:0]`. The core remains in this state until the user asserts `m_axis_rx_tready` to signal the acceptance of the data presented on `m_axis_rx_tdata[31:0]`. At that point, the core presents subsequent TLP DWORDs as long as `m_axis_rx_tready` remains asserted. If the user deasserts `m_axis_rx_tready` during the middle of a transfer, the core stalls the transfer of data until the user asserts `m_axis_rx_tready` again. There is no limit to the number of cycles the user can keep `m_axis_rx_tready` deasserted. The core pauses until the user is again ready to receive TLPs.

[Figure 5-17](#) illustrates the core asserting `m_axis_rx_tvalid` along with presenting data on `m_axis_rx_tdata[31:0]`. The user application logic inserts wait states by deasserting `m_axis_rx_tready`. The core does not present the next TLP DWORD until it detects `m_axis_rx_tready` assertion. The user application logic can assert or deassert `m_axis_rx_tready` as required to balance receipt of new TLP transfers with the rate of TLP data processing inside the application logic.



[Figure 5-17: User Application Throttling Receive TLP](#)

Receiving Back-to-Back Transactions on the Receive AXI4-Stream Interface

The user application logic must be designed to handle presentation of back-to-back TLPs on the receive interface AXI4-Stream interface by the core. The core can assert `m_axis_rx_tvalid` for a new TLP at the clock cycle after `m_axis_rx_tlast` assertion for the previous TLP. [Figure 5-18](#) illustrates back-to-back TLPs presented on the receive interface.

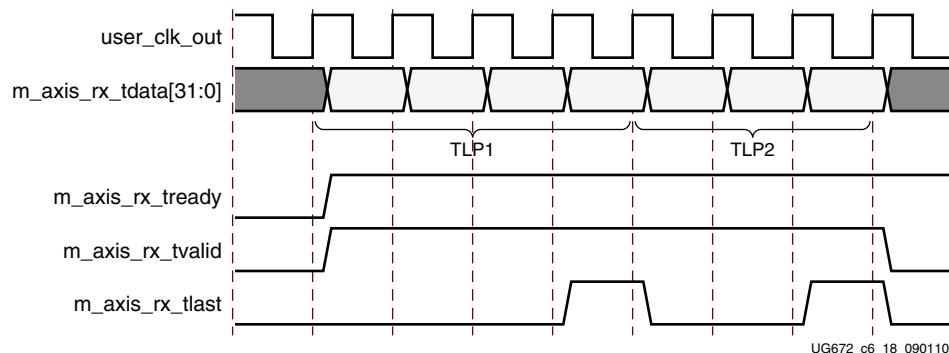


Figure 5-18: Receive Back-to-Back Transactions

If the user application cannot accept back-to-back packets, it can stall the transfer of the TLP by deasserting `m_axis_rx_tready` as discussed in the previous section. [Figure 5-19](#) shows an example of using `m_axis_rx_tready` to pause the acceptance of the second TLP.

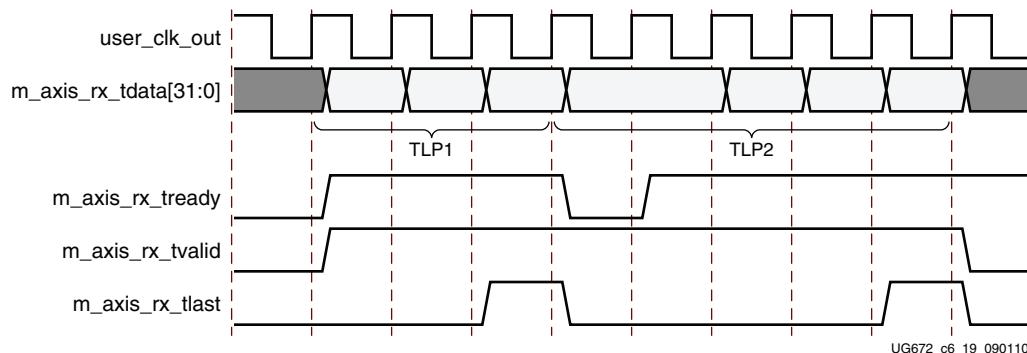


Figure 5-19: User Application Throttling of Back-to-Back TLPs

Packet Re-ordering on Receive AXI4-Stream Interface

Transaction processing in the core receiver is fully compliant with the PCI transaction ordering rules, described in Chapter 2 of the *PCI Express Base Specification*. The transaction ordering rules allow for Posted and Completion TLPs to bypass blocked Non-Posted TLPs.

The user application can deassert `rx_np_ok` if it is not ready to accept Non-Posted Transactions from the core, (as shown in [Figure 5-20](#)) but can receive Posted and Completion Transactions. The user application must deassert `rx_np_ok` at least two clock cycles before `m_axis_rx_tlast` of the second-to-last Non-Posted packet the user can accept. While `rx_np_ok` is deasserted, received Posted and Completion Transactions pass Non-Posted Transactions. After the user application is ready to accept Non-Posted Transactions, it must reassert `rx_np_ok`. Previously bypassed Non-Posted Transactions are presented to the user application before other received TLPs. There is no limit as to how long `rx_np_ok` can be deasserted, however users must take care to not deassert `rx_np_ok` for extended

periods, because this can cause a completion time-out in the Requester. See the *PCI Express Base Specification* for more information on the Completion Timeout Mechanism.

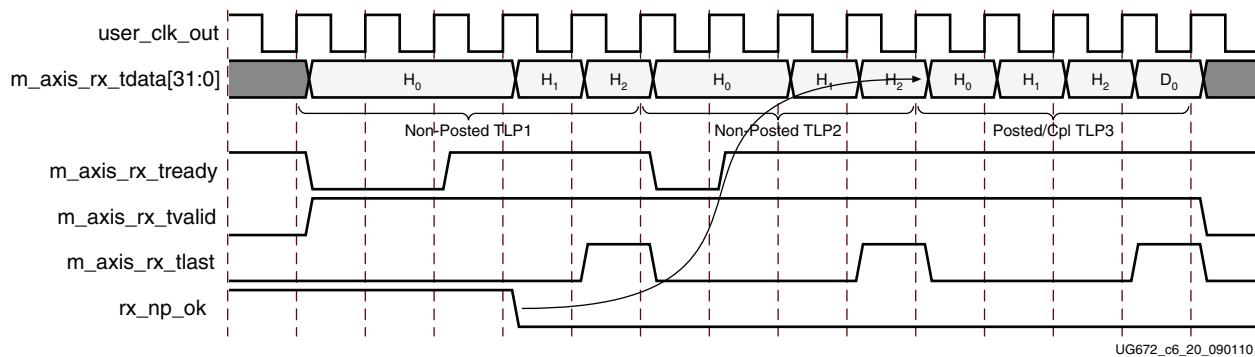


Figure 5-20: Packet Re-ordering on Receive AXI4-Stream Interface

Packet re-ordering allows the user application to optimize the rate at which Non-Posted TLPs are processed, while continuing to receive and process Posted and Completion TLPs in a non-blocking fashion. The *rx_np_ok* signaling restrictions require that the user application be able to receive and buffer at least three Non-Posted TLPs. The following algorithm describes the process of managing the Non-Posted TLP buffers.

Consider that *Non-Posted_Buffers_Available* denotes the size of Non-Posted buffer space available to user application. The size of the Non-Posted buffer space is greater than three Non-Posted TLPs. *Non-Posted_Buffers_Available* is decremented when a Non-Posted TLP is accepted for processing from the core, and is incremented when Non-Posted TLP is drained for processing by the user application.

```

For every clock cycle, do    {
    if (Valid transaction Start-Of-Frame accepted by user application) {
        Extract TLP Format and Type from the 1st TLP DW
        if (TLP type == Non Posted) {
            if (Non-Posted_Buffers_Available <= 2) // Accounts for the
                current and possibly the next NP TLP
                Deassert rx_np_ok on the following clock cycle.
            else if (Other optional user policies to stall Non-Posted
                transactions)
                Deassert rx_np_ok on the following clock cycle.
            else // (Non-Posted_Buffers_Available > 2)
                Assert rx_np_ok on the following clock cycle.
                Decrement Non-Posted_Buffers_Available in User Application
        } else { // Posted and Completion TLPs
            Process the received TLPs
        }
    }
}

```

Packet Data Poisoning and TLP Digest on Receive AXI4-Stream Interface

To simplify logic within the user application, the core performs automatic pre-processing based on values of TLP Digest (TD) and Data Poisoning (EP) header bit fields on the received TLP.

All received TLPs with the Data Poisoning bit in the header set ($EP = 1$) are presented to the user. The core asserts the `rerr_fwd` signal for the duration of each poisoned TLP, as illustrated in [Figure 5-21](#).

If the TLP Digest bit field in the TLP header is set ($TD = 1$), the TLP contains an End-to-End CRC (ECRC). The core performs the following operations based on how the user configured the core during core generation:

- If the Trim TLP Digest option is on, the core removes and discards the ECRC field from the received TLP and clears the TLP Digest bit in the TLP header.
- If the Trim TLP Digest option is off, the core does not remove the ECRC field from the received TLP and presents the entire TLP including TLP Digest to the user application receiver interface.

See [Chapter 4, Generating and Customizing the Core](#), for more information about how to enable the Trim TLP Digest option during core generation.

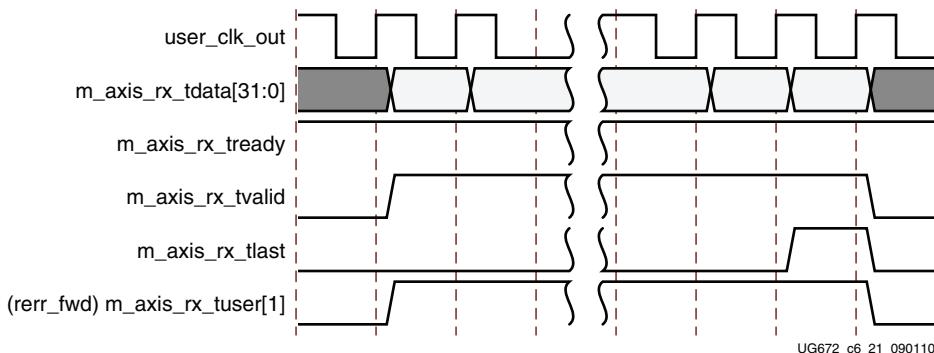


Figure 5-21: Receive Transaction Data Poisoning

Packet Base Address Register Hit on Receive AXI4-Stream Interface

The core decodes incoming Memory and I/O TLP request addresses to determine which Base Address Register (BAR) in the core's Type0 configuration space is being targeted, and indicates the decoded base address on `bar_hit[6:0]`. For each received Memory or I/O TLP, a minimum of one and a maximum of two (adjacent) bit(s) are set to 1b. If the received TLP targets a 32-bit Memory or I/O BAR, only one bit is asserted. If the received TLP targets a 64-bit Memory BAR, two adjacent bits are asserted. If the core receives a TLP that is not decoded by one of the BARs, then the core drops it without presenting it to the user and an Unsupported Request message is automatically generated. Even if the core is configured for a 64-bit BAR, the system might not always allocate a 64-bit address, in which case only one `bar_hit[6:0]` signal is asserted.

[Table 5-2](#) illustrates mapping between bar_hit[6:0] and the BARs, and the corresponding byte offsets in the core Type0 configuration header.

Table 5-2: bar_hit to Base Address Register Mapping

bar_hit[x]	m_axis_rx_tuser[x]	BAR	Byte Offset
0	2	0	10h
1	3	1	14h
2	4	2	18h
3	5	3	1Ch
4	6	4	20h
5	7	5	24h
6	8	Expansion ROM BAR	30h

For a Memory or I/O TLP Transaction on the receive interface, bar_hit[6:0] is valid for the entire TLP, starting with the assertion of m_axis_rx_tvalid, as shown in [Figure 5-22](#). When receiving non-Memory and non-I/O transactions, the signal bar_hit[6:0] is undefined.

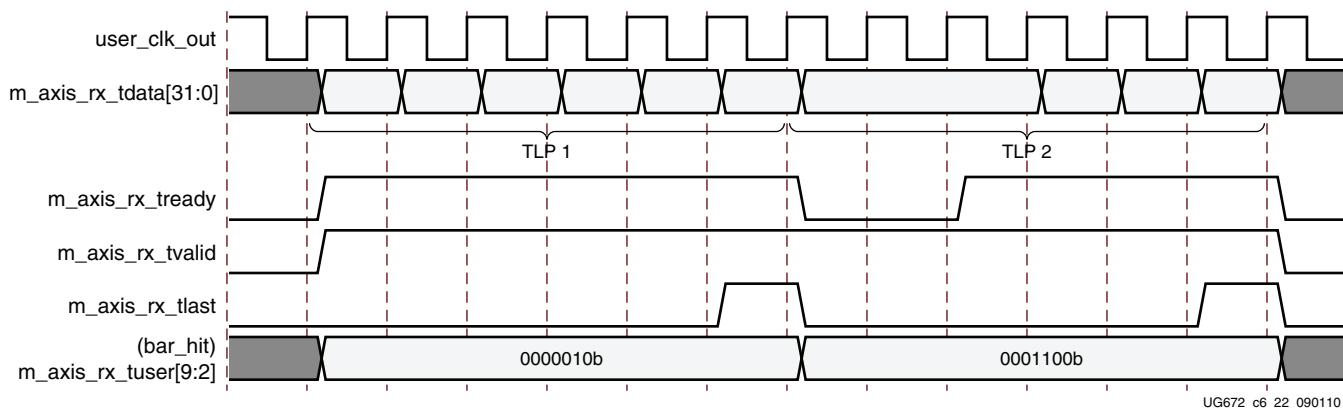


Figure 5-22: BAR Target Determination Using bar_hit

The signal bar_hit[6:0] enables received Memory and I/O transactions to be directed to the appropriate destination apertures within the user application. By utilizing bar_hit[6:0], application logic can inspect only the lower order Memory and I/O address bits within the address aperture to simplify decoding logic.

Packet Transfer During Link-Down Event on Receive AXI4-Stream Interface

The loss of communication with the link partner is signaled by deassertion of `user_lnk_up`. When `user_lnk_up` is deasserted, it effectively acts as a *Hot Reset* to the entire core. For this reason, all TLPs stored inside the core or being presented to the receive interface are irrecoverably lost. A TLP in progress on the Receive AXI4-Stream interface will be presented to its correct length, according to the Length field in the TLP header. However, the TLP is corrupted and should be discarded by the user application. Figure 5-23 illustrates packet transfer discontinue scenario.

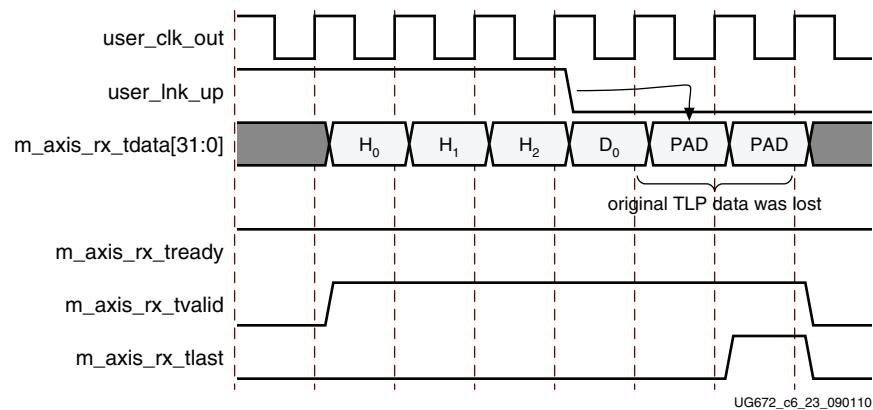


Figure 5-23: **Receive Transaction Discontinue**

Receiver Flow Control Credits Available

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express provides the user application information about the state of the receiver buffer pool queues. This information represents the current space available for the Posted, Non-Posted, and Completion queues.

One Header Credit is equal to either a 3 or 4 DWORD TLP Header and one Data Credit is equal to 16 bytes of payload data. [Table 5-3](#) provides values on credits available immediately after user_lnk_up assertion but before the reception of any TLP. If space available for any of the above categories is exhausted, the corresponding credit available signals indicate a value of zero. Credits available return to initial values after the receiver has drained all TLPs.

Table 5-3: Transaction Receiver Credits Available Initial Values

Credit Category	Performance Level	128 byte Capability MPS	256 byte Capability MPS	512 byte Capability MPS
Non-Posted Header	Good	8		
	High			
Posted Header	Good	16	24	32
	High	30	32	32
Posted Data	Good	41	96	211
	High	89	211	467
Completion Header	Good	16	24	40
	High	30	40	40
Completion Data	Good	41	96	211
	High	89	211	467

The user application can use the fc_ph[7:0], fc_pd[11:0], fc_nph[7:0], fc_npd[11:0], fc_cplh[7:0], fc_cpld[11:0], and fc_sel[2:0] signals to efficiently utilize and manage receiver buffer space available in the core and the core application. For additional information, see [Flow Control Credit Information, page 93](#).

Endpoint cores for PCI Express have a unique requirement where the user application must use advanced methods to prevent buffer overflows while requesting Non-Posted Read Requests from an upstream component. According to the specification, a PCI Express Endpoint is required to advertise infinite storage credits for Completion Transactions in its receivers. This means that endpoints must internally manage Memory Read Requests transmitted upstream and not overflow the receiver when the corresponding Completions are received. The user application transmit logic must use Completion credit information presented to modulate the rate and size of Memory Read requests, to stay within the instantaneous Completion space available in the core receiver. For additional information, see [Appendix E, Managing Receive-Buffer Space for Inbound Completions](#).

Design with Configuration Space Registers and Configuration Interface

This section describes the use of the Configuration Interface for accessing the PCI Express Configuration Space Type 0 registers that are part of the integrated Endpoint block core. The Configuration Interface includes a read Configuration Port for accessing the registers. In addition, some commonly used registers are mapped directly on the Configuration Interface for convenience.

Registers Mapped Directly onto the Configuration Interface

The integrated Endpoint block core provides direct access to select command and status registers in its Configuration Space. Values in these registers are modified by Configuration Writes received from the Root Complex and cannot be modified by the user application. [Table 5-4](#) defines the command and status registers mapped to the configuration port.

Table 5-4: Command and Status Registers Mapped to the Configuration Port

Port Name	Direction	Description
cfg_bus_number[7:0]	Output	Bus Number: Default value after reset is 00h. Refreshed whenever a Type 0 Configuration Write packet is received.
cfg_device_number[4:0]	Output	Device Number: Default value after reset is 00000b. Refreshed whenever a Type 0 Configuration Write packet is received.
cfg_function_number[2:0]	Output	Function Number: Function number of the core, hard wired to 000b.
cfg_status[15:0]	Output	Status Register: Status register from the Configuration Space Header.
cfg_command[15:0]	Output	Command Register: Command register from the Configuration Space Header.
cfg_dstatus[15:0]	Output	Device Status Register: Device status register from the PCI Express Extended Capability Structure.
cfg_dcommand[15:0]	Output	Device Command Register: Device control register from the PCI Express Extended Capability Structure.
cfg_lstatus[15:0]	Output	Link Status Register: Link status register from the PCI Express Extended Capability Structure.
cfg_lcommand[15:0]	Output	Link Command Register: Link control register from the PCI Express Extended Capability Structure.

Device Control and Status Register Definitions

`cfg_bus_number[7:0], cfg_device_number[4:0], cfg_function_number[2:0]`

Together, these three values comprise the core ID, which the core captures from the corresponding fields of inbound Type 0 Configuration Write accesses. The user application is responsible for using this core ID as the Requestor ID on any requests it originates, and using it as the Completer ID on any Completion response it sends. This core supports only one function; for this reason, the function number is hardwired to 000b.

`cfg_status[15:0]`

This bus allows the user application to read the Status register in the PCI Configuration Space Header. [Table 5-5](#) defines these bits. See the *PCI Express Base Specification* for detailed information.

Table 5-5: Bit Mapping on Header Status Register

Bit	Name
<code>cfg_status[15]</code>	Detected Parity Error
<code>cfg_status[14]</code>	Signaled System Error
<code>cfg_status[13]</code>	Received Master Abort
<code>cfg_status[12]</code>	Received Target Abort
<code>cfg_status[11]</code>	Signaled Target Abort
<code>cfg_status[10:9]</code>	DEVSEL Timing (hardwired to 00b)
<code>cfg_status[8]</code>	Master Data Parity Error
<code>cfg_status[7]</code>	Fast Back-to-Back Transactions Capable (hardwired to 0)
<code>cfg_status[6]</code>	Reserved
<code>cfg_status[5]</code>	66 MHz Capable (hardwired to 0)
<code>cfg_status[4]</code>	Capabilities List Present (hardwired to 1)
<code>cfg_status[3]</code>	Interrupt Status
<code>cfg_status[2:0]</code>	Reserved

cfg_command[15:0]

This bus reflects the value stored in the Command register in the PCI Configuration Space Header. [Table 5-6](#) provides the definitions for each bit in this bus. See the *PCI Express Base Specification* for detailed information.

Table 5-6: Bit Mapping on Header Command Register

Bit	Name
cfg_command[15:11]	Reserved
cfg_command[10]	Interrupt Disable
cfg_command[9]	Fast Back-to-Back Transactions Enable (hardwired to 0)
cfg_command[8]	SERR Enable
cfg_command[7]	IDSEL Stepping/Wait Cycle Control (hardwired to 0)
cfg_command[6]	Parity Error Enable (hardwired to 0)
cfg_command[5]	VGA Palette Snoop (hardwired to 0)
cfg_command[4]	Memory Write and Invalidate (hardwired to 0)
cfg_command[3]	Special Cycle Enable (hardwired to 0)
cfg_command[2]	Bus Master Enable
cfg_command[1]	Memory Address Space Decoder Enable
cfg_command[0]	I/O Address Space Decoder Enable

The user application must monitor the Bus Master Enable bit (cfg_command[2]) and refrain from transmitting requests while this bit is not set. This requirement applies only to requests; completions can be transmitted regardless of this bit.

cfg_dstatus[15:0]

This bus reflects the value stored in the Device Status register of the PCI Express Extended Capabilities Structure. [Table 5-7](#) defines each bit in the cfg_dstatus bus. See the *PCI Express Base Specification* for detailed information.

Table 5-7: Bit Mapping on PCI Express Device Status Register

Bit	Name
cfg_dstatus[15:6]	Reserved
cfg_dstatus[5]	Transaction Pending
cfg_dstatus[4]	AUX Power Detected
cfg_dstatus[3]	Unsupported Request Detected
cfg_dstatus[2]	Fatal Error Detected
cfg_dstatus[1]	Non-Fatal Error Detected
cfg_dstatus[0]	Correctable Error Detected

cfg_dcommand[15:0]

This bus reflects the value stored in the Device Control register of the PCI Express Extended Capabilities Structure. [Table 5-8](#) defines each bit in the cfg_dcommand bus. See the *PCI Express Base Specification* for detailed information.

Table 5-8: Bit Mapping of PCI Express Device Control Register

Bit	Name
cfg_dcommand[15]	Reserved
cfg_dcommand[14:12]	Max_Read_Request_Size
cfg_dcommand[11]	Enable No Snoop
cfg_dcommand[10]	Auxiliary Power PM Enable
cfg_dcommand[9]	Phantom Functions Enable
cfg_dcommand[8]	Extended Tag Field Enable
cfg_dcommand[7:5]	Max_Payload_Size
cfg_dcommand[4]	Enable Relaxed Ordering
cfg_dcommand[3]	Unsupported Request Reporting Enable
cfg_dcommand[2]	Fatal Error Reporting Enable
cfg_dcommand[1]	Non-Fatal Error Reporting Enable
cfg_dcommand[0]	Correctable Error Reporting Enable

cfg_lstatus[15:0]

This bus reflects the value stored in the Link Status register in the PCI Express Extended Capabilities Structure. [Table 5-9](#) defines each bit in the cfg_lstatus bus. See the *PCI Express Base Specification* for details.

Table 5-9: Bit Mapping of PCI Express Link Status Register

Bit	Name
cfg_lstatus[15:13]	Reserved
cfg_lstatus[12]	Slot Clock Configuration
cfg_lstatus[11]	Reserved
cfg_lstatus[10]	Reserved
cfg_lstatus[9:4]	Negotiated Link Width
cfg_lstatus[3:0]	Link Speed

cfg_lcommand[15:0]

This bus reflects the value stored in the Link Control register of the PCI Express Extended Capabilities Structure. [Table 5-10](#) provides the definition of each bit in cfg_lcommand bus. See the *PCI Express Base Specification* for more details.

Table 5-10: Bit Mapping of PCI Express Link Control Register

Bit	Name
cfg_lcommand[15:8]	Reserved
cfg_lcommand [7]	Extended Synch
cfg_lcommand [6]	Common Clock Configuration
cfg_lcommand [5]	Retrain Link (Reserved for an endpoint device)
cfg_lcommand [4]	Link Disable
cfg_lcommand [3]	Read Completion Boundary
cfg_lcommand[2]	Reserved
cfg_lcommand [1:0]	Active State Link PM Control

Accessing Additional Registers through the Configuration Port

Configuration registers that are not directly mapped to the user interface can be accessed by configuration-space address using the ports shown in [Table 2-9, page 23](#).

The user application must supply the read address as a DWORD address, not a byte address. To calculate the DWORD address for a register, divide the byte address by four. For example:

- The DWORD address of the Command/Status Register in the PCI Configuration Space Header is 01h. (The byte address is 04h.)
- The DWORD address for BAR0 is 04h. (The byte address is 10h.)

To read any register in the configuration space shown in [Table 2-2, page 13](#), the user application drives the register DWORD address onto cfg_dwaddr[9:0]. The core drives the content of the addressed register onto cfg_do[31:0]. The value on cfg_do[31:0] is qualified by signal assertion on cfg_rd_wr_done. [Figure 5-24](#) illustrates an example with two consecutive reads from the Configuration Space.

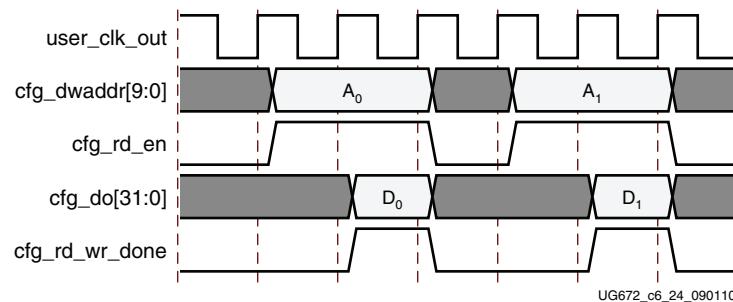


Figure 5-24: Example Configuration Space Access

User Implemented Configuration Space

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express enables users to optionally implement registers in the PCI Configuration Space, the PCI Express Extended Configuration Space, or both, in the user application. The user application is required to return Config Completions for all address within this space. For more information about enabling and customizing this feature, see [Chapter 4, Generating and Customizing the Core](#).

PCI Configuration Space

If the user chooses to implement registers within 0x6C to 0xFF in the PCI Configuration Space, the start address of the address region they wish to implement can be defined during the core generation process.

The user application is responsible for generating all Completions to Configuration Reads and Writes from the user-defined start address to the end of PCI Configuration Space (0xFF). Configuration Reads to unimplemented registers within this range should be responded to with a Completion with 0x00000000 as the data, and configuration writes should be responded to with a successful Completion.

For example, to implement address range 0xC0 to 0xCF, there are several address ranges defined that should be treated differently depending on the access. [Table 5-11](#) shows more details on this example.

Table 5-11: Example: User Implemented Space 0xC0 to 0xCF

	Configuration Writes	Configuration Reads
0x00 to 0xBF	Core responds automatically	Core responds automatically
0xC0 to 0xCF	User application responds with Successful Completion	User application responds with register contents
0xD0 to 0xFF	User application responds with Successful Completion	User application responds with 0x00000000

PCI Express Extended Configuration Space

The starting address of the region in the PCI Express Extended Configuration Space that is optionally available for users to implement depends on the PCI Express Extended Capabilities the user has enabled in the Spartan-6 FPGA Integrated Endpoint Block for PCI Express, as shown in [Table 5-12](#).

Table 5-12: Min Start Addresses of the User Implemented Extended Capabilities

	No Capabilities Selected	DSN
Starting byte address available	100h	10Ch

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express allows the user to select the start address of the user implemented PCI Express Extended Configuration Space. This space must be implemented in the user application. The user application is required to generate a CplD with 0x00000000 for Configuration Read and successful Cpl for Configuration Write to addresses in this selected range not implemented in the user application. The user can choose to implement a Configuration Space with a start address other than that allowed by the integrated Endpoint block for PCI Express. In such a case, the core returns a completion with 0x00000000 for configuration accesses to the region that the user has chosen to not implement. [Table 5-13](#) illustrates this scenario.

Table 5-13: Example: User Defined Start Address for Configuration Space

Configuration Space	Byte Address
DSN Capability	100h - 108h
Reserved Extended Configuration Space (Core Returns Successful Completion with 0x00000000)	10Ch - 164h
User Implemented PCI Express Extended Configuration Space	168h - 47Ch
User Implemented Reserved PCI Express Extended Configuration Space (User application Returns Successful Completion with 0x00000000)	480h - FFFh

Table 5-13 illustrates an example Configuration of the PCI Express Extended Configuration Space, with these settings:

- DSN Capability Enabled
- User Implemented PCI Express Extended Configuration Space Enabled
- User Implemented PCI Express Extended Configuration Space Start Address 168h

In this configuration, the DSN Capability occupies the registers at 100h-108h. The remaining PCI Express Extended Configuration Space, starting at address 10Ch is available to the user to implement. For this example, the user has chosen to implement registers in the address region starting 168h.

In this scenario, the core returns successful Completions with 0x00000000 for Configuration accesses to registers 10Ch-164h. **Table 5-13** also illustrates a case where the user only implements the registers from 168h to 47Ch. In this case, the user is responsible for returning successful Completions with 0x00000000 for configuration accesses to 480h-FFFh.

Additional Packet Handling Requirements

The user application must manage the following mechanisms to ensure protocol compliance, because the core does not manage them automatically.

Generation of Completions

The integrated Endpoint block core does not generate Completions for Memory Reads or I/O requests made by a remote device. The user is expected to service these completions according to the rules specified in the *PCI Express Base Specification*.

Tracking Non-Posted Requests and Inbound Completions

The Integrated Endpoint Block for PCIe does not track transmitted I/O requests or Memory Reads that have yet to be serviced with inbound Completions. The user application is required to keep track of such requests using the Tag ID or other information.

Keep in mind that one Memory Read request can be answered by several Completion packets. The user application must accept all inbound Completions associated with the original Memory Read until all requested data has been received.

The *PCI Express Base Specification* requires that an endpoint advertise infinite Completion Flow Control credits as a receiver; the endpoint can only transmit Memory Reads and I/O requests if it has enough space to receive subsequent Completions.

The integrated Endpoint block core does not keep track of receive-buffer space for Completion. Rather, it sets aside a fixed amount of buffer space for inbound Completions. The user application must keep track of this buffer space to know if it can transmit requests requiring a Completion response. See [Appendix E, Managing Receive-Buffer Space for Inbound Completions](#) for more information.

Reporting User Error Conditions

The user application must report errors that occur during Completion handling using dedicated error signals on the core interface, and must observe the Device Power State before signaling an error to the core. If the user application detects an error (for example, a Completion Timeout) while the device has been programmed to a non-D0 state, the user application is responsible to signal the error after the device is programmed back to the D0 state.

After the user application signals an error, the core reports the error on the PCI Express Link and also sets the appropriate status bit(s) in the Configuration Space. Because status bits must be set in the appropriate Configuration Space register, the user application cannot generate error reporting packets on the transmit interface. The type of error-reporting packets transmitted depends on whether or not the error resulted from a Posted or Non-Posted Request. User-reported Posted errors cause Message packets to be sent to the Root Complex if enabled to do so through the Device Control Error Reporting bits and/or the Status SERR Enable bit. User-reported non-Posted errors cause Completion packets with non-successful status to be sent to the Root Complex unless the error is regarded as an Advisory Non-Fatal Error. For more information about Advisory Non-Fatal Errors, see Chapter 6 of the *PCI Express Base Specification*. Errors on Non-Posted Requests can result in either Messages to the Root Complex or Completion packets with non-Successful status sent to the original Requester.

Error Types

The user application triggers six types of errors using the signals defined in [Table 2-9, page 23](#).

- End-to-end CRC ECRC Error
- Unsupported Request Error
- Completion Timeout Error
- Unexpected Completion Error
- Completer Abort Error
- Correctable Error

Multiple errors can be detected in the same received packet; for example, the same packet can be an Unsupported Request and have an ECRC error. If this happens, only one error should be reported. Because all user-reported errors have the same severity, the user application design can determine which error to report. The cfg_err_posted signal, combined with the appropriate error reporting signal, indicates what type of error-reporting packets are transmitted. The user can signal only one error per clock cycle. See [Figure 5-25](#), [Figure 5-26](#), and [Figure 5-27](#), and [Table 5-14](#) and [Table 5-15](#).

Table 5-14: User-Indicated Error Signaling

Reported Error	<code>cfg_err_posted</code>	Action
None	Don't care	No Action Taken
<code>cfg_err_ur</code>	0 or 1	0: If enabled, a Non-Fatal Error Message is sent. 1: A Completion with a "status=unsupported request" is sent.
<code>cfg_err_cpl_abort</code>	0 or 1	0: If enabled, a Non-Fatal Error message is sent. 1: A Completion with a "status=unsupported request" is sent.
<code>cfg_err_cpl_timeout</code>	Don't care	If enabled, a Non-Fatal Error Message is sent.
<code>cfg_err_ecrc</code>	Don't care	If enabled, a Non-Fatal Error Message is sent.
<code>cfg_err_cor</code>	Don't care	If enabled, a Correctable Error Message is sent.

Table 5-15: Possible Error Conditions for TLPs Received by the User Application

Received TLP Type	Possible Error Condition					Error Qualifying Signal Status	
	Unsupported Request (<code>cfg_err_ur</code>)	Completion Abort (<code>cfg_err_cpl_abort</code>)	Correctable Error (<code>cfg_err_cor</code>)	ECRC Error (<code>cfg_err_ecrc</code>)	Value to Drive on (<code>cfg_err_posted</code>)	Drive Data on (<code>cfg_err_tlp_cpl_header[47:0]</code>)	
Memory Write	✓	X	N/A	✓	0		No
Memory Read	✓	✓	N/A	✓	1		Yes
I/O	✓	✓	N/A	✓	1		Yes
Completion	X	X	N/A	✓	0		No

Notes:

1. A checkmark indicates a possible error condition for a given TLP type. For example, users can signal Unsupported Request or ECRC Error for a Memory Write TLP, if these errors are detected. An X indicates not a valid error condition for a given TLP type. For example, users should never signal Completion Abort in response to a Memory Write TLP.

Whenever an error is detected in a Non-Posted Request, the user application deasserts `cfg_err_posted` and provides header information on `cfg_err_tlp_cpl_header[47:0]` during the same clock cycle the error is reported, as illustrated in Figure 5-25. The additional header information is necessary to construct the required Completion with non-Successful status. Additional information about when to assert or deassert `cfg_err_posted` is provided in the following sections.

If an error is detected on a Posted Request, the user application instead asserts `cfg_err_posted`, but otherwise follows the same signaling protocol. This results in a Non-Fatal Message to be sent, if enabled.

The core's ability to generate error messages can be disabled by the Root Complex issuing a configuration write to the Endpoint core's Device Control register and the PCI Command register setting the appropriate bits to 0. For more information about these registers, see Chapter 7 of the *PCI Express Base Specification*. However, error-reporting status bits are always set in the Configuration Space whether or not their Messages are disabled.

If several non-Posted errors are signaled on cfg_err_ur or cfg_err_cpl_abort in a short amount of time, it is possible for the core to be unable to buffer them all. If that occurs, then cfg_err_cpl_rdy is deasserted, and the user must cease signaling those types of errors on the same cycle. In addition, the user must not resume signaling those types of errors until cfg_err_cpl_rdy is reasserted.

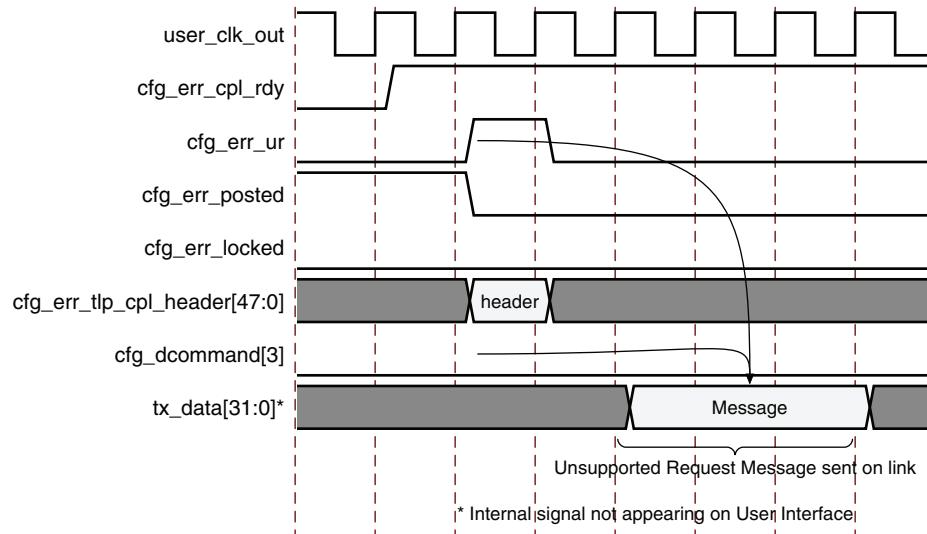


Figure 5-25: Signaling Unsupported Request for Non-Posted TLP

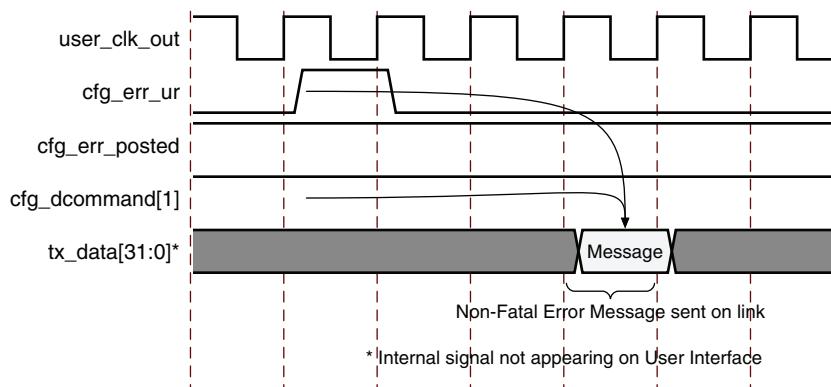


Figure 5-26: Signaling Unsupported Request for Posted TLP

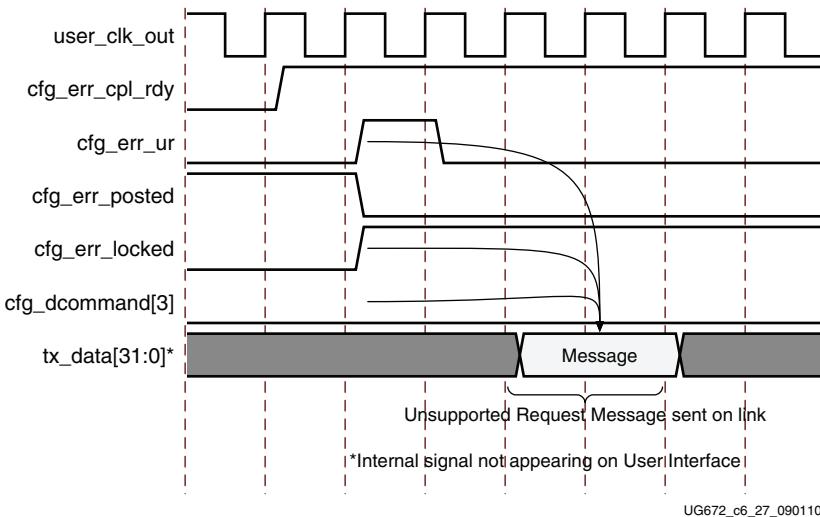


Figure 5-27: Signaling Locked Unsupported Request for Locked Non-Posted TLP

Completion Timeouts

The integrated Endpoint block core does not implement Completion timers; for this reason, the user application must track how long its pending Non-Posted Requests have each been waiting for a Completion and trigger timeouts on them accordingly. The core has no method of knowing when such a timeout has occurred, and for this reason does not filter out inbound Completions for expired requests.

If a request times out, the user application must assert `cfg_err_cpl_timeout`, which causes an error message to be sent to the Root Complex. If a Completion is later received after a request times out, the user application must treat it as an Unexpected Completion.

Unexpected Completions

The integrated Endpoint block core automatically reports Unexpected Completions in response to inbound Completions whose Requestor ID is different than the Endpoint ID programmed in the Configuration Space. These completions are not passed to the user application. The current version of the core regards an Unexpected Completion to be an Advisory Non-Fatal Error (ANFE), and no message is sent. Other types of unexpected completions are passed to the user application, and the user determines how to handle these.

Completer Abort

If the user application is unable to transmit a normal Completion in response to a Non-Posted Request it receives, it must signal `cfg_err_cpl_abort`. The `cfg_err_posted` signal can also be set to 1 simultaneously to indicate Non-Posted and the appropriate request information placed on `cfg_err_tlp_cpl_header[47:0]`. This sends a Completion with non-Successful status to the original Requester, but does not send an Error Message. When in Legacy mode if the `cfg_err_locked` signal is set to 0 (to indicate the transaction causing the error was a locked transaction), a Completion Locked with Non-Successful status is sent. If the `cfg_err_posted` signal is set to 1b (to indicate a Posted transaction), no Completion is sent, but a Non-Fatal Error Message is sent (if enabled).

Unsupported Request

If the user application receives an inbound Request it does not support or recognize, it must assert `cfg_err_ur` to signal an Unsupported Request. The `cfg_err_posted` signal must also be asserted or deasserted depending on whether the packet in question is a Posted or Non-Posted Request. If the packet is Posted, a Non-Fatal Error Message is sent out (if enabled); if the packet is Non-Posted, a Completion with a non-Successful status is sent to the original Requester. When in Legacy mode if the `cfg_err_locked` signal is set to 1b (to indicate the transaction causing the error was a locked transaction), a Completion Locked with Unsupported Request status is sent.

The Unsupported Request condition can occur for several reasons, including:

- An inbound Memory Write packet violates the user application's programming model, for example, if the user application has been allotted a 4 KB address space but only uses 3 KB, and the inbound packet addresses the unused portion. (Note: If this occurs on a Non-Posted Request, the user application should use `cfg_err_cpl_abort` to flag the error.)
- An inbound packet uses a packet Type not supported by the user application, for example, an I/O request to a memory-only device.

ECRC Error

The integrated Endpoint block core does not check the ECRC field for validity. If the user application chooses to check this field, and finds the CRC is in error, it can assert `cfg_err_ecrc`, causing a Non-Fatal Error Message to be sent.

Flow Control Credit Information

Using the Flow Control Credit Signals

The integrated Endpoint block provides the user application with information about the state of the Transaction Layer transmit and receive buffer credit pools. This information represents the current space available, as well as the credit "limit" and "consumed" information for the Posted, Non-Posted, and Completion pools.

[Table 2-8, page 20](#) defines the Flow Control Credit signals. Credit status information is presented on these signals:

- `fc_ph[7:0]`
- `fc_pd[11:0]`
- `fc_nph[7:0]`
- `fc_npdl[11:0]`
- `fc_cplh[7:0]`
- `fc_cpld[11:0]`

Collectively, these signals are referred to as `fc_*`.

The `fc_*` signals provide information about each of the six credit pools defined in the *PCI Express Base Specification: Header and Data Credits for Each of Posted, Non-Posted, and Completion*.

Six different types of flow control information can be read by the user application. The fc_sel[2:0] input selects the type of flow control information represented by the fc_* outputs. The Flow Control Information Types are shown in [Table 5-16](#).

Table 5-16: Flow Control Information Types

fc_sel[2:0]	Flow Control Information Type
000	Receive Credits Available Space
001	Receive Credits Limit
010	Receive Credits Consumed
011	Reserved
100	Transmit Credits Available Space
101	Transmit Credit Limit
110	Transmit Credits Consumed
111	Reserved

The fc_sel[2:0] signals can be changed on every clock cycle to indicate a different Flow Control Information Type. There is a two clock-cycle delay between the value of fc_sel[2:0] changing and the corresponding Flow Control Information Type being presented on the fc_* outputs. [Figure 5-28](#) illustrates the timing of the Flow Control Credits signals.

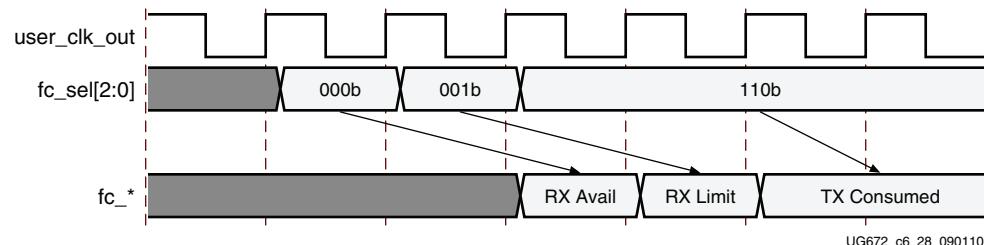


Figure 5-28: Flow Control Credits

The output values of the fc_* signals represent credit values as defined in the *PCI Express Base Specification*. One Header Credit is equal to either a 3 or 4 DWORD TLP Header and one Data Credit is equal to 16 bytes of payload data. Initial credit information is available immediately after user_Lnk_up assertion, but before the reception of any TLP. [Table 5-17](#) defines the possible values presented on the fc_* signals. Initial credit information varies depending on the size of the receive buffers within the integrated Endpoint block and the Link Partner.

Table 5-17: fc_* Value Definition

Header Credit Value	Data Credit Value	Meaning
00 – 7F	000 – 7FF	User credits
FF-80	FFF-800	Negative credits available ⁽¹⁾
7F	7FF	Infinite credits available ⁽¹⁾

Notes:

- Only Transmit Credits Available Space indicate Negative or Infinite credits available.

Receive Credit Flow Control Information

Receive Credit Flow Control Information can be obtained by setting `fc_sel[2:0]` to 000b, 001b, or 010b. The Receive Credit Flow Control information indicates the current status of the receive buffers within the integrated Endpoint block.

Receive Credits Available Space: `fc_sel[2:0] = 000b`

Receive Credits Available space shows the credit space available in the integrated Endpoint block's Transaction Layer local receive buffers for each credit pool. If space available for any of the credit pools is exhausted, the corresponding `fc_*` signal indicates a value of zero. Receive Credits Available Space returns to its initial values after the user application has drained all TLPs from the integrated Endpoint block.

In the case where infinite credits have been advertised to the Link Partner for a specific Credit pool, such as Completion Credits for Endpoints, the user application should use this value along with the methods described in [Appendix E, Managing Receive-Buffer Space for Inbound Completions](#) to avoid completion buffer overflow.

Receive Credits Limit: `fc_sel[2:0] = 001b`

Receive Credits Limit show the credits granted to the link partner. The `fc_*` values are initialized with the values advertised by the integrated Endpoint block during Flow Control initialization and are updated as a cumulative count as TLPs are read out of the Transaction Layer's receive buffers via the AXI4-Stream interface. This value is referred to as CREDITS_ALLOCATED within the *PCI Express Base Specification*.

In the case where infinite credits have been advertised for a specific credit pool, the Receive Buffer Credits Limit for that pool will always indicate zero credits.

Receive Credits Consumed: `fc_sel[2:0] = 010b`

Receive Buffer Credits Consumed show the credits consumed by the link partner (and received by the integrated Endpoint block). The initial `fc_*` values are always zero and are updated as a cumulative count, as packets are received by the Transaction Layers receive buffers. This value is referred to as CREDITS RECEIVED in the *PCI Express Base Specification*.

Transmit Credit Flow Control Information

Transmit Credit Flow Control Information can be obtained by setting `fc_sel[2:0]` to 100b, 101b, or 110b. The Transmit Credit Flow Control information indicates the current status of the receive buffers within the Link Partner.

Transmit Credits Available Space: `fc_sel[2:0] = 100b`

Transmit Credits Available Space indicates the available credit space within the receive buffers of the Link Partner for each credit pool. If space available for any of the credit pools is exhausted, the corresponding `fc_*` signal indicates a value of zero or negative. Transmit Credits Available Space returns to its initial values after the integrated Endpoint block has successfully sent all TLPs to the Link Partner.

If the value is negative, more header or data has been written into the integrated Endpoint block's local transmit buffers than the Link Partner can currently consume. Because the block does not allow posted packets to pass completions, a posted packet that is written is not transmitted if there is a completion ahead of it waiting for credits (as indicated by a zero or negative value). Similarly, a completion that is written is not transmitted if a posted packet is ahead of it waiting for credits. The user application can monitor the Transmit

Credits Available Space to ensure that these temporary blocking conditions do not occur, and that the bandwidth of the PCI Express Link is fully utilized by only writing packets to the integrated Endpoint block that have sufficient space within the Link Partner's Receive buffer. Non-Posted packets can always be bypassed within the integrated Endpoint block; so, any Posted or Completion packet written will pass Non-Posted packets waiting for credits.

The Link Partner can advertise infinite credits for one or more of the three traffic types. Infinite credits are indicated to the user by setting the Header and Data credit outputs to their maximum value as indicated in [Table 5-17](#).

Transmit Credits Limit: $fc_sel[2:0] = 101b$

Transmit Credits Limit shows the receive buffer limits of the Link Partner for each credit pool. The fc_* values are initialized with the values advertised by the Link Partner during Flow Control initialization and are updated as a cumulative count as Flow Control updates are received from the Link Partner. This value is referred to as CREDITS_LIMIT in the *PCI Express Base Specification*.

In the case where infinite credits have been advertised for a specific Credit pool, the Transmit Buffer Credits Limit always indicates zero credits for that pool.

Transmit Credits Consumed: $fc_sel[2:0] = 110b$

Transmit Credits Consumed show the credits consumed of the Receive Buffer of the Link Partner by the integrated Endpoint block. The initial value is always zero and is updated as a cumulative count, as packets are transmitted to the Link Partner. This value is referred to as CREDITS_CONSUMED in the *PCI Express Base Specification*.

Power Management

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express core supports these power management modes:

- Active State Power Management (ASPM)
- Programmed Power Management (PPM)

Implementing these power management functions as part of the PCI Express design enables the PCI Express hierarchy to seamlessly exchange power-management messages to save system power. All power management message identification functions are implemented. The sections below describe the user logic definition to support the ASPM and PPM modes of power management.

For additional information on ASPM and PPM implementation, see the *PCI Express Base Specification*.

Active State Power Management

The Active State Power Management (ASPM) functionality is autonomous and transparent from a user-logic function perspective. The core supports the conditions required for ASPM.

Programmed Power Management

To achieve considerable power savings on the PCI Express hierarchy tree, the core supports these link states of Programmed Power Management (PPM):

- L0: Active State (data exchange state)
- L1: Higher Latency, lower power standby state
- L3: Link Off State

All PPM messages are always initiated by an upstream link partner. Programming the core to a non-D0 state, results in PPM messages being exchanged with the upstream link-partner. The PCI Express Link transitions to a lower power state after completing a successful exchange.

PPM L0 State

The L0 state represents *normal* operation and is transparent to the user logic. The core reaches the L0 (active state) after a successful initialization and training of the PCI Express Link(s) as per the protocol.

PPM L1 State

These steps outline the transition of the core to the PPM L1 state:

1. The transition to a lower power PPM L1 state is always initiated by an upstream device, by programming the PCI Express device power state to D3-hot (or to D1 or D2 if they are supported).
2. The core then throttles/stalls the user logic from initiating any new transactions on the user interface by deasserting s_axis_tx_tready. Any pending transactions on the user interface are however accepted fully and can be completed later.
3. The core exchanges appropriate power management messages with its link partner to successfully transition the link to a lower power PPM L1 state. This action is transparent to the user logic.
4. All user transactions are stalled for the duration of time when the device power state is non-D0.
5. The device power state is communicated to the user logic through the user configuration port interface. The user logic is responsible for performing a successful read operation to identify the device power state.
6. The user logic, after identifying the device power state as non-D0, can initiate a request through cfg_pm_wake to the upstream link partner to configure the device back to the D0 power state.
7. The user logic must poll the PME_Status bit of the PMCSR (via the Configuration Interface). If a PME message is not acknowledged by the host within 100 ms (+50%/-5%) by the host clearing the PME_Status bit, the Endpoint is required to retransmit. This functionality is not provided by the Spartan-6 FPGA Integrated Endpoint Block for PCI Express. For more information, see section 5.3.3.3.1 of the *PCI Express Base Specification v1.1*.

Note: If the upstream link partner has not configured the device to allow the generation of PM_PME messages (PME_En bit of PMCSR = 0), the assertion of cfg_pm_wake is ignored by the core.

PPM L3 State

These steps outline the transition of the Integrated Endpoint Block for PCIe core to the PPM L3 state:

1. The core negotiates a transition to the L23 Ready Link State upon receiving a PME_Turn_Off message from the upstream link partner.

2. Upon receiving a PME_Turn_Off message, the Endpoint core initiates a handshake with the user logic through cfg_to_turnoff (see [Table 2-8, page 20](#)) and expects a cfg_turnoff_ok back from the user logic.
3. A successful handshake results in a transmission of the Power Management Turn-off Acknowledge (PME-turnoff_ack) Message by the Endpoint core to its upstream link partner.
4. The Endpoint core closes all its interfaces, disables the Physical/Data-Link/Transaction layers and is ready for removal of power to the core.

Power-down negotiation follows these steps:

1. Before power and clock are turned off, the Root Complex or the Hot-Plug controller in a downstream switch issues a PME_Turn_Off broadcast message.
2. When the Endpoint PIPE for PCIe core receives this TLP, it asserts cfg_to_turnoff to the user application and starts polling the cfg_turnoff_ok input.
3. When the user application detects the assertion of cfg_to_turnoff, it must complete any packet in progress and stop generating any new packets. After the user application is ready to be turned off, it asserts cfg_turnoff_ok to the core. After assertion of cfg_turnoff_ok, the user application has committed to being turned off.
4. The Endpoint core sends a PME_TO_Ack when it detects assertion of cfg_turnoff_ok.

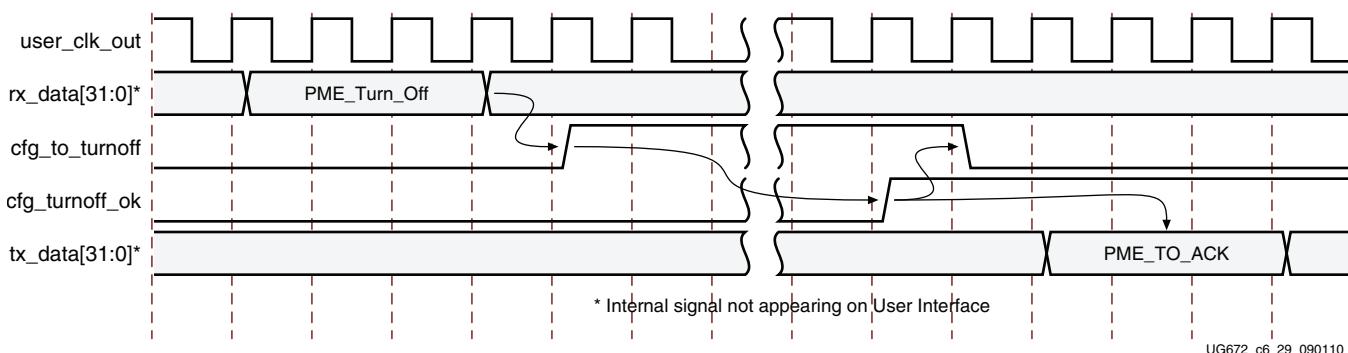


Figure 5-29: Power Management Handshaking

Generating Interrupt Requests

The integrated Endpoint block supports sending interrupt requests as either legacy interrupts or Message Signaled Interrupts (MSI). The mode is programmed using the MSI Enable bit in the Message Control Register of the MSI Capability Structure. For more information on the MSI capability structure, refer to section 6.8 of the *PCI Local Base Specification* v3.0. The state of the MSI Enable bit is reflected by the cfg_interrupt_msienable output:

- cfg_interrupt_msienable = 0: Legacy Interrupt (INTx) mode
- cfg_interrupt_msienable = 1: MSI mode

If the MSI Enable bit is set to a 1, the core generates MSI requests by sending Memory Write TLPs. If the MSI Enable bit is set to 0, the core generates legacy interrupt messages as long as the Interrupt Disable bit in the PCI Command Register is set to 0:

- cfg_command[10] = 0: interrupts enabled
- cfg_command[10] = 1: interrupts disabled (request are blocked by the core)

The user application requests interrupt service in one of two ways, each of which are described below. The user application must determine which method to use based on the value of the cfg_interrupt_msienable output. When 0, the Legacy Interrupt method must be used; when 1, the MSI method.

The MSI Enable bit in the MSI control register and the Interrupt Disable bit in the PCI Command register are programmed by the Root Complex. The user application has no direct control over these bits. Regardless of the interrupt type used, the user initiates interrupt requests through the use of cfg_interrupt and cfg_interrupt_rdy as shown in [Table 5-18](#).

Table 5-18: Interrupt Signalling

Port Name	Direction	Description
cfg_interrupt	Input	Assert to request an interrupt. Leave asserted until the interrupt is serviced.
cfg_interrupt_rdy	Output	Asserted when the core accepts the signaled interrupt request.

The user application requests interrupt service in one of two ways, each of which are described below.

MSI Mode

- As shown in [Figure 5-30](#), the user application first asserts cfg_interrupt. Additionally the user application supplies a value on cfg_interrupt_di[7:0] if Multi-Vector MSI is enabled (see below).
- The core asserts cfg_interrupt_rdy to signal that the interrupt has been accepted and the core sends a MSI Memory Write TLP. On the following clock cycle, the user application deasserts cfg_interrupt if no further interrupts are to be sent.

The MSI request is either a 32-bit addressable Memory Write TLP or a 64-bit addressable Memory Write TLP. The address is taken from the Message Address and Message Upper Address fields of the MSI Capability Structure, while the payload is taken from the Message Data field. These values are programmed by system software through configuration writes to the MSI Capability structure. When the core is configured for Multi-Vector MSI, system software can permit Multi-Vector MSI messages by programming a non-zero value to the Multiple Message Enable field.

The type of MSI TLP sent (32-bit addressable or 64-bit addressable) depends on the value of the Upper Address field in the MSI capability structure. By default, MSI messages are sent as 32-bit addressable Memory Write TLPs. MSI messages use 64-bit addressable Memory Write TLPs only if the system software programs a non-zero value into the Upper Address register.

When Multi-Vector MSI messages are enabled, the user application can override one or more of the lower-order bits in the Message Data field of each transmitted MSI TLP to differentiate between the various MSI messages sent upstream. The number of lower-order bits in the Message Data field available to the user application is determined by the lesser of the value of the Multiple Message Capable field, as set in the CORE Generator software, and the Multiple Message Enable field, as set by system software and available as the cfg_interrupt_mmenable[2:0] core output. The core masks any bits in cfg_interrupt_di[7:0] which are not configured by system software via Multiple Message Enable.

This pseudo-code shows the processing required:

```
// Value MSI_Vector_Num must be in range: 0 ≤ MSI_Vector_Num ≤  
(2^cfg_interrupt_mmenable)-1  
  
if (cfg_interrupt_msienable) {           // MSI Enabled  
    if (cfg_interrupt_mmenable > 0) {   // Multi-Vector MSI Enabled  
        cfg_interrupt_di[7:0] = {Padding_0s, MSI_Vector_Num};  
    } else {                           // Single-Vector MSI Enabled  
        cfg_interrupt_di[7:0] = Padding_0s;  
    }  
} else {  
    // Legacy Interrupts Enabled  
}
```

For example:

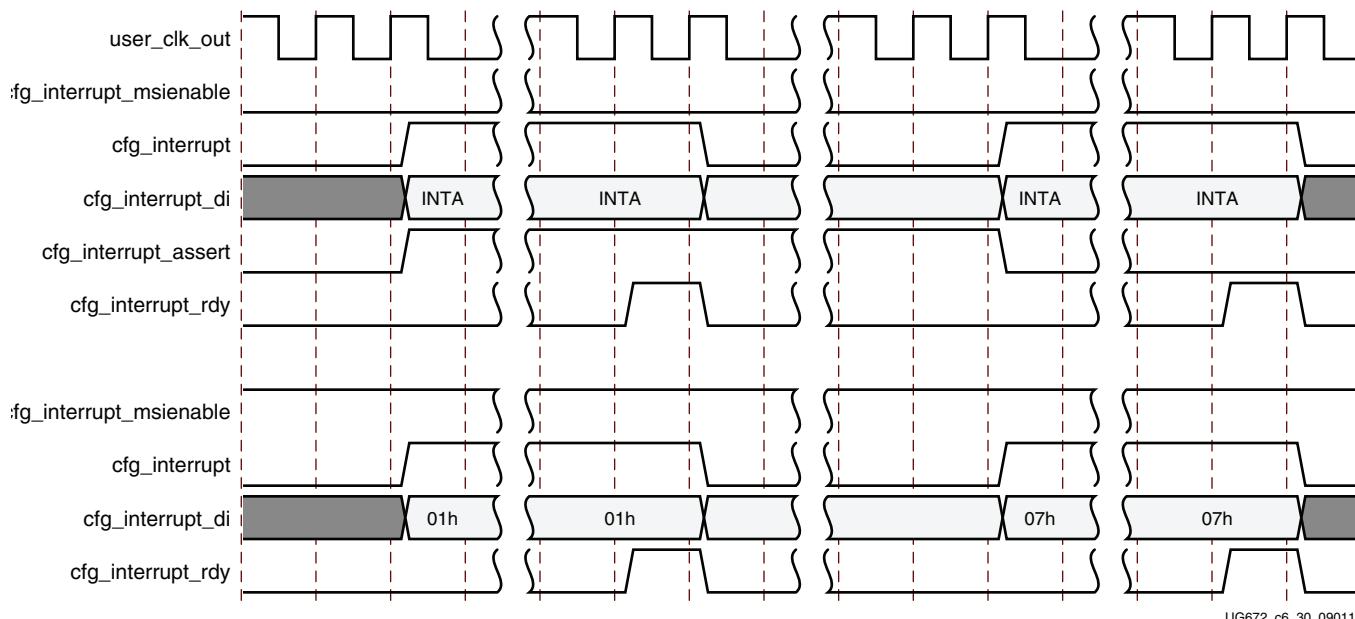
1. If cfg_interrupt_mmenable[2:0] == 000b, i.e., 1 MSI Vector Enabled, then cfg_interrupt_di[7:0] = 00h;
2. if cfg_interrupt_mmenable[2:0] == 101b, i.e., 32 MSI Vectors Enabled, then cfg_interrupt_di[7:0] = {{000b}, {MSI_Vector#}};

where MSI_Vector# is a 5-bit value and is allowed to be 00000b ≤ MSI_Vector# ≤ 11111b

Legacy Interrupt Mode

- As shown in [Figure 5-30](#), the user application first asserts cfg_interrupt and cfg_interrupt_assert to assert the interrupt. The user application should select a specific interrupt (INTA, INTB, INTC, or INTD) using cfg_interrupt_di[7:0] as shown in [Table 5-19](#).
- The core then asserts cfg_interrupt_rdy to indicate the interrupt has been accepted. On the following clock cycle, the user application deasserts cfg_interrupt and, if the Interrupt Disable bit in the PCI Command register is set to 0, the core sends an assert interrupt message (Assert_INTA, Assert_INTB, and so forth).
- After the user application has determined that the interrupt has been serviced, it asserts cfg_interrupt while deasserting cfg_interrupt_assert to deassert the interrupt. The appropriate interrupt must be indicated via cfg_interrupt_di[7:0].
- The core then asserts cfg_interrupt_rdy to indicate the interrupt deassertion has been accepted. On the following clock cycle, the user application deasserts cfg_interrupt and the core sends a deassert interrupt message (Deassert_INTA, Deassert_INTB, and so forth).

Note: The Interrupt Status bit in the Status register in the PCI Configuration Space Header is not set to reflect the state of a pending interrupt in the device/function.



[Figure 5-30: Requesting Interrupt Service: MSI and Legacy Mode](#)

[Table 5-19: Legacy Interrupt Mapping](#)

cfg_interrupt_di[7:0] value	Legacy Interrupt
00h	INTA
01h	INTB
02h	INTC
03h	INTD

Clocking and Reset of the Integrated Endpoint Block Core

Reset

The Integrated Endpoint Block for PCI Express core uses sys_reset to reset the system. This is an asynchronous reset signal asserted during the PCI Express Fundamental Reset.

Asserting this signal causes a hard reset of the entire core, including the transceivers. After the reset is released, the core attempts to link train and resume normal operation. In a typical endpoint application, for example, an add-in card, a sideband reset signal is normally present and should be connected to sys_reset. For endpoint applications that do not have a sideband system reset signal, the initial hardware reset should be generated locally. Three reset events can occur in PCI Express:

- **Cold Reset.** A Fundamental Reset that occurs at the application of power. The signal sys_reset is asserted to cause the cold reset of the core.
- **Warm Reset.** A Fundamental Reset triggered by hardware without the removal and re-application of power. The sys_reset signal is asserted to cause the warm reset to the core.
- **Hot Reset:** In-band propagation of a reset across the PCI Express Link through the protocol. In this case, sys_reset is not used. In the case of Hot Reset, the received_hot_reset signal is asserted to indicate the source of the reset.

The user application interface of the core has an output signal called user_reset_out. This signal is deasserted synchronously with respect to user_clk_out. user_reset_out is asserted as a result of any of these conditions:

- **Fundamental Reset:** Occurs (cold or warm) due to assertion of sys_reset.
- **PLL within the core:** Loses lock, indicating an issue with the stability of the clock input.
- **Loss of Transceiver PLL Lock:** The Lane 0 transceiver loses lock, indicating an issue with the PCI Express link.

The user_reset_out signal deasserts synchronously with user_clk_out after all of the above reasons are resolved, allowing the core to attempt to train and resume normal operation.

Important Notes:

1. Systems designed to the PCI Express electro-mechanical specification provide a sideband reset signal, which uses 3.3V signaling levels—see the FPGA device data sheet to understand the requirements for interfacing to such signals.
2. During the simulation starting from time 0, the integrated Endpoint block simulation model will model the behavior of the integrated Endpoint block as it comes out of power-on reset. In part of the power-on-reset modeling, all block outputs are driven High. It might appear initially that the sys_reset signal is not causing the device to reset because of the modeling of power-on reset. After the power-on-reset sequence is complete, the integrated Endpoint block should respond to all signals as specified.

Clocking

The integrated Endpoint block core input system clock signal is called sys_clk. The core accepts a 100 or 125 MHz clock input. The clock frequency used must match the clock frequency selection in the CORE Generator software GUI. For more information, see [Answer Record 18329](#). Standard PCI Express slots provide a compliant 100 MHz reference clock that can be routed directly to the FPGA without the need for any external devices.

In a typical PCI Express solution, the PCI Express reference clock is a spread spectrum clock (SSC), provided at 100 MHz. In most commercial PCI Express systems, SSC cannot be disabled. For more information regarding SSC and PCI Express, see section 4.3.1.1.1 of the *PCI Express Base Specification*.

Synchronous and Non-Synchronous Clocking

There are two ways to clock the PCI Express system:

- Using synchronous clocking, where a shared clock source is used for all devices.
- Using non-synchronous clocking, where each device has its own clock source.

Important Note: Xilinx recommends that designers use synchronous clocking when using the core. All add-in card designs must use synchronous clocking due to the characteristics of the provided reference clock. See the *Spartan-6 FPGA GTP Transceivers User Guide* and device data sheet for additional information regarding reference clock requirements.

For synchronous clocked systems, each link partner device shares the same clock source. When using the 125 MHz reference clock option, an external PLL must be used to do a multiply of 5/4 to convert the 100 MHz clock to 125 MHz, as illustrated in [Figure 5-32](#) and [Figure 5-33](#). See [Answer Record 18329](#) for more information about clocking requirements.

Further, even if the device is part of an embedded system, if the system uses commercial PCI Express root complexes or switches along with typical mother board clocking schemes, synchronous clocking should still be used as shown in [Figure 5-32](#).

[Figure 5-31](#), [Figure 5-32](#), and [Figure 5-33](#) illustrate high-level representations of the board layouts. Designers must ensure that proper coupling, termination, and so forth are used when laying out the board.

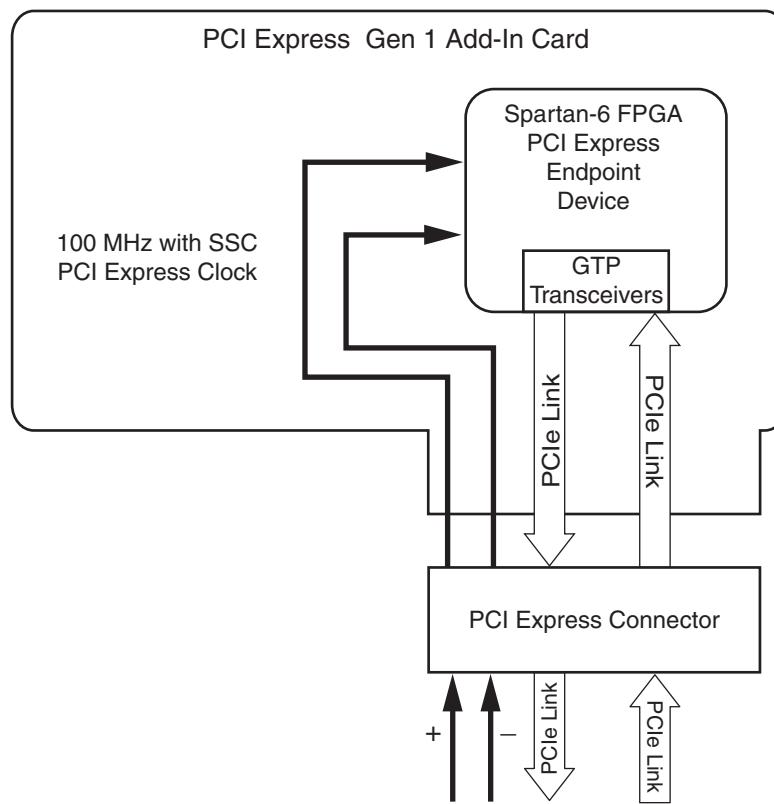


Figure 5-31: Spartan-6 FPGA PCI Express Gen 1 Using 100 MHz Reference Clock

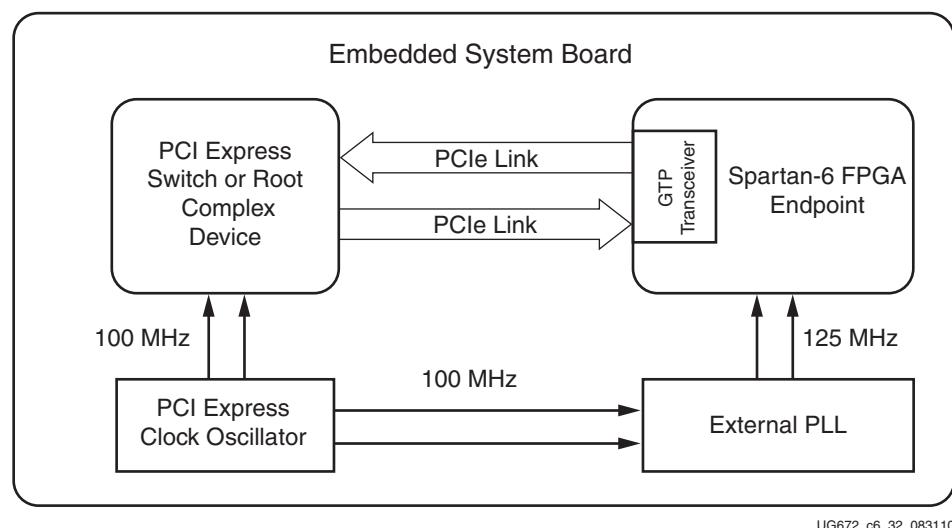
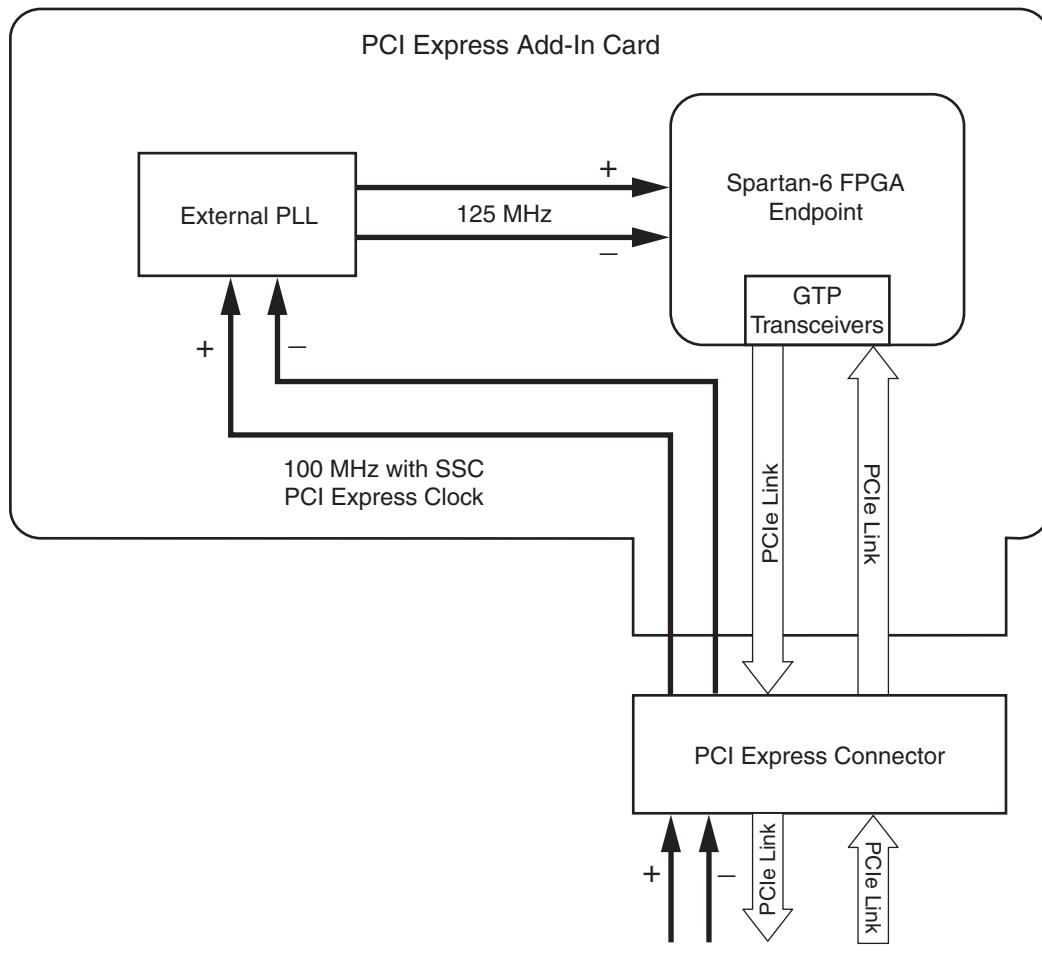


Figure 5-32: Embedded System Using 125 MHz Reference Clock



UG672_c6_33_083110

Figure 5-33: Open System Add-In Card Using 125 MHz Reference Clock

Core Constraints

The Spartan®-6 FPGA Integrated Endpoint Block for PCI Express® solution requires the specification of timing and other physical implementation constraints to meet specified performance requirements for PCI Express. These constraints are provided with the Endpoint Solution in a User Constraints File (UCF). Pinouts and hierarchy names in the generated UCF correspond to the provided example design.

To achieve consistent implementation results, a UCF containing these original, unmodified constraints must be used when a design is run through the Xilinx tools. For additional details on the definition and use of a UCF or specific constraints, see the Xilinx Libraries Guide and/or Development System Reference Guide.

Constraints provided with the integrated Endpoint block solution have been tested in hardware and provide consistent results. Constraints can be modified, but modifications should only be made with a thorough understanding of the effect of each constraint. Additionally, support is not provided for designs that deviate from the provided constraints.

Contents of the User Constraints File

Although the UCF delivered with each core shares the same overall structure and sequence of information, the content of each core's UCF varies. The sections that follow define the structure and sequence of information in a generic UCF file.

Part Selection Constraints: Device, Package, and Speed Grade

The first section of the UCF specifies the exact device for the implementation tools to target, including the specific part, package, and speed grade. In some cases, device-specific options can be included. The device in the UCF reflects the device chosen in the CORE Generator software project.

User Timing Constraints

The User Timing constraints section is not populated; it is a placeholder for the designer to provide timing constraints on user-implemented logic.

User Physical Constraints

The User Physical constraints section is not populated; it is a placeholder for the designer to provide physical constraints on user-implemented logic.

Core Pinout and I/O Constraints

The Core Pinout and I/O constraints section contains constraints for I/Os belonging to the core's System (SYS) and PCI Express (PCI_EXP) interfaces. It includes location constraints for pins and I/O logic as well as I/O standard constraints.

Core Physical Constraints

Physical constraints are used to limit the core to a specific area of the device and to specify locations for clock buffering and other logic instantiated by the core.

Core Timing Constraints

This Core Timing constraints section defines clock frequency requirements for the core and specifies which nets the timing analysis tool should ignore.

Required Modifications

Some constraints provided in the UCF utilize hierarchical paths to elements within the integrated Endpoint block core. These constraints assume an instance name of *core* for the core. If a different instance name is used, replace *core* with the actual instance name in all hierarchical constraints.

For example:

Using *xilinx_pcie_ep* as the instance name, the physical constraint

```
INST core/GT_i/tile0_gtpa1_dual_wrapper_i/gtpa1_dual_i
LOC = GTPA1_DUAL_X0Y0
```

becomes

```
INST xilinx_pcie_ep/GT_i/tile0_gtpa1_dual_wrapper_i/gtpa1_dual_i
LOC = GTPA1_DUAL_X0Y0
```

The provided UCF includes a line specifying attributes for the sys_reset pin, but it is up to the user to un-comment that line and provide a pin location. In addition, the UCF includes blank sections for constraining user-implemented logic. While the constraints provided adequately constrain the integrated Endpoint block core itself, they cannot adequately constrain user-implemented logic interfaced to the core. Additional constraints must be implemented by the designer.

Device Selection

The device selection portion of the UCF informs the implementation tools which part, package, and speed grade to target for the design. Because integrated Endpoint block cores are designed for specific part and package combinations, this section should not be modified by the designer.

The device selection section always contains a part selection line, but can also contain part or package-specific options. An example part selection line:

```
CONFIG PART = xc6slx45t-fgg484-2;
```

Core I/O Assignments

This section controls the placement and options for I/Os belonging to the core's System (SYS) interface and PCI Express (PCI_EXP) interface. NET constraints in this section control the pin location and I/O options for signals in the SYS group. Locations and options vary depending on which derivative of the core is used and should not be changed without fully understanding the system requirements.

For example:

```
NET sys_clk_p LOC = Y4;  
NET sys_clk_n LOC = Y3;
```

See [Clocking and Reset of the Integrated Endpoint Block Core, page 102](#) for detailed information about reset and clock requirements.

Each GTPA1_DUAL tile contains two transceivers. Any GTPA1_DUAL tile along the top edge of the device can be used with the integrated Endpoint block for PCIe. Either of the two transceivers in the GTPA1_DUAL tile can be used. For GTPA1_DUAL pinout information, see the *Spartan-6 FPGA GTP Transceivers User Guide*.

INST constraints are used to control placement of signals that belong to the PCI_EXP group. These constraints control the location of the transceiver(s) used, which implicitly controls pin locations for the transmit and receive differential pair. The provided transceiver wrapper file consumes *both* transceivers in a tile even though only one is used.

For example:

```
INST core/GT_i/tile0_gtpa1_dual_wrapper_i/gtpa1_dual_i LOC =  
GTPA1_DUAL_X0Y0;
```

Core Physical Constraints

Physical constraints can be included in the constraints file to control the location of clocking and memory elements. Specific physical constraints are chosen to match each supported device and package combination—it is very important to leave these constraints unmodified except for changing the hierarchical name, as described above.

Core Timing Constraints

Timing constraints are provided for all integrated Endpoint block solutions, although they differ based on core configuration. In all cases they are crucial and must not be modified, except to specify the top-level hierarchical name. Timing constraints are divided into two categories:

- **TIG constraints.** Used on paths where specific delays are unimportant, to instruct the timing analysis tools to refrain from issuing *Unconstrained Path* warnings.
- **Frequency constraints.** Group clock nets into time groups and assign properties and requirements to those groups.

TIG constraints example:

```
NET sys_reset TIG;
```

Clock constraints example:

First, the input reference clock period is specified, which can be either 100 or 125 MHz (selected in the CORE Generator™ software GUI).

```
NET sys_clk_c PERIOD = 8ns;
```

Next, the internally generated clock net and period is specified, which can be 100 or 125 MHz. (*Both* clock constraints must be specified as having the same period.)

```
NET core/gt_refclk_out(0) TNM_NET = GT_REFCLK_OUT ;
TIMESPEC TS_GT_REFCLK_OUT = PERIOD GT_REFCLK_OUT 8ns HIGH 50 % ;
```

Relocating the Integrated Endpoint Block

While Xilinx does not provide technical support for designs whose system clock input, GTP transceivers, or block RAM locations are different from the provided examples, it is possible to relocate the core within the FPGA. The locations selected in the provided examples are the recommended pinouts. These locations have been chosen based on the proximity to the Endpoint Block, which enables meeting timing, and because they are conducive to layout requirements for add-in card design. If the core is moved, the relative location of all transceivers and clocking resources should be maintained to ensure timing closure.

Supported Core Pinouts

Table 6-1 defines the supported core pinouts for the available LXT part and package combinations. The CORE Generator software provides a UCF for the selected part and package that matches the content of this table.

Table 6-1: Spartan-6 FPGA LXT Pinout

Package	Part	GTPA1_DUAL	Channel	sys_clk_n	sys_clk_p	pci_exp_txn	pci_exp_txp	pci_exp_rxn	pci_exp_rxp
CSG324	XC6SLX25T	X0Y0	0	A8	B8	A4	B4	C5	D5
	XC6SLX45T		1	C9	D9	A6	B6	C7	D7
	XC6SLX25T	X1Y0	0	A10	B10	A12	B12	C11	D11
	XC6SLX45T		1	E10	F10	A14	B14	C13	D13
	XC6SLX45T	X0Y0	0	A10	B10	A6	B6	C7	D7
	XC6SLX45T		1	C11	D11	A8	B8	C9	D9
CSG484	XC6SLX45T	X1Y0	0	A12	B12	A14	B14	C13	D13
	XC6SLX45T		1	E14	F14	A16	B16	C15	D15

Table 6-1: Spartan-6 FPGA LXT Pinout (Cont'd)

Package	Part	GTPA1_DUAL	Channel	sys_clk_n	sys_clk_p	pci_exp_txn	pci_exp_txp	pci_exp_rxn	pci_exp_rxp
CSG484	XC6SLX75T	X0Y1	0	A10	B10	A6	B6	C7	D7
	XC6SLX100T								
	XC6SLX150T								
	XC6SLX75T		1	C11	D11	A8	B8	C9	D9
	XC6SLX100T								
	XC6SLX150T	X1Y1	0	A12	B12	A14	B14	C13	D13
	XC6SLX75T								
	XC6SLX100T		1	E14	F14	A16	B16	C15	D15
	XC6SLX150T								
FGG484	XC6SLX25T	X0Y0	0	B10	A10	A6	B6	C7	D7
	XC6SLX45T								
	XC6SLX25T	X1Y0	0	D11	C11	A8	B8	C9	D9
	XC6SLX45T								
	XC6SLX45T	X1Y0	1	F12	E12	A16	B16	C15	D15
	XC6SLX45T								
FGG484	XC6SLX75T	X0Y1	0	B10	A10	A6	B6	C7	D7
	XC6SLX100T								
	XC6SLX150T		1	D11	C11	A8	B8	C9	D9
	XC6SLX75T								
	XC6SLX100T		X1Y1	0	B12	A12	A14	B14	C13
	XC6SLX150T								
	XC6SLX75T			1	F12	E12	A16	B16	C15
	XC6SLX100T								
	XC6SLX150T								

Table 6-1: Spartan-6 FPGA LXT Pinout (Cont'd)

Package	Part	GTPA1_DUAL	Channel	sys_clk_n	sys_clk_p	pci_exp_txn	pci_exp_txp	pci_exp_rxn	pci_exp_rxp
FGG676	XC6SLX75T	X0Y1	0	A10	B10	A6	B6	C7	D7
	XC6SLX100T								
	XC6SLX150T								
	XC6SLX75T		1	C11	D11	A8	B8	C9	D9
	XC6SLX100T								
	XC6SLX150T	X1Y1	0	C15	D15	A18	B18	C17	D17
	XC6SLX75T								
	XC6SLX100T		1	A16	B16	A20	B20	C19	D19
	XC6SLX150T								
FGG900	XC6SLX100T	X0Y1	0	A13	B13	A9	B9	C10	D10
	XC6SLX150T								
	XC6SLX100T		1	C14	D14	A11	B11	C12	D12
	XC6SLX150T								
	XC6SLX100T	X1Y1	0	C18	D18	A21	B21	C20	D20
	XC6SLX150T								
	XC6SLX100T		1	A19	B19	A23	B23	C22	D22
	XC6SLX150T								

FPGA Configuration

This chapter discusses how to configure the Spartan®-6 FPGA so that the device can link up and be recognized by the system. This information is provided so the user can choose the correct FPGA configuration method for the system and verify that it works as expected.

This chapter discusses how specific requirements of the *PCI Express Base Specification* and *PCI Express Card Electromechanical Specification* apply to FPGA configuration. Where appropriate, Xilinx recommends that the user read the actual specifications for detailed information. This chapter is divided into four sections:

- **Configuration Terminology.** Defines terms used in this chapter.
- **Configuration Access Time.** Several specification items govern when an Endpoint device needs to be ready to receive configuration accesses from the host (Root Complex).
- **Board Power in Real-World Systems.** Understanding real-world system constraints related to board power and how they affect the specification requirements.
- **Recommendations.** Describes methods for FPGA configuration and includes sample problem analysis for FPGA configuration timing issues.

Configuration Terminology

In this chapter, these terms are used to differentiate between FPGA configuration and configuration of the PCI Express device:

- **Configuration of the FPGA.** *FPGA configuration* is used.
- **Configuration of the PCI Express device.** After the link is active, *configuration* is used.

Configuration Access Time

In standard systems for PCI Express, when the system is powered up, configuration software running on the processor starts scanning the PCI Express bus to discover the machine topology.

The process of scanning the PCI Express hierarchy to determine its topology is referred to as the *enumeration process*. The root complex accomplishes this by initiating configuration transactions to devices as it traverses and determines the topology.

All PCI Express devices are expected to have established the link with their link partner and be ready to accept configuration requests during the enumeration process. As a result, there are requirements as to when a device needs to be ready to accept configuration requests after power up; if the requirements are not met, the following occurs:

- If a device is not ready and does not respond to configuration requests, the root complex does not discover it and treats it as non-existent.
- The operating system does not report the device's existence and the user's application is not able to communicate with the device.

Choosing the appropriate FPGA configuration method is key to ensuring the device is able to communicate with the system in time to achieve link up and respond to the configuration accesses.

Configuration Access Specification Requirements

Two PCI Express specification items are relevant to configuration access:

1. Section 6.6 of *PCI Express Base Specification*, rev 1.1 states “A system must guarantee that all components intended to be software visible at boot time are ready to receive Configuration Requests within 100 ms of the end of Fundamental Reset at the Root Complex.” For detailed information about how this is accomplished, see the specification; it is beyond the scope of this discussion.

Xilinx compliance to this specification is validated by the PCI Express-CV tests. The [PCI Special Interest Group \(PCI-SIG\)](#) provides the PCI Express Configuration Test Software to verify the device meets the requirement of being able to receive configuration accesses within 100 ms of the end of the fundamental reset. The software, available to any member of the PCI-SIG, generates several resets using the in-band reset mechanism and PERST# toggling to validate robustness and compliance to the specification.

2. Section 6.6 of *PCI Express Base Specification*, rev 1.1 defines three parameters necessary “where power and PERST# are supplied.” The parameter T_{PVPERL} applies to FPGA configuration timing and is defined as:

T_{PVPERL} - PERST# must remain active at least this long after power becomes valid.

The *PCI Express Base Specification* does not give a specific value for T_{PVPERL} – only its meaning is defined. The most common form factor used by designers with the integrated Endpoint block core is an ATX-based form factor. The *PCI Express Card Electromechanical Specification* focuses on requirements for ATX-based form factors. This applies to most designs targeted to standard desktop or server type motherboards. Figure 7-1 shows the relationship between Power Stable and PERST#. (This figure is based on Figure 2-10 from section 2.1 of *PCI Express Card Electromechanical Specification*, rev 1.1.)

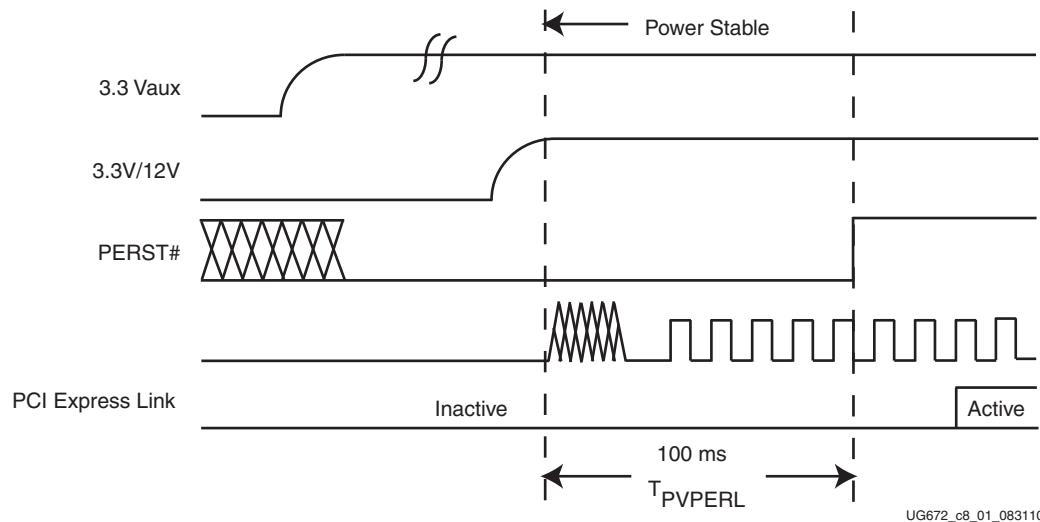


Figure 7-1: Power Up

Section 2.6.2 of the *PCI Express Card Electromechanical Specification* defines T_{PVPERL} as a minimum of 100 ms, indicating that from the time power is stable the system reset is asserted for at least 100 ms (as shown in Table 7-1).

Table 7-1: T_{PVPERL} Specification

Symbol	Parameter	Min	Max	Units
T_{PVPERL}	Power stable to PERST# inactive	100		ms

From Figure 7-1 and Table 7-1, it is possible to obtain a simple equation to define the FPGA configuration time as follows:

$$\text{FPGA Configuration Time} \leq T_{PWRVLD} + T_{PVPERL} \quad \text{Equation 7-1}$$

Given that T_{PVPERL} is defined as 100 ms minimum, this becomes:

$$\text{FPGA Configuration Time} \leq T_{PWRVLD} + 100 \text{ ms} \quad \text{Equation 7-2}$$

Note: Although T_{PWRVLD} is included in Equation 7-2, it has yet to be defined in this discussion because it depends on the type of system in use. The [Board Power in Real-World Systems](#) section defines T_{PWRVLD} for both ATX-based and non ATX-based systems.

FPGA configuration time is only relevant at cold boot; subsequent warm or hot resets do not cause reconfiguration of the FPGA. If the design appears to be having problems due to FPGA configuration, the user should issue a warm reset as a simple test, which resets the system, including the PCI Express link, but keeps the board powered. If the problem does not appear, the issue could be FPGA configuration time related.

Board Power in Real-World Systems

Several boards are used in PCI Express systems. The *ATX Power Supply Design* specification, endorsed by Intel, is used as a guideline and for this reason followed in the majority of mother boards and 100% of the time if it is an Intel-based motherboard. The relationship between power rails and power valid signaling is described in the [ATX 12V Power Supply Design Guide](#). Figure 7-2, redrawn here and simplified to show the information relevant to FPGA configuration, is based on the information and diagram found in section 3.3 of the *ATX 12V Power Supply Design Guide*. For the entire diagram and definition of all parameters, see the *ATX 12V Power Supply Design Guide*.

Figure 7-2 shows that power stable indication from Figure 7-1 for the PCI Express system is indicated by the assertion of PWR_OK. PWR_OK is asserted High after some delay once the power supply has reached 95% of nominal.

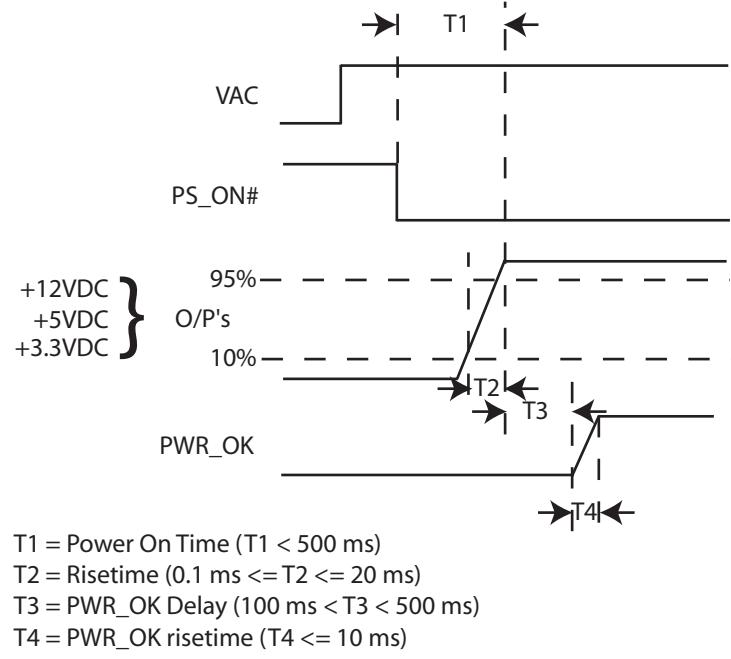


Figure 7-2: ATX Power Supply

Figure 7-2 shows that power is actually valid before PWR_OK is asserted High. This is represented by T3 and is the PWR_OK delay. The *ATX 12V Power Supply Design Guide* defines PWR_OK as $100 \text{ ms} < T3 < 500 \text{ ms}$, indicating the following: From the point at which the power level reaches 95% of nominal, there is a minimum of at least 100 ms but no more than 500 ms of delay before PWR_OK is asserted. Remember, according to the *PCI Express Card Electromechanical Specification*, the PERST# is guaranteed to be asserted a minimum of 100 ms from when power is stable indicated in an ATX system by the assertion of PWR_OK.

Again, the FPGA configuration time equation is:

$$\text{FPGA Configuration Time} \leq T_{\text{PWRVLD}} + 100 \text{ ms} \quad \text{Equation 7-3}$$

T_{PWRVLD} is defined as PWR_OK delay period, that is, T_{PWRVLD} represents the amount of time that power is valid in the system before PWR_OK is asserted. This time can be added to the amount of time the FPGA has to configure. The minimum values of T2 and T4 are negligible and considered zero for purposes of these calculations. For ATX-based

motherboards, which represent the majority of real-world motherboards in use, T_{PWRVLD} can be defined as:

$$100 \text{ ms} \leq T_{PWRVLD} \leq 500 \text{ ms}$$

This provides the following requirement for FPGA configuration time in both ATX and non-ATX-based motherboards:

- FPGA Configuration Time \leq 200 ms (for ATX based motherboard)
- FPGA Configuration Time \leq 100 ms (for non-ATX based motherboard)

The second equation for the non-ATX based motherboards assumes a T_{PWRVLD} value of 0 ms because it is not defined in this context. Designers with non-ATX based motherboards should evaluate their own power supply design to obtain a value for T_{PWRVLD} .

This chapter assumes that the FPGA power (V_{CCINT}) is stable before or at the same time as PWR_OK is asserted. If this is not the case, additional time must be subtracted from the available time for FPGA configuration. Xilinx recommends to avoid designing add-in cards that have staggered voltage regulators with long delays.

Hot-Plug Systems

Hot-plug systems generally employ the use of a hot-plug power controller located on the system motherboard. Many discrete hot-plug power controllers extend T_{PVPERL} beyond the minimum 100 ms. Add-in card designers should consult the hot-plug power controller data sheet to determine the value of T_{PVPERL} . If the hot-plug power controller is unknown, a T_{PVPERL} value of 100 ms should be assumed.

Recommendations

Xilinx recommends using a Quad-SPI Flash device in Master Serial/SPI mode with a CCLK frequency of 33 MHz, which allows time for the FPGA configuration of the Spartan-6 FPGA in ATX-based motherboards. Configuration options are shown as green cells in [Table 7-2](#) and [Table 7-3](#) depending on the type of system in use. This section discusses these recommendations and includes sample analysis of potential problems that might arise during FPGA configuration.

FPGA Configuration Times for Spartan-6 Devices

During power up, the FPGA configuration sequence is performed in four steps:

1. Wait for POR (Power on Reset) for all voltages (V_{CCINT} , V_{CCAUX} , and V_{CCO}) in the FPGA to trip, referred to as POR Trip Time
2. Wait for completion (deassertion) of INIT to allow the FPGA to initialize before accepting a bitstream transfer.

Note: As a general rule, steps 1 and 2 require \leq 50ms

3. Wait for assertion of DONE, the actual time required for a bitstream to transfer, and depends on:
 - Bitstream size
 - Clock frequency
 - Transfer mode used in the Flash Device
 - SPI = Serial Peripheral Interface
 - BPI = Byte Peripheral Interface

- PFP = Platform Flash PROMs

For detailed information about the configuration process, see the *Spartan-6 FPGA Configuration User Guide* (UG380).

Table 7-2 and **Table 7-3** show the comparative data for all Spartan-6 FPGA LXT devices with respect to a variety of flash devices and programming modes. The default clock rate for configuring the device is always 2 MHz. Any reference to a different clock rate implies a change in the settings of the device being used to program the FPGA. The configuration clock (CCLK), when driven by the FPGA, has variation and is not exact. See the *Spartan-6 FPGA Configuration User Guide* for more information on CCLK tolerances.

Configuration Time Matrix: ATX Motherboards

Table 7-2 shows the configuration methods that allow the device to be configured before the end of the fundamental reset in ATX-based systems. The table values represent the bitstream transfer time only. The matrix is color-coded to show which configuration methods allow the device to configure within 200 ms once the FPGA initialization time is included. Choose a configuration method shaded in green when using ATX-based systems to ensure that the device is recognized.

Table 7-2: Configuration Time Matrix (ATX Motherboards): Spartan-6 FPGA Bitstream Transfer Time in Milliseconds

Spartan-6 FPGA	Bitstream (Bits)	SPIx4 ⁽¹⁾	XCF32P ⁽²⁾ (Slave-SMAPx8)
XC6SLX25T	6,411,440	36	25
XC6SLX45T	11,875,104	66	45
XC6SLX75T	19,624,608	110	75
XC6SLX100T	26,543,136	148	101
XC6SLX150T	33,761,568	188	128
GREEN: Bitstream Transfer Time + FPGA INIT Time (50 ms) <= 200 ms			
YELLOW: Bitstream Transfer Time + FPGA INIT Time (50 ms) > 200 ms			
RED: Bitstream Transfer Time > 200 ms			

Notes:

1. CCLK assumptions: 45 MHz
2. CCLK assumptions: 33 MHz

Configuration Time Matrix: Non-ATX-Based Motherboards

Table 7-3 shows the configuration methods that allow the device to be configured before the end of the fundamental reset in non-ATX-based systems. This assumes T_{PWRVLD} is zero. The table values represent the bitstream transfer time only. The matrix is color-coded to show which configuration methods allow the device to configure within 100 ms once the FPGA initialization time is included. Choose a configuration method shaded in green when using non-ATX-based systems to ensure that the device is recognized.

For some of the larger FPGAs, it might not be possible to configure within the 100 ms window. In these cases, the user system needs to be evaluated to see if any margin is available that can be assigned to T_{PWRVLD} .

Table 7-3: Configuration Time Matrix (Generic Platforms: Non-ATX Motherboards): Spartan-6 FPGA Bitstream Transfer Time in Milliseconds

Spartan-6 FPGA	Bitstream (Bits)	SPIx4 ⁽¹⁾	XCF32P ⁽²⁾ (Slave-SMAPx8)
XC6SLX25T	6,411,440	36	25
XC6SLX45T	11,875,104	66	45
XC6SLX75T	19,624,608	110	75
XC6SLX100T	26,543,136	148	101
XC6SLX150T	33,761,568	188	128
GREEN: Bitstream Transfer Time + FPGA INIT Time (50 ms) <= 100 ms			
YELLOW: Bitstream Transfer Time + FPGA INIT Time (50 ms) > 100 ms			
RED: Bitstream Transfer Time > 100 ms			

Notes:

1. CCLK assumptions: 45 MHz
2. CCLK assumptions: 33 MHz

Sample Problem Analysis

This section presents data from an ASUS PL5 system to demonstrate the relationships between Power Valid, FPGA Configuration, and PERST#. [Figure 7-3](#) shows a case where the endpoint failed to be recognized due to a FPGA configuration time issue. [Figure 7-4](#) shows a successful FPGA configuration with the endpoint being recognized by the system.

Failed FPGA Recognition

[Figure 7-3](#) illustrates a failed cold boot test using the default configuration time on an **XC6SLX50T** FPGA. In this example, the host failed to recognize the Xilinx FPGA. Although a second PERST# pulse assists in allowing more time for the FPGA to configure, the slowness of the FPGA configuration clock (2 MHz) causes configuration to complete well after this second deassertion. During this time, the system enumerated the bus and did not recognize the FPGA.

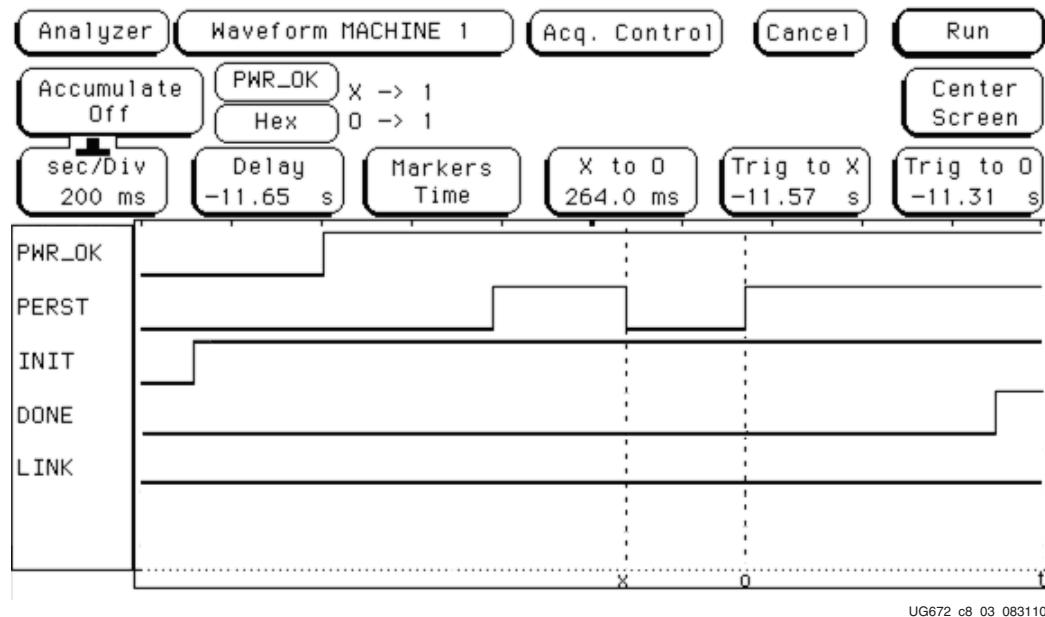


Figure 7-3: Default Configuration Time on LX50T Device (2 MHz Clock)

Successful FPGA Recognition

Figure 7-4 illustrates a successful cold boot test on the same system. In this test, the CCLK was running at 50 MHz, allowing the FPGA to configure in time to be enumerated and recognized. The figure shows that the FPGA began initialization approximately 250 ns before PWR_OK. DONE going High shows that the FPGA was configured even before PWR_OK was asserted.

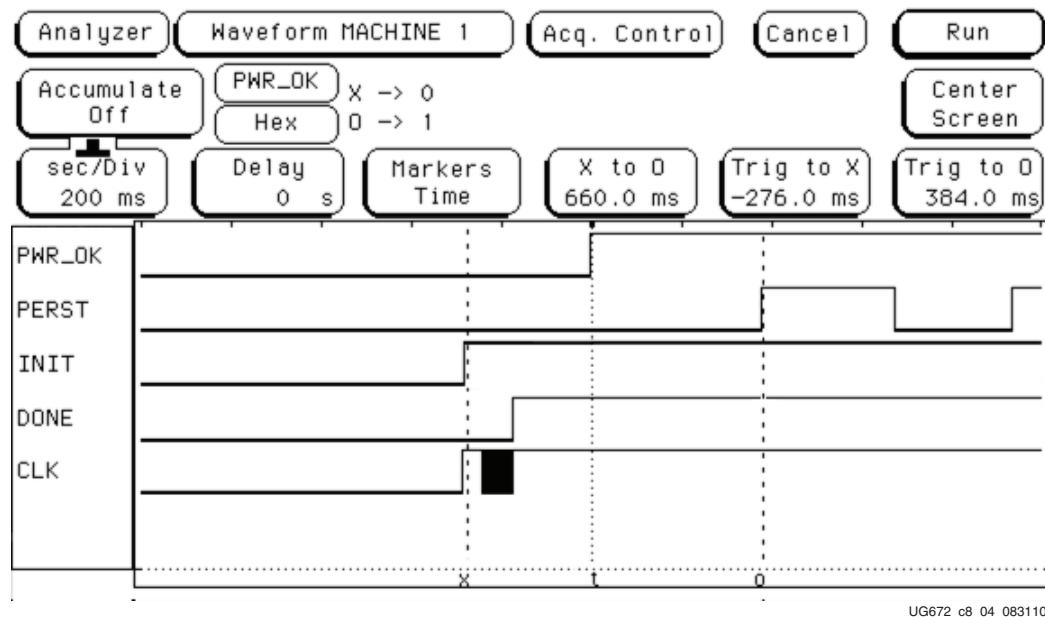


Figure 7-4: Fast Configuration Time on LX50T Device (50 MHz Clock)

Workarounds for Closed Systems

For failing FPGA configuration combinations, as represented by pink cells and yellow cells in [Table 7-2](#) and [Table 7-3](#), designers might be able to work around the problem in closed systems or systems where they can guarantee behavior. These options are not recommended for products where the targeted end system is unknown.

1. Check if the motherboard and BIOS generate multiple PERST# pulses at start-up. This can be determined by capturing the signal on the board using an oscilloscope. (This is similar to what is shown in [Figure 7-3](#). If multiple PERST#s are generated, this typically adds extra time for FPGA configuration.

Define $T_{PERSTPERIOD}$ as the total sum of the pulse width of PERST# and deassertion period before the next PERST# pulse arrives. Because the FPGA is not power cycled or reconfigured with additional PERST# assertions, the $T_{PERSTPERIOD}$ number can be added to the FPGA configuration equation.

$$\text{FPGA Configuration Time} \leq T_{PWRVLD} + T_{PERSTPERIOD} + 100 \text{ ms}$$

2. In closed systems, it might be possible to create scripts to force the system to perform a warm reset after the FPGA is configured, after the initial power up sequence. This resets the system along with the PCI Express sub-system allowing the device to be recognized by the system.

Known Restrictions

This chapter describes restrictions or issues where the integrated Endpoint block deviates from the *PCI Express Base Specification*, Rev.1.1, or in cases where the specification is ambiguous. All issues listed in this chapter are considered low impact and are not a concern for most applications. The Comments sections describe where the associated problem might occur so that designers can decide quickly if further investigation is needed.

Master Data Parity Error Bit Set Incorrectly

The Master Data Parity Error bit of the Status Register is erroneously set when the Error Poisoned status bit is set in Completion with Data TLPs.

Area of Impact

Configuration Space

Detailed Description

Transmitting a Completion with Data TLP (CplID) with the Error Poisoned (EP) bit set to 1b causes the Master Data Parity Error bit (bit 8) in the Status Register (PCI Configuration Space Header address 0x006) to be set. This is not allowed for Endpoints.

Comments

There are no hardware-related side effects to setting the Master Data Parity Error bit.

System software effects are system dependent; however, it is unlikely that software will react to this bit being set in an Endpoint.

Non-Posted UpdateFC During PPM Transition

A Non-Posted UpdateFC DLLP is not sent immediately after the link transitions from power state L1 to L0 (due to PPM non-D0).

Area of Impact

Programmed Power Management (PPM)

Detailed Description

When the link partner has used up all of the integrated Endpoint block's Non-Posted credits, and then places the integrated Endpoint block into a non-D0 PPM state, the integrated Endpoint block must immediately send a Non-Posted UpdateFC DLLP following an exit from state L1 to L0 because of a PPM change from the non-D0 state. The integrated Endpoint block does not send the Non-Posted UpdateFC immediately upon entry to D0. It waits for an internal timer to time-out, which leads to temporary reduced performance.

Comments

The probability of this error occurring is extremely low. The link partner would have to deplete all non-posted credits in the integrated Endpoint block, and then immediately put the integrated Endpoint block into a non-D0 PPM state. It is unlikely that users would want to change to a non-D0 PPM state while there are outstanding non-posted requests.

To avoid this temporary condition of reduced performance, users should ensure there are no outstanding non-posted requests before moving to a non-D0 PPM state.

Additional Resources

Xilinx Resources

For support resources such as Answers, Documentation, Downloads, and Forums, see the Xilinx Support website at:

<http://www.xilinx.com/support>.

For a glossary of technical terms used in Xilinx documentation, see:

<http://www.xilinx.com/company/terms.htm>.

Solution Centers

See the [Xilinx Solution Centers](#) for support on devices, software tools, and intellectual property at all stages of the design cycle. Topics include design assistance, advisories, and troubleshooting tips.

The Solution Center specific to the Spartan®-6 FPGA Integrated Block for PCI Express core is located at [Xilinx Solution Center for PCI Express](#).

Additional Core Resources

For additional documentation specific to the Spartan-6 FPGA Integrated Block, refer to the [Spartan-6 FPGA Integrated Block for PCI Express documentation page](#).

Additional information and resources related to the PCI Express technology are available from these websites:

- [PCI Express at PCI-SIG](#)
- [Developer Network for PCI Express](#)

The AMBA® AXI4-Stream specification is available at:

- [AMBA AXI4-Stream Protocol Specification](#)

Programmed Input/Output Example Design

Programmed Input/Output (PIO) transactions are generally used by a PCI Express® system host CPU to access Memory Mapped Input/Output (MMIO) and Configuration Mapped Input/Output (CMIO) locations in the PCI Express fabric. Endpoints for PCI Express accept Memory and I/O Write transactions and respond to Memory and I/O Read transactions with Completion with Data transactions.

The PIO example design (PIO design) is included with the Endpoint for PCI Express generated by the CORE Generator™ software, which allows users to bring up their system board with a known established working design to verify the link and functionality of the board.

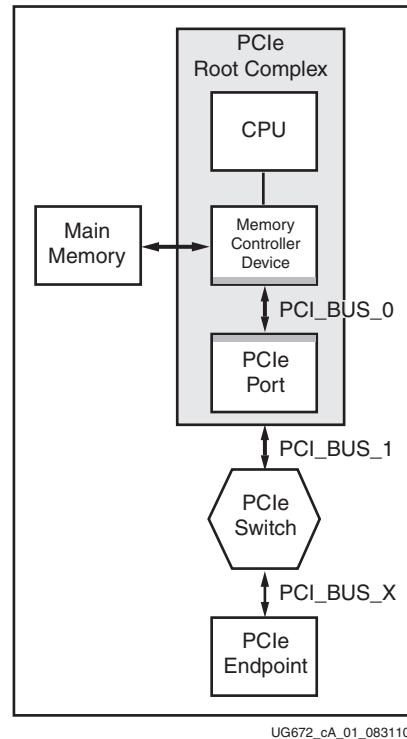
Note: The PIO design Port Model is shared by the Spartan®-6 FPGA Integrated Endpoint Block for PCI Express, Endpoint Block Plus for PCI Express, and Endpoint PIPE for PCI Express solutions. This appendix represents all the solutions generically using the name Endpoint for PCI Express (or Endpoint for PCIe®).

System Overview

The PIO design is a simple target-only application that interfaces with the Endpoint for PCIe core's Transaction (AXI4-Stream) interface and is provided as a starting point for customers to build their own designs. These features are included:

- Four transaction-specific 2 KB target regions using the internal Xilinx FPGA block RAMs, providing a total target space of 8192 bytes
- Supports single DWORD payload Read and Write PCI Express transactions to 32-/64-bit address memory spaces and I/O space with support for completion TLPs
- Utilizes the core's bar_hit[6:0] signals to differentiate between TLP destination Base Address Registers
- Provides separate implementations optimized for 32-bit, 64-bit, and 128-bit AXI4-Stream interfaces

Figure B-1 illustrates the PCI Express system architecture components, consisting of a Root Complex, a PCI Express switch device, and an Endpoint for PCIe. PIO operations move data *downstream* from the Root Complex (CPU register) to the Endpoint, and/or *upstream* from the Endpoint to the Root Complex (CPU register). In either case, the PCI Express protocol request to move the data is initiated by the host CPU.



UG672_cA_01_083110

Figure B-1: System Overview

Data is moved downstream when the CPU issues a store register to a MMIO address command. The Root Complex typically generates a Memory Write TLP with the appropriate MMIO location address, byte enables and the register contents. The transaction terminates when the Endpoint receives the Memory Write TLP and updates the corresponding local register.

Data is moved upstream when the CPU issues a load register from a MMIO address command. The Root Complex typically generates a Memory Read TLP with the appropriate MMIO location address and byte enables. The Endpoint generates a Completion with Data TLP after it receives the Memory Read TLP. The Completion is steered to the Root Complex and payload is loaded into the target register, completing the transaction.

PIO Hardware

The PIO design implements a 8192 byte target space in FPGA block RAM, behind the Endpoint for PCIe. This 32-bit target space is accessible through single DWORD I/O Read, I/O Write, Memory Read 64, Memory Write 64, Memory Read 32, and Memory Write 32 TLPs.

The PIO design generates a completion with 1 DWORD of payload in response to a valid Memory Read 32 TLP, Memory Read 64 TLP, or I/O Read TLP request presented to it by the core. In addition, the PIO design returns a completion without data with successful status for I/O Write TLP request.

The PIO design processes a Memory or I/O Write TLP with 1 DWORD payload by updating the payload into the target address in the FPGA block RAM space.

Base Address Register Support

The PIO design supports four discrete target spaces, each consisting of a 2 KB block of memory represented by a separate Base Address Register (BAR). Using the default parameters, the CORE Generator software produces a core configured to work with the PIO design defined in this section, consisting of:

- One 64-bit addressable Memory Space BAR
- One 32-bit Addressable Memory Space BAR

Users can change the default parameters used by the PIO design; however, in some cases they might need to change the back-end user application depending on their system. See [Changing CORE Generator Tool Default BAR Settings](#) for information about changing the default CORE Generator software parameters and the effect on the PIO design.

Each of the four 2 KB address spaces represented by the BARs corresponds to one of four 2 KB address regions in the PIO design. Each 2 KB region is implemented using a 2 KB dual-port block RAM. As transactions are received by the core, the core decodes the address and determines which of the four regions is being targeted. The core presents the TLP to the PIO design and asserts the appropriate bits of `bar_hit[6:0]`, as defined in [Table B-1](#).

Table B-1: TLP Traffic Types

Block RAM	TLP Transaction Type	Default BAR	<code>bar_hit[6:0]</code>
ep_io_mem	I/O TLP transactions	Disabled	Disabled
ep_mem_32	32-bit address Memory TLP transactions	2	111_1011b
ep_mem64	64-bit address Memory TLP transactions	0-1	111_1100b
ep_mem_erm	32-bit address Memory TLP transactions destined for EROM	Exp. ROM	011_1111b

Changing CORE Generator Tool Default BAR Settings

Users can change the CORE Generator tool parameters and continue to use the PIO design to create customized Verilog or VHDL source to match the selected BAR settings.

However, because the PIO design parameters are more limited than the core parameters, consider these example design limitations when changing the default CORE Generator tool parameters:

- The example design supports one I/O space BAR, one 32-bit Memory space (that cannot be the Expansion ROM space), and one 64-bit Memory space. If these limits are exceeded, only the first space of a given type is active—accesses to the other spaces do not result in completions.
- Each space is implemented with a 2 KB memory. If the corresponding BAR is configured to a wider aperture, accesses beyond the 2 KB limit wrap around and overlap the 2 KB memory space.
- The PIO design supports one I/O space BAR, which by default is disabled, but can be changed if desired.

Although there are limitations to the PIO design, Verilog or VHDL source code is provided so the user can tailor the example design to their specific needs.

TLP Data Flow

This section defines the data flow of a TLP successfully processed by the PIO design. For detailed information about the interface signals within the sub-blocks of the PIO design, see [Receive Path, page 134](#) and [Transmit Path, page 136](#).

The PIO design successfully processes single DWORD payload Memory Read and Write TLPs and I/O Read and Write TLPs. Memory Read or Memory Write TLPs of lengths larger than one DWORD are not processed correctly by the PIO design; however, the core *does* accept these TLPs and passes them along to the PIO design. If the PIO design receives a TLP with a length of greater than 1 DWORD, the TLP is received completely from the core and discarded. No corresponding completion is generated.

Memory and I/O Write TLP Processing

When the Endpoint for PCIe receives a Memory or I/O Write TLP, the TLP destination address and transaction type are compared with the values in the core BARs. If the TLP passes this comparison check, the core passes the TLP to the Receive AXI4-Stream interface of the PIO design. The PIO design handles Memory writes and I/O TLP writes in different ways: the PIO design responds to *I/O writes* by generating a Completion Without Data (cpl), a requirement of the PCI Express specification.

Along with the start of packet, end of packet, and ready handshaking signals, the Receive AXI4-Stream interface also asserts the appropriate bar_hit[6:0] signal to indicate to the PIO design the specific destination BAR that matched the incoming TLP. On reception, the PIO design's RX State Machine processes the incoming Write TLP and extracts the TLPs data and relevant address fields so that it can pass this along to the PIO design's internal block RAM write request controller.

Based on the specific bar_hit[6:0] signal asserted, the RX State Machine indicates to the internal write controller the appropriate 2 KB block RAM to use prior to asserting the write enable request. For example, if an I/O Write Request is received by the core targeting BAR0, the core passes the TLP to the PIO design and asserts bar_hit[0]. The RX State machine extracts the lower address bits and the data field from the I/O Write TLP and instructs the internal Memory Write controller to begin a write to the block RAM.

In this example, the assertion of bar_hit[0] instructed the PIO memory write controller to access ep_mem0 (which by default represents 2 KB of I/O space). While the write is being carried out to the FPGA block RAM, the PIO design RX state machine deasserts the m_axis_rx_tready, causing the Receive AXI4-Stream interface to stall receiving any further TLPs until the internal Memory Write controller completes the write to the block RAM. Deasserting m_axis_rx_tready in this way is not required for all designs using the core—the PIO design uses this method to simplify the control logic of the RX state machine.

Memory and I/O Read TLP Processing

When the Endpoint for PCIe receives a Memory or I/O Read TLP, the TLP destination address and transaction type are compared with the values programmed in the core BARs. If the TLP passes this comparison check, the core passes the TLP to the Receive AXI4-Stream interface of the PIO design.

Along with the start of packet, end of packet, and ready handshaking signals, the Receive AXI4-Stream interface also asserts the appropriate bar_hit[6:0] signal to indicate to the PIO design the specific destination BAR that matched the incoming TLP. On reception, the PIO design's state machine processes the incoming Read TLP and extracts the relevant TLP

information and passes it along to the PIO design's internal block RAM read request controller.

Based on the specific bar_hit[6:0] signal asserted, the RX state machine indicates to the internal read request controller the appropriate 2 KB block RAM to use before asserting the read enable request. For example, if a Memory Read 32 Request TLP is received by the core targeting the default MEM32 BAR2, the core passes the TLP to the PIO design and asserts bar_hit[2]. The RX State machine extracts the lower address bits from the Memory 32 Read TLP and instructs the internal Memory Read Request controller to start a read operation.

In this example, the assertion of bar_hit[2] instructs the PIO memory read controller to access the Mem32 space, which by default represents 2 KB of memory space. A notable difference in handling of memory write and read TLPs is the requirement of the receiving device to return a Completion with Data TLP in the case of memory or I/O read request.

While the read is being processed, the PIO design RX state machine deasserts m_axis_rx_tready, causing the Receive AXI4-Stream interface to stall receiving any further TLPs until the internal Memory Read controller completes the read access from the block RAM and generates the completion. Deasserting m_axis_rx_tready in this way is not required for all designs using the core. The PIO design uses this method to simplify the control logic of the RX state machine.

PIO File Structure

Table B-2 defines the PIO design file structure. Based on the specific core targeted, not all files delivered by CORE Generator software are necessary, and some files might not be delivered. The major difference is that some of the Endpoint for PCIe solutions use a 32-bit user datapath, others use a 64-bit datapath, and the PIO design works with both. The width of the datapath depends on the specific core being targeted.

Table B-2: PIO Design File Structure

File	Description
PIO.[v vhd]	Top-level design wrapper
PIO_EP.[v vhd]	PIO application module
PIO_TO_CTRL.[v vhd]	PIO turn-off controller module
PIO_32.v	32b interface macro define
PIO_64.v	64b macro define
PIO_32_RX_ENGINE.[v vhd]	32b Receive engine
PIO_32_TX_ENGINE.[v vhd]	32b Transmit engine
PIO_64_RX_ENGINE.[v vhd]	64b Receive engine
PIO_64_TX_ENGINE.[v vhd]	64b Transmit engine
PIO_EP_MEM_ACCESS.[v vhd]	Endpoint memory access module
PIO_EP_MEM.[v vhd]	Endpoint memory

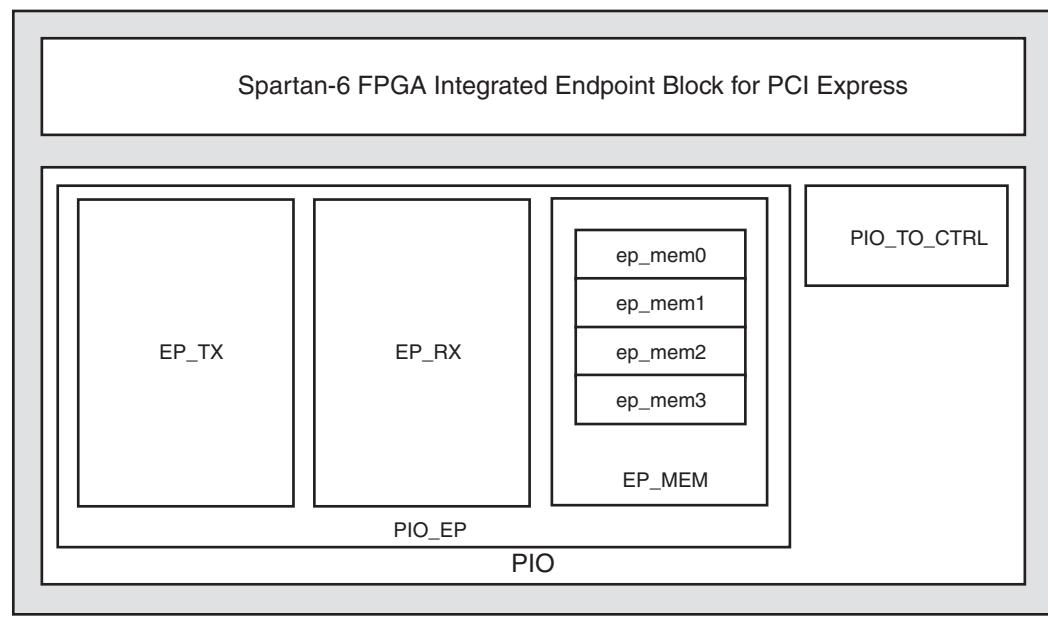
Two configurations of the PIO Design are provided: PIO_32 and PIO_64, with 32 and 64-bit AXI4-Stream interfaces, respectively. The PIO configuration generated depends on the selected endpoint type (that is, Spartan-6 FPGA integrated Endpoint block, PIPE, PCI

Express, and Block Plus) as well as the number of PCI Express lanes selected by the user. **Table B-3** identifies the PIO configuration generated based on the user's selection.

Table B-3: PIO Configuration

Core	x1	x2	x4	x8
Endpoint for PIPE	PIO_32	NA	NA	NA
Endpoint for PCI Express (Soft IP)	PIO_32	NA	PIO_64	PIO_64
Endpoint for PCI Express Block Plus	PIO_64	NA	PIO_64	PIO_64
Spartan-6 FPGA Integrated Endpoint Block	PIO_32	NA	NA	NA
Virtex®-6 FPGA Integrated Block	PIO_64	PIO_64	PIO_64	PIO_64

Figure B-2 shows the various components of the PIO design, which is separated into four main parts: the TX Engine, RX Engine, Memory Access Controller, and Power Management Turn-Off Controller.



UG672_cA_02_032111

Figure B-2: PIO Design Components

PIO Application

[Figure B-3](#) and [Figure B-4](#) depict 64-bit and 32-bit PIO application top-level connectivity, respectively. The datapath width, either 32 bits or 64 bits, depends on which Endpoint for PCIe core is used. The PIO_EP module contains the PIO FPGA block RAM memory modules and the transmit and receive engines. The PIO_TO_CTRL module is the Endpoint Turn-Off controller unit, which responds to power turn-off message from the host CPU with an acknowledgment.

The PIO_EP module connects to the Endpoint AXI4-Stream and Configuration (cfg) interfaces.

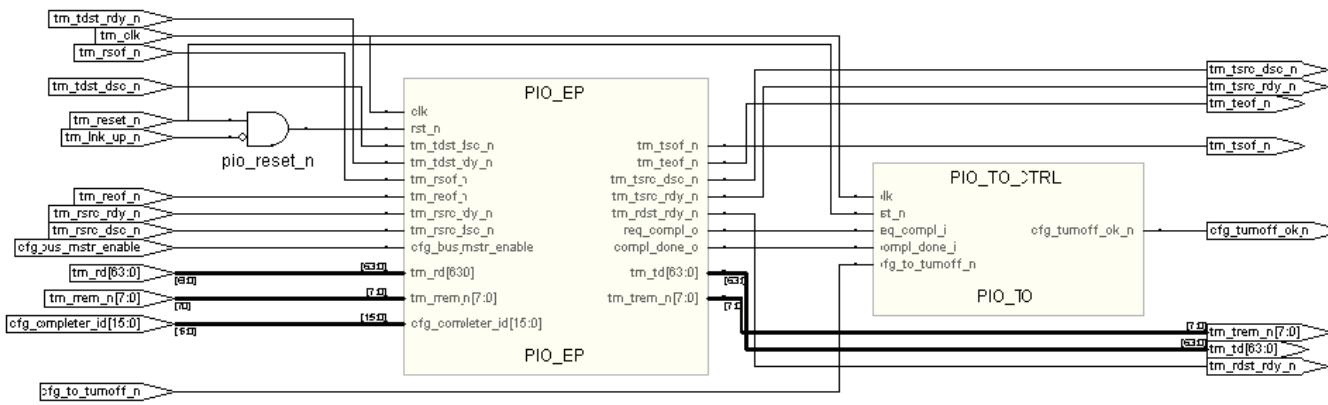


Figure B-3: PIO 64-Bit Application

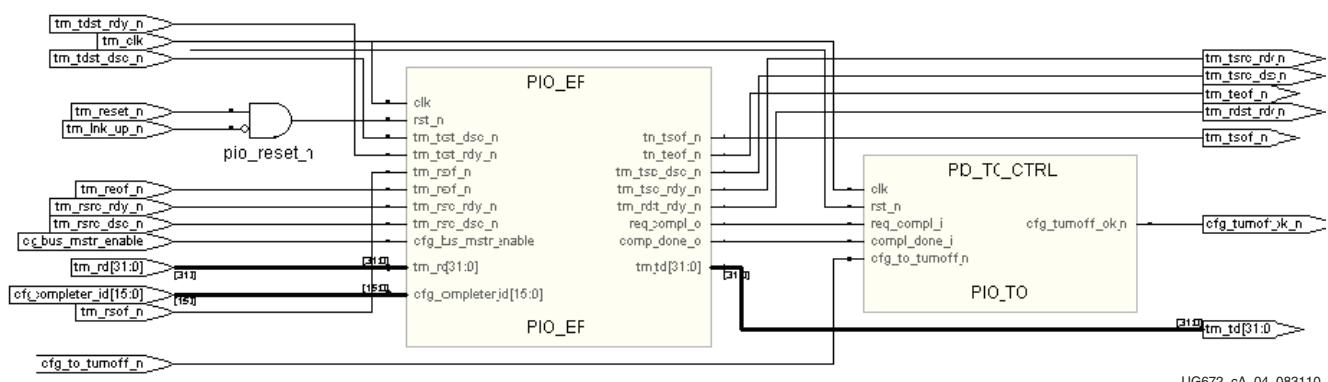
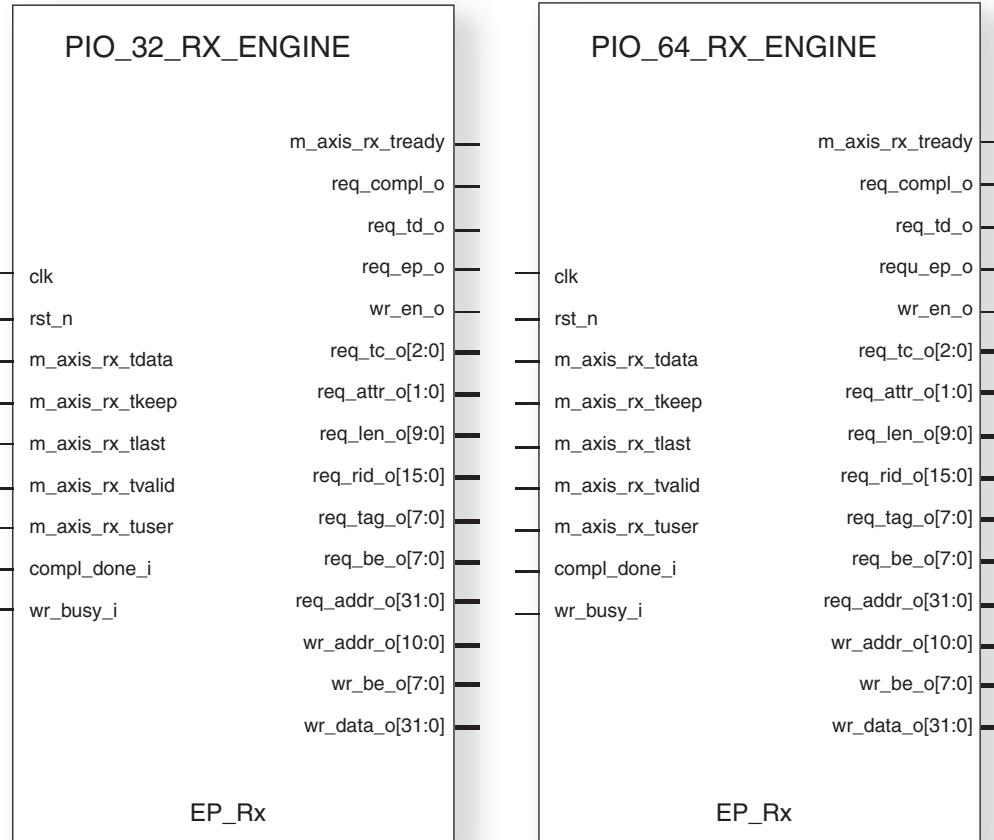


Figure B-4: PIO 32-Bit Application

Receive Path

[Figure B-5](#) illustrates the PIO_32_RX_ENGINE and PIO_64_RX_ENGINE modules. The datapath of the module must match the datapath of the core being used. These modules connect with Endpoint for PCIe Receive interface.



UG671_aA_05_032111

Figure B-5: RX Engines

The PIO_32_RX_ENGINE and PIO_64_RX_ENGINE modules receive and parse incoming read and write TLPs.

The RX engine parses one DWORD 32 and 64-bit addressable memory and I/O read requests. The RX state machine extracts needed information from the TLP and passes it to the memory controller, as defined in [Table B-4](#).

Table B-4: RX Engine: Read Outputs

Port	Description
req_compl_o	Completion request (active High)
req_td_o	Request TLP Digest bit
req_ep_o	Request Error Poisoning bit
req_tc_o[2:0]	Request Traffic Class
req_attr_o[1:0]	Request Attributes

Table B-4: RX Engine: Read Outputs (Cont'd)

Port	Description
req_len_o[9:0]	Request Length
req_rid_o[15:0]	Request Requester Identifier
req_tag_o[7:0]	Request Tag
req_be_o[7:0]	Request Byte Enable
req_addr_o[10:0]	Request Address

The RX Engine parses one DWORD 32- and 64-bit addressable memory and I/O write requests. The RX state machine extracts needed information from the TLP and passes it to the memory controller, as defined in [Table B-5](#).

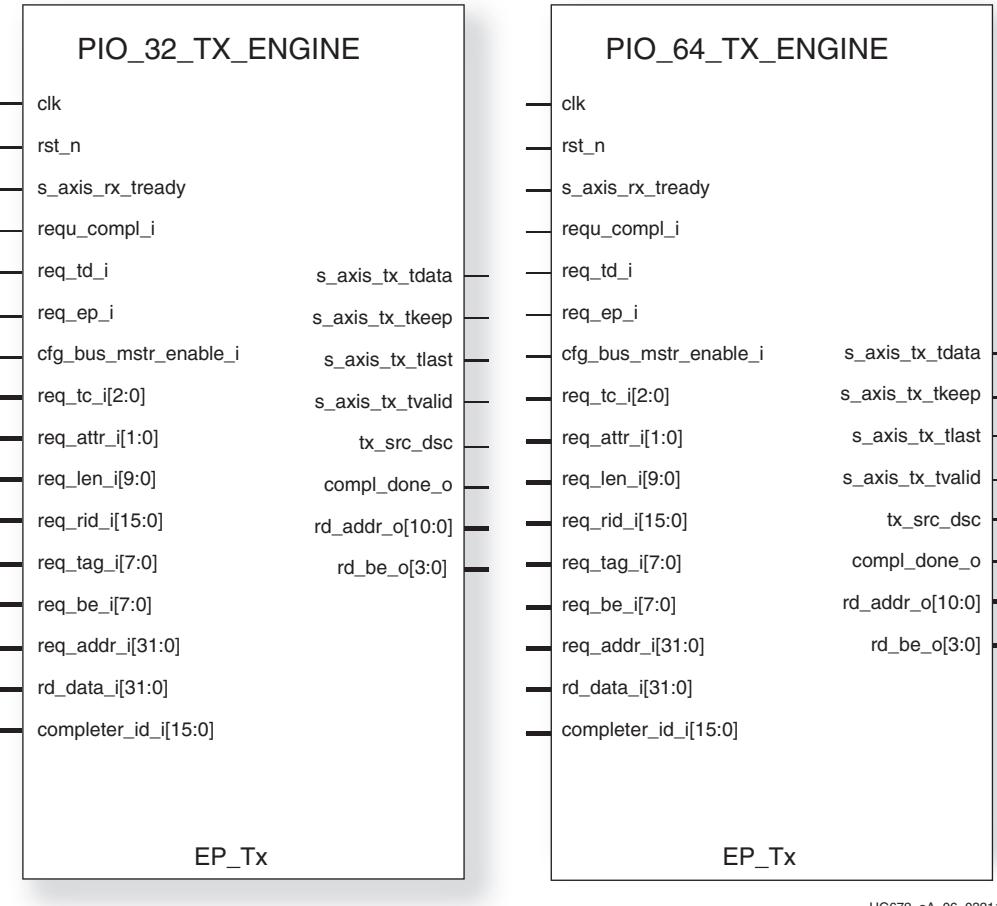
Table B-5: RX Engine: Write Outputs

Port	Description
wr_en_o	Write received
wr_addr_o[10:0]	Write address
wr_be_o[7:0]	Write byte enable
wr_data_o[31:0]	Write data

The read datapath stops accepting new transactions from the core while the application is processing the current TLP. This is accomplished by m_axis_rx_tready deassertion. For an ongoing Memory or I/O Read transaction, the module waits for compl_done_i input to be asserted before it accepts the next TLP, while an ongoing Memory or I/O Write transaction is deemed complete after wr_busy_i is deasserted.

Transmit Path

Figure B-6 shows the PIO_32_TX_ENGINE and PIO_64_TX_ENGINE modules. The datapath of the module must match the datapath of the core being used. These modules connect with the core Transaction interface.



UG672_aA_06_032111

Figure B-6: TX Engines

The PIO_32_TX_ENGINE and PIO_64_TX_ENGINE modules generate completions for received memory and I/O read TLPs. The PIO design does not generate outbound read or write requests. However, users can add this functionality to further customize the design.

The PIO_32_TX_ENGINE and PIO_64_TX_ENGINE modules generate completions in response to one DWORD 32 and 64-bit addressable memory and I/O read requests. Information necessary to generate the completion is passed to the TX Engine, as defined in Table B-6.

Table B-6: TX Engine Inputs

Port	Description
req_compl_i	Completion request (active High)
req_td_i	Request TLP Digest bit
req_ep_i	Request Error Poisoning bit
req_tc_i[2:0]	Request Traffic Class

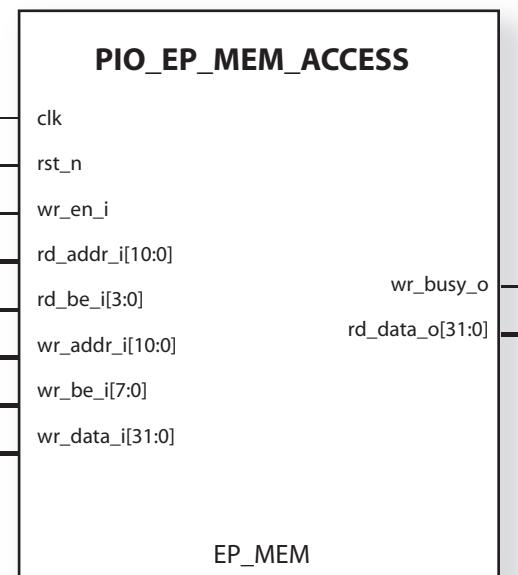
Table B-6: TX Engine Inputs (Cont'd)

Port	Description
req_attr_i[1:0]	Request Attributes
req_len_i[9:0]	Request Length
req_rid_i[15:0]	Request Requester Identifier
req_tag_i[7:0]	Request Tag
req_be_i[7:0]	Request Byte Enable
req_addr_i[10:0]	Request Address

After the completion is sent, the TX engine asserts the compl_done_i output indicating to the RX engine that it can assert m_axis_rx_tready and continue receiving TLPs.

Endpoint Memory

Figure B-7 displays the PIO_EP_MEM_ACCESS module. This module contains the Endpoint memory space.



UG672_cA_07_083110

Figure B-7: EP Memory Access

The PIO_EP_MEM_ACCESS module processes data written to the memory from incoming Memory and I/O Write TLPs and provides data read from the memory in response to Memory and I/O Read TLPs.

The EP_MEM module processes 1 DWORD 32- and 64-bit addressable Memory and I/O Write requests based on the information received from the RX Engine, as defined in Table B-7. While the memory controller is processing the write, it asserts the wr_busy_o output indicating it is busy.

Table B-7: EP Memory: Write Inputs

Port	Description
wr_en_i	Write received
wr_addr_i[10:0]	Write address
wr_be_i[7:0]	Write byte enable
wr_data_i[31:0]	Write data

Both 32 and 64-bit Memory and I/O Read requests of one DWORD are processed based on the inputs defined in [Table B-8](#). After the read request is processed, the data is returned on rd_data_o[31:0].

Table B-8: EP Memory: Read Inputs

Port	Description
req_be_i[7:0]	Request Byte Enable
req_addr_i[31:0]	Request Address

PIO Operation

PIO Read Transaction

Figure B-8 depicts a Back-To-Back Memory Read request to the PIO design. The receive engine deasserts m_axis_rx_tready as soon as the first TLP is completely received. The next Read transaction is accepted only after compl_done_o is asserted by the transmit engine, indicating that Completion for the first request was successfully transmitted.

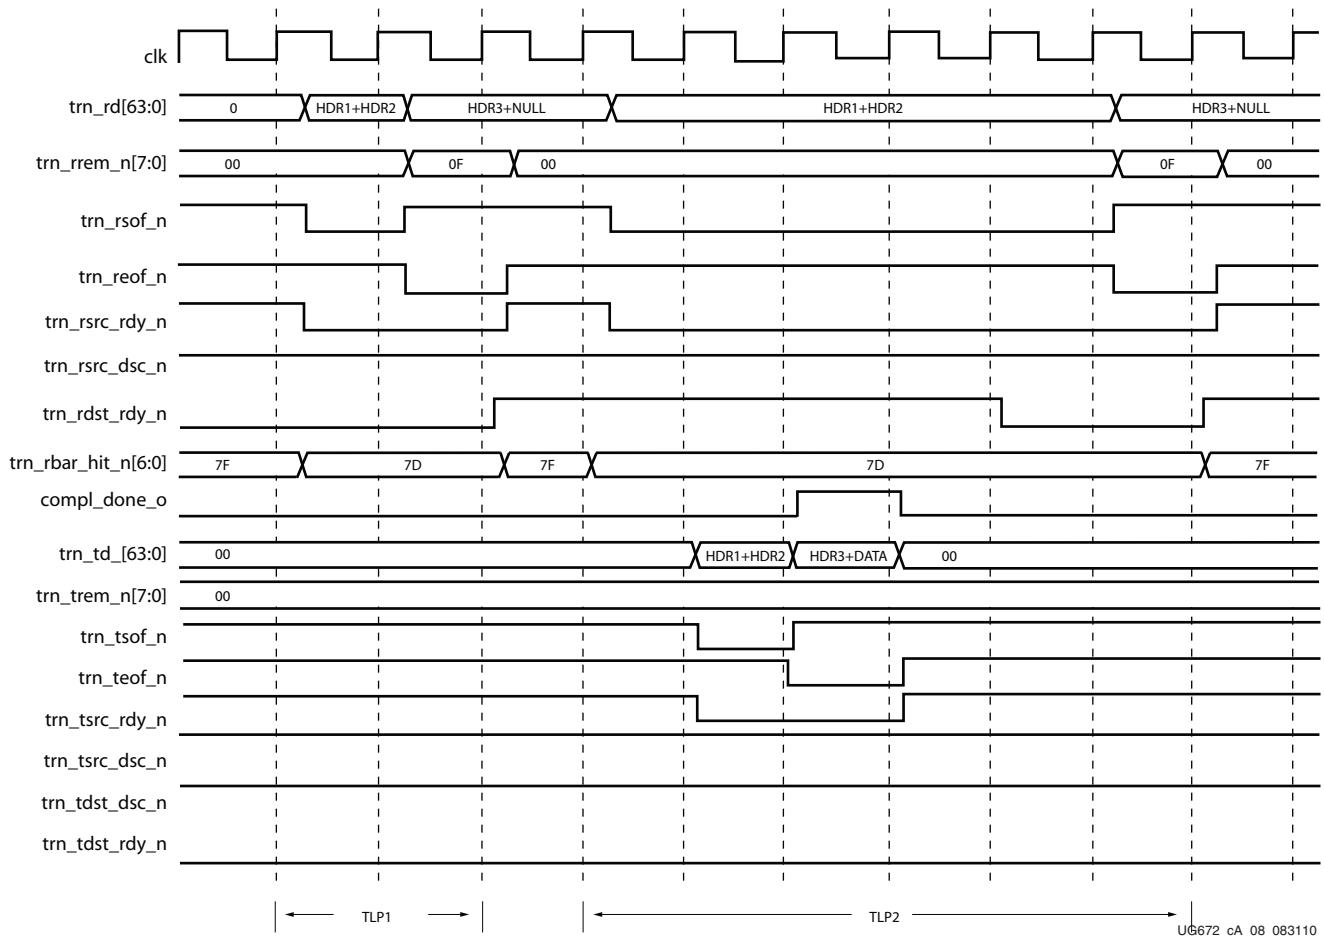


Figure B-8: Back-to-Back Read Transactions

PIO Write Transaction

Figure B-9 depicts a back-to-back Memory Write to the PIO design. The next Write transaction is accepted only after wr_busy_o is deasserted by the memory access unit, indicating that data associated with the first request was successfully written to the memory aperture.

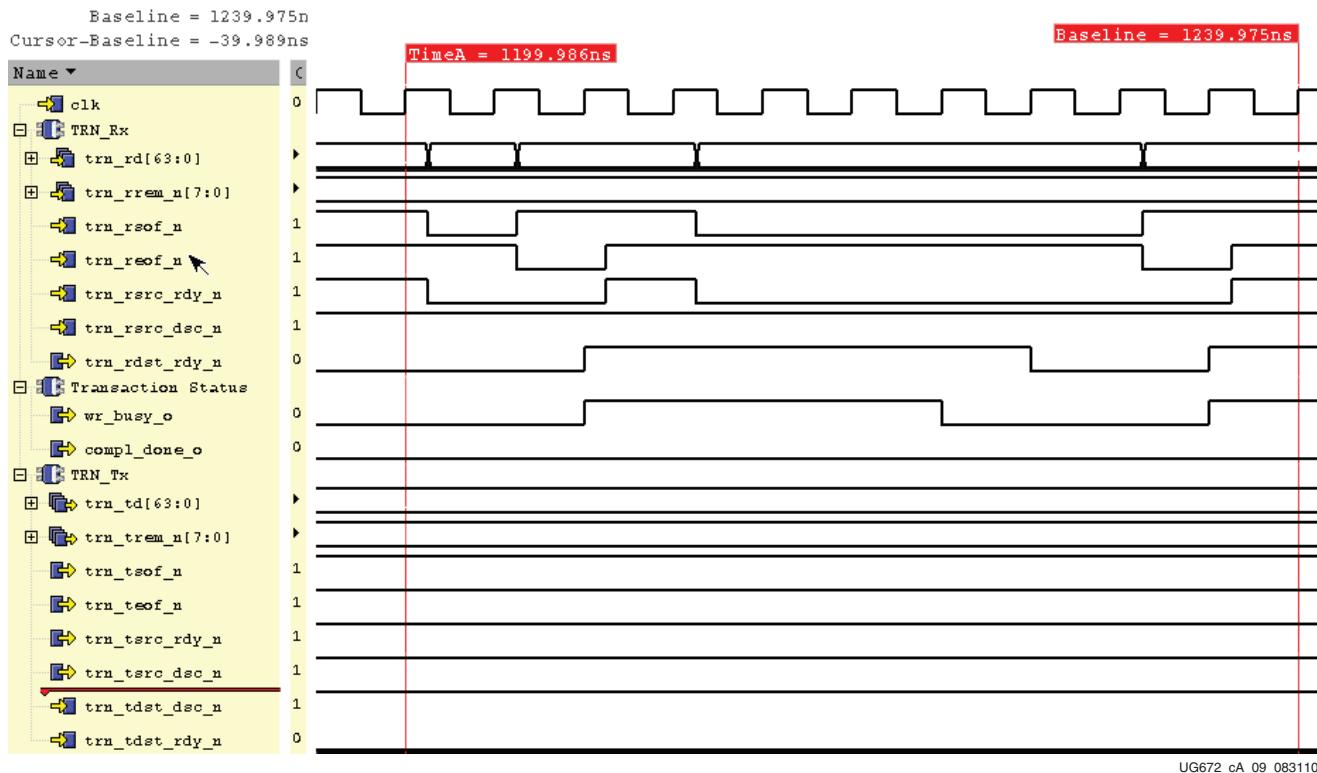


Figure B-9: Back-to-Back Write Transactions

Device Utilization

Table B-9 shows the PIO design FPGA resource utilization.

Table B-9: PIO Design FPGA Resources

Resources	Utilization
LUTs	300
Flip-Flops	500
Block RAMs	4

Summary

The PIO design demonstrates the Endpoint for PCIe and its interface capabilities. In addition, it enables rapid bring-up and basic validation of end user Endpoint add-in card FPGA hardware on PCI Express platforms. Users can leverage standard operating system utilities that enable generation of read and write transactions to the target space in the reference design.

Root Port Model Test Bench for Endpoint

The PCI Express Root Port Model is a robust test bench environment that provides a test program interface that can be used with the provided PIO design or with the user's design. The purpose of the Root Port Model is to provide a source mechanism for generating downstream PCI Express TLP traffic to stimulate the customer design, and a destination mechanism for receiving upstream PCI Express TLP traffic from the customer design in a simulation environment.

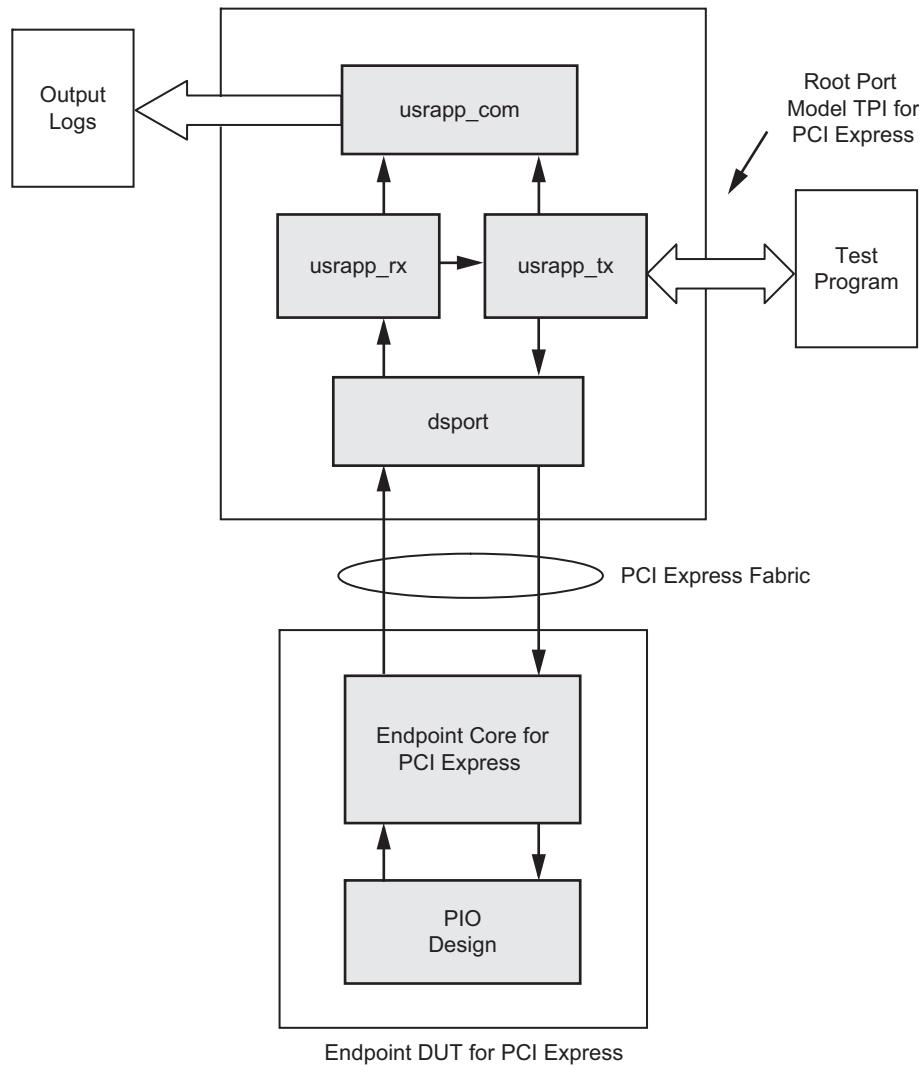
Note: The Root Port Model is shared by the Spartan-6 FPGA Integrated Block and Endpoint Block Plus for PCI Express solutions. This appendix represents these solutions.

Source code for the Root Port Model is included to provide the model for a starting point for the user test bench. All the significant work for initializing the core's configuration space, creating TLP transactions, generating TLP logs, and providing an interface for creating and verifying tests are complete, allowing the user to dedicate efforts to verifying the correct functionality of the design rather than spending time developing an Endpoint core test bench infrastructure.

The Root Port Model consists of:

- Test Programming Interface (TPI), which allows the user to stimulate the Endpoint device for the PCI Express
- Example tests that illustrate how to use the test program TPI
- Verilog or VHDL source code for all Root Port Model components, which allow the user to customize the test bench

Figure B-10 illustrates the Root Port Model coupled with the PIO design.



UG671_aA_05_081710

Figure B-10: Root Port Model and Top-Level Endpoint

Architecture

The Root Port Model consists of these blocks, illustrated in Figure B-10:

- dsport (Root Port)
- usrapp_tx
- usrapp_rx
- usrapp_com (Verilog only)

The usrapp_tx and usrapp_rx blocks interface with the dsport block for transmission and reception of TLPs to/from the Endpoint Design Under Test (DUT). The Endpoint DUT consists of the Endpoint for PCIe and the PIO design (displayed) or customer design.

The usrapp_tx block sends TLPs to the dsport block for transmission across the PCI Express Link to the Endpoint DUT. In turn, the Endpoint DUT device transmits TLPs across the PCI Express Link to the dsport block, which are subsequently passed to the

usrapp_rx block. The dsport and core are responsible for the data link layer and physical link layer processing when communicating across the PCI Express fabric. Both usrapp_tx and usrapp_rx utilize the usrapp_com block for shared functions, for example, TLP processing and log file outputting. Transaction sequences or test programs are initiated by the usrapp_tx block to stimulate the endpoint device's fabric interface. TLP responses from the endpoint device are received by the usrapp_rx block. Communication between the usrapp_tx and usrapp_rx blocks allow the usrapp_tx block to verify correct behavior and act accordingly when the usrapp_rx block has received TLPs from the endpoint device.

Simulating the Design

Four simulation script files are provided with the model to facilitate simulation with Synopsys VCS and VCS MX, Cadence IES, Mentor Graphics ModelSim, and Xilinx ISE Simulator (ISim):

- simulate_vcs.sh (Verilog Only)
- simulate_ncsim.sh (Verilog Only)
- simulate_mti.do
- simulate_isim.bat/simulate_isim.sh

The example simulation script files are located in the following directory:

```
<project_dir>/<component_name>/simulation/functional
```

Instructions for simulating the PIO design using the Root Port Model are provided in [Chapter 3, Getting Started Example Design](#).

Note: For Cadence IES users, the work construct must be manually inserted into the cds.lib file:
DEFINE WORK WORK.

Scaled Simulation Timeouts

The simulation model of the Virtex-6 FPGA Integrated Block for PCI Express uses scaled down times during link training to allow for the link to train in a reasonable amount of time during simulation. According to the *PCI Express Specification, rev. 2.0*, there are various timeouts associated with the link training and status state machine (LTSSM) states. The Virtex-6 FPGA Integrated Block for PCI Express scales these timeouts by a factor of 256 in simulation, except in the Recovery Speed_1 LTSSM state, where the timeouts are not scaled.

Test Selection

Table B-10 describes the tests provided with the Root Port Model, followed by specific sections for VHDL and Verilog test selection.

Table B-10: Root Port Model Provided Tests

Test Name	Test in VHDL/Verilog	Description
sample_smoke_test0	Verilog and VHDL	Issues a PCI Type 0 Configuration Read TLP and waits for the completion TLP; then compares the value returned with the expected Device/Vendor ID value.
sample_smoke_test1	Verilog	Performs the same operation as sample_smoke_test0 but makes use of expectation tasks. This test uses two separate test program threads: one thread issues the PCI Type 0 Configuration Read TLP and the second thread issues the Completion with Data TLP expectation task. This test illustrates the form for a parallel test that uses expectation tasks. This test form allows for confirming reception of any TLPs from the customer's design. Additionally, this method can be used to confirm reception of TLPs when ordering is unimportant.

VHDL Test Selection

Test selection is implemented in the VHDL Downstream Port Model by means of overriding the `test_selector` generic within the `tests` entity. The `test_selector` generic is a string with a one-to-one correspondence to each test within the `tests` entity.

The user can modify the generic mapping of the instantiation of the `tests` entity within the `pci_exp_usrapp_tx` entity. Currently, there is one test defined inside the `tests` entity, `sample_smoke_test0`. Additional customer-defined tests should be added inside `tests.vhd`. Currently, specific tests cannot be selected from the VHDL simulation scripts.

Verilog Test Selection

The Verilog test model used for the Root Port Model lets the user specify the name of the test to be run as a command line parameter to the simulator. For example, the `simulate_ncsim.sh` script file, used to start the Cadence IES simulator, explicitly specifies the test `sample_smoke_test0` to be run using this command line syntax:

```
ncsim work.boardx01 +TESTNAME=sample_smoke_test0
```

To change the test to be run, change the value provided to `TESTNAME` defined in the test files `sample_tests1.v` and `pio_tests.v`. The same mechanism is used for VCS and ModelSim. ISim uses the `-testplusarg` options to specify `TESTNAME`, for example:
`demo_tb.exe -gui -view wave.wcfg -wdb wave_isim -tclbatch
isim_cmd.tcl -testplusarg TESTNAME=sample_smoke_test0`.

VHDL and Verilog Root Port Model Differences

The following subsections identify differences between the VHDL and Verilog Root Port Model.

Verilog Expectation Tasks

The most significant difference between the Verilog and the VHDL test bench is that the Verilog test bench has Expectation Tasks. Expectation tasks are API calls used in

conjunction with a bus mastering customer design. The test program issues a series of expectation task calls, that is, the task calls expect a memory write TLP and a memory read TLP. If the customer design does not respond with the expected TLPs, the test program fails. This functionality was implemented using the fork-join construct in Verilog, which is not available in VHDL and subsequently not implemented.

Verilog Command Line versus VHDL tests.vhd Module

The Verilog test bench allows test programs to be specified at the command line, while the VHDL test bench specifies test programs within the `tests.vhd` module.

Generating Wave Files

- The Verilog test bench uses recordvars and dumpfile commands within the code to generate wave files.
- The VHDL test bench leaves the generating wave file functionality up to the simulator.

Speed Differences

The VHDL test bench is slower than the Verilog test bench, especially when testing the x8 core. For initial design simulation and speed enhancement, the user might want to use the x1 core, identify basic functionality issues, and then move to x4 or x8 simulation when testing design performance.

Waveform Dumping

[Table B-11](#) describes the available simulator waveform dump file formats, each of which is provided in the simulator's native file format. The same mechanism is used for VCS and ModelSim.

Table B-11: Simulator Dump File Format

Simulator	Dump File Format
Synopsys VCS	.vpd
ModelSim	.vcd
Cadence IES	.trn
ISim	.wdb

VHDL Flow

Waveform dumping in the VHDL flow does not use the `+dump_all` mechanism described in the [Verilog Flow](#) section. Because the VHDL language itself does not provide a common interface for dumping waveforms, each VHDL simulator has its own interface for supporting waveform dumping. For both the supported ModelSim and IUS flows, dumping is supported by invoking the VHDL simulator command line with a command line option that specifies the respective waveform command file, `wave.do` (ModelSim), `wave.sv` (IES), and `wave.wcfg` (ISim). This command line can be found in the respective simulation script files `simulate_mti.do`, `simulate_ncsim.sh`, and `simulate_isim.bat[sh]`.

ModelSim

This command line initiates waveform dumping for the ModelSim flow using the VHDL test bench:

```
>vsim +notimingchecks -do wave.do -L unisim -L work work.board
```

IES

This command line initiates waveform dumping for the IES flow using the VHDL test bench:

```
>ncsim -gui work.board -input @"simvision -input wave.sv"
```

Verilog Flow

The Root Port Model provides a mechanism for outputting the simulation waveform to file by specifying the `+dump_all` command line parameter to the simulator.

For example, the script file `simulate_ncsim.sh` (used to start the Cadence IES simulator) can indicate to the Root Port Model that the waveform should be saved to a file using this command line:

```
ncsim work.boardx01 +TESTNAME=sample_smoke_test0 +dump_all
```

Output Logging

When a test fails on the example or customer design, the test programmer debugs the offending test case. Typically, the test programmer inspects the wave file for the simulation and cross-reference this to the messages displayed on the standard output. Because this approach can be very time consuming, the Root Port Model offers an output logging mechanism to assist the tester with debugging failing test cases to speed the process.

The Root Port Model creates three output files (`tx.dat`, `rx.dat`, and `error.dat`) during each simulation run. Log files `rx.dat` and `tx.dat` each contain a detailed record of every TLP that was received and transmitted, respectively, by the Root Port Model. With an understanding of the expected TLP transmission during a specific test case, the test programmer can more easily isolate the failure.

The log file `error.dat` is used in conjunction with the expectation tasks. Test programs that utilize the expectation tasks will generate a general error message to standard output. Detailed information about the specific comparison failures that have occurred due to the expectation error is located within `error.dat`.

Parallel Test Programs

There are two classes of tests supported by the Root Port Model:

- Sequential tests. Tests that exist within one process and behave similarly to sequential programs. The test depicted in [Test Program: pio_writeReadBack_test0, page 148](#) is an example of a sequential test. Sequential tests are very useful when verifying behavior that have events with a known order.
- Parallel tests. Tests involving more than one process thread. The test `sample_smoke_test1` is an example of a parallel test with two process threads. Parallel tests are very useful when verifying that a specific set of events have occurred, however the order of these events are not known.

A typical parallel test uses the form of one command thread and one or more expectation threads. These threads work together to verify a device's functionality. The role of the command thread is to create the necessary TLP transactions that cause the device to receive and generate TLPs. The role of the expectation threads is to verify the reception of an expected TLP. The Root Port Model TPI has a complete set of expectation tasks to be used in conjunction with parallel tests.

Because the example design is a target-only device, only Completion TLPs can be expected by parallel test programs while using the PIO design. However, the full library of expectation tasks can be used for expecting any TLP type when used in conjunction with the customer's design (which can include bus-mastering functionality). Currently, the VHDL version of the Root Port Model Test Bench does not support Parallel tests.

Test Description

The Root Port Model provides a Test Program Interface (TPI). The TPI provides the means to create tests by simply invoking a series of Verilog tasks. All Root Port Model tests should follow the same six steps:

1. Perform conditional comparison of a unique test name
2. Set up master timeout in case simulation hangs
3. Wait for Reset and link-up
4. Initialize the configuration space of the endpoint
5. Transmit and receive TLPs between the Root Port Model and the Endpoint DUT
6. Verify that the test succeeded

Test Program: pio_writeReadBack_test0

```

1.    else if(testname == "pio_writeReadBack_test1"
2. begin
3. // This test performs a 32 bit write to a 32 bit Memory space and performs a read back
4. TSK_SIMULATION_TIMEOUT(10050);
5. TSK_SYSTEM_INITIALIZATION;
6. TSK_BAR_INIT;
7. for (ii = 0; ii <= 6; ii = ii + 1) begin
8.     if (BAR_INIT_P_BAR_ENABLED[ii] > 2'b00) // bar is enabled
9.         case(BAR_INIT_P_BAR_ENABLED[ii])
10.             2'b01 : // IO SPACE
11.                 begin
12.                     $display("[%t] : NOTHING: to IO 32 Space BAR %x", $realtime, ii);
13.                 end
14.             2'b10 : // MEM 32 SPACE
15.                 begin
16.                     $display("[%t] : Transmitting TLPs to Memory 32 Space BAR %x",
17.                             $realtime, ii);
18. //-----
19. // Event : Memory Write 32 bit TLP
20. //-----
21.             DATA_STORE[0] = 8'h04;
22.             DATA_STORE[1] = 8'h03;
23.             DATA_STORE[2] = 8'h02;
24.             DATA_STORE[3] = 8'h01;
25.             P_READ_DATA = 32'hffff_ffff; // make sure P_READ_DATA has known initial value
26.             TSK_TX_MEMORY_WRITE_32(DEFAULT_TAG, DEFAULT_TC, 10'd1, BAR_INIT_P_BAR[ii][31:0], 4'hF,
27.             4'hF, 1'b0);
28.             TSK_TX_CLK_EAT(10);
29.             DEFAULT_TAG = DEFAULT_TAG + 1;
30. //-----
31. // Event : Memory Read 32 bit TLP
32. //-----
33.             TSK_TX_MEMORY_READ_32(DEFAULT_TAG, DEFAULT_TC, 10'd1, BAR_INIT_P_BAR[ii][31:0], 4'hF,
34.             4'hF);
35.             TSK_WAIT_FOR_READ_DATA;
36.             if (P_READ_DATA != {DATA_STORE[3], DATA_STORE[2], DATA_STORE[1], DATA_STORE[0] })
37.                 begin
38.                     $display("[%t] : Test FAILED --- Data Error Mismatch, Write Data %x != Read Data %x",
39.                         $realtime, {DATA_STORE[3], DATA_STORE[2], DATA_STORE[1], DATA_STORE[0]}, P_READ_DATA);
40.                 end
41.             else
42.                 begin
43.                     $display("[%t] : Test PASSED --- Write Data: %x successfully received", $realtime,
44.                         P_READ_DATA);
45.                 end

```

Expanding the Root Port Model

The Root Port Model was created to work with the PIO design, and for this reason is tailored to make specific checks and warnings based on the limitations of the PIO design. These checks and warnings are enabled by default when the Root Port Model is generated by the CORE Generator software. However, these limitations can easily be disabled so that they do not affect the customer's design.

Because the PIO design was created to support at most one I/O BAR, one Mem64 BAR, and two Mem32 BARs (one of which must be the EROM space), the Root Port Model by default makes a check during device configuration that verifies that the core has been configured to meet this requirement. A violation of this check causes a warning message to be displayed as well as for the offending BAR to be gracefully disabled in the test bench. This check can be disabled by setting the `pio_check_design` variable to zero in the `pci_exp_usrapp_tx.v` file.

Root Port Model TPI Task List

The Root Port Model TPI tasks include the following, which are further defined in [Tables B-12 through B-16](#).

- [Test Setup Tasks](#)
- [TLP Tasks](#)
- [BAR Initialization Tasks](#)
- [Example PIO Design Tasks](#)
- [Expectation Tasks](#)

Table B-12: Test Setup Tasks

Name	Input(s)		Description
TSK_SYSTEM_INITIALIZATION	None		Waits for transaction interface reset and link-up between the Root Port Model and the Endpoint DUT. This task must be invoked prior to the Endpoint core initialization.
TSK_USR_DATA_SETUP_SEQ	None		Initializes global 4096 byte DATA_STORE array entries to sequential values from zero to 4095.
TSK_TX_CLK_EAT	clock count	31:30	Waits clock_count transaction interface clocks.
TSK_SIMULATION_TIMEOUT	timeout	31:0	Sets master simulation timeout value in units of transaction interface clocks. This task should be used to ensure that all DUT tests complete.

Table B-13: TLP Tasks

Name	Input(s)			Description
TSK_TX_TYPE0_CONFIGURATION_READ	tag_ reg_addr_ first_dw_be_	7:0 11:0 3:0		Waits for transaction interface reset and link-up between the Root Port Model and the Endpoint DUT. This task must be invoked prior to Endpoint core initialization.
TSK_TX_TYPE1_CONFIGURATION_READ	tag_ reg_addr_ first_dw_be_	7:0 11:0 3:0		Sends a Type 1 PCI Express Config Read TLP from Root Port Model to reg_addr_ of Endpoint DUT with tag_ and first_dw_be_ inputs. CplID returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.

Table B-13: TLP Tasks (*Cont'd*)

Name	Input(s)	Description				
TSK_TX_TYPE0_CONFIGURATION_WRITE	tag_ reg_addr_ reg_data_ first_dw_be_	7:0 11:0 31:0 3:0	Sends a Type 0 PCI Express Config Write TLP from Root Port Model to reg_addr_ of Endpoint DUT with tag_ and first_dw_be_ inputs. Cpl returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.			
TSK_TX_TYPE1_CONFIGURATION_WRITE	tag_ reg_addr_ reg_data_ first_dw_be_	7:0 11:0 31:0 3:0	Sends a Type 1 PCI Express Config Write TLP from Root Port Model to reg_addr_ of Endpoint DUT with tag_ and first_dw_be_ inputs. Cpl returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.			
TSK_TX_MEMORY_READ_32	tag_ tc_ len_ addr_ last_dw_be_ first_dw_be_	7:0 2:0 9:0 31:0 3:0 3:0	Sends a PCI Express Memory Read TLP from Root Port to 32-bit memory address addr_ of Endpoint DUT. CplID returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.			
TSK_TX_MEMORY_READ_64	tag_ tc_ len_ addr_ last_dw_be_ first_dw_be_	7:0 2:0 9:0 63:0 3:0 3:0	Sends a PCI Express Memory Read TLP from Root Port Model to 64-bit memory address addr_ of Endpoint DUT. CplID returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.			
TSK_TX_MEMORY_WRITE_32	tag_ tc_ len_ addr_ last_dw_be_ first_dw_be_ ep_	7:0 2:0 9:0 31:0 3:0 3:0 -	Sends a PCI Express Memory Write TLP from Root Port Model to 32-bit memory address addr_ of Endpoint DUT. CplID returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID. The global DATA_STORE byte array is used to pass write data to task.			
TSK_TX_MEMORY_WRITE_64	tag_ tc_ len_ addr_ last_dw_be_ first_dw_be_ ep_	7:0 2:0 9:0 63:0 3:0 3:0 -	Sends a PCI Express Memory Write TLP from Root Port Model to 64-bit memory address addr_ of Endpoint DUT. CplID returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID. The global DATA_STORE byte array is used to pass write data to task.			

Table B-13: TLP Tasks (Cont'd)

Name	Input(s)	Description	
TSK_TX_COMPLETION	tag_ tc_ len_ comp_status_	7:0 2:0 9:0 2:0	Sends a PCI Express Completion TLP from Root Port Model to the Endpoint DUT using global COMPLETE_ID_CFG as the completion ID.
TSK_TX_COMPLETION_DATA	tag_ tc_ len_ byte_count lower_addr comp_status ep_	7:0 2:0 9:0 11:0 6:0 2:0 —	Sends a PCI Express Completion with Data TLP from Root Port Model to the Endpoint DUT using global COMPLETE_ID_CFG as the completion ID. The global DATA_STORE byte array is used to pass completion data to task.
TSK_TX_MESSAGE	tag_ tc_ len_ data message_rtg message_code	7:0 2:0 9:0 63:0 2:0 7:0	Sends a PCI Express Message TLP from Root Port Model to Endpoint DUT. Completion returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.
TSK_TX_MESSAGE_DATA	tag_ tc_ len_ data message_rtg message_code	7:0 2:0 9:0 63:0 2:0 7:0	Sends a PCI Express Message with Data TLP from Root Port Model to Endpoint DUT. The global DATA_STORE byte array is used to pass message data to task. Completion returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.
TSK_TX_IO_READ	tag_ addr_ first_dw_be_	7:0 31:0 3:0	Sends a PCI Express I/O Read TLP from Root Port Model to I/O address addr_[31:2] of the Endpoint DUT. CplD returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.
TSK_TX_IO_WRITE	tag_ addr_ first_dw_be_ data	7:0 31:0 3:0 31:0	Sends a PCI Express I/O Write TLP from Root Port Model to I/O address addr_[31:2] of the Endpoint DUT. CplD returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.

Table B-13: TLP Tasks (*Cont'd*)

Name	Input(s)			Description
TSK_TX_BAR_READ	bar_index byte_offset tag_ tc_	2:0 31:0 7:0 2:0		Sends a PCI Express 1 DWORD Memory 32, Memory 64, or I/O Read TLP from the Root Port Model to the target address corresponding to offset byte_offset from BAR bar_index of the Endpoint DUT. This task sends the appropriate Read TLP based on how BAR bar_index has been configured during initialization. This task can only be called after TSK_BAR_INIT has successfully completed. CplD returned from the Endpoint DUT will use the contents of global COMPLETE_ID_CFG as the completion ID.
TSK_TX_BAR_WRITE	bar_index byte_offset tag_ tc_ data_	2:0 31:0 7:0 2:0 31:0		Sends a PCI Express 1 DWORD Memory 32, Memory 64, or I/O Write TLP from the Root Port to the target address corresponding to offset byte_offset from BAR bar_index of the Endpoint DUT. This task sends the appropriate Write TLP based on how BAR bar_index has been configured during initialization. This task can only be called after TSK_BAR_INIT has successfully completed.
TSK_WAIT_FOR_READ_DATA	None			Waits for the next completion with data TLP that was sent by the Endpoint DUT. On successful completion, the first DWORD of data from the CplD will be stored in the global P_READ_DATA. This task should be called immediately following any of the read tasks in the TPI that request Completion with Data TLPs to avoid any race conditions. By default this task will locally time out and terminate the simulation after 1000 transaction interface clocks. The global cpld_to_finish can be set to zero so that local time out returns execution to the calling test and does not result in simulation timeout. For this case test programs should check the global cpld_to, which when set to one indicates that this task has timed out and that the contents of P_READ_DATA are invalid.

Table B-14: BAR Initialization Tasks

Name	Input(s)	Description
TSK_BAR_INIT	None	<p>Performs a standard sequence of Base Address Register initialization tasks to the Endpoint device using the PCI Express fabric. Performs a scan of the Endpoint's PCI BAR range requirements, performs the necessary memory and I/O space mapping calculations, and finally programs the Endpoint so that it is ready to be accessed.</p> <p>On completion, the user test program can begin memory and I/O transactions to the device. This function displays to standard output a memory and I/O table that details how the Endpoint has been initialized. This task also initializes global variables within the Root Port Model that are available for test program usage. This task should only be called after TSK_SYSTEM_INITIALIZATION.</p>
TSK_BAR_SCAN	None	<p>Performs a sequence of PCI Type 0 Configuration Writes and Configuration Reads using the PCI Express fabric in order to determine the memory and I/O requirements for the Endpoint.</p> <p>The task stores this information in the global array BAR_INIT_P_BAR_RANGE[]. This task should only be called after TSK_SYSTEM_INITIALIZATION.</p>
TSK_BUILD_PCIE_MAP	None	<p>Performs memory and I/O mapping algorithm and allocates Memory 32, Memory 64, and I/O space based on the Endpoint requirements.</p> <p>This task has been customized to work in conjunction with the limitations of the PIO design and should only be called after completion of TSK_BAR_SCAN.</p>
TSK_DISPLAY_PCIE_MAP	None	Displays the memory mapping information of the Endpoint core's PCI Base Address Registers. For each BAR, the BAR value, the BAR range, and BAR type is given. This task should only be called after completion of TSK_BUILD_PCIE_MAP.

Table B-15: Example PIO Design Tasks

Name	Input(s)		Description
TSK_TX_READBACK_CONFIG	None		<p>Performs a sequence of PCI Type 0 Configuration Reads to the Endpoint device's Base Address Registers, PCI Command Register, and PCIe Device Control Register using the PCI Express fabric.</p> <p>This task should only be called after TSK_SYSTEM_INITIALIZATION.</p>
TSK_MEM_TEST_DATA_BUS	bar_index	2:0	<p>Tests whether the PIO design FPGA block RAM data bus interface is correctly connected by performing a 32-bit walking ones data test to the I/O or memory address pointed to by the input bar_index.</p> <p>For an exhaustive test, this task should be called four times, once for each block RAM used in the PIO design.</p>
TSK_MEM_TEST_ADDR_BUS	bar_index nBytes	2:0 31:0	<p>Tests whether the PIO design FPGA block RAM address bus interface is accurately connected by performing a walking ones address test starting at the I/O or memory address pointed to by the input bar_index.</p> <p>For an exhaustive test, this task should be called four times, once for each block RAM used in the PIO design.</p> <p>Additionally, the nBytes input should specify the entire size of the individual block RAM.</p>
TSK_MEM_TEST_DEVICE	bar_index nBytes	2:0 31:0	<p>Tests the integrity of each bit of the PIO design FPGA block RAM by performing an increment/decrement test on all bits starting at the block RAM pointed to by the input bar_index with the range specified by input nBytes.</p> <p>For an exhaustive test, this task should be called four times, once for each block RAM used in the PIO design.</p> <p>Additionally, the nBytes input should specify the entire size of the individual block RAM.</p>

Table B-16: Expectation Tasks

Name	Input(s)		Output	Description
TSK_EXPECT_CPLD	traffic_class td ep attr length completer_id completer_status bcm byte_count requester_id tag address_low	2:0 - - 1:0 9:0 15:0 2:0 - 11:0 15:0 7:0 6:0	Expect status	Waits for a Completion with Data TLP that matches traffic_class, td, ep, attr, length, and payload. Returns a 1 on successful completion; 0 otherwise.
TSK_EXPECT_CPL	traffic_class td ep attr completer_id completer_status bcm byte_count requester_id tag address_low	2:0 - - 1:0 15:0 2:0 - 11:0 15:0 7:0 6:0	Expect status	Waits for a Completion without Data TLP that matches traffic_class, td, ep, attr, and length. Returns a 1 on successful completion; 0 otherwise.
TSK_EXPECT_MEMRD	traffic_class td ep attr length requester_id tag last_dw_be first_dw_be address	2:0 - - 1:0 9:0 15:0 7:0 3:0 3:0 29:0	Expect status	Waits for a 32-bit Address Memory Read TLP with matching header fields. Returns a 1 on successful completion; 0 otherwise. This task can only be used in conjunction with Bus Master designs.

Table B-16: Expectation Tasks (Cont'd)

Name	Input(s)		Output	Description
TSK_EXPECT_MEMRD64	traffic_class td ep attr length requester_id tag last_dw_be first_dw_be address	2:0 - - 1:0 9:0 15:0 7:0 3:0 3:0 61:0	Expect status	Waits for a 64-bit Address Memory Read TLP with matching header fields. Returns a 1 on successful completion; 0 otherwise. This task can only be used in conjunction with Bus Master designs.
TSK_EXPECT_MEMWR	traffic_class td ep attr length requester_id tag last_dw_be first_dw_be address	2:0 - - 1:0 9:0 15:0 7:0 3:0 3:0 29:0	Expect status	Waits for a 32-bit Address Memory Write TLP with matching header fields. Returns a 1 on successful completion; 0 otherwise. This task can only be used in conjunction with Bus Master designs.
TSK_EXPECT_MEMWR64	traffic_class td ep attr length requester_id tag last_dw_be first_dw_be address	2:0 - - 1:0 9:0 15:0 7:0 3:0 3:0 61:0	Expect status	Waits for a 64-bit Address Memory Write TLP with matching header fields. Returns a 1 on successful completion; 0 otherwise. This task can only be used in conjunction with Bus Master designs.
TSK_EXPECT_IOWR	td ep requester_id tag first_dw_be address data	- - 15:0 7:0 3:0 31:0 31:0	Expect status	Waits for an I/O Write TLP with matching header fields. Returns a 1 on successful completion; 0 otherwise. This task can only be used in conjunction with Bus Master designs.

Migration Considerations

This appendix describes the changes in the interfaces and signals that are necessary when migrating from the LogiCORE™ IP Spartan®-3 FPGA Endpoint PIPE core to a LogiCORE IP Spartan-6 FPGA integrated Endpoint block core for revisions 1.x and earlier. Starting with revision 2.0, the Spartan-6 FPGA Integrated Endpoint Block has changed from the TRN interface to the AXI4-Stream interface, please see [Appendix J, TRN to AXI Interface Migration Considerations](#) for additional migration considerations.

Integrated PHY

The first and most notable change between the Spartan-3 FPGA Endpoint PIPE core and the Spartan-6 FPGA integrated Endpoint block core is that the external SerDes PHY device that previously resided outside of the device has been integrated into the Spartan-6 architecture. This means that all of the PXPIPE signals (33 ports) are replaced with the serial interface (4 ports).

System Clocking and Reset

For the Spartan-6 FPGA integrated Endpoint block, the sys_clk signal was added.

In the Spartan-3 FPGA design, the system clock was provided by the external PHY on the port, RXCLK. For more information about sys_clk and how to set it up, see [Clocking in Chapter 5](#).

Interface Changes

Streaming Signal Added

The trn_tstr_n signal was added to the integrated Endpoint block to allow packets to be streamed in the transmit direction. For more information on trn_tstr_n, see [Table 2-6, page 16](#).

TRN Transmit Destination Discontinue Removed

The trn_tdst_dsc_n signal has been replaced with trn_tx_terr_drop_n. The signal serves the same purpose to signal when a packet has been dropped. However, the signal now is asserted 1 to 2 clock cycles after end of the packet that was dropped. The User Application is not required to do anything in response to trn_terr_drop_n; it is intended for diagnosing problems when bringing up new designs. This makes timing significantly easier to meet.

TRN Buffer Available Size Change

The Spartan-3 FPGA Endpoint PIPE signal `trn_tbuf_av[4:0]` is one bit wider. The signal is now `trn_tbuf_av[5:0]`. This change reflects the increased number of transmit buffers supported by the Spartan-6 FPGA core.

CMM Arbitration

The Spartan-6 FPGA integrated Endpoint block design has two signals that allow for the user to control the arbitration between the CMM and the TRN interfaces for transmitted packets. These signals are `trn_tcfc_req_n` and `trn_tcfc_gnt_n`.

To maintain the same behavior as the Spartan-3 FPGA Endpoint PIPE, assign the signal `trn_tcfc_gnt_n` asserted (`1'b0`).

TRN Credit Buses Additional Functionality

The TRN credit buses have different names in the Spartan-6 FPGA integrated Endpoint block design. The letter “r” has been removed from the signal names. [Table C-1](#) shows the old and new names.

Table C-1: Credit Bus Name Change from Spartan-3 to Spartan-6 Devices

Spartan-3 FPGAs	Spartan-6 FPGAs
<code>trn_rfc_nph</code>	<code>trn_fc_nph</code>
<code>trn_rfc_npd</code>	<code>trn_fc_npd</code>
<code>trn_rfc_ph</code>	<code>trn_fc_ph</code>
<code>trn_rfc_pd</code>	<code>trn_fc_pd</code>
<code>trn_rfc_cplh</code>	<code>trn_fc_cplh</code>
<code>trn_rfc_cpld</code>	<code>trn_fc_cpld</code>

There is also a signal named `trn_fc_sel[2:0]` that controls what values are placed on the `trn_*` bus.

To maintain the same behavior as the Spartan-3 FPGA Endpoint PIPE, set the `trn_fc_sel[2:0]` signal to `3'b000`. For more information, see [Flow Control Credit Information in Chapter 5](#).

Configuration Error Completion Ready

A signal named `cfg_err_cpl_rdy_n` has been added for the Spartan-6 FPGA integrated Endpoint block. For more information, see [Table 2-9, page 23](#).

Configuration Error Locked

A signal named `cfg_err_locked_n` has been added for the Spartan-6 FPGA integrated Endpoint block. For more information, see [Table 2-9, page 23](#).

Removed Configuration Signals

These signals were removed because they were either unused or not needed:

- cfg_di
- cfg_wr_en_n
- cfg_byte_en_n
- cfg_err_cpl_unexpected_n

Hot Reset

A signal named received_hot_reset has been added for the Spartan-6 FPGA integrated Endpoint block. For more information, see [Table 2-3, page 14](#).

Block RAM Settings

The block RAM settings can now be customized. See the CORE Generator™ software GUI for supported settings.

Signal Change Summary

These signals have been added for the Spartan-6 FPGA integrated Endpoint block:

- sys_clk
- trn_tstr_n
- trn_tcfg_req_n
- trn_tcfg_gnt_n
- cfg_err_cpl_rdy_n
- cfg_err_locked_n
- received_hot_reset
- trn_fc_sel[2:0]

These signals have been removed from the Spartan-6 FPGA integrated Endpoint block:

- cfg_di
- cfg_wr_en_n
- cfg_byte_en_n
- cfg_err_cpl_unexpected_n

[Table C-2](#) shows the signal name changes.

Table C-2: Signal Name Change from Spartan-3 to Spartan-6 Devices

Spartan-3 FPGAs	Spartan-6 FPGAs
trn_rfc_nph	trn_fc_nph
trn_rfc_npd	trn_fc_npdu
trn_rfc_ph	trn_fc_ph
trn_rfc_pd	trn_fc_pd
trn_rfc_cplh	trn_fc_cplh

Table C-2: Signal Name Change from Spartan-3 to Spartan-6 Devices (Cont'd)

Spartan-3 FPGAs	Spartan-6 FPGAs
trn_rfc_cpld	trn_fc_cpld
trn_tbuf_av[4:0]	trn_tbuf_av[5:0]
trn_tdst_dsn_n	trn_tx_terr_drop_n

Debugging Designs

This appendix provides information on using resources available on the Xilinx Support website, available debug tools, and a step-by-step process for debugging designs that use the Spartan®-6 Integrated Endpoint Block for PCI Express®. This appendix uses flow diagrams to guide the user through the debug process.

The following information is found in this appendix:

- [Finding Help on Xilinx.com](#)
- [Contacting Xilinx Technical Support](#)
- [Debug Tools](#)
- [Hardware Debug](#)
- [Simulation Debug](#)

Finding Help on Xilinx.com

To help in the design and debug process when using the Integrated Endpoint Block for PCI Express, the Xilinx Support webpage (www.xilinx.com/support) contains key resources such as Product documentation, Release Notes, Answer Records, and links to opening a Technical Support case.

Documentation

The Data Sheet and User Guide are the main documents associated with the Spartan-6 FPGA integrated Endpoint block, as shown in [Table D-1](#).

Table D-1: Spartan-6 FPGA Integrated Endpoint Block for PCI Express Documentation

Designation	Description
DS	Data Sheet: provides a high-level description of the integrated Endpoint block and key features. It includes information on which ISE software version is supported by the current LogiCORE™ IP version used to instantiate the integrated Endpoint block.
UG	User Guide: provides information on generating an integrated Endpoint block design, detailed descriptions of the interface and how to use the product. The User Guide contains waveforms to show interactions with the block and other important information needed to design with the product.

These Integrated Endpoint Block for PCI Express documents along with documentation related to all products that aid in the design process can be found on the Xilinx Support

webpage. Documentation is sorted by product family at the main support page or by solution at the Documentation Center.

To see the available documentation by device family:

- Navigate to www.xilinx.com/support.
- Select **Spartan-6** from the **Device List** drop-down menu.
- This will sort all available Spartan-6 FPGA documentation by Hardware Documentation, Configuration Solutions Documentation, Related Software Documentation, Tools, IP, and Data Files.

To see the available documentation by solution:

- Navigate to www.xilinx.com/support.
- Select the **Documentation** tab located at the top of the webpage.
- This is the Documentation Center where Xilinx documentation is sorted by Devices, Boards, IP, Design Tools, Doc Type, and Topic.

Release Notes and Known Issues

Known issues for all cores, including the Spartan-6 FPGA Integrated Endpoint Block for PCI Express, are described in the [IP Release Notes Guide](#).

Answer Records

Answer Records include information on commonly encountered problems, helpful information on how to resolve these problems, and any known issues with a product. Answer Records are created and maintained daily ensuring users have access to the most up-to-date information on Xilinx products. Answer Records can be found by searching the Answers Database.

To use the Answers Database Search:

- Navigate to www.xilinx.com/support. The Answers Database Search is located at the top of this webpage.
- Enter keywords in the provided search field and select **Search**.
 - Examples of searchable keywords are product names, error messages, or a generic summary of the issue encountered.
 - To see all answer records directly related to the Spartan-6 FPGA Integrated Endpoint Block for PCI Express, search for the phrase *Spartan-6 Integrated Endpoint Block for PCI Express*.

Contacting Xilinx Technical Support

Xilinx provides premier technical support for customers encountering issues that require additional assistance.

To contact Technical Support:

- Navigate to www.xilinx.com/support.
- Open a WebCase by selecting the WebCase link located under **Support Quick Links**.

When opening a WebCase, include:

- Target FPGA including package and speed grade
- All applicable software versions of the ISE tool, synthesis (if not XST), and simulator
- The xco file created during generation of the LogiCORE IP wrapper
 - This file is located in the directory targeted for the CORE Generator™ software project

Additional files might be required based on the specific issue. See the relevant sections in this debug guide for further information on specific files to include with the WebCase.

Debug Tools

There are many tools available to debug PCI Express design issues. It is important to know which tools would be useful for debugging for the various situations encountered. This chapter references these tools:

Example Design

Xilinx Endpoint for PCI Express products come with a synthesizable back-end application called the PIO design that has been tested and is proven to be interoperable in available systems. The design appropriately handles all incoming 1 DWORD read and write transactions. It returns completions for non-posted transactions and updates the target memory space for writes. For more information, see [Appendix B, Programmed Input/Output Example Design](#).

ChipScope Pro Tool

The ChipScope™ Pro tool inserts logic analyzer, bus analyzer, and virtual I/O software cores directly into the user design. The ChipScope Pro tool allows the user to set trigger conditions to capture application and integrated Endpoint block port signals in hardware. Captured signals can then be analyzed through the ChipScope Pro Logic Analyzer tool. For detailed information on the ChipScope Pro tool, visit www.xilinx.com/chipscope.

Link Analyzers

Third-party link analyzers show link traffic in a graphical or text format. Lecroy, Agilent, and Vmetro are companies that make common analyzers available today. These tools greatly assist in debugging link issues and allow users to capture data which Xilinx support representatives can view to assist in interpreting link behavior.

Third-Party Software Tools

This section describes third-party software tools that can be useful in debugging.

LSPCI (Linux)

LSPCI is available on Linux platforms and allows users to view the PCI Express device configuration space. LSPCI is usually found in the /sbin directory. LSPCI displays a list of devices on the PCI buses in the system. See the LSPCI manual for all command options. Some useful commands for debugging include:

- `lspci -x -d [<vendor>] : [<device>]`

This displays the first 64 bytes of configuration space in hexadecimal form for the device with vendor and device ID specified (omit the -d option to display information for all devices). The default Vendor/Device ID for Xilinx cores is 10EE:6012. Here is a sample of a read of the configuration space of a Xilinx device:

```
> lspci -x -d 10EE:6012
81:00.0 Memory controller: Xilinx Corporation: Unknown device 6012
 00: ee 10 12 60 07 00 10 00 00 00 80 05 10 00 00 00
 10: 00 00 80 fa 00 00 00 00 00 00 00 00 00 00 00 00 00
 20: 00 00 00 00 00 00 00 00 00 00 00 00 ee 10 6f 50
 30: 00 00 00 00 40 00 00 00 00 00 00 00 00 05 01 00 00
```

Included in this section of the configuration space are the Device ID, Vendor ID, Class Code, Status and Command registers, and Base Address Registers.

- `lspci -xxxx -d [<vendor>] : [<device>]`

This displays the extended configuration space of the device. It might be useful to read the extended configuration space on the root and look for the Advanced Error Reporting (AER) registers. These registers provide more information on why the device has flagged an error (for example, it might show that a correctable error was issued because of a replay timer time-out).

- `lspci -k`

Shows kernel drivers handling each device and kernel modules capable of handling it (works with kernel 2.6 or later).

PCItree (Windows)

PCItree can be downloaded at www.pcitree.de and allows the user to view the PCI Express device configuration space and perform 1 DWORD memory writes and reads to the aperture.

The configuration space is displayed by default in the lower right corner when the device is selected, as shown in [Figure D-1](#).

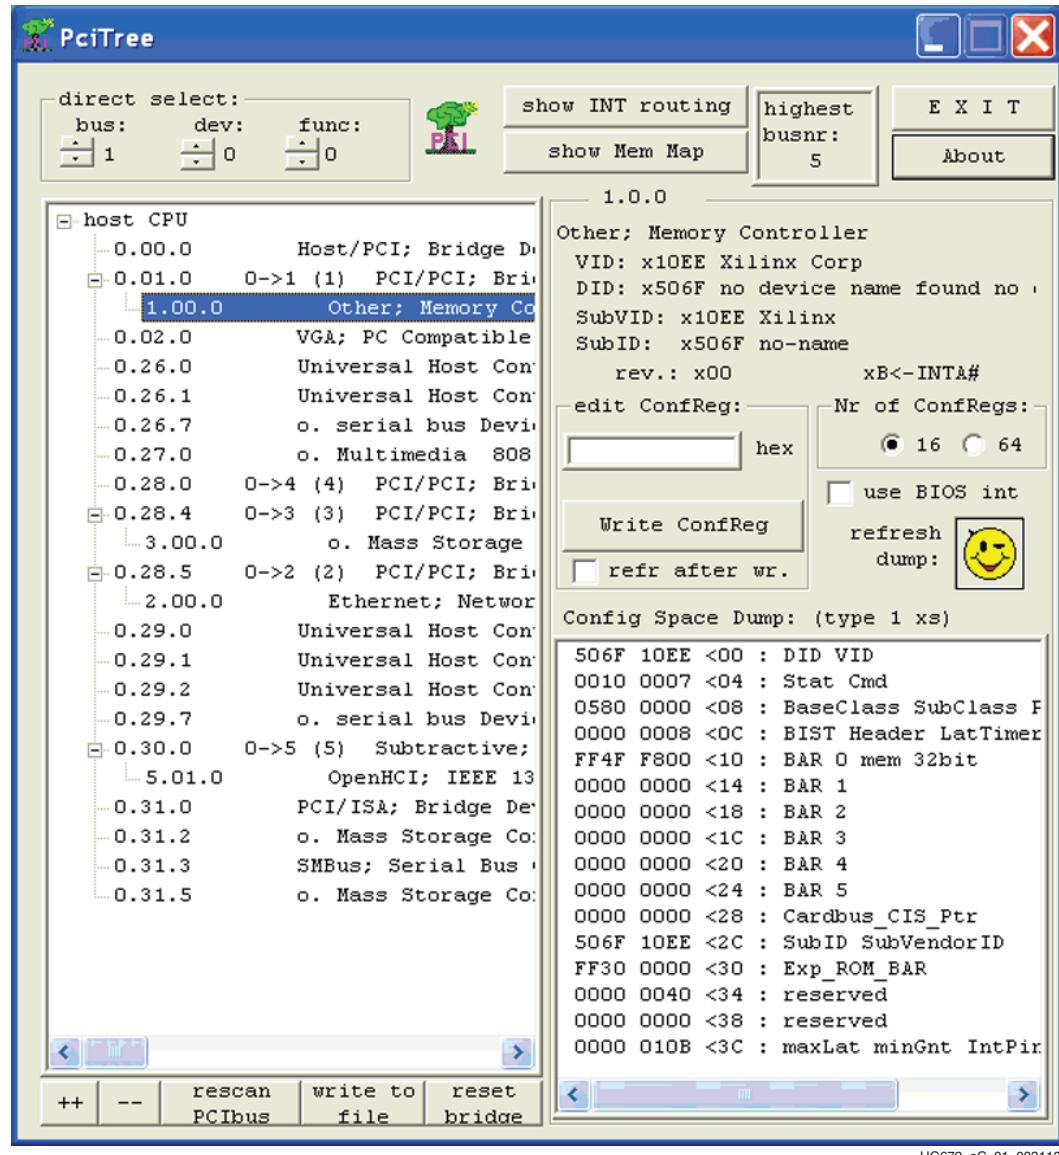


Figure D-1: PCI tree with Read of Configuration Space

UG672_aC_01_083110

HWDIRECT (Windows)

HWDIRECT can be purchased at www.e protek.com and allows the user to view the PCI Express device configuration space as well as the extended configuration space (including the AER registers on the root).

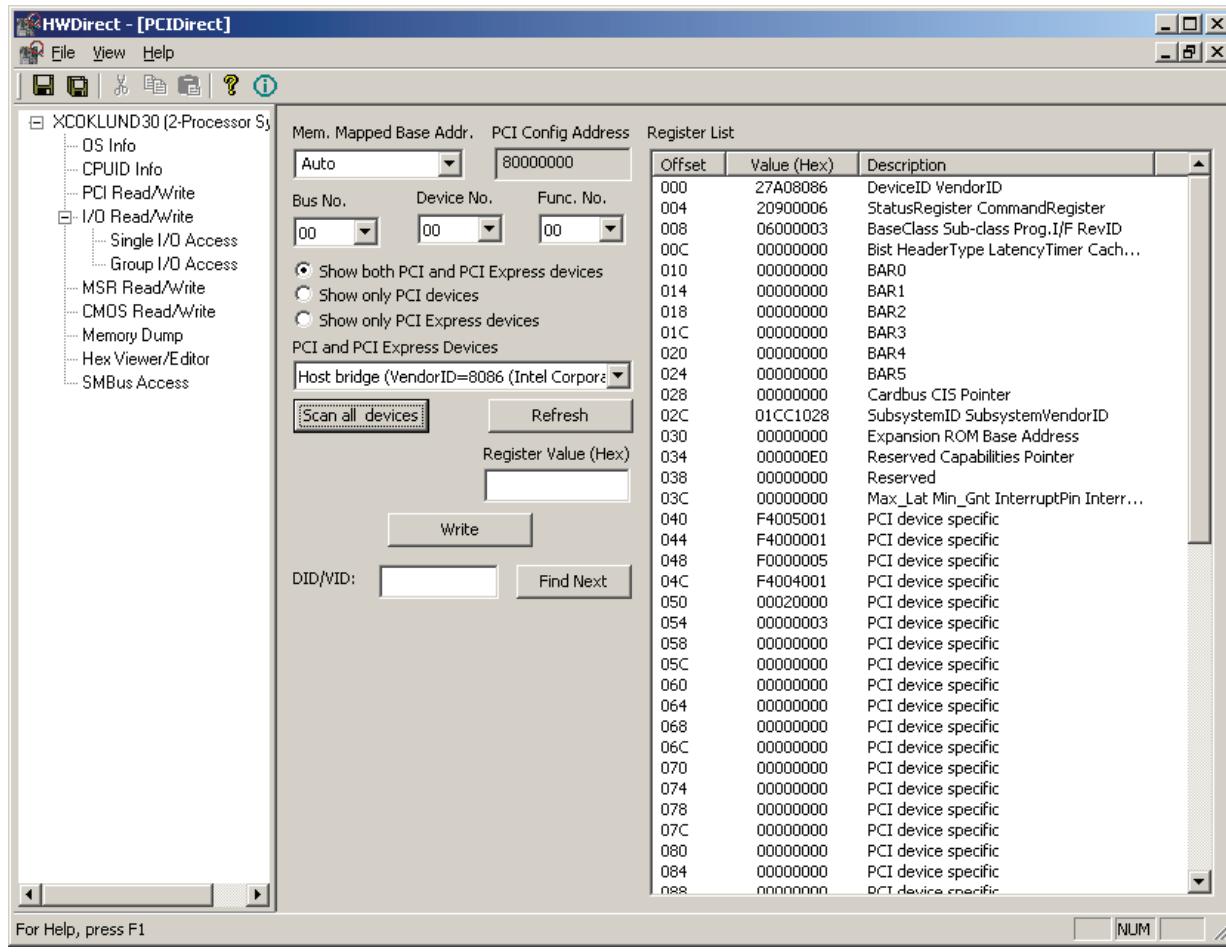


Figure D-2: HWDIRECT with Read of Configuration Space

PCI-SIG Software Suites

PCI-SIG® software suites such as PCIE-CV can be used to test compliance with the specification. This software can be downloaded at www.pcisig.com.

Debug Ports

The Spartan-6 FPGA Integrated Endpoint Block for PCI Express has debug ports described in [Table D-3](#) providing insight to why the different error conditions occur. The receiver might detect different problems that result in either a Fatal, Non-fatal, or correctable error. Also, the receiver might detect an unsupported request. Four of the debug signals shown in [Table D-2](#) mirror the lower four bits of the PCI Express device status register. When one of these conditions occurs, another signal is asserted for one clock cycle to show the reason causing the error.

Table D-2: Device Status Register Debug Ports

Name	Device Status Bit
dbg_reg_detected_correctable	Bit 0 - correctable
dbg_reg_detected_non_fatal	Bit 1 - Non-Fatal
dbg_reg_detected_fatal	Bit 2 - Fatal
dbg_reg_detected_unsupported	Bit 3 - Unsupported Request

[Table D-3](#) defines the debug port signals.

Table D-3: Spartan-6 FPGA Integrated Block for PCI Express Debug Ports

Port	Direction	Clock Domain	Description
dbg_bad_dllp_status	Output	USERCLK	This signal pulses High for one USERCLK cycle when a DLLP CRC error is detected.
dbg_bad_tlp_lcrc	Output	USERCLK	This signal pulses High for one USERCLK cycle when a TLP with an LCRC error is detected.
dbg_bad_tlp_seq_num	Output	USERCLK	This signal pulses High for one USERCLK cycle when a TLP with an invalid sequence number is detected.
dbg_bad_tlp_status	Output	USERCLK	This signal pulses High for one USERCLK cycle when a bad TLP is detected, for reasons other than a bad LCRC or a bad sequence number.
dbg_dl_protocol_status	Output	USERCLK	This signal pulses High for one USERCLK cycle if an out-of-range ACK or NAK is received.
dbg_fc_protocol_err_status	Output	USERCLK	This signal pulses High for one USERCLK cycle if there is a protocol error with the received flow control updates.
dbg_mlfrmd_length	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had a length that did not match what was in the TLP header.
dbg_mlfrmd_mps	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had a length in violation of the negotiated MPS.
dbg_mlfrmd_tcvc	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had an invalid TC or VC value.
dbg_mlfrmd_tlp_status	Output	USERCLK	This signal pulses High for one USERCLK cycle when a malformed TLP is received. See the other DBGMLFRMD* signals for further clarification. Note: There is skew between DBGMLFRMD* and DBGMLFRMDTLPSTATUS.

Table D-3: Spartan-6 FPGA Integrated Block for PCI Express Debug Ports (Cont'd)

Port	Direction	Clock Domain	Description
dbg_mlfrmd_unrec_type	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had an invalid/unrecognized type field value.
dbg_poistlpstatus	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a TLP was received with the EP (poisoned) status bit set.
dbg_rcvr_overflow_status	Output	USERCLK	This signal pulses High for one USERCLK cycle if a received TLP violates the advertised credit.
dbg_reg_detected_correctable	Output	USERCLK	This signal is a mirror of the internal signal used to indicate a correctable error is detected. The error is cleared upon a read by the Root Complex (RC).
dbg_reg_detected_fatal	Output	USERCLK	This signal is a mirror of the internal signal used to indicate that a fatal error has been detected. The error is cleared upon a read by the RC.
dbg_reg_detected_non_fatal	Output	USERCLK	This signal is a mirror of the internal signal used to indicate that a non-fatal error has been detected. The error is cleared upon a read by the RC.
dbg_reg_detected_unsupported	Output	USERCLK	This signal is a mirror of the internal signal used to indicate that an unsupported request has been detected. The error is cleared upon a read by the RC.
dbg_rply_rollover_status	Output	USERCLK	This signal pulses High for one USERCLK cycle when the rollover counter expires.
dbg_rply_timeout_status	Output	USERCLK	This signal pulses High for one USERCLK cycle when the replay time-out counter expires.
dbg_ur_no_bar_hit	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received read or write request did not match any configured BAR.
dbg_ur_pois_cfg_wr	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a CfgWr TLP with the Error/Poisoned bit (EP) = 1 was received.
dbg_ur_status	Output	USERCLK	This signal pulses High for one USERCLK cycle when an unsupported request is received. See the DBGUR* signals for further clarification. Note: There is skew between DBGUR* and DBGURSTATUS.
dbg_ur_unsup_msg	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that an Msg or MsgD TLP with an unsupported type was received.

Using the Debug Ports

The debug ports are outputs on the integrated Endpoint block and users can access them by opening the wrapper source file in the source directory. This file is named <corename>.v [hd] where <corename> represents core name entered in the CORE Generator tool. Signals are defined for each of these ports as either wires in the Verilog version or signals in the VHDL version.

The debug ports can be used in both simulation and in hardware to debug problems. In simulation, these signals can easily be added to the waveform viewer without any changes to the code because they are already defined in the wrapper file. In hardware, users might want to use the ChipScope Pro tool to monitor these signals or probe them to external ports or add additional logic such as counters to enable more in depth analysis. Users might need to bring these signals to upper layers in the design and this can be done by modifying the port description and instantiations of the files.

Figure D-3 shows a common problem faced by many users. This illustrates the behavior of the debug ports when a non-fatal error condition occurs. The scenario is a memory write is sent from the downstream port to the Spartan-6 FPGA Endpoint. This memory write does not correctly target any of the available BARs in the design. This results in a non-fatal error condition. However, there are other reasons that cause a non-fatal error, so monitoring the cfg_dev_status_nonfatal_err_detected output of the core might not be sufficient to debug the problem. By using the debug ports, the user can see that the non-fatal error was caused by a BAR miss due to dbg_ur_no_bar_hit asserted for one cycle.

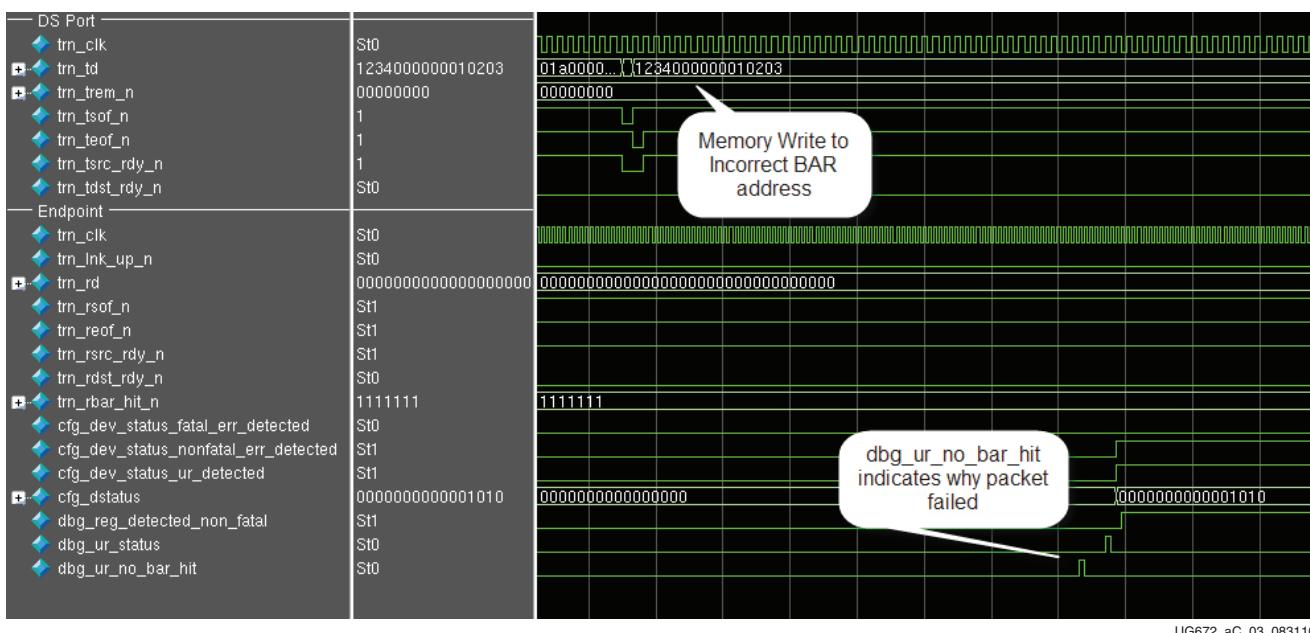
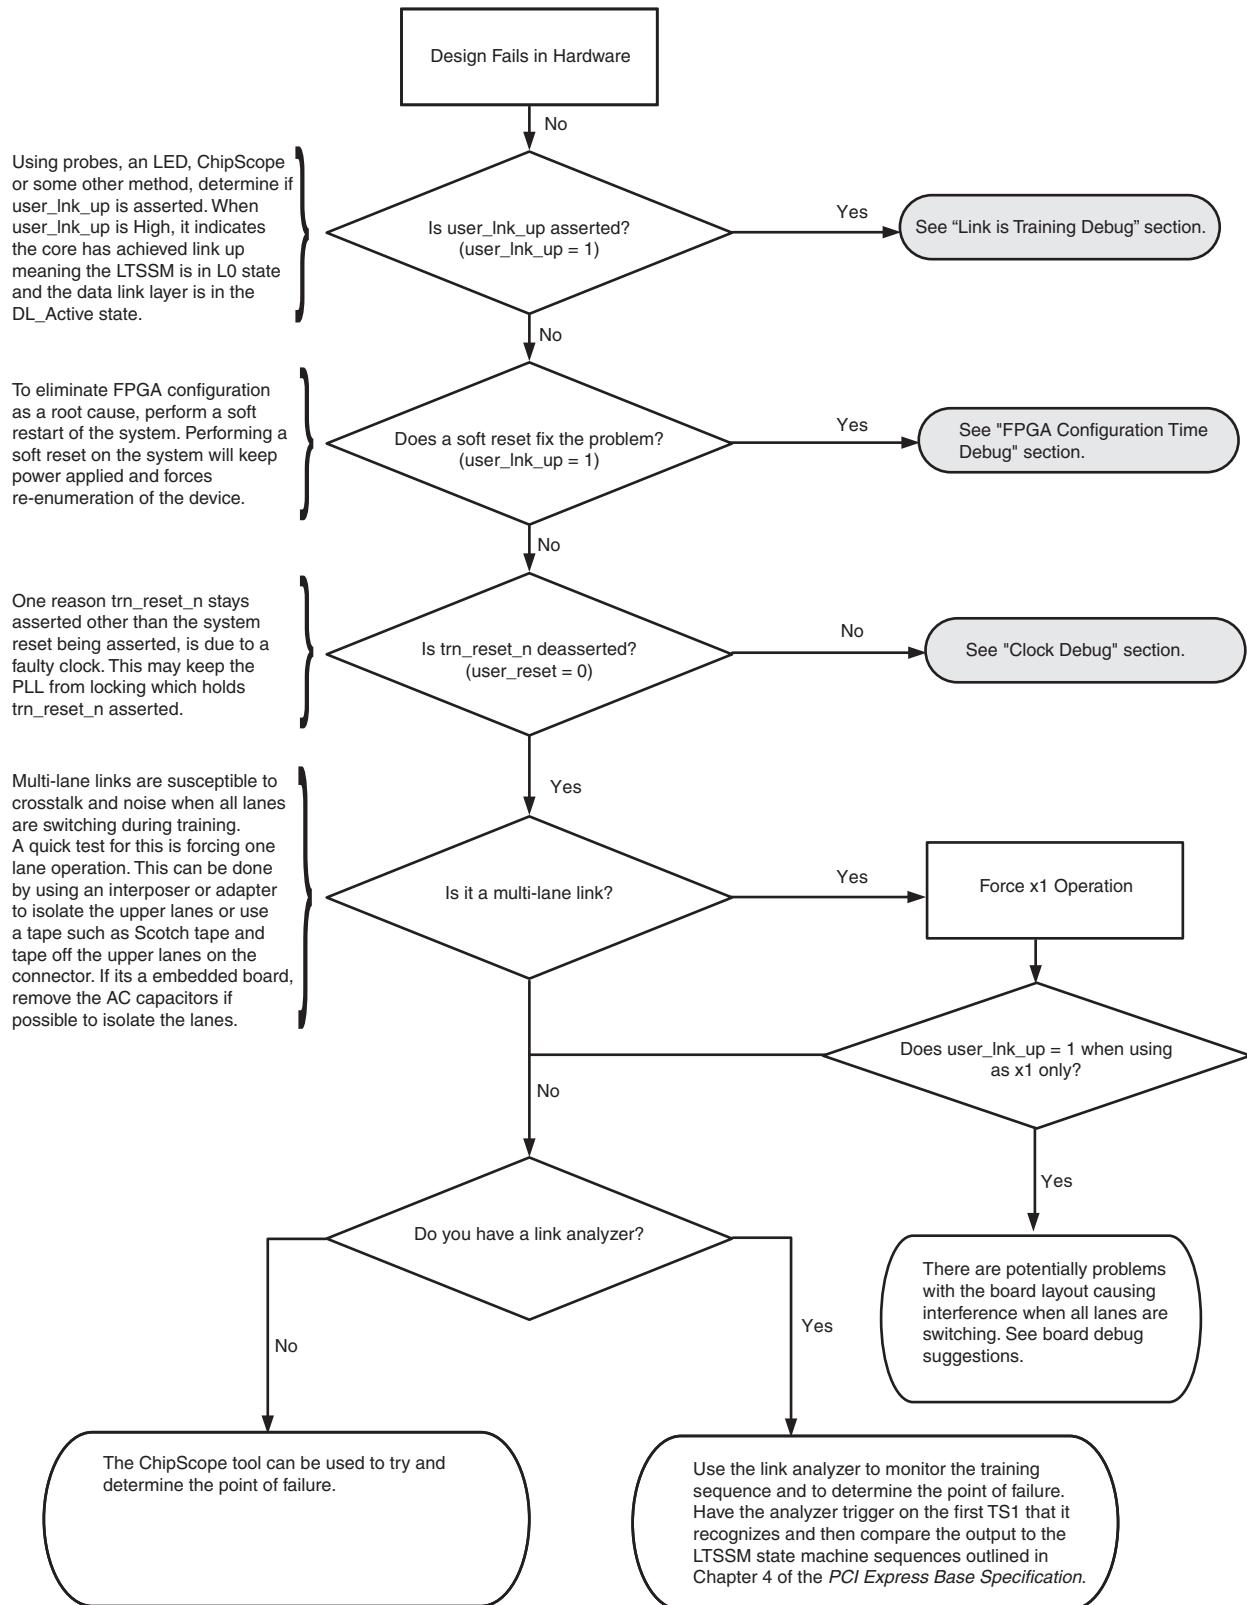


Figure D-3: Debug Wave Screenshot

In the ChipScope tool, the user needs to decide how best to trigger the ChipScope tool to capture these problems. There are various ways to do this, but one suggestion would be to use the signals in [Table D-3](#) as triggers. At least one of these signals is asserted. When the ChipScope tool triggers, the rest of the signals can be analyzed to find out exactly what caused the error condition.

Hardware Debug

Hardware issues can range from device recognition issues to problems seen after hours of testing. This section provides debug flow diagrams for some of the most common issues experienced by users. Endpoints that are shaded gray indicate that more information can be found in sections after [Figure D-4](#).



UG672_aC_04_092010

Figure D-4: Design Fails in Hardware Debug Flow Diagram

FPGA Configuration Time Debug

Device initialization and configuration issues can be caused by not having the FPGA configured fast enough to enter link training and be recognized by the system. Section 6.6 of *PCI Express Base Specification, v1.1* states two rules that might be impacted by FPGA Configuration Time:

- A component must enter the LTSSM Detect state within 20 ms of the end of the Fundamental reset.
- A system must guarantee that all components intended to be software visible at boot time are ready to receive Configuration Requests within 100 ms of the end of Conventional Reset at the Root Complex.

These statements basically mean there is a finite time in which the FPGA must be configured by, and not meeting these requirements could cause problems with link training and device recognition.

Configuration can be accomplished using an onboard PROM or dynamically using JTAG. When using JTAG to configure the device, configuration typically occurs after the Chipset has enumerated each peripheral. After configuring the FPGA, a soft reset is required to restart enumeration and configuration of the device. A soft reset on a Windows based PC is performed by going to **Start -> Shut Down** and then selecting **Restart**.

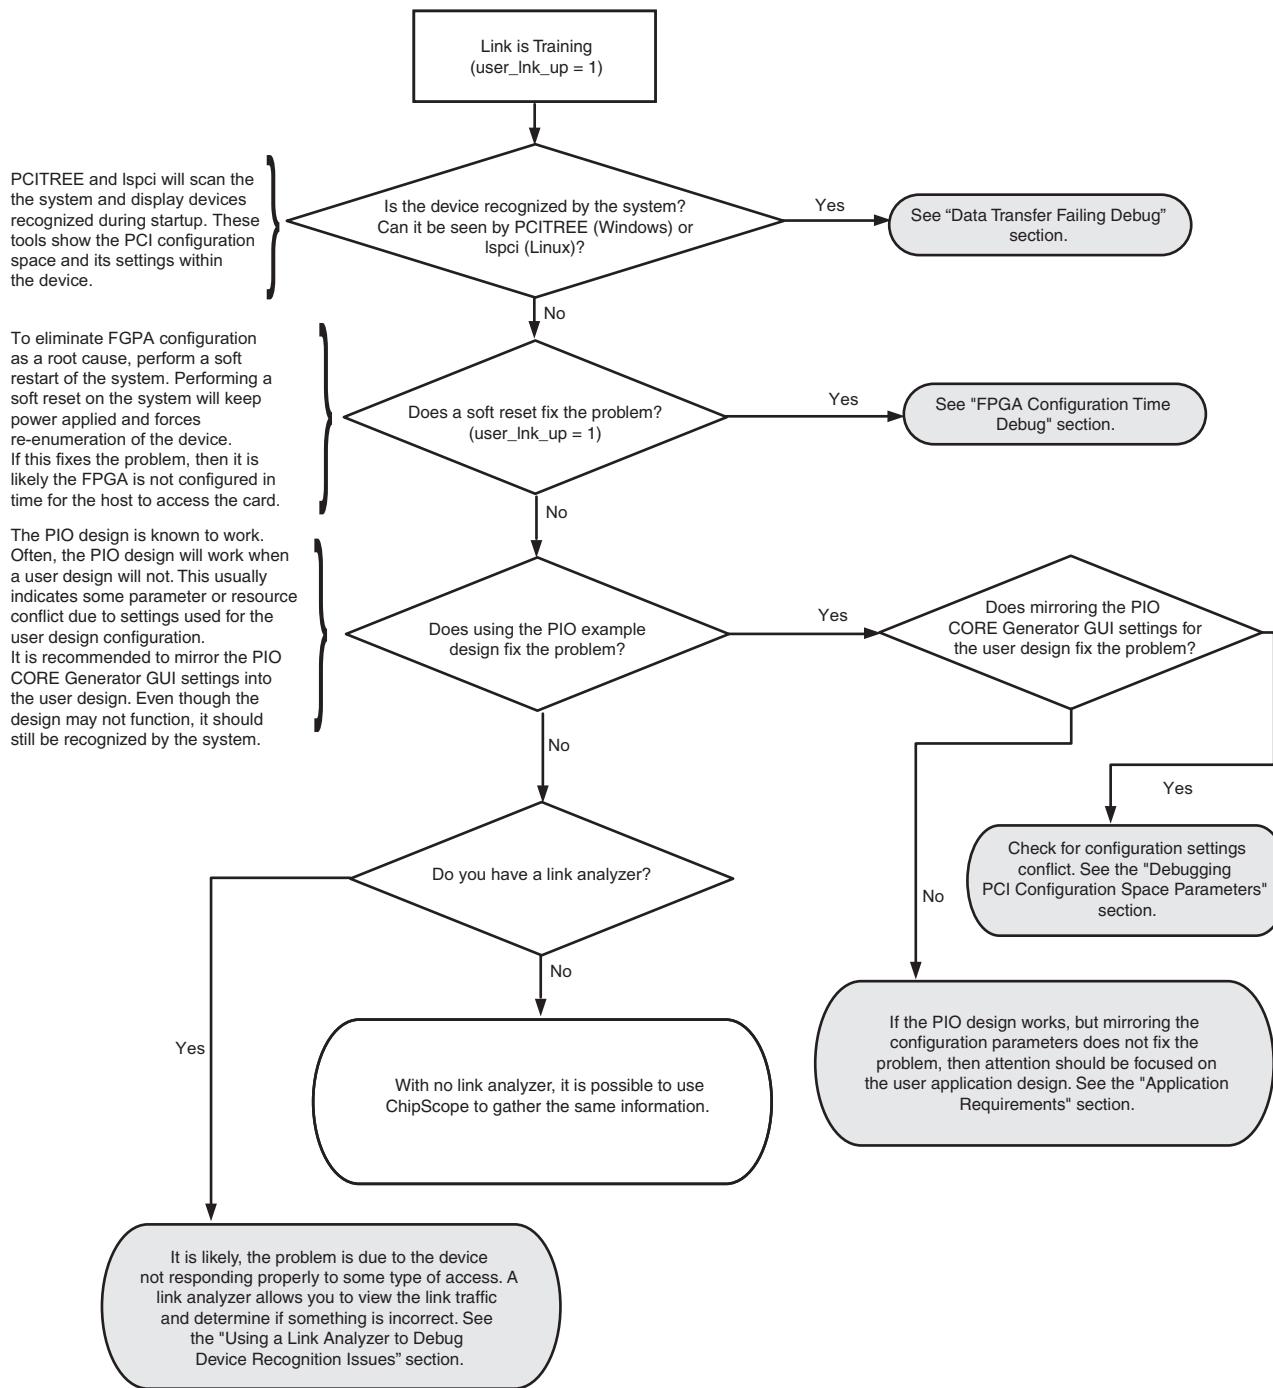
To eliminate FPGA configuration as a root cause, perform a soft restart of the system. Performing a soft reset on the system keeps power applied and forces re-enumeration of the device. If the device links up and is recognized after a soft reset is performed, then FPGA configuration is most likely the problem. Most typical systems use ATX power supplies which provides some margin on this 100 ms window as the power supply is normally valid before the 100 ms window starts. For more information on FPGA configuration, see [Chapter 7, FPGA Configuration](#).

Clock Debug

One reason to not deassert the user_reset_out signal is that the fabric PLL and transceiver PLL have not locked to the incoming clock. To verify lock, monitor the transceiver's RXPLLKDET output and the fabric PLL's LOCK output. If the PLLs are not locking, you must ensure the incoming reference clock meets the requirements discussed in the *Spartan-6 FPGA GTP Transceivers User Guide* (UG386). The REFCLK signal should be routed to the dedicated reference clock input pins on the serial transceiver, and the user design should instantiate the IBUFDS primitive in the user design. See the *Spartan-6 FPGA GTP Transceivers User Guide* for more information on PCB layout requirements including reference clock requirements.

Reference clock jitter has the potential to close both the TX and RX eyes, depending on the frequency content of the phase jitter. Therefore, as clean a reference clock as possible must be maintained. Reduce crosstalk on the REFCLK signal by isolating the clock signal from nearby high-speed traces. Maintain a separation of at least 25 mils from the nearest aggressor signals. See [Appendix F, Board Design Guidelines](#) for more information on reference clock considerations. The PCI Special Interest Group website provides other tools for ensuring the reference clocks are compliant to the requirements of the *PCI Express Specification*: www.pcisig.com/specifications/pciexpress/compliance/compliance_library.

Link is Training Debug



UG672_aC_05_092010

Figure D-5: Link Trained Debug Flow Diagram

Debugging PCI Configuration Space Parameters

Often, a user application fails to be recognized by the system, but the Xilinx PIO Example design works. In these cases, the user application is often using a PCI configuration space

setting that is interfering with the system's ability to recognize and allocate resources to the card.

Xilinx PCI Express solutions handle all configuration transactions internally and generate the correct responses to incoming configuration requests. Chipsets have limits as to the amount of system resources they can allocate, and the core must be configured to adhere to these limitations.

The resources requested by the endpoint are identified by the BAR settings within the Endpoint configuration space. Verify that the resources requested in each BAR can be allocated by the chipset. I/O BARs are especially limited so configuring a large I/O BAR typically prevents the chipset from configuring the device. Generate a core that implements a small amount of Memory (approximately 2 KB) to identify if this is the root cause.

The Class Code setting selected in the CORE Generator software GUI can also affect configuration. The Class Code informs the Chipset as to what type of device the Endpoint is. Chipsets might expect a certain type of device to be plugged into the PCI Express slot and configuration might fail if it reads an unexpected Class Code. The BIOS could be configurable to workaround this issue.

Use the PIO design with default settings to rule out any device allocation issues. The PIO design default settings have proven to work in all systems encountered when debugging problems. If the default settings allow the device to be recognized, then change the PIO design settings to match the intended user application by changing the PIO configuration in the CORE Generator software GUI. Trial and error might be required to pinpoint the problem if a link analyzer is not available.

Using a link analyzer, it is possible to monitor the link traffic and possibly determine when during the enumeration and configuration process problems occur.

Application Requirements

During enumeration, it is possible for the chipset to issue TLP traffic that is passed from the core to the backend application. A common oversight when designing custom backend applications is to not have logic which handles every type incoming request. As a result, no response is created and problems arise. The PIO design has the necessary backend functions to respond correctly to any incoming request. It is the responsibility of the application to generate the correct response. These packet types are presented to the application:

- Requests targeting the Expansion ROM (if enabled)
- Message TLPs
- Memory or I/O requests targeting a BAR
- All completion packets

The PIO design, can be used to rule out any of these types of concerns, as the PIO design responds to all incoming transactions to the user application in some way to ensure the host receives the proper response allowing the system to progress. If the PIO design works, but the custom application does not, this means that some transaction is not being handled properly.

The ChipScope analyzer should be implemented on the wrapper AXI4-Stream Receive interface to identify if requests targeting the backend application are drained and completed successfully. The AXI4-Stream interface signals that should be probed in the ChipScope analyzer are defined in [Table E-4, page 185](#).

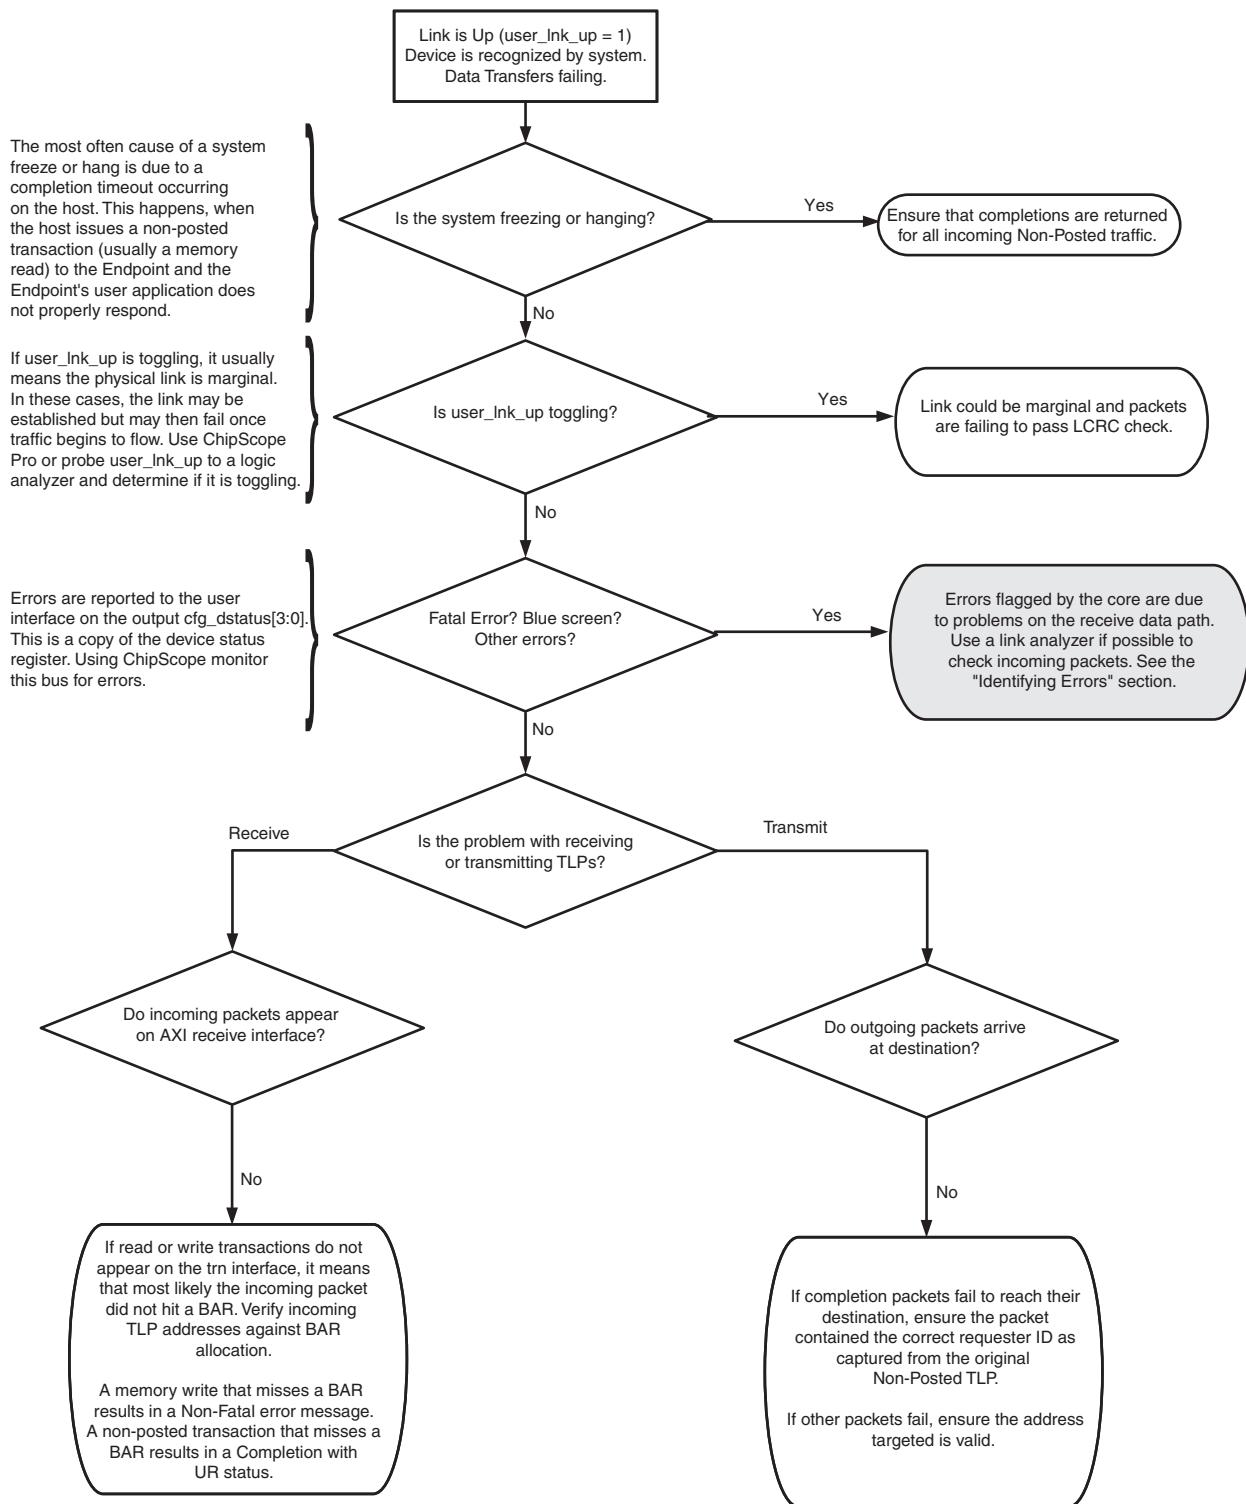
Using a Link Analyzer to Debug Device Recognition Issues

In cases where the link is up (`user_lnk_up = 1`), but the device is not recognized by the system, a link analyzer can help solve the problem. It is likely the FPGA is not responding properly to some type of access. Use the link view to analyze the traffic and see if anything looks out of place.

To focus on the problem, it might be necessary to try different triggers. Here are some trigger examples:

- Trigger on the first `INIT_FC1` and/or `UPDATE_FC` in either direction. This allows the analyzer to begin capture after link up.
- The first TLP normally transmitted to an endpoint is the Set Slot Power Limit Message. This usually occurs before Configuration traffic begins. This might be a good trigger point.
- Trigger on Configuration TLPs.
- Trigger on Memory Read or Memory Write TLPs.

Data Transfer Failing Debug



UG672_aC_06_092010

Figure D-6: Data Transfer Debug Flow Diagram

Identifying Errors

Hardware symptoms of system lock up issues are indicated when the system hangs or a blue screen appears (PC systems). The *PCI Express Base Specification v2.0* requires that error detection be implemented at the receiver. A system lock up or hang is commonly the result of a Fatal Error and is reported in bit 2 of the receivers Device Status register. Using the ChipScope tool, monitor the core's device status register to see if a fatal error is being reported.

A fatal error reported at the Root complex implies an issue on the transmit side of the EP. The Root Complex Device Status register can often times be seen using PCITree (Windows) or LSPCI (Linux). If a fatal error is detected, refer to the [Transmit](#) section. A Root Complex can often implement Advanced Error Reporting (AER) which further distinguishes the type of error reported. AER provides valuable information as to why a certain error was flagged and is provided as an extended capability within a devices configuration space. Section 7.10 of the *PCI Express Base Specification v2.0* provides more information on AER registers.

Transmit

Fatal Error Detected on Root or Link Partner

Check to make sure the TLP is correctly formed and that the payload (if one is attached) matches what is stated in the header length field. The Endpoints device status register does not report errors created by traffic on the transmit channel.

The following signals should be monitored on the Transmit interface to verify all traffic being initiated is correct. Refer to [Table 2-6](#) for signal descriptions.

- user_lnk_up
- s_axis_tx_tlast
- s_axis_tx_tdata
- s_axis_tx_trb
- s_axis_tx_tvalid
- s_axis_tx_tready

Fatal Error Not Detected

Ensure that the address provided in the TLP header is valid. The kernel mode driver attached to the device is responsible for obtaining the system resources allocated to the device. In a Bus Mastering design, the driver is also responsible for providing the application with a valid address range. System hangs or blue screens might occur if a TLP contains an address which does not target the designated address range for that device.

Receive

Xilinx solutions for PCI Express provide the Device Status register to the application on CFG_DSTATUS[3:0]. Debug ports are available to help users determine the exact cause of errors on the Endpoint's receiver. See [Debug Ports, page 167](#) for information on these ports.

Table D-4: Description of CFG_DSTATUS[3:0]

CFG_DSTATUS[3:0]	Description
CFG_DSTATUS[0]	Correctable Error Detected
CFG_DSTATUS[1]	Non-Fatal Error Detected

Table D-4: Description of CFG_DSTATUS[3:0]

CFG_DSTATUS[3:0]	Description
CFG_DSTATUS[2]	Fatal Error Detected
CFG_DSTATUS[3]	UR Detected

System lock up conditions due to issues on the receive channel of the PCI Express core are often result of an error message being sent upstream to the root. Error messages are only sent when error reporting is enabled in the Device Control register.

A fatal condition is reported if any of these occur:

- Training Error
- DLL Protocol Error
- Flow Control Protocol Error
- Malformed TLP
- Receiver Overflow

The first four bullets are not common in hardware because both Xilinx PCI Express solutions and connected components have been thoroughly tested in simulation and hardware. However, a receiver overflow is a possibility. Users must ensure they follow requirements discussed in [Receiver Flow Control Credits Available in Chapter 5](#) when issuing memory reads.

Debug ports are available to help users determine the exact cause of errors on the endpoint's receiver. See [Debug Ports, page 167](#) for information on these ports.

Non-Fatal Errors

This section lists conditions that are reported as Non-Fatal errors. See the *PCI Express Base Specification* for more details.

If the error is being reported by the root, the Advanced Error Reporting (AER) registers can be read to determine the condition that led to the error. Use a tool such as HWDIRECT, discussed in [Third-Party Software Tools, page 164](#), to read the root's AER registers. Chapter 7 of the *PCI Express Base Specification* defines the AER registers. If the error is signaled by the endpoint, debug ports are available to help determine the specific cause of the error.

Correctable Non-Fatal errors are:

- Receiver Error
- Bad TLP
- Bad DLLP
- Replay Timeout
- Replay NUM Rollover

The first three errors listed above are detected by the receiver and are not common in hardware systems. The replay error conditions are signaled by the transmitter. If an ACK is not received for a packet within the allowed time, it is replayed by the transmitter. Throughput can be reduced if many packets are being replayed, and the source can usually be determined by examining the link analyzer or ChipScope tool captures.

Uncorrectable Non-Fatal errors are:

- Poisoned TLP
- Received ECRC Check Failed
- Unsupported Request (UR)
- Completion Timeout
- Completer Abort
- Unexpected Completion

An unsupported request usually indicates that the address in the TLP did not fall within the address space allocated to the BAR. This often points to a problem with the address translation performed by the driver. Ensure also that the BAR has been assigned correctly by the root at start-up. LSPCI or PCItree discussed in [Third-Party Software Tools, page 164](#) can be used to read the BAR values for each device.

A completion timeout indicates that no completion was returned for a transmitted TLP and is reported by the requester. This can cause the system to hang (could include a blue screen on Windows) and is usually caused when one of the devices locks up and stops responding to incoming TLPs. If the root is reporting the completion timeout, the ChipScope analyzer can be used to investigate why the User Application did not respond to a TLP (for example, the User Application is busy, there are no transmit buffers available, or `ts_axis_tx_tready` is deasserted). If the endpoint is reporting the Completion timeout, a link analyzer would show the traffic patterns during the time of failure and would be useful in determining the root cause.

Next Steps

If the debug suggestions listed above do not resolve the issue, open a support case to have the appropriate Xilinx expert assist with the issue.

To create a technical support case in Webcase, see the Xilinx website at:

www.xilinx.com/support/clearexpress/websupport.htm

Items to include when opening a case:

- Detailed description of the issue and results of the steps listed above.
- Attach ChipScope analyzer VCD captures taken in the steps above.

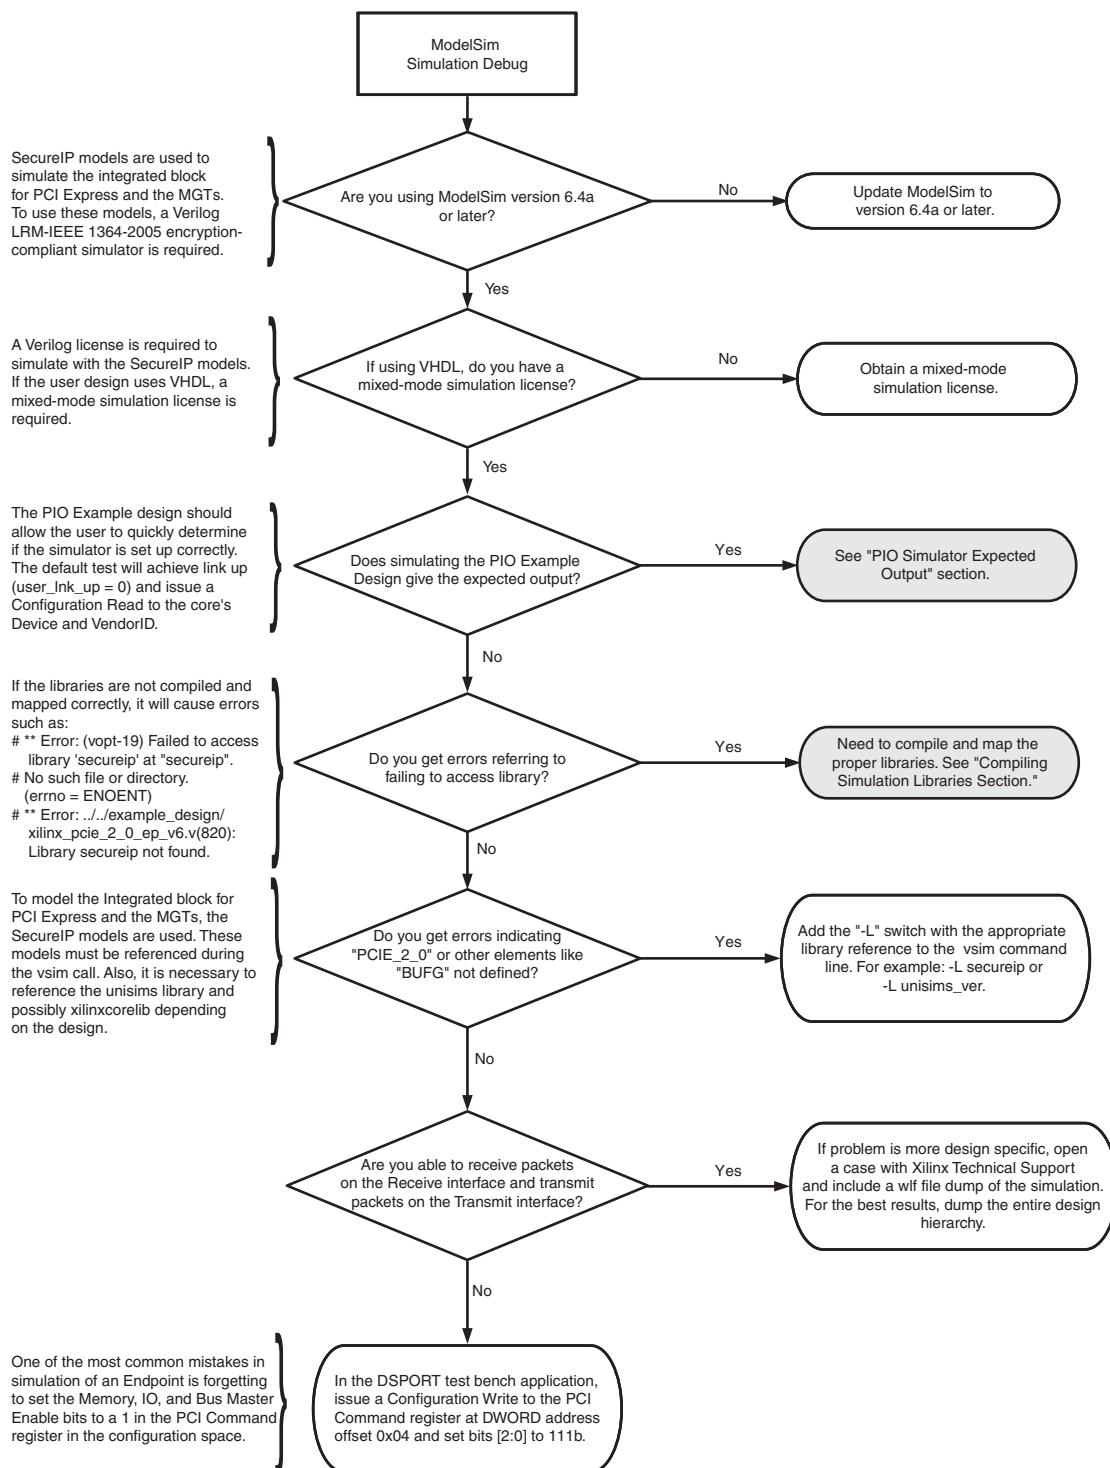
To discuss possible solutions, use the Xilinx User Community:

forums.xilinx.com/xlnx/

Simulation Debug

This section provides simulation debug flow diagrams for some of the most common issues experienced by users. Endpoints that are shaded gray indicate that more information can be found in sections below the figure.

ModelSim Debug



UG672_aC_07_092210

Figure D-7: ModelSim Debug Flow Diagram

PIO Simulator Expected Output

The PIO design simulation should give the output as follows:

```
# Loading work.board(fast)
# Loading unisims_ver.IBUFDS_GTXE1(fast)
# Loading work.pcie_clocking_v6(fast)
# Loading unisims_ver.PCIE_2_0(fast)
# Loading work.pcie_gtx_v6(fast)
# Loading unisims_ver.GTXE1(fast)
# Loading unisims_ver.RAMB36(fast)
# Loading unisims_ver.RAMB16_S36_S36(fast)
# Loading unisims_ver.PCIE_2_0(fast_1)
# Loading work.glbl(fast)
# [          0] board.EP.core.pcie_2_0_i.pcie_bram_i ROWS_TX 1 COLS_TX 2
# [          0] board.EP.core.pcie_2_0_i.pcie_bram_i ROWS_RX 1 COLS_RX 2
# [          0] board.EP.core.pcie_2_0_i.pcie_bram_i.pcie_brams_tx NUM_BRAMS 2
DOB_REG 1 WIDTH 36 RAM_WRITE_LATENCY 0 RAM_RADDR_LATENCY 0 RAM_RDATA_LATENCY 2
# [          0] board.EP.core.pcie_2_0_i.pcie_bram_i.pcie_brams_rx NUM_BRAMS 2
DOB_REG 1 WIDTH 36 RAM_WRITE_LATENCY 0 RAM_RADDR_LATENCY 0 RAM_RDATA_LATENCY 2
# [          0] board.RP.rport.pcie_2_0_i.pcie_bram_i ROWS_TX 1 COLS_TX 2
# [          0] board.RP.rport.pcie_2_0_i.pcie_bram_i ROWS_RX 1 COLS_RX 2
# [          0] board.RP.rport.pcie_2_0_i.pcie_bram_i.pcie_brams_tx NUM_BRAMS 2
DOB_REG 1 WIDTH 36 RAM_WRITE_LATENCY 0 RAM_RADDR_LATENCY 0 RAM_RDATA_LATENCY 2
# [          0] board.RP.rport.pcie_2_0_i.pcie_bram_i.pcie_brams_rx NUM_BRAMS 2
DOB_REG 1 WIDTH 36 RAM_WRITE_LATENCY 0 RAM_RADDR_LATENCY 0 RAM_RDATA_LATENCY 2
# Running test {sample_smoke_test0}.....
# [          0] : System Reset Asserted...
# [ 4995000] : System Reset De-asserted...
# [ 64069100] : Transaction Reset Is De-asserted...
# [ 73661100] : Transaction Link Is Up...
# [ 73661100] : Expected Device/Vendor ID = 000710ee
# [ 73661100] : Reading from PCI/PCI-Express Configuration Register 0x00
# [ 73673000] : TSK_PARSE_FRAME on Transmit
# [ 74941000] : TSK_PARSE_FRAME on Receive
# [ 75273000] : TEST PASSED --- Device/Vendor ID 000710ee successfully received
# ** Note: $finish : ./tests/sample_tests1.v(29)
#     Time: 75273 ns  Iteration: 3  Instance: /board/RP/tx_usrapp
```

Compiling Simulation Libraries

Use the compxlib command to compile simulation libraries. This tool is delivered as part of the Xilinx software. For more information see the ISE Software Manuals and specifically the *Development System Reference Guide* under the section titled compxlib.

Assuming the Xilinx and ModelSim environments are set up correctly, this is an example of compiling the SecureIP and UniSims libraries for Verilog into the current directory

```
compxlib -s mti_se -arch spartan6 -l verilog -lib secureip -lib unisims
-dir ./
```

There are many other options available for compxlib described in the *Development System Reference Guide*.

Compxlib produces a modelsim.ini file containing the library mappings. In ModelSim, to see the current library mappings type **vmap** at the prompt. The mappings can be updated in the INI file or to map a library at the ModelSim prompt type:

```
vmap [<logical_name>] [<path>]
```

For example:

```
Vmap unisims_ver C:\my_unisim_lib
```

Next Step

If the debug suggestions listed above do not resolve the issue, a support case should be opened to have the appropriate Xilinx expert assist with the issue.

To create a technical support case in Webcase, see the Xilinx website at:

www.xilinx.com/support/clearexpress/websupport.htm

Items to include when opening a case:

- Detailed description of the issue and results of the steps listed above.
- Attach a VCD or WLF dump of the simulation.

To discuss possible solutions, use the Xilinx User Community:

forums.xilinx.com/xlnx/

Managing Receive-Buffer Space for Inbound Completions

The PCI Express Base Specification requires all Endpoints to advertise infinite Flow Control credits for received Completions to their link partners. This means that an Endpoint must only transmit Non-Posted Requests for which it has space to accept Completion responses. This appendix describes how a User Application can manage the receive-buffer space in the PCI Express Endpoint core to fulfill this requirement.

General Considerations and Concepts

Completion Space

Table E-1 defines the completion space reserved in the receive buffer by the core. The values differ for different versions of the core, and also differ based on whether the designer chooses to have TLP Digests (ECRC) removed from the incoming packet stream. Values are credits, expressed in decimal.

Table E-1: Receiver-Buffer Completion Space

Capability Max Payload Size (bytes)	Performance Level : Good		Performance Level : High	
	Cpl. Hdr. (Total_CpIH)	Cpl. Data (Total_CpID)	Cpl. Hdr. (Total_CpIH)	Cpl. Data (Total_CpID)
128	8	64	16	128
256	16	128	32	256
512	32	256	32	256

Maximum Request Size

A Memory Read cannot request more than the value stated in Max_Request_Size, which is given by Configuration bits cfg_dcommand[14:12] as defined in [Table E-2](#). If the User Application chooses not to read the Max_Request_Size value, it must use the default value of 128 bytes.

Table E-2: Max Request Size Settings

cfg_dcommand[14:12]	Max_Request_Size			
	Bytes	DW	QW	Credits
000b	128	32	16	8
001b	256	64	32	16
010b	512	128	64	32
011b	1024	256	128	64
100b	2048	512	256	128
101b	4096	1024	512	256
110b-111b	Reserved			

Read Completion Boundary

A Memory Read can be answered with multiple Completions, which when put together return all requested data. To make room for packet-header overhead, the User Application must allocate enough space for the maximum number of Completions that might be returned.

To make this process easier, the *Base Specification* quantizes the length of all Completion packets such that each must start and end on a naturally aligned Read Completion Boundary (RCB), unless it services the starting or ending address of the original request. The value of RCB is determined by Configuration bit cfg_lcommand[3] as defined in [Table E-3](#). If the User Application chooses not to read the RCB value, it must use the default value of 64 bytes.

Table E-3: Read Completion Boundary Settings

cfg_lcommand[3]	Read Completion Boundary			
	Bytes	DW	QW	Credits
0	64	16	8	4
1	128	32	16	8

When calculating the number of Completion credits a Non-Posted Request requires, the user must determine how many RCB-bounded blocks the Completion response might require; this is the same as the number of Completion Header credits required.

Methods of Managing Completion Space

A User Application can choose one of four methods to manage receive-buffer Completion space, as listed in [Table E-4](#). For convenience, this discussion refers to these methods as LIMIT_FC, PACKET_FC, RCB_FC, and DATA_FC. Each has advantages and disadvantages that the designer needs to consider when developing the user application.

Table E-4: Managing Receive Completion Space Methods

Method	Description	Advantage	Disadvantage
LIMIT_FC	Limit the total number of outstanding NP Requests	Simplest method to implement in user logic	Much Completion capacity goes unused
PACKET_FC	Track the number of outstanding CplH and CplD credits; allocate and deallocate on a per-packet basis	Relatively simple user logic; finer allocation granularity means less wasted capacity than LIMIT_FC	As with LIMIT_FC, credits for an NP are still tied up until the Request is completely satisfied
RCB_FC	Track the number of outstanding CplH and CplD credits; allocate and deallocate on a per-RCB basis	Ties up credits for less time than PACKET_FC	More complex user logic than LIMIT_FC or PACKET_FC
DATA_FC	Track the number of outstanding CplH and CplD credits; allocate and deallocate on a per-RCB basis	Lowest amount of wasted capacity	Most complex user logic

The LIMIT_FC Method

The LIMIT_FC method is the simplest to implement. The User Application assesses the maximum number of outstanding Non-Posted Requests allowed at one time, MAX_NP. To calculate this value, perform these steps:

1. Determine the number of CplH credits required by a Max_Request_Size packet:

$$\text{Max_Header_Count} = \text{ceiling}(\text{Max_Request_Size} / \text{RCB})$$
2. Determine the greatest number of maximum-sized Completions supported by the CplD credit pool:

$$\text{Max_Packet_Count_CplD} = \text{floor}(\text{CplD} / \text{Max_Request_Size})$$
3. Determine the greatest number of maximum-sized Completions supported by the CplH credit pool:

$$\text{Max_Packet_Count_CplH} = \text{floor}(\text{CplH} / \text{Max_Header_Count})$$
4. Use the *smaller* of the two quantities from steps 2 and 3 to obtain the maximum number of outstanding Non-Posted requests:

$$\text{MAX_NP} = \min(\text{Max_Packet_Count_CplH}, \text{Max_Packet_Count_CplD})$$

With knowledge of MAX_NP, the User Application can load a register NP_PENDING with zero at reset and make sure it always stays with the range 0 to MAX_NP. When a Non-Posted Request is transmitted, NP_PENDING decrements by one. When *all* Completions for an outstanding NP Request are received, NP_PENDING increments by one.

Although this method is the simplest to implement, it potentially wastes the most receiver space because an entire Max_Request_Size block of Completion credit is allocated for each Non-Posted Request, regardless of actual request size. The amount of waste becomes greater when the User Application issues a larger proportion of short Memory Reads (on the order of a single DWORD), I/O Reads and I/O Writes.

The PACKET_FC Method

The PACKET_FC method allocates blocks of credit in finer granularities than LIMIT_FC, using the receive Completion space more efficiently with a small increase in user logic.

Start with two registers, CPLH_PENDING and CPLD_PENDING, (loaded with zero at reset), and then perform these steps:

- When the User Application needs to send an NP request determine the potential number of CplH and CplD credits, it might require:

$$\text{NP_CplH} = \text{ceiling}[(\text{Start_Address mod RCB}) + \text{Request_Size}) / \text{RCB}]$$

$$\text{NP_CplD} = \text{ceiling}[(\text{Start_Address mod 16 bytes}) + \text{Request_Size}) / 16 \text{ bytes}] \\ (\text{except I/O Write, which returns zero data})$$

The modulo and ceiling functions ensure that any fractional RCB or credit blocks are rounded up. For example, if a Memory Read requests 8 bytes of data from address 7Ch, the returned data can potentially be returned over two Completion packets (7Ch-7Fh, followed by 80h-83h). This would require two RCB blocks and two data credits.

- Check the following:

$$\text{CPLH_PENDING} + \text{NP_CplH} \leq \text{Total_CplH} \text{ (from Table E-1)}$$

$$\text{CPLD_PENDING} + \text{NP_CplD} \leq \text{Total_CplD} \text{ (from Table E-1)}$$

- If both inequalities are true, transmit the Non-Posted Request, increase CPLH_PENDING by NP_CplH and CPLD_PENDING by NP_CplD. For each NP Request transmitted, keep NP_CplH and NP_CplD for later use.
- When all Completion data is returned for an NP Request, decrement CPLH_PENDING and CPLD_PENDING accordingly.

This method is less wasteful than LIMIT_FC but still ties up all of an NP Request's Completion space until the *entire* request is satisfied. RCB_FC and DATA_FC provide finer de-allocation granularity at the expense of more logic.

The RCB_FC Method

The RCB_FC method allocates and de-allocates blocks of credit in RCB granularity. Credit is freed on a per-RCB basis.

As with PACKET_FC, start with two registers, CPLH_PENDING and CPLD_PENDING (loaded with zero at reset).

- Calculate the number of data credits per RCB:

$$\text{CplD_PER_RCB} = \text{RCB} / 16 \text{ bytes}$$

- When the User Application needs to send an NP request, determine the potential number of CplH credits it might require. Use this to allocate CplD credits with RCB granularity:

$$\text{NP_CplH} = \text{ceiling}[(\text{Start_Address mod RCB}) + \text{Request_Size}) / \text{RCB}]$$

$$\text{NP_CplD} = \text{NP_CplH} \times \text{CplD_PER_RCB}$$

3. Check the following:

$$\text{CPLH_PENDING} + \text{NP_CplH} \leq \text{Total_CplH}$$

$$\text{CPLD_PENDING} + \text{NP_CplD} \leq \text{Total_CplD}$$

4. If both inequalities are true, transmit the Non-Posted Request, increase CPLH_PENDING by NP_CplH and CPLD_PENDING by NP_CplD.
5. At the start of each incoming Completion, or when that Completion begins at or crosses an RCB without ending at that RCB, decrement CPLH_PENDING by 1 and CPLD_PENDING by CplD_PER_RCB. Any Completion can cross more than one RCB. The number of RCB crossings can be calculated by:

$$\text{RCB_CROSSED} = \text{ceiling}[(\text{Lower_Address mod RCB}) + \text{Length}) / \text{RCB}]$$

Lower_Address and Length are fields that can be parsed from the Completion header. Alternatively, a designer can load a register CUR_ADDR with Lower_Address at the start of each incoming Completion, increment per DW or QW as appropriate, then count an RCB whenever CUR_ADDR rolls over.

This method is less wasteful than PACKET_FC but still gives us an RCB granularity. If a User Application transmits I/O requests, the User Application could adopt a policy of only allocating one CplD credit for each I/O Read and zero CplD credits for each I/O Write. The User Application would have to match each incoming Completion's Tag with the Type (Memory Write, I/O Read, I/O Write) of the original NP Request.

The DATA_FC Method

The DATA_FC method provides the finest allocation granularity at the expense of logic.

As with PACKET_FC and RCB_FC, start with two registers, CPLH_PENDING and CPLD_PENDING (loaded with zero at reset).

1. When the User Application needs to send an NP request, determine the potential number of CplH and CplD credits it might require:

$$\text{NP_CplH} = \text{ceiling}[(\text{Start_Address mod RCB}) + \text{Request_Size}) / \text{RCB}]$$

$$\text{NP_CplD} = \text{ceiling}[(\text{Start_Address mod 16 bytes}) + \text{Request_Size}) / 16 \text{ bytes}] \\ (\text{except I/O Write, which returns zero data})$$

2. Check the following:

$$\text{CPLH_PENDING} + \text{NP_CplH} \leq \text{Total_CplH}$$

$$\text{CPLD_PENDING} + \text{NP_CplD} \leq \text{Total_CplD}$$

3. If both inequalities are true, transmit the Non-Posted Request, increase CPLH_PENDING by NP_CplH and CPLD_PENDING by NP_CplD.
4. At the start of each incoming Completion, or when that Completion begins at or crosses an RCB without ending at that RCB, decrement CPLH_PENDING by 1. The number of RCB crossings can be calculated by:

$$\text{RCB_CROSSED} = \text{ceiling}[(\text{Lower_Address mod RCB}) + \text{Length}) / \text{RCB}]$$

Lower_Address and Length are fields that can be parsed from the Completion header. Alternatively, a designer can load a register CUR_ADDR with Lower_Address at the start of each incoming Completion, increment per DW or QW as appropriate, then count an RCB whenever CUR_ADDR rolls over.

5. At the start of each incoming Completion, or when that Completion begins at or crosses at a naturally aligned credit boundary, decrement CPLD_PENDING by 1. The number of credit-boundary crossings is given by:

$$\text{DATA_CROSSED} = \text{ceiling}[(\text{Lower_Address mod } 16 \text{ B}) + \text{Length}) / 16 \text{ B}]$$

Alternatively, a designer can load a register CUR_ADDR with Lower_Address at the start of each incoming Completion, increment per DW or QW as appropriate, then count an RCB whenever CUR_ADDR rolls over each 16-byte address boundary.

This method is the least wasteful but requires the greatest amount of user logic. If even finer granularity is desired, the user can scale the Total_CplID value by 2 or 4 to get the number of Completion QWORDs or DWORDs, respectively, and adjust the data calculations accordingly.

Board Design Guidelines

Overview

This appendix discusses topics related to implementing a PCI Express® design that uses the Spartan®-6 FPGA on a printed circuit board (PCB). Optimal performance requires an understanding of the functionality of the device pins and needs to address issues such as device interfacing, protocol specifications, and signal integrity.

Recommendations made in this chapter are guidelines and do not guarantee a working design.

The information presented here discusses PCB considerations specific to the PCI Express specifications. This chapter should be used in conjunction with these documents for a comprehensive understanding of PCB design with Xilinx FPGAs.

- [UG386, Spartan-6 FPGA GTP Transceivers User Guide](#) - Specifically, see the “Board Design Guidelines” chapter.
- [UG393, Spartan-6 FPGA PCB Design Guide](#).

The PCI-SIG maintains multiple specifications that can apply depending on the form factor of the design. This document only considers the subset of these specifications focused on chip-to-chip and add-in card implementations. [Table F-1](#) shows the specifications that correlate to the applicable form factors.

Table F-1: PCI-SIG Specifications and Form Factor

Specification Name	Form-factor
PCI Express Base Specification Revision 1.1	Chip-to-chip on a single PCB
PCI Express Card Electromechanical Specification (CEM) Revision 1.1	ATX: desktop/server consisting of System card and Add-in card

Example PCB Reference

Xilinx delivers the SP605 board with an x1 PCI Express add-in card connection. This chapter uses this board as an example for certain recommendations.

For documentation such as schematics, gerbers, and a bill-of-material for the SP605 board, see the Spartan-6 FPGA SP605 Evaluation Kit product page:

[www.xilinx.com /sp605](http://www.xilinx.com/sp605)

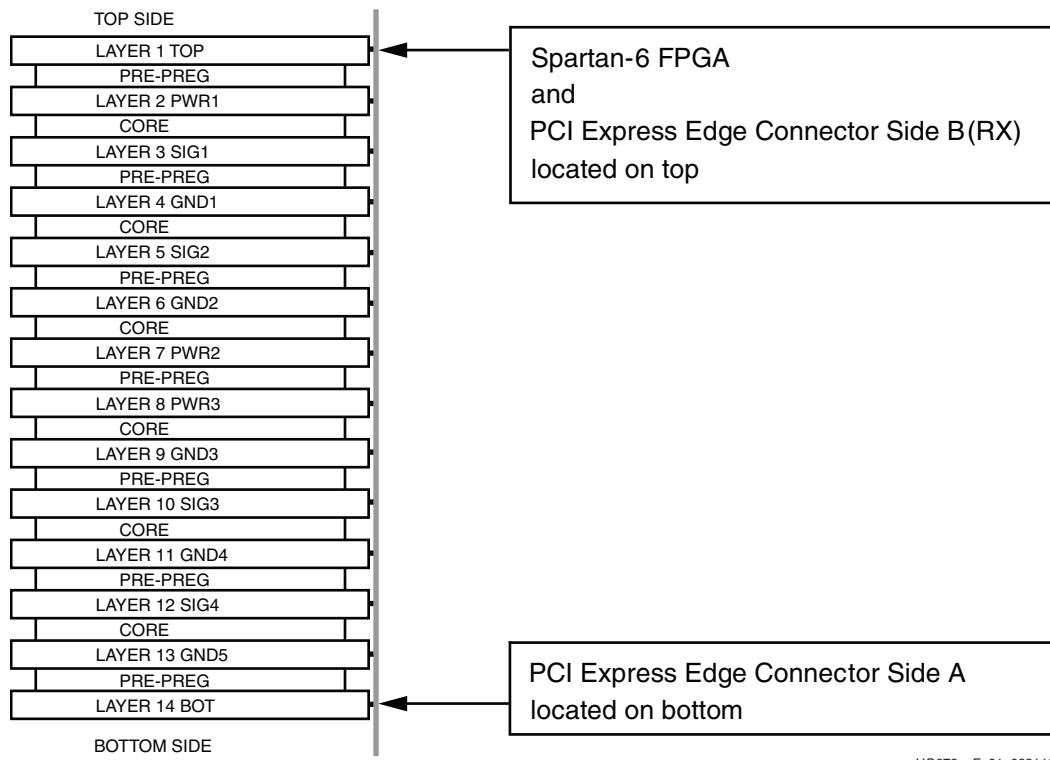
Board Stackup

Board stackup design is dependent on many variables, including design, manufacturing, and cost constraints. See the information on board stackup design in the *Spartan-6 FPGA PCB Design Guide* and the *Spartan-6 FPGA GTP Transceivers User Guide*.

Generally speaking, signal layers for high-speed signals such as PCI Express data signals should be sandwiched between ground planes. It is also preferable to use the layers closest to the top or bottom of the device so that via stubs are minimized.

SP605 Example

Figure F-1 shows the stackup that the SP605 Add-in Card reference board employs. All internal signal layers are sandwiched between (uninterrupted) ground and power planes.

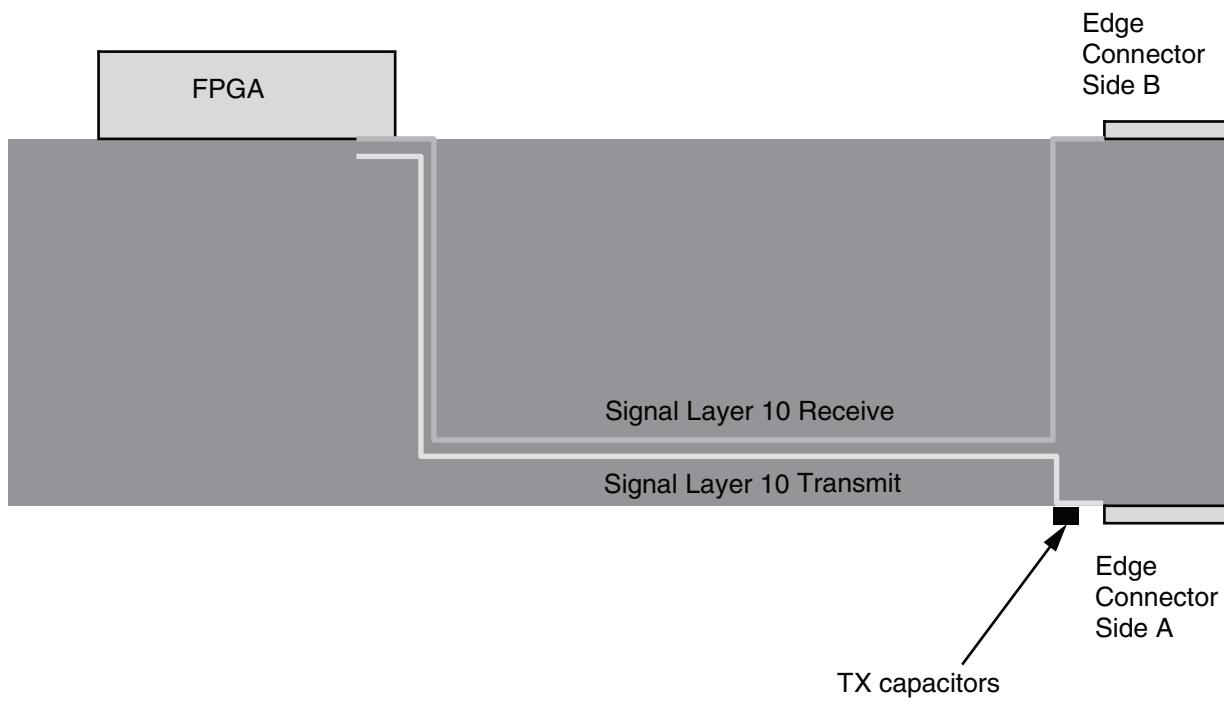


UG672_aF_01_083110

Figure F-1: SP605 Board Stackup

Transmit (TX) data lines initiate from the FPGA on the top layer, immediately drop to SIG3 (Layer 10) for routing across the PCB, and then terminate at the PCI Express edge connector side A on the bottom layer.

Receive (RX) data lines initiate from the FPGA on the top layer, immediately drop to SIG3 (Layer 10) for routing across the PCB, and then terminate at the PCI Express edge connector side B on the top layer.



UG672_aE_02_083110

Figure F-2: Transmit and Receive Data Lines

Power Supply Design

The *Spartan-6 FPGA PCB Design Guide* discusses general Power Distribution System (PDS) design for the FPGA, including the required decoupling capacitors for the VCCINT, VCCO, and VCCAUX supplies.

It is also imperative to ensure a clean power supply on MGTAVCC and MGTAVTT power supplies. Consult the *Spartan-6 FPGA GTP Transceivers User Guide* for more details on GTP transceiver power supply layout and other requirements for filtering and design.

Data Routing Guidelines

Breakout from FPGA BGA

The *Spartan-6 FPGA GTP Transceivers User Guide* discusses how to break out the high-speed GTP transceiver signals from the BGA and provides examples of such. Design constraints might require microstrips for the BGA exit path or from via to the PCI Express edge connector launch or SMT pads. In such cases, the microstrip trace must be kept as short as possible.

An example Receive and Transmit breakout pattern from the SP605 board are shown in [Figure F-3](#). Transmit lines are shown in green, and receive lines are shown in red.

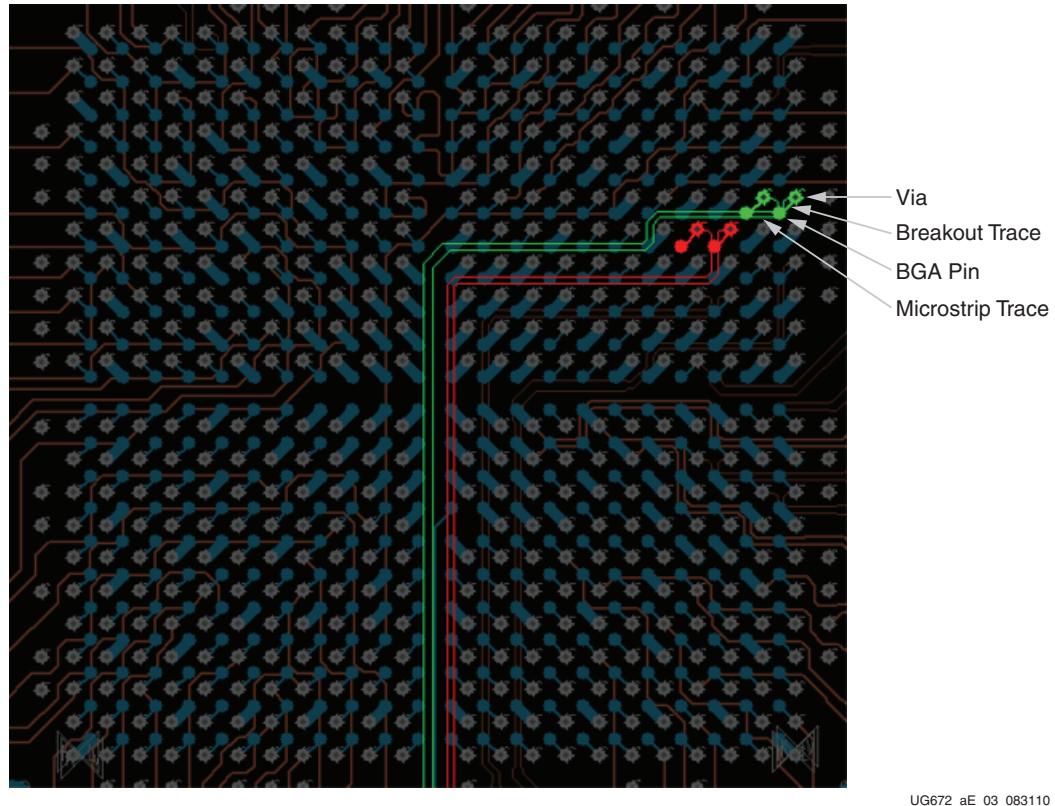


Figure F-3: Receive Breakout Pattern

Microstrip vs. stripline

Striplines are to be used whenever possible, as are the uppermost and lowermost stripline layers to minimize via stubs. When the stackup is being planned, these layers should be placed as close to the top and bottom layers whenever possible.

Plane Reference and Splits

Ground planes should be used as reference planes for signals, as opposed to noisier power planes. Each reference plane should be contiguous for the length of the trace, because routing over plane-splits creates an impedance discontinuity. In this case, the impedance of the trace changes because its coupling to the reference plane is changed abruptly at the plane split.

Bends

Follow the recommendations in the *Spartan-6 FPGA PCB Design Guide* regarding microstrip and stripline bends. Tight bends (such as 90 degrees) should be avoided; only mitered, 45-degree or less, bends are recommended.

Propagation Delay

PCI Express generally does not specify a maximum propagation delay for data signals, with the exception of add-in cards. Add-in card designs should meet the propagation delay specification in the CEM specification for data traces. The delay from the edge finger to the GTP transceiver must not exceed 750 ps.

Intrapair Skew

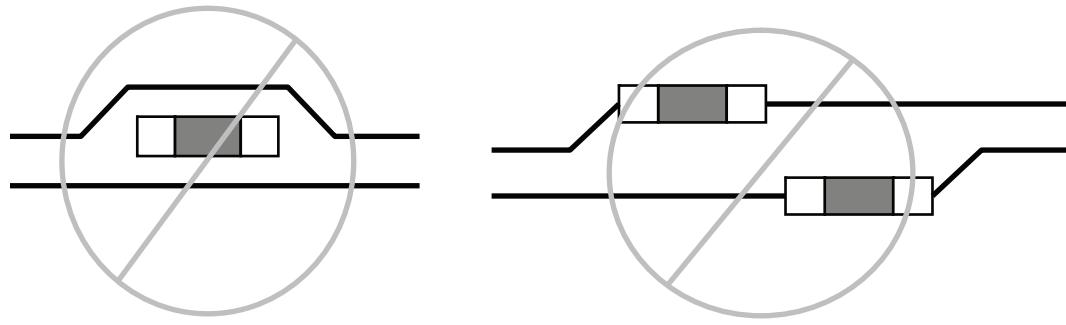
Intrapair skew refers to the skew between a P and N leg of a differential pair. Skew can introduce common-mode effects which lead to increased EMI, crosstalk and other DC effects. It is important to match the skew for differential pairs as close as possible.

Xilinx recommends intrapair trace length-matching to within 5 mils to minimize these effects.

Symmetrical Routing

Always use symmetrical routing. This prevents common-mode effects, such as EMI, from being introduced into the system.

[Figure F-4](#) illustrates two examples of non-symmetrical routing, which should be avoided.



UG672_aE_04_083110

Figure F-4: Non-Symmetrical Routing Examples

Vias

Users should follow the recommendations in the *Spartan-6 FPGA PCB Design Guide* for differential vias. Specifically, wherever high-speed signals must transition signal layers, a Ground-Signal-Signal-Ground (GSSG) type via should be used if possible. This provides a low inductance return current path.

All vias for a differential pair should employ symmetrical routing rules.

Trace Impedance

Differential data-line trace impedance was not specified in the Rev 1.0, 1.0a, or 1.1 (1.x) of the PCI Express Base and PCI Express CEM Specifications. The transmitters and receivers were specified to have 100Ω nominal differential impedance; therefore, most 1.x designs opt for a default 100Ω differential trace impedance for all PCI Express differential connections.

Xilinx recommends using simulation techniques to determine the optimum trace impedance. Simulation using HSPICE or Hyperlynx can help determine the optimum trace impedance to reduce signal loss.

PCB dielectric material, board stack up, microstrip, and strip line traces affect signal impedance. It is important that all of these factors are taken into consideration together.

If a simulator is not available, Xilinx recommends these basic guidelines for differential data-line trace impedance targets:

- $100\Omega \pm 10\%$ for 2.5 Gb/s only links

Trace Separation

Generally, simulation or post-layout analysis tools should be used to determine the optimum spacing required to reduce crosstalk from nearby aggressor signals. In the absence of these tools, Xilinx suggests that spacing between differential pairs and other non-PCI Express signals should be at least three times the dielectric height above the reference planes to minimize crosstalk. Exceptions to this are allowed in the break-out area of the FPGA; however, these sections should be kept as short as possible.

Lane Polarity Inversion

The *PCI Express Base Specification (1.x)* requires that all PCI Express receivers support polarity inversion. This gives the PCB designer flexibility to avoid having to cross P and N lines within a given differential pair.

GTP receivers support lane polarity inversion on a per transceiver basis.

AC Coupling

System and Add-in Cards

AC coupling capacitors should be placed on the TX pairs. Place the capacitors either near the edge connector or the FPGA—not in the middle of the interconnect.

Chip-to-Chip

AC coupling capacitors can be placed anywhere on the interconnect, except in the very middle.

General Guidelines

Capacitors for coupled traces should always be located at the same relative place as its partner, that is, symmetrical routing guidelines apply for differential pairs.

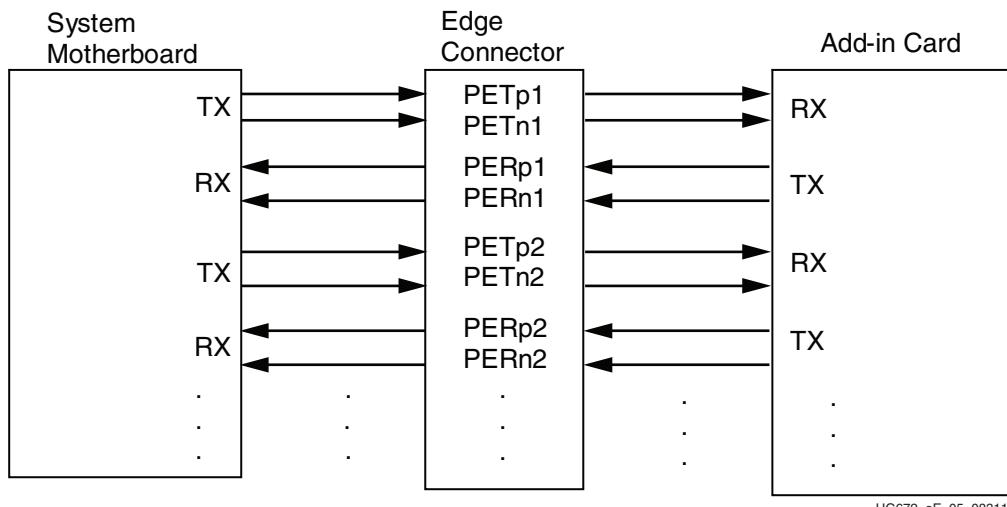
Use 0.1 μF ceramic chip capacitors in the smallest package possible.

Data Signal Termination

No external resistor terminators are required with the exception of a precision 50Ω resistor connected to the RCAL circuitry for the GTP transceiver column. Make sure the trace length and geometry to both legs of the resistor are equal. See the *Spartan-6 FPGA GTP Transceivers User Guide* for more information.

Additional Considerations for Add-In Card Designs

1. Board thickness for add-in cards should not exceed 0.062 inches.
2. Care must be taken when connecting the RX and TX data lines to the edge connector. The edge connector pin names for the TX and RX data lines as defined in the CEM specification are named from the view of the system board. That is, the RX (PERxx) lines are connected to the receiver on the system board and the transmitter on the add-in card. Similarly the TX (PETxx) lines are connected to the transmitter on the system board and the receiver on the add-in card. That means the add-in card should route the edge connector PERxx pins to the transmitter and the PETxx pins to the receiver on an Endpoint configured FPGA. [Figure F-5](#) illustrates how to connect the data lines for an add-in card design.



UG672_aE_05_083110

Figure F-5: Add-In Card Design Connections

Reference Clock Considerations

Jitter

Reference clock jitter has the potential to close both the TX and RX eyes, depending on the frequency content of the phase jitter. Therefore, it is very important to maintain as clean a reference clock as possible.

Reduce crosstalk on the REFCLK signal by isolating the clock signal from nearby high-speed traces. Maintain a separation of at least 25 mils from the nearest aggressor signals.

Ensure a clean power supply on MGTAVCC power supply. See the *Spartan-6 FPGA GTP Transceivers User Guide* for more details on GTP transceiver power supply layout and design.

In some cases where the designer has no control over the clock source, it might be desirable to add a jitter attenuator device.

If an external PLL or jitter attenuator device is used, ensure that it meets the specifications for PLL bandwidth as defined in the *PCI Express Base Specification*. The PLL bandwidth specification is different for 1.x and 2.0 versions of the specification.

Trace Impedance

The reference clock should use a 100Ω differential trace impedance.

Termination

The REFCLK signal should be routed to the dedicated reference clock input pins on the GTP transceiver, and the user design should instantiate an IBUFDS primitive in the user design. An internal 100Ω differential termination biased to $2/3$ MGTAVCC is automatically included on these input pins when the IBUFDS is used, and no external termination is required or needed for Spartan-6 devices. This is true for both HSCL and LVDS clocks.

See the *Spartan-6 FPGA GTP Transceivers User Guide* (UG386) for more information on GTP transceiver reference clock termination.

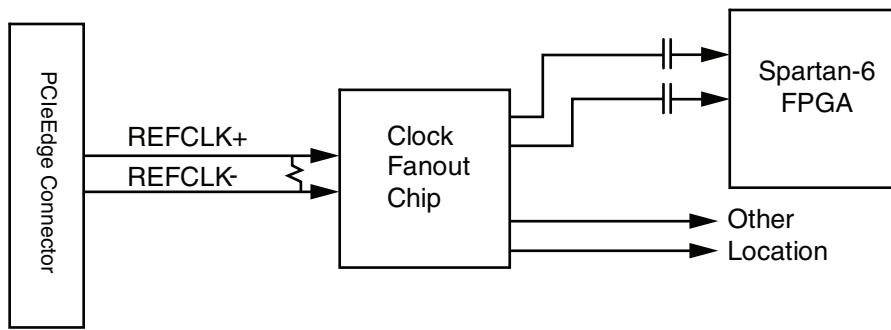
AC Coupling

The REFCLK signal should be AC coupled at the input to the FPGA. Xilinx recommends $0.1\ \mu\text{F}$ ceramic-chip capacitors for this purpose. See *Spartan-6 FPGA GTP Transceivers User Guide* for more information

Fanout

If the reference clock needs to be routed to more than one location, then a dedicated clock fanout device should be used. Make sure to follow the specifications for the fanout device. For instance, 100Ω termination might be required on the input to the fanout device.

Figure F-6 shows an example of a clock fanout device used to route the reference clock to multiple locations. The Spartan-6 FPGA requires no external resistive termination (just AC coupling capacitors). The fanout device is shown with a single resistor terminator at its clock input pins.



UG672_aE_06_083110

Figure F-6: Fanout Block Diagram

Sideband PCI Express Signals

PERST#

The PERST# signal must be routed to the FPGA for add-in cards. This 3.3V signal should be routed to an 3.3 V I/O bank (that is, V_{CCO} connected to 3.3V). If a non-3.3V I/O bank is used, an external circuit is necessary to interface with the Spartan-6 FPGA inputs. This external circuit could consist of a level translator such as the ST Micro ST2378E, a resistor

network, or other transistor-based circuit. There is no termination required for this signal, although the integrated Endpoint Block core implements a pull-up on the input from within the example UCF file.

PRSNT#

The PRSNT# pins should be connected as recommended in the CEM specification. Also see the SP605 board for an example

Summary Checklist

Table F-2 provides a checklist which summarizes the items discussed in this chapter.

Table F-2: Board Design Checklist

Item	
Board Stackup	
	Follow guidelines in UG393 and UG386.
Power Supply Design	
	Follow guidelines in UG393 and UG386.
High-Speed Data Signal Routing	
	Use stripline routing when possible.
	Avoid routing over reference plane splits or voids.
	Bends < 45 degrees.
	Add-in cards must not exceed 750 ps propagation delay.
	Length match intrapair skew to within 3 ps.
	Use Ground-Signal-Signal-Ground (GSSG) type vias when possible.
	Limit the number of vias.
	100Ω differential trace impedance for 2.5 Gb/s data signals.
	20 mil trace separation between differential pairs (exception in breakout area).
	AC coupling 0.1 μF ceramic chip capacitors on all TX lines.
	50Ω precision resistor connected to the RCAL circuit for GTP transceivers (see UG386).
	Add-in cards must not exceed 0.062 inches in thickness.
Reference Clock (REFCLK)	
	100Ω differential trace impedance.
	Maintain separation of at least 25 mils from nearby aggressor signals.
	Ensure clean power supply on MGTAVCC.
	No external termination required at input to FPGA (however, user must instantiate IBUFDS primitive).
	AC coupling 0.1 μF ceramic chip capacitors.

Table F-2: Board Design Checklist (*Cont'd*)

	Item
Sideband Signals for Add-In Cards	
	PERST# routes directly to 3.3V I/O bank. Use external circuitry if routing to non-3.3V I/O bank.
	PRSNT# connects as recommended in CEM specification.

PCIE_A1 Port Descriptions

This appendix describes the physical interfaces visible on the Spartan®-6 FPGA integrated Endpoint block's software primitive, PCIE_A1.

This appendix contains these sections:

- [Clock and Reset Interface](#)
- [Transaction Layer Interface](#)
- [Block RAM Interface](#)
- [GTP Transceiver Interface](#)
- [Configuration Management Interface](#)
- [Debug Interface Ports](#)

Clock and Reset Interface

[Table G-1](#) defines the ports in the Clock and Reset interface.

Table G-1: Clock and Reset Interface Port Descriptions

Port	Direction	Clock Domain	Description
CLOCKLOCKED	Input	USERCLK	LOCKED signal from the PLL.
MGTCLK	Input	MGTCLK	PIPE interface clock.
RECEIVEDHOTRST	Output	MGTCLK	Received hot reset. When asserted, this output indicates when an in-band hot reset has been received.
SYSRESETN	Input	NONE	Asynchronous system reset (active Low). When this input is asserted, the integrated Endpoint block is held in reset until PLL LOCK; thus it can be used to reset the integrated Endpoint block.
USERCLK	Input	USERCLK	User interface clock.
USERRSTN	Output	USERCLK	User interface reset (active Low). This output should be used to reset the user design logic (it is asserted when the integrated Endpoint block is reset).

Transaction Layer Interface

Packets are presented to and received from the integrated Endpoint block's Transaction Layer through the Transaction Layer interface. [Table G-2](#) defines the ports in the Transaction Layer interface.

Table G-2: Transaction Layer Interface Port Descriptions

Port	Direction	Clock Domain	Description
TRNFCCPLD[11:0]	Output	USERCLK	Completion Data Flow Control Credits. This output contains the number of Completion Data FC credits for the selected flow control type.
TRNFCCPLH[7:0]	Output	USERCLK	Completion Header Flow Control Credits. This output contains the number of Completion Header FC credits for the selected flow control type.
TRNFCNPD[11:0]	Output	USERCLK	Non-Posted Data Flow Control Credits. This output contains the number of Non-Posted Data FC credits for the selected flow control type.
TRNFCNPH[7:0]	Output	USERCLK	Non-Posted Header Flow Control Credits. This output contains the number of Non-Posted Header FC credits for the selected flow control type.
TRNFCPD[11:0]	Output	USERCLK	Posted Data Flow Control Credits. This output contains the number of Posted Data FC credits for the selected flow control type.
TRNFCPH[7:0]	Output	USERCLK	Posted Header Flow Control Credits. This output contains the number of Posted Header FC credits for the selected flow control type.
TRNFCSEL[2:0]	Input	USERCLK	Flow Control Informational Select. This input selects the type of flow control information presented on the TRNFC* signals. Valid values are: 000b: Receive buffer available space 001b: Receive credits granted to the link partner 010b: Receive credits consumed 100b: Transmit user credits available 101b: Transmit credit limit 110b: Transmit credits consumed
TRNLNKUPN	Output	USERCLK	Transaction Link Up (active Low). This output is asserted when the core and the connected upstream link partner port are ready and able to exchange data packets. It is deasserted when the core and link partner are attempting to establish communication, and when communication with the link partner is lost due to errors on the transmission channel. When the core is driven to the Hot Reset and Link Disable states by the link partner, TRNLNKUPN is deasserted and all TLPs stored in the core are lost.

Table G-2: Transaction Layer Interface Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
TRNRBARHITN[6:0]	Output	USERCLK	<p>Receive BAR Hit (active Low). This output indicates the BAR(s) targeted by the current receive transaction:</p> <ul style="list-style-type: none"> TRNRBARHITN[0]: BAR0 TRNRBARHITN[1]: BAR1 TRNRBARHITN[2]: BAR2 TRNRBARHITN[3]: BAR3 TRNRBARHITN[4]: BAR4 TRNRBARHITN[5]: BAR5 TRNRBARHITN[6]: Expansion ROM Address <p>If two BARs are configured into a single 64-bit address, both corresponding TRNRBARHITN bits are asserted.</p>
TRNRD[31:0]	Output	USERCLK	Receive Data. This bus contains the packet data being received.
TRNRDSTRDYN	Input	USERCLK	Receive Destination Ready (active Low). This input is asserted to indicate that the user application is ready to accept data on TRNRD. Simultaneous assertion of TRRSRCRDYN and TRNRDSTRDYN marks the successful transfer of data on TRNRD.
TRNREOFN	Output	USERCLK	Receive End-of-Frame (active Low). When asserted, this output indicates the end of a packet.
TRNRERRFWDN	Output	USERCLK	Receive Error Forward (active Low). This output marks the current packet in progress as error-poisoned. It is asserted by the integrated Endpoint block for the entire length of the packet.
TRNRNPOKN	Input	USERCLK	Receive Non-Posted OK (active Low). The user application asserts this input whenever it is ready to accept a Non-Posted Request packet. This allows Posted and Completion packets to bypass Non-Posted packets in the inbound queue if necessitated by the user application. When the user application approaches a state where it is unable to service Non-Posted Requests, it must deassert TRNRNPOKN one clock cycle before the integrated Endpoint block presents TRNREOFN of the last Non-Posted TLP the user application can accept.
TRNRSOFN	Output	USERCLK	Receive Start-of-Frame (active Low). When asserted, this output indicates the start of a packet.
TRRSRCDSN	Output	USERCLK	Receive Source Discontinue (active Low). When asserted, this output indicates that the integrated Endpoint block is aborting the current packet transfer. It is asserted when the physical link is going into reset.
TRRSRCRDYN	Output	USERCLK	Receive Source Ready (active Low). When asserted, this output indicates that the integrated Endpoint block is presenting valid data on TRNRD.
TRNTBUFAV[5:0]	Output	USERCLK	Transmit Buffers Available. This output provides the number of transmit buffers available in the integrated Endpoint block. The maximum number is 32. Each transmit buffer can accommodate one TLP up to the supported Maximum Payload Size.

Table G-2: Transaction Layer Interface Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
TRNTCFGGNTN	Input	USERCLK	Transmit Configuration Grant (active Low). The user application asserts this input in response to TRNTCFGREQN, to allow the integrated Endpoint block to transmit an internally generated TLP. If the user does not need to postpone internally generated TLPs, this signal can be continuously asserted.
TRNTCFGREQN	Output	USERCLK	Transmit Configuration Request (active Low). This output is asserted when the integrated Endpoint block is ready to transmit a Configuration Completion or other internally generated TLP.
TRNTD[31:0]	Input	USERCLK	Transmit Data. This bus contains the packet data to be transmitted.
TRNTDSTRDYN	Output	USERCLK	Transmit Destination Ready (active Low). When asserted, this output indicates that the integrated Endpoint block is ready to accept data on TRNTD. Simultaneous assertion of TRNTSRCRDYN and TRNTDSTRDYN marks a successful transfer of data on TRNTD.
TRNTEOFN	Input	USERCLK	Transmit End-of-Frame (active Low). This input signals the end of a packet.
TRNTERRDROPN	Output	USERCLK	Transmit Error Drop (active Low). When asserted, this output indicates that the integrated Endpoint block discarded a packet because of a length violation or, when streaming, data was not presented on consecutive clock cycles. Length violations only include packets longer than the supported maximum payload size and do not include packets whose payload does not match the payload advertised in the TLP header length field.
TRNTERRFWDN	Input	USERCLK	Transmit Error Forward (active Low). This input marks the current packet in progress as error-poisoned. If TRNTSTRN is deasserted, TRNTERRFWDN can be asserted any time between start of frame (SOF) and end of frame (EOF), inclusive. If TRNTSTRN is asserted, TRNTERRFWDN can only be asserted at SOF.
TRNTSOFN	Input	USERCLK	Transmit Start-of-Frame (active Low). When asserted, this input indicates the start of a packet.
TRNTSRCRDYN	Input	USERCLK	Transmit Source Ready (active Low). When asserted, this input indicates that the user application is presenting valid data on TRNTD.
TRNTSTRN	Input	USERCLK	Transmit Streamed (active Low). When asserted, this input indicates a packet is presented on consecutive clock cycles and transmission on the link can begin before the entire packet has been written to the integrated Endpoint block.

Block RAM Interface

The Transmit (TX) and Receive (RX) buffers are implemented with block RAM. [Table G-3](#) defines the TX buffer and RX buffer ports for the Block RAM interface.

Table G-3: Block RAM Interface Port Descriptions

Port	Direction	Clock Domain	Description
MIMRXRADDR[11:0]	Output	USERCLK	RX buffer read address
MIMRXRDATA[34:0]	Input	USERCLK	RX buffer read data
MIMRXREN	Output	USERCLK	RX buffer read enable
MIMRXWADDR[11:0]	Output	USERCLK	RX buffer write address
MIMRXWDATA[34:0]	Output	USERCLK	RX buffer write data
MIMRXWEN	Output	USERCLK	RX buffer write enable
MIMTXRADDR[11:0]	Output	USERCLK	TX buffer read address
MIMTXRDATA[35:0]	Input	USERCLK	TX buffer read data
MIMTXREN	Output	USERCLK	TX buffer read enable
MIMTXWADDR[11:0]	Output	USERCLK	TX buffer write address
MIMTXWDATA[35:0]	Output	USERCLK	TX buffer write data
MIMTXWEN	Output	USERCLK	TX buffer write enable

GTP Transceiver Interface

[Table G-4](#) defines the PIPE per Lane ports within the GTP Transceiver interface. There are two copies of the PIPE per Lane ports, one for each port ($n = A$ or B). Depending on which GTP transceiver is used, the LogiCORE IP core for PCI Express selects the correct port to use for the design.

Table G-4: PIPE per Lane Port Descriptions for the GTP Transceiver Interface

Port	Direction	Clock Domain	Description
PIPEGTRESETDONE n	Input	MGTCLK	When asserted, this input indicates that the GTP transceiver has finished reset and is ready for use.
PIPEPHYSTATUS n	Input	MGTCLK	PIPEPHYSTATUS n is asserted for a single cycle to indicate completion of GTP transceiver functions such as Power Management state transitions and receiver detection on lane n .
PIPERXCHARISK n [1:0]	Input	MGTCLK	This output defines the control bit(s) for received data: 0b: Data byte 1b: Control byte The lower bit corresponds to the lower byte of PIPERXDATA n [7:0] while the upper bit describes of PIPERXDATA n [15:8].
PIPERXDATA n [15:0]	Input	MGTCLK	This input contains the received data.
PIPERXENTERELECIDLE n	Input	MGTCLK	This input indicates an electrical idle on the Receiver.

Table G-4: PIPE per Lane Port Descriptions for the GTP Transceiver Interface (Cont'd)

Port	Direction	Clock Domain	Description
PIPERXPOLARITY n	Output	MGTCLK	When High, this output instructs the GTP transceiver to invert polarity (on the RX differential pair).
PIPERXRESET n	Output	MGTCLK	When asserted, this output resets the receive portion of the GTP transceiver.
PIPERXSTATUS $n[2:0]$	Input	MGTCLK	This input encodes the receiver status and error codes for the received data stream and receiver detection on lane n : 000b: Data received OK 001b: Reserved 010b: Reserved 011b: Receiver Detected 100b: 8B/10B decode error 101b: Elastic Buffer overflow 110b: Elastic Buffer underflow 111b: Receive disparity error
PIPETXCHARDISPMODE $n[1:0]$	Output	MGTCLK	PIPETXCHARDISPMODE and PIPETXCHARDISPVAL allow the 8B/10B disparity of outgoing data to be controlled when 8B/10B encoding is enabled. PIPETXCHARDISPMODE[1] corresponds to TXDATA[15:8] and PIPETXCHARDISPMODE[0] corresponds to PIPETXDATA[7:0]. For PCI Express operation, PIPETXCHARDISPMODE maps to the PIPE signal TXCOMPLIANCE given that PIPETXCHARDISPVAL is Low. When PIPETXCHARDISPMODE is High and PIPETXCHARDISPVAL is Low, the running disparity is set to negative. This functionality is used when transmitting the compliance pattern.
PIPETXCHARDISPVAL $n[1:0]$	Output	MGTCLK	PIPETXCHARDISPMODE and PIPETXCHARDISPVAL allow the 8B/10B disparity of outgoing data to be controlled when 8B/10B encoding is enabled. TXCHARDISPVAL[1] corresponds to TXDATA[15:8] and TXCHARDISPVAL[0] corresponds to TXDATA[7:0]. For PCI Express operation, PIPETXCHARDISPVAL should always be Low.
PIPETXCHARISK $n[1:0]$	Output	MGTCLK	This output determines the control bit(s) for received data: 0b: Data byte 1b: Control byte The lower bit corresponds to the lower byte of PIPETXDATA $n[7:0]$ while the upper bit describes PIPETXDATA $n[15:8]$.

Table G-4: PIPE per Lane Port Descriptions for the GTP Transceiver Interface (Cont'd)

Port	Direction	Clock Domain	Description
PIPETXDATA n [15:0]	Output	MGTCLK	This output contains the transmit data.
PIPETXELECIDLE n	Output	MGTCLK	This output forces the transmit output to electrical idle in all power states.
PIPETXPOWERDOWN n [1:0]	Output	MGTCLK	This output is the Power Management signal for the transmitter for lane n : 00b: P0 (Normal operation) 01b: P0s (Low recovery time power-saving state) 10b: P1 (Longer recovery time power state) 11b: Reserved
PIPETXRCVRDET n	Output	MGTCLK	When asserted, this output enables the GTP transceiver to begin either a receiver detection operation or loopback.

Configuration Management Interface

The Configuration Management Interface contains these signal groupings:

- [Management Interface Ports](#)
- [Error Reporting Ports](#)
- [Interrupt Generation and Status Ports](#)
- [Power Management Ports](#)
- [Configuration Specific Register Ports](#)
- [Miscellaneous Configuration Management Ports](#)

Management Interface Ports

[Table G-5](#) defines the Management Interface ports within the Configuration Management interface. These ports are used when reading and writing the Configuration Space Registers.

Table G-5: Management Interface Port Descriptions

Port	Direction	Clock Domain	Description
CFGDO[31:0]	Output	USERCLK	Management Data Out. This 32-bit data output obtains read data from the configuration space inside the integrated Endpoint block.
CFGDWADDR[9:0]	Input	USERCLK	Management DWORD Address. This 10-bit address input provides a configuration register DWORD address during configuration register accesses.
CFGRDENN	Input	USERCLK	Management Read Enable (active Low). This input is the read-enable for configuration register accesses.
CFGWRDWDONEN	Output	USERCLK	Management Read or Write Done (active Low). The read-write done signal indicates successful completion of the user configuration register access operation. For a user configuration register read operation, this signal validates the value of the CFGDO[31:0] data bus. The integrated Endpoint block does not support write operations.

Error Reporting Ports

[Table G-6](#) defines the Error Reporting ports within the Configuration Management interface.

Table G-6: Error Reporting Port Descriptions

Port	Direction	Clock Domain	Description
CFGERRCORN	Input	USERCLK	Configuration Error Correctable Error (active Low). The user asserts this signal to report a Correctable Error.
CFGERRCPLABORTN	Input	USERCLK	Configuration Error Completion Aborted (active Low). The user asserts this signal to report a completion was aborted. This signal is ignored if CFGERRCPLRDYN is deasserted.

Table G-6: Error Reporting Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
CFGERRCPLRDYN	Output	USERCLK	Configuration Error Completion Ready (active Low). When asserted, this signal indicates that the core can accept assertions on CFGERRURN and CFGERRCPLABORTN for Non-Posted Transactions. Assertions on CFGERRURN and CFGERRCPLABORTN are ignored when CFGERRCPLRDYN is deasserted.
CFGERRCPLTIMEOUTN	Input	USERCLK	Configuration Error Completion Time-out (active Low). The user asserts this signal to report a completion timed out.
CFGERRECRCN	Input	USERCLK	ECRC Error Report (active Low). The user asserts this signal to report an end-to-end CRC (ECRC) error.
CFGERRLOCKEDN	Input	USERCLK	Configuration Error Locked (active Low). This input is used to further qualify the CFGERRURN or CFGERRCPLABORTN input signal. When this input is asserted concurrently with one of those two signals, it indicates that the transaction that caused the error was an MRdLk transaction and not an MRd. The integrated Endpoint block generates a CplLk instead of a Cpl if the appropriate response is to send a Completion.
CFGERRPOSTEDN	Input	USERCLK	Configuration Error Posted (active Low). This input is used to further qualify any of the CFGERR* input signals. When this input is asserted concurrently with one of the other signals, it indicates that the transaction that caused the error was a posted transaction.
CFGERRTLPCLHEADER[47:0]	Input	USERCLK	Configuration Error TLP Completion Header. This 48-bit input accepts the header information from the user when an error is signaled. This information is required so that the integrated Endpoint block can issue a correct completion, if required. This information should be extracted from the received error TLP and presented in the listed format: [47:41] Lower Address [40:29] Byte Count [28:26] TC [25:24] Attr [23:8] Requester ID [7:0] Tag
CFGERRURN	Input	USERCLK	Configuration Error Unsupported Request (active Low). The user asserts this signal to report that an Unsupported Request (UR) was received. This signal is ignored if CFGERRCPLRDYN is deasserted.

Interrupt Generation and Status Ports

Table G-7 defines the Interrupt Generation and Status ports within the Configuration Management interface.

Table G-7: Interrupt Generation and Status Port Descriptions

Port	Direction	Clock Domain	Description
CFGINTERRUPTASSERTN	Input	USERCLK	Configuration Legacy Interrupt Assert/Deassert Select. This input selects between Assert and Deassert messages for Legacy interrupts when CFGINTERRUPTN is asserted. It is not used for MSI interrupts. Value Message Type: 0b: Assert 1b: Deassert
CFGINTERRUPTDI[7:0]	Input	USERCLK	Configuration Interrupt Data In. For Message Signaling Interrupts (MSI), this input provides the portion of the Message Data that the Endpoint must drive to indicate MSI vector number, if Multi-Vector Interrupts are enabled. The value indicated by CFGINTERRUPTMMENABLE[2:0] determines the number of lower-order bits of Message Data that the Endpoint provides; the remaining upper bits of CFGINTERRUPTDI[7:0] are not used. For Single-Vector Interrupts, CFGINTERRUPTDI[7:0] is not used. For Legacy Interrupt Messages (ASSERTINTX, DEASSERTINTX), this input indicates which message type is sent, where Value Legacy Interrupt is: 00h: INTA 01h: INTB 02h: INTC 03h: INTD
CFGINTERRUPTDO[7:0]	Output	USERCLK	Configuration Interrupt Data Out. This output is the value of the lowest eight bits of the Message Data field in the Endpoint's MSI capability structure. This value is used in conjunction with CFGINTERRUPTMMENABLE[2:0] to drive CFGINTERRUPTDI[7:0].
CFGINTERRUPTMMENABLE[2:0]	Output	USERCLK	Configuration Interrupt Multiple Message Enabled. This output has the value of the Multiple Message Enable field, where values range from 000b to 101b. A value of 000b indicates that single vector MSI is enabled. Other values indicate the number of bits that can be used for multi-vector MSI.

Table G-7: Interrupt Generation and Status Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
CFGINTERRUPTMSIENABLE	Output	USERCLK	Configuration Interrupt MSI Enabled. 0: Only Legacy (INTx) interrupts can be sent 1: The Message Signaling Interrupt (MSI) messaging is enabled
CFGINTERRUPTN	Input	USERCLK	Configuration Interrupt Request (active Low). When asserted, this input causes the selected interrupt message type to be transmitted by the integrated Endpoint block. The signal should be asserted until CFGINTERRUPTRDYN is asserted.
CFGINTERRUPTRDYN	Output	USERCLK	Configuration Interrupt Ready (active Low). This output is the interrupt grant signal. The simultaneous assertion of CFGINTERRUPTRDYN and CFGINTERRUPTN indicates that the integrated Endpoint block has successfully transmitted the requested interrupt message.

Power Management Ports

Table G-8 defines the Power Management ports within the Configuration Management interface.

Table G-8: Power Management Port Descriptions

Port	Direction	Clock Domain	Description
CFGPMWAKEN	Input	USERCLK	Send PMPME Message (active Low). A one-clock cycle assertion of this input signals the integrated Endpoint block to send a Power Management Wake Event (PMPME) Message TLP to the upstream link partner.
CFGTOTURNOFFN	Output	USERCLK	Configuration To Turnoff: This output signal notifies the user that a PME_TURN_Off message has been received, and the Configuration and Capabilities Module (CCM) starts polling the CFGTURNOFFOKN input coming in from the user. When CFGTURNOFFOKN is asserted, the CMM sends a PME_To_Ack message to the upstream device.
CFGTURNOFFOKN	Input	USERCLK	Configuration Turnoff OK (active Low). This input is the power turn-off ready signal. The user application can assert this input to notify the Endpoint that it is safe for power to be turned off.

Configuration Specific Register Ports

Table G-9 defines the Configuration Specific Register ports within the Configuration Management interface. These ports directly mirror the contents of commonly used registers located within the PCI Express Configuration Space.

Table G-9: Configuration Specific Register Port Descriptions

Port	Direction	Clock Domain	Description
CFGCOMMANDBUSMASTERENABLE	Output	USERCLK	Configuration Command, Bus Master Enable, Command[2]. The integrated Endpoint block takes no action based on this setting; the user logic must. When this output is asserted, the user logic is allowed to issue Memory or I/O Requests (including MSI interrupts); otherwise, the user logic must not issue those requests.
CFGCOMMANDINTERRUPTDISABLE	Output	USERCLK	Configuration Command, Interrupt Disable, Command[10]. When this output is asserted, the integrated Endpoint block is prevented from asserting INTx interrupts.
CFGCOMMANDIOENABLE	Output	USERCLK	Configuration Command, I/O Space Enable, Command[0]. 0: The integrated Endpoint block filters these accesses and responds with a UR. 1: Allows the device to receive I/O Space accesses.
CFGCOMMANDMEMENABLE	Output	USERCLK	Configuration Command, Memory Space Enable, Command[1]. 0: The integrated Endpoint block filters these accesses and responds with a UR. 1: Allows the device to receive Memory Space accesses.
CFGCOMMANDSERREN	Output	USERCLK	Configuration Command, SERR Enable (active Low), Command[8]. When this output is asserted, reporting of Non-fatal and Fatal errors is enabled. If enabled, errors are reported either through this bit or through the PCI Express specific bits in the Device Control Register.
CFGDEVCONTROLAUXPOWEREN	Output	USERCLK	Not used.
CFGDEVCONTOLCORRERRREPORTINGEN	Output	USERCLK	Configuration Device Control, Correctable Error Reporting Enable, DEVICECTRL[0]. This bit, in conjunction with other bits, controls sending ERRCOR messages.

Table G-9: Configuration Specific Register Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
CFGDEVCONTROLENABLER0	Output	USERCLK	Configuration Device Control, Enable Relaxed Ordering, DEVICECTRL[4]. When this output is asserted, the user logic is permitted to set the Relaxed Ordering bit in the Attributes field of transactions it initiates that do not require strong write ordering.
CFGDEVCONTROLEXTTAGEN	Output	USERCLK	Configuration Device Control, Tag Field Enable, DEVICECTRL[8]. When this output is asserted, the user logic can use an 8-bit Tag field as a Requester. When this output is deasserted, the user logic is restricted to a 5-bit Tag field. The integrated Endpoint block does not enforce the number of Tag bits used, either in outgoing request TLPs or incoming Completions.
CFGDEVCONTROLFATALERRREPORTINGEN	Output	USERCLK	Configuration Device Control, Fatal Error Reporting Enable, DEVICECTRL[2]. This bit, in conjunction with other bits, controls sending ERRFATAL messages.
CFGDEVCONTROLMAXPAYLOAD[2:0]	Output	USERCLK	Configuration Device Control, MAXPAYLOADSIZE, DEVICECTRL[7:5]. This field sets the maximum TLP payload size. As a Receiver, the user logic must handle TLPs as large as the set value. As a Transmitter, the user logic must not generate TLPs exceeding the set value. 000b: 128-byte maximum payload size 001b: 256-byte maximum payload size 010b: 512-byte maximum payload size
CFGDEVCONTROLMAXREADREQ[2:0]	Output	USERCLK	Configuration Device Control, MAXREADREQUESTSIZE, DEVICECTRL[14:12]. This field sets the maximum Read Request size for the user logic as a Requester. The user logic must not generate Read Requests with size exceeding the set value. 000b: 128-byte maximum Read Request size 001b: 256-byte maximum Read Request size 010b: 512-byte maximum Read Request size

Table G-9: Configuration Specific Register Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
CFGDEVCONTROLNONFATALREPORTINGEN	Output	USERCLK	Configuration Device Control, Non-Fatal Error Reporting Enable, DEVICECTRL[1]. This bit, in conjunction with other bits, controls sending ERRNONFATAL messages.
CFGDEVCONTROLNOSNOOPEN	Output	USERCLK	Configuration Device Control, Enable No Snoop, DEVICECTRL[11]. When this output is asserted, the user logic is permitted to set the No Snoop bit in TLPs it initiates that do not require hardware-enforced cache coherency.
CFGDEVCONTROLPHANTOMEN	Output	USERCLK	Configuration Device Control, Phantom Functions Enable, DEVICECTRL[9]. When this output is asserted, the user logic can use unclaimed Functions as Phantom Functions to extend the number of outstanding transaction identifiers. If this output is deasserted, the user logic is not allowed to use Phantom Functions.
CFGDEVCONTROLURERRREPORTINGEN	Output	USERCLK	Configuration Device Control, UR Reporting Enable, DEVICECTRL[3]. This bit, in conjunction with other bits, controls the signaling of URs by sending Error messages.
CFGDEVSTATUSCORRERRDETECTED	Output	USERCLK	Configuration Device Status, Correctable Error Detected, DEVICESTATUS[0]. This output indicates the status of correctable errors detected. Errors are logged in this register regardless of whether error reporting is enabled or not in the Device Control Register.
CFGDEVSTATUSFATALERRDETECTED	Output	USERCLK	Configuration Device Status, Fatal Error Detected, DEVICESTATUS[2]. This output indicates the status of Fatal errors detected. Errors are logged in this register regardless of whether error reporting is enabled or not in the Device Control Register.
CFGDEVSTATUSNONFATALERRDETECTED	Output	USERCLK	Configuration Device Status, Non-Fatal Error Detected, DEVICESTATUS[1]. This output indicates the status of Non-fatal errors detected. Errors are logged in this register regardless of whether error reporting is enabled or not in the Device Control Register.

Table G-9: Configuration Specific Register Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
CFGDEVSTATUSURDETECTED	Output	USERCLK	Configuration Device Status, Unsupported Request Detected, DEVICESTATUS[3]. This output indicates that the integrated Endpoint block received a UR. Errors are logged in this register regardless of whether error reporting is enabled or not in the Device Control Register.
CFGLINKCONTROLASPMCONTROL[1:0]	Output	USERCLK	Configuration Link Control, ASPM Control, LINKCTRL[1:0]. This 2-bit output indicates the level of ASPM supported, where: 00b: Disabled 01b: L0s Entry Enabled 10b: Not used 11b: Not used
CFGLINKCONTROLCOMMONCLOCK	Output	USERCLK	Configuration Link Control, Common Clock Configuration, LINKCTRL[6]. When this output is asserted, this component and the component at the opposite end of this Link are operating with a distributed common reference clock. When this output is deasserted, the components are operating with an asynchronous reference clock.
CFGLINKCONTROLEXTENDEDSYNC	Output	USERCLK	Configuration Link Control, Extended Sync, LINKCTRL[7]. When this output is asserted, the transmission of additional Ordered Sets is forced when exiting the L0s state and when in the Recovery state.
CFGLINKCONTROLRCB	Output	USERCLK	Configuration Link Control, RCB, LINKCTRL[3]. This output indicates the Read Completion Boundary value, where: 0: 64B 1: 128B
CFGTRNPENDINGN	Input	USERCLK	User Transaction Pending (active Low). When asserted, this input sets the Transactions Pending bit in the Device Status Register (DEVICESTATUS[5]). Note: The user is required to assert this input if the User Application has not received a completion to a request.

Miscellaneous Configuration Management Ports

Table G-10 defines the Miscellaneous ports within the Configuration Management interface.

Table G-10: Miscellaneous Configuration Management Port Descriptions

Port	Direction	Clock Domain	Description
CFGBUSNUMBER[7:0]	Output	USERCLK	Configuration Bus Number. This 8-bit output provides the assigned bus number for the device. The user application must use this information in the Bus Number field of outgoing TLP requests. The default value after reset is 00h. This output is refreshed whenever a Type 0 Configuration Write packet is received.
CFGDEVICENUMBER[4:0]	Output	USERCLK	Configuration Device Number: This 5-bit output provides the assigned device number for the device. The user application must use this information in the Device Number field of outgoing TLP requests. The default value after reset is 00000b. This output is refreshed whenever a Type 0 Configuration Write packet is received.
CFGDEVID[15:0]	Input	USERCLK	Device ID value. This 16-bit input must be stable when SYSRESETN is deasserted.
CFGDSN[63:0]	Input	USERCLK	Configuration Device Serial Number. This 64-bit input indicates the value that should be transferred to the Device Serial Number Capability.
CFGFUNCTIONNUMBER[2:0]	Output	USERCLK	Configuration Function Number. This 3-bit output provides the function number for the device. The user application must use this information in the Function Number field of outgoing TLP requests. The function number is hardwired to 000b.

Table G-10: Miscellaneous Configuration Management Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
CFGLTSSMSTATE[4:0]	Output	MGTCLK	<p>This 5-bit output is a mirror of the LTSSM state machine bits:</p> <ul style="list-style-type: none"> 00000b: Detect.Quiet 00001b: Detect.Active 00010b: Polling.Active 00011b: Polling.Config 00100b: Polling.Compliance 00101b: Configuration.Linkwidth.Start 00110b: Configuration.Linkwidth.Start 00111b: Configuration.Linkwidth.Accept 01000b: Configuration.Linkwidth.Accept 01001b: Configuration.Lanenum.Wait 01010b: Configuration.Lanenum.Accept 01011b: Configuration.Complete 01100b: Configuration.Idle 01101b: L0 01110b: L1.Entry 01111b: L1.Entry 10000b: L1.Entry 10001b: L1.Idle 10010b: L1.Exit-to-recovery 10011b: Recovery.RcvrLock 10100b: Recovery.RcvrCfg 10101b: Recovery.Idle 10110b: Hot Reset 10111b: Disabled 11000b: Disabled 11001b: Disabled 11010b: Disabled 11011b: Detect.Quiet
CFGPCIELINKSTATEN[2:0]	Output	USERCLK	<p>PCI Express Link State. This encoded bus reports the PCI Express Link State Information to the user:</p> <ul style="list-style-type: none"> 110b: L0 state 101b: L0s state 011b: L1 state 111b: Under transition
CFGREVID[7:0]	Input	USERCLK	Revision ID Value. This input must be stable when SYSRESETN is deasserted.
CFGSUBSYSID[15:0]	Input	USERCLK	Subsystem ID Value. This input must be stable when SYSRESETN is deasserted.

Table G-10: Miscellaneous Configuration Management Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
CFGSUBSYSVENID[15:0]	Input	USERCLK	Subsystem Vendor ID Reset Value. This input must be stable when SYSRESETN is deasserted.
CFGVENID[15:0]	Input	USERCLK	Vendor ID Value. This input must be stable when SYSRESETN is deasserted.

Debug Interface Ports

Table G-11 describes the Debug Interface ports.

Table G-11: Debug Interface Port Descriptions

Port	Direction	Clock Domain	Description
DBGBADDLLPSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle when a DLLP CRC error is detected.
DBGBADTLPPLCRC	Output	USERCLK	This signal pulses High for one USERCLK cycle when a TLP with an LCRC error is detected.
DBGBADTLPSEQNUM	Output	USERCLK	This signal pulses High for one USERCLK cycle when a TLP with an invalid sequence number is detected.
DBGBADTLPSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle when a bad TLP is detected, for reasons other than a bad LCRC or a bad sequence number.
DBGDLPROTOCOLSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle if an out-of-range ACK or NAK is received.
DBGFCPROTOCOLERRSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle if there is a protocol error with the received flow control updates.
DBGMLFRMDLENGTH	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had a length that did not match what was in the TLP header.
DBGMLFRMDMPS	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had a length in violation of the negotiated MPS.
DBGMLFRMDTCVC	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had an invalid TC or VC value.
DBGMLFRMDTLPSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle when a malformed TLP is received. See the other DBGMLFRMD* signals for further clarification. Note: There is skew between DBGMLFRMD* and DBGMLFRMDTLPSTATUS.
DBGMLFRMDUNRECTYPE	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received TLP had an invalid/unrecognized type field value.

Table G-11: Debug Interface Port Descriptions (Cont'd)

Port	Direction	Clock Domain	Description
DBGPOISTLPSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle if a poisoned TLP is received.
DBGRCVROVERFLOWSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle if a received TLP violates the advertised credit.
DBGREGDETECTEDCORRECTABLE	Output	USERCLK	This signal is a mirror of the internal signal used to indicate a correctable error is detected. The error is cleared upon a read by the Root Complex (RC).
DBGREGDETECTEDFATAL	Output	USERCLK	This signal is a mirror of the internal signal used to indicate that a fatal error has been detected. The error is cleared upon a read by the RC.
DBGREGDETECTEDNONFATAL	Output	USERCLK	This signal is a mirror of the internal signal used to indicate that a non-fatal error has been detected. The error is cleared upon a read by the RC.
DBGREGDETECTEDUNSUPPORTED	Output	USERCLK	This signal is a mirror of the internal signal used to indicate that an unsupported request has been detected. The error is cleared upon a read by the RC.
DBGRPLYROLLOVERSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle when the rollover counter expires.
DBGRPLYTIMEOUTSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle when the replay time-out counter expires.
DBGURNOBARHIT	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a received read or write request did not match any configured BAR.
DBGURPOISCFGWR	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that a CfgWr TLP with the Error/Poisoned bit (EP) = 1 was received.
DBGURSTATUS	Output	USERCLK	This signal pulses High for one USERCLK cycle when an unsupported request is received. See the DBGUR* signals for further clarification. Note: There is skew between DBGUR* and DBGURSTATUS.
DBGURUNSUPMSG	Output	USERCLK	This signal pulses High for one USERCLK cycle to indicate that an Msg or MsgD TLP with an unsupported type was received.

PCIE_A1 Attribute Descriptions

Table H-1 defines the attributes on the PCIE_A1 library primitive for the Spartan®-6 FPGA Integrated Endpoint Block for PCI Express® designs. All attributes are set in the LogiCORE™ IP; they are documented in this chapter for reference. Users should not change the attribute settings as set in the CORE Generator™ software GUI for proper operation of the design.

Table H-1: PCIE_A1 Attributes

Attribute Name	Type	Description
BAR0	32-bit Hex	<p>This attribute specifies the mask/settings for Base Address Register (BAR) 0. If BAR is not to be implemented, this attribute is set to 32'h00000000. Bits are defined as follows:</p> <ul style="list-style-type: none"> • Memory Space BAR: <ul style="list-style-type: none"> 0: Mem Space Indicator (set to 0) [2:1]: Type field (10 for 64-bit, 00 for 32-bit) 3: Prefetchable (0 or 1) • I/O Space BAR: <ul style="list-style-type: none"> 0: I/O Space Indicator (set to 1) 1: Reserved (set to 0) [31:2]: Mask for writable bits of BAR. The uppermost 31:n bits are set to 1, where 2^n = memory aperture size in bytes. For a 64-bit BAR, the uppermost 63:n bits of {BAR1, BAR0} are set to 1.
BAR1	32-bit Hex	<p>This attribute specifies the mask/settings for BAR1 if BAR0 is a 32-bit BAR, or the upper bits of {BAR1, BAR0} if BAR0 is a 64-bit BAR. If BAR is not to be implemented, this attribute is set to 32'h00000000. See the BAR0 description if this attribute functions as the upper bits of a 64-bit BAR. Bits are defined as follows:</p> <ul style="list-style-type: none"> • Memory Space BAR (not the upper bits of BAR0): <ul style="list-style-type: none"> 0: Mem Space Indicator (set to 0) [2:1]: Type field (10 for 64-bit, 00 for 32-bit) 3: Prefetchable (0 or 1) • I/O Space BAR: <ul style="list-style-type: none"> 0: I/O Space Indicator (set to 1) 1: Reserved (set to 0) [31:2]: Mask for writable bits of BAR. The uppermost 31:n bits are set to 1, where 2^n = I/O aperture size in bytes

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
BAR2	32-bit Hex	<p>For an Endpoint, this attribute specifies the mask/settings for BAR2 if BAR1 is a 32-bit BAR, or the upper bits of {BAR2, BAR1} if BAR1 is the lower part of a 64-bit BAR. If BAR is not to be implemented, this attribute is set to 32'h00000000. See the BAR1 description if this attribute functions as the upper bits of a 64-bit BAR.</p> <p>For an Endpoint, bits are defined as follows:</p> <ul style="list-style-type: none"> Memory Space BAR (not upper bits of BAR1): <ul style="list-style-type: none"> 0: Mem Space Indicator (set to 0) [2:1]: Type field (10 for 64-bit, 00 for 32-bit) 3: Prefetchable (0 or 1) [31:4]: Mask for writable bits of BAR. For a 32-bit BAR, the uppermost 31:n bits are set to 1, where 2^n = memory aperture size in bytes. For a 64-bit BAR, the uppermost 63:n bits of {BAR3, BAR2} are set to 1. I/O Space BAR: <ul style="list-style-type: none"> 0: I/O Space Indicator (set to 1) 1: Reserved (set to 0) [31:2]: Mask for writable bits of BAR. The uppermost 31:n bits are set to 1, where 2^n = I/O aperture size in bytes
BAR3	32-bit Hex	<p>For an Endpoint, this attribute specifies the mask/settings for BAR3 if BAR2 is a 32-bit BAR, or the upper bits of {BAR3, BAR2} if BAR2 is the lower part of a 64-bit BAR. If BAR is not to be implemented, this attribute is set to 32'h00000000. See the BAR2 description if this functions as the upper bits of a 64-bit BAR.</p> <p>For an Endpoint, bits are defined as follows:</p> <ul style="list-style-type: none"> Memory Space BAR (not upper bits of BAR2): <ul style="list-style-type: none"> 0: Mem Space Indicator (set to 0) [2:1]: Type field (10 for 64-bit, 00 for 32-bit) 3: Prefetchable (0 or 1) [31:4]: Mask for writable bits of BAR. For a 32-bit BAR, the uppermost 31:n bits are set to 1, where 2^n = memory aperture size in bytes. For a 64-bit BAR, the uppermost 63:n bits of {BAR4, BAR3} are set to 1. I/O Space BAR: <ul style="list-style-type: none"> 0: I/O Space Indicator (set to 1) 1: Reserved (set to 0) [31:2]: Mask for writable bits of BAR. The uppermost 31:n bits are set to 1, where 2^n = I/O aperture size in bytes

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
BAR4	32-bit Hex	<p>For an Endpoint, this attribute specifies mask/settings for Base Address Register (BAR) 4 if BAR3 is a 32-bit BAR, or the upper bits of {BAR4, BAR3}, if BAR3 is the lower part of a 64-bit BAR. If BAR is not to be implemented, this attribute is set to 32'h00000000. See the BAR3 description if this functions as the upper bits of a 64-bit BAR.</p> <p>For an Endpoint, bits are defined as follows:</p> <ul style="list-style-type: none"> Memory Space BAR (not upper bits of BAR3): <ul style="list-style-type: none"> 0: Mem Space Indicator (set to 0) [2:1]: Type field (10 for 64-bit, 00 for 32-bit) 3: Prefetchable (0 or 1) [31:4]: Mask for writable bits of BAR. For a 32-bit BAR, the uppermost 31:n bits are set to 1, where 2^n = memory aperture size in bytes. For a 64-bit BAR, the uppermost 63:n bits of {BAR5, BAR4} to 1. I/O Space BAR: <ul style="list-style-type: none"> 0: I/O Space Indicator (set to 1) 1: Reserved (set to 0) [31:2]: Mask for writable bits of BAR. The uppermost 31:n bits are set to 1, where 2^n = I/O aperture size in bytes
BAR5	32-bit Hex	<p>For an Endpoint, this attribute specifies mask/settings for BAR5 if BAR4 is a 32-bit BAR or the upper bits of {BAR5, BAR4} if BAR4 is the lower part of a 64-bit BAR. If BAR is not to be implemented, this attribute is set to 32'h00000000. See the BAR4 description if this functions as the upper bits of a 64-bit BAR.</p> <p>For an Endpoint, bits are defined as follows:</p> <ul style="list-style-type: none"> Memory Space BAR (not upper bits of BAR4): <ul style="list-style-type: none"> 0: Mem Space Indicator (set to 0) [2:1]: Type field (00 for 32-bit; BAR5 cannot be the lower part of a 64-bit BAR) 3: Prefetchable (0 or 1) [31:4]: Mask for writable bits of BAR. The uppermost 31:n bits are set to 1, where 2^n = memory aperture size in bytes I/O Space BAR: <ul style="list-style-type: none"> 0: I/O Space Indicator (set to 1) 1: Reserved (set to 0) [31:2]: Mask for writable bits of BAR. The uppermost 31:n bits are set to 1, where 2^n = I/O aperture size in bytes
CARDBUS_CIS_POINTER	32-bit Hex	Pointer to the Cardbus data structure. This value is transferred to the Cardbus CIS Pointer Register. It is set to 0 if the Cardbus pointer is not implemented.

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
CLASS_CODE	24-bit Hex	Code identifying basic function, subclass, and applicable programming interface. This value is transferred to the Class Code Register.
DEV_CAP_ENDPOINT_L0S_LATENCY	3-bit Binary	Endpoint L0s Acceptable Latency. This attribute records the latency that the Endpoint can withstand on transitions from the L0s state to the L0 state. Valid settings are: 000b: Less than 64 ns 001b: 64 ns to 128 ns 010b: 128 ns to 256 ns 011b: 256 ns to 512 ns 100b: 512 ns to 1 μ s 101b: 1 μ s to 2 μ s 110b: 2 μ s to 4 μ s 111b: More than 4 μ s
DEV_CAP_ENDPOINT_L1_LATENCY	3-bit Binary	Endpoint L1 Acceptable Latency. Records the latency that the endpoint can withstand on transitions from the L1 state to the L0 state (if the L1 state is supported). Valid settings are: 000b: Less than 1 μ s 001b: 1 μ s to 2 μ s 010b: 2 μ s to 4 μ s 011b: 4 μ s to 8 μ s 100b: 8 μ s to 16 μ s 101b: 16 μ s to 32 μ s 110b: 32 μ s to 64 μ s 111b: More than 64 μ s
DEV_CAP_EXT_TAG_SUPPORTED	Boolean	Extended Tags support. FALSE: 5-bit tag TRUE: 8-bit tag
DEV_CAP_MAX_PAYLOAD_SUPPORTED	3-bit Binary	This attribute specifies the maximum payload supported. Valid (supported) settings are: 000b: 128 bytes 001b: 256 bytes 010b: 512 bytes This value is transferred to the Device Capabilities Register.

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
DEV_CAP_PHANTOM_FUNCTIONS_SUPPORT	2-bit Binary	Phantom Function Support. This attribute indicates the number of functions re-allocated as Tag bits. Valid settings are: 00b: 0 01b: 1 10b: 2 11b: 3
DEV_CAP_ROLE_BASED_ERROR	Boolean	When this attribute is set to TRUE, compliant error reporting is supported.
DISABLE_BAR_FILTERING	Boolean	When this attribute is set to TRUE, BAR filtering is disabled. This setting does not change the behavior of the BAR hit outputs.
DISABLE_ID_CHECK	Boolean	When this attribute is set to TRUE, checking for Requester ID of received completions is disabled.
DISABLE_SCRAMBLING	Boolean	When this attribute is TRUE, Scrambling of transmit data is turned off.
ENABLE_RX_TD_ECRC_TRIM	Boolean	When this attribute is set to TRUE, received TLPs have their td bit set to 0 and the ECRC is removed.
EXPANSION_ROM	22-bit Hex	This attribute specifies the mask/settings for the Expansion ROM BAR. If the BAR is not to be implemented, this attribute is set to 22'h00000000. Bits are defined as follows: 0: Expansion ROM implemented (set to 1 to implement ROM) [21:1]: Mask for writable bits of BAR. The uppermost 21:n bits are set to 1, where 2^n = ROM aperture size in bytes
FAST_TRAIN	Boolean	When this attribute is set to TRUE, the timers in the LTSSM state machine are shortened to reduce simulation time. Specifically, the transition out of Polling.Active requires sending 16 TS1s and receiving 8 TS1s. The LTSSM timer values of 1 ms, 2 ms, 12 ms, 24 ms, and 48 ms are reduced to 3.9 µs, 7.81 µs, 46.8 µs, 93.75 µs, and 187.5 µs, respectively (reduced by a factor of 256). This attribute must be set to FALSE for silicon designs.
GTP_SEL	Boolean	This attribute indicates which port interface is used: FALSE: Transceiver A port interface TRUE: Transceiver B port interface

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
LINK_CAP_ASPM_SUPPORT	2-bit Binary	<p>Active State PM Support. This attribute indicates the level of active state power management supported by the selected PCI Express Link:</p> <ul style="list-style-type: none"> 00b: Reserved 01b: L0s entry supported 10b: Reserved 11b: Reserved
LINK_CAP_L0S_EXIT_LATENCY	3-bit Binary	<p>This attribute sets the exit latency from the L0s state to be applied (at 2.5 Gb/s) where a common clock is used. This value is transferred to the Link Capabilities Register.</p> <p>Valid settings are:</p> <ul style="list-style-type: none"> 000b: Less than 64 ns 001b: 64 ns to less than 128 ns 010b: 128 ns to less than 256 ns 011b: 256 ns to less than 512 ns 100b: 512 ns to less than 1 μs 101b: 1 μs to less than 2 μs 110b: 2 μs to 4 μs 111b: More than 4 μs
LINK_CAP_L1_EXIT_LATENCY	3-bit Binary	<p>This attribute sets the exit latency from the L1 state to be applied (at 2.5 Gb/s) where a common clock is used. This value is transferred to the Link Capabilities Register.</p> <p>Valid settings are:</p> <ul style="list-style-type: none"> 000b: Less than 1 μs 001b: 1 μs to less than 2 μs 010b: 2 μs to less than 4 μs 011b: 4 μs to less than 8 μs 100b: 8 μs to less than 16 μs 101b: 16 μs to less than 32 μs 110b: 32 μs to 64 μs 111b: More than 64 μs
LINK_STATUS_SLOT_CLOCK_CONFIG	Boolean	Slot Clock Configuration. This attribute indicates where the component uses the same physical reference clock that the platform provides on the connector. For a port that connects to the slot, this attribute indicates that it uses a clock with a common source to that used by the slot. For an adaptor inserted in the slot, this attribute indicates that it uses the same clock source as the slot, not a locally derived clock source. This value is transferred to the Link Status Register, bit 12.

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
LL_ACK_TIMEOUT	15-bit Hex	This attribute sets an ACK time-out counter override value. The value is in increments of USERCLK periods. It should be set to 0 unless the user wishes to override the default (internal) setting.
LL_ACK_TIMEOUT_EN	Boolean	When set to TRUE, the value specified by LL_ACK_TIMEOUT is added to the internal value, increasing the ACK Timeout delay. When set to FALSE, the value provided on LL_ACK_TIMEOUT is subtracted from the internal value, decreasing the ACK Timeout delay
LL_REPLAY_TIMEOUT	15-bit Hex	This attribute sets a replay timer override value. The value is in increments of USERCLK periods. It should be set to 0 unless the user wishes to override the default (internal) setting.
LL_REPLAY_TIMEOUT_EN	Boolean	When set to TRUE, the value specified by LL_REPLAY_TIMEOUT is added to the internal value, increasing the Replay Timeout delay. When set to FALSE, the value provided on LL_REPLAY_TIMEOUT is subtracted from the internal value, decreasing the Replay Timeout delay
MSI_CAP_MULTIMSG_EXTENSION	Boolean	Multiple Message Capable Extension. When set to TRUE, this attribute allows 256 unique messages to be sent by the user regardless of the setting of MSI_CAP_MULTIMSGCAP). Note: Enabling this feature (TRUE) violates the PCI Express Base Specification and should only be used in closed systems.
MSI_CAP_MULTIMSGCAP	3-bit Binary	Multiple Message Capable. Each MSI function can request up to 32 unique messages. System software can read this field to determine the number of messages requested. The number of messages requested are encoded as follows: 000b: 1 vector 001b: 2 vectors 010b: 4 vectors 011b: 8 vectors 100b: 16 vectors 101b: 32 vectors 110b - 111b: Reserved
PCIE_CAP_CAPABILITY_VERSION	4-bit Hex	This attribute indicates the version number of the PCI-SIG defined PCI Express capability structure. It must be set to 0001b.

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
PCIE_CAP_DEVICE_PORT_TYPE	4-bit Hex	<p>This attribute identifies the type of device/port. Valid settings are (all other values are unsupported):</p> <ul style="list-style-type: none"> 0000b: PCI Express Endpoint device 0001b: Legacy PCI Express Endpoint device <p>This value is transferred to the PCI Express Capabilities Register.</p>
PCIE_CAP_INT_MSG_NUM	5-bit Hex	Interrupt Message Number. This value is transferred to the PCI Express Cap Register [13:9]. It is not used internally by the integrated Endpoint block.
PCIE_CAP_SLOT_IMPLEMENTED	Boolean	This attribute must be set to FALSE.
PCIE_GENERIC	12-bit Hex	<p>The 12 bits are assigned as follows:</p> <ul style="list-style-type: none"> 11: This bit must be 0. 10: <ul style="list-style-type: none"> • 0: Electrical idle is not received until an Electrical Idle Ordered Set (EIOS) is received, if no EIOS core enters the LTSSM RECOVERY state • 1: An electrical idle can occur without an EIOS (the EIOS is assumed). This is the default and recommended setting. [9:7]: These bits drive the Interrupt Pin Register in the PCI Configuration Space. A value of 0 indicates no Legacy interrupts are implemented. Values of 1, 2, 3, and 4 indicate INTA, INTB, INTC, and INTD, respectively. Other values are not permitted. 6: <ul style="list-style-type: none"> • 0: The DSN Extended Capability is not implemented • 1: The DSN Extended Capability is implemented 5: <ul style="list-style-type: none"> • 0: 8B/10B Not_in_table is not inferred • 1: 8B/10B Not_in_table from the GTP transceiver is inferred from RXSTATUS. This is the default and recommended setting. 4: <ul style="list-style-type: none"> • 0: A read to an unimplemented config space returns completion with data of zero. This is the default and recommended setting. • 1: A read to an unimplemented config space returns a UR (legacy behavior of PIPE) [3:0]: These bits drive nFTS[7:4]. The lower bits of nFTS are set to Fh. The default value is 0xF.
PLM_AUTO_CONFIG	Boolean	This attribute must be set to FALSE.

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
PM_CAP_AUXCURRENT	3-bit Binary	AUX Current. Requested auxiliary current allocation. This value is transferred to the PM Capabilities Register, bits [24:22]. The integrated Endpoint block does not support AUX power, so this field should be set to 000b.
PM_CAP_D1SUPPORT	Boolean	D1 Support. This value is transferred to the PM Capabilities Register, bit 25.
PM_CAP_D2SUPPORT	Boolean	D2 Support. This value is transferred to the PM Capabilities Register, bit 26.
PM_CAP_DSI	Boolean	Device Specific Initialization (DSI). This value is transferred to the PM Capabilities Register, bit 21.
PM_CAP_PME_CLOCK	Boolean	When this attribute is set to TRUE, a PCI™ clock is required for PME generation. This attribute must be set to FALSE per the specification. The value is transferred to the PM Capabilities Register, bit 19.
PM_CAP_PMESUPPORT	5-bit Hex	PME Support. These five bits indicate support for D3cold, D3hot, D2, D1, and D0, respectively. This value is transferred to the PM Capabilities Register, bits [31:27].
PM_CAP_VERSION	3-bit Binary	The version of Power Management specification followed. This value is transferred to the PM Capabilities Register, bits [18:16]. This attribute must be set to 3.
PM_DATA_SCALE0	2-bit Hex	Power Management Data Scale Register 0. This attribute specifies the scale applied to PM_DATA0. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are: 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x
PM_DATA_SCALE1	2-bit Hex	Power Management Data Scale Register 1. This attribute specifies the scale applied to PM_DATA1. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are: 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
PM_DATA_SCALE2	2-bit Hex	<p>Power Management Data Scale Register 2. This attribute specifies the scale applied to PM_DATA2. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are:</p> <ul style="list-style-type: none"> 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x
PM_DATA_SCALE3	2-bit Hex	<p>Power Management Data Scale Register 3. This attribute specifies the scale applied to PM_DATA3. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are:</p> <ul style="list-style-type: none"> 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x
PM_DATA_SCALE4	2-bit Hex	<p>Power Management Data Scale Register 4. This attribute specifies the scale applied to PM_DATA4. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are:</p> <ul style="list-style-type: none"> 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x
PM_DATA_SCALE5	2-bit Hex	<p>Power Management Data Scale Register 5. This attribute specifies the scale applied to PM_DATA5. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are:</p> <ul style="list-style-type: none"> 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
PM_DATA_SCALE6	2-bit Hex	<p>Power Management Data Scale Register 6. This attribute specifies the scale applied to PM_DATA6. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are:</p> <ul style="list-style-type: none"> 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x
PM_DATA_SCALE7	2-bit Hex	<p>Power Management Data Scale Register 7. This attribute specifies the scale applied to PM_DATA7. The power consumption of the device is determined by multiplying the contents of the Base Power Data Register field with the value corresponding to the encoding returned by this field. Defined encodings are:</p> <ul style="list-style-type: none"> 00b: 1.0x 01b: 0.1x 10b: 0.01x 11b: 0.001x
PM_DATA0	8-bit Hex	<p>Power Management Data Register 0 (D0 Power Consumed). This value appears in the Data field of the PM Status Register if the host has written the value 0000b to the Data Select field of the PM Control Register.</p>
PM_DATA1	8-bit Hex	<p>Power Management Data Register 1 (D1 Power Consumed). This value appears in the Data field of the PM Status Register if the host has written the value 0001b to the Data Select field of the PM Control Register.</p>
PM_DATA2	8-bit Hex	<p>Power Management Data Register 2 (D2 Power Consumed). This value appears in the Data field of the PM Status Register if the host has written the value 0010b to the Data Select field of the PM Control Register.</p>
PM_DATA3	8-bit Hex	<p>Power Management Data Register 3 (D3 Power Consumed). This value appears in the Data field of the PM Status Register if the host has written the value 0011b to the Data Select field of the PM Control Register.</p>
PM_DATA4	8-bit Hex	<p>Power Management Data Register 4 (D0 Power Dissipated). This value appears in the Data field of the PM Status Register if the host has written the value 0100b to the Data Select field of the PM Control Register.</p>

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
PM_DATA5	8-bit Hex	Power Management Data Register 5 (D1 Power Dissipated). This value appears in the Data field of the PM Status Register if the host has written the value 0101b to the Data Select field of the PM Control Register.
PM_DATA6	8-bit Hex	Power Management Data Register 6 (D2 Power Dissipated). This value appears in the Data field of the PM Status Register if the host has written the value 0110b to the Data Select field of the PM Control Register.
PM_DATA7	8-bit Hex	Power Management Data Register 7 (D3 Power Dissipated). This value appears in the Data field of the PM Status Register if the host has written the value 0111b to the Data Select field of the PM Control Register.
SLOT_CAP_ATT_BUTTON_PRESENT	Boolean	Attention Button Present. When this attribute is TRUE, an Attention Button is implemented on the chassis for this slot. This value is transferred to the Slot Capabilities Register. This attribute must be set to FALSE for Endpoints.
SLOT_CAP_ATT_INDICATOR_PRESENT	Boolean	Attention Indicator Present. When this attribute is TRUE, an Attention Indicator is implemented on the chassis for this slot. This value is transferred to the Slot Capabilities Register. This attribute must be set to FALSE for Endpoints.
SLOT_CAP_POWER_INDICATOR_PRESENT	Boolean	Power Indicator Present. When this attribute is TRUE, a Power Indicator is implemented on the chassis for this slot. This value is transferred to the Slot Capabilities Register. This attribute must be set to FALSE for Endpoints.
TL_RX_RAM_RADDR_LATENCY	Boolean	This attribute specifies the read address latency for RX RAMs in terms of USER_CLK cycles. FALSE: No fabric pipeline register is on the read address and enable block RAM inputs TRUE: A fabric pipeline register is on the read address and enable block RAM inputs
TL_RX_RAM_RDATA_LATENCY	2-bit Binary	This attribute specifies the read data latency for RX RAMs in terms of USER_CLK cycles. 01b: The block RAM output register is disabled 10b: The block RAM output register is enabled 11b: The block RAM output register is enabled and a fabric pipeline register is added to the block RAM data output

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
TL_RX_RAM_WRITE_LATENCY	Boolean	<p>This attribute specifies the write latency for RX RAMs in terms of cycles of USER_CLK.</p> <p>FALSE: No fabric pipeline register is on the write address and enable block RAM inputs</p> <p>TRUE: A fabric pipeline register is on the write address and enable block RAM inputs</p>
TL_TFC_DISABLE	Boolean	When this attribute is set to TRUE, checking of flow control values and transmit packets in the order they were presented on the TRN TX interface is disabled.
TL_TX_CHECKS_DISABLE	Boolean	When this attribute is set to TRUE, all TLM checks of incoming data are disabled.
TL_TX_RAM_RADDR_LATENCY	Boolean	<p>This attribute specifies the read address latency for TX RAMs in terms of USER_CLK cycles.</p> <p>FALSE: No fabric pipeline register on the read address and enable block RAM inputs</p> <p>TRUE: A fabric pipeline register is on the read address and enable block RAM inputs</p>
TL_TX_RAM_RDATA_LATENCY	2-bit Binary	<p>This attribute specifies the read data latency for TX RAMs in terms of USER_CLK cycles.</p> <p>01b: The block RAM output register is disabled</p> <p>01b: The block RAM output register is enabled</p> <p>11b: The block RAM output register is enabled and a fabric pipeline register is added to the block RAM data output</p>
USR_CFG	Boolean	When this attribute is set to TRUE, the user application is permitted to add or implement PCI Legacy capability registers beyond address BFh. This option should be selected when the user application implements such a legacy capability configuration space, starting at C0h.
USR_EXT_CFG	Boolean	When this attribute is set to TRUE, the user application is permitted to add or implement PCI Express extended capability registers beyond address 1FFh. This box should be checked when the user application implements such an extended capability configuration space starting at 200h.
VC0_CPL_INFINITE	Boolean	<p>When this attribute is set to TRUE, the block advertises infinite completions.</p> <p>Note: For Endpoints, this attribute must be set to TRUE for compliance.</p>
VC0_RX_RAM_LIMIT	12-bit Hex	This attribute must be set to RX buffer bytes/4.

Table H-1: PCIE_A1 Attributes (Cont'd)

Attribute Name	Type	Description
VC0_TOTAL_CREDITS_CD	11-bit Hex	<p>Number of credits that should be advertised for Completion data received on Virtual Channel 0. The bytes advertised must be less than or equal to the block RAM bytes available.</p> <p>The equation to calculate bytes advertised is:</p> $(ph * (rx_td_ecrc_trim ? 16 : 20)) + (pd * 16) + (nph * 20) + (ch * 16) + (cd * 16)$ <p>The equation to calculate block RAM bytes available is:</p> $(vc0_rx_ram_limit + 1) * 4$ <p>See Table H-2, page 234 for valid settings.</p>
VC0_TOTAL_CREDITS_CH	7-bit Hex	<p>Number of credits that should be advertised for Completion headers received on Virtual Channel 0. The sum of the Posted, Non-Posted, and Completion header credits must be ≤ 80.</p> <p>See Table H-2, page 234 for valid settings.</p>
VC0_TOTAL_CREDITS_NPH	7-bit Hex	<p>Number of credits that should be advertised for Non-Posted headers received on Virtual Channel 0. The number of Non-Posted data credits advertised by the block is equal to the number of Non-Posted header credits. The sum of the Posted, Non-Posted, and Completion header credits must be ≤ 80.</p> <p>This attribute must be set to 8.</p>
VC0_TOTAL_CREDITS_PD	11-bit Hex	<p>Number of credits that should be advertised for Posted data received on Virtual Channel 0. The bytes advertised must be less than or equal to the block RAM bytes available.</p> <p>The equation to calculate bytes advertised is:</p> $(ph * (rx_td_ecrc_trim ? 16 : 20)) + (pd * 16) + (nph * 20) + (ch * 16) + (cd * 16)$ <p>The equation to calculate block RAM bytes available is:</p> $(vc0_rx_ram_limit + 1) * 4$ <p>See Table H-2, page 234 for valid settings.</p>
VC0_TOTAL_CREDITS_PH	7-bit Hex	<p>Number of credits that should be advertised for Posted headers received on Virtual Channel 0. The sum of the Posted, Non-Posted, and Completion header credits must be ≤ 80.</p>
VC0_TX_LASTPACKET	5-bit Hex	<p>Index of the last packet buffer used by TX TLM (that is, the number of buffers - 1). This value is calculated from the maximum payload size supported and the number of block RAMs configured for transmit.</p> <p>The equation is:</p> $((TX\ buffer\ bytes) / (MPS_in_bytes + 20) - 1)$ <p>See Table H-2 for valid settings.</p>

Table H-2: Valid Data Credit Combinations

Parameter Name	Valid Combinations					
DEV_CAP_MAX_PAYLOAD_SUPPORTED	0	0	1	1	2	2
VC0_TOTAL_CREDITS_PD	36	92	92	204	204	716
VC0_TOTAL_CREDITS_CD	64	128	128	256	256	256
VC0_TOTAL_CREDITS_CH	8	16	16	32	32	32
VC0_TX_LAST_PACKET	12	26	13	28	14	29

PCIE_A1 Timing Parameter Descriptions

This appendix lists the timing parameter names and descriptions related to the Spartan®-6 FPGA Integrated Endpoint Block for PCI Express® designs. This information is useful for debugging timing issues. Values for these timing parameters can be obtained by running the Speedprint tool. Usage of Speedprint is documented in the *Development System Reference Guide*.

The timing parameters on the integrated Endpoint block consist of either Setup/Hold or Clock-to-Out parameters. [Table I-1](#) lists the timing parameter names, descriptions, signal grouping, and related clock domain for a given parameter. In the table, parameter Tpcicck_XXX is a setup time (before the clock edge), and parameter Tpcickc_XXX is a hold time (after clock edge).

Table I-1: PCIE_A1 Timing Parameters

Name	Clock Domain	Signal Grouping
Sequential Setup and Hold Times for Integrated Endpoint Block Inputs		
Tpcicck_CFG / Tpcickc_CFG	USERCLK	CFGDEVID[15:0] CFGDSN[63:0] CFGDWADDR[9:0] CFGERRCORN CFGERRCPLTIMEOUTN CFGERRECRCN CFGERRLOCKEDN CFGERRPOSTEDN CFGERRTLPCPLHEADER[47:0] CFGERRURN CFGINTERRUPTASSERTN CFGINTERRUPTDI[7:0] CFGINTERRUPTN CFGRDENN CFGREVID[7:0] CFGSUBSYSID[15:0] CFGSUBSYSENID[15:0] CFGTRNPENDINGN CFGVENID[15:0] TRNTCFGGNTN
Tpcicck_ERR / Tpcickc_ERR	USERCLK	CFGERRCPLABORTN

Table I-1: PCIE_A1 Timing Parameters (Cont'd)

Name	Clock Domain	Signal Grouping
Tpcicck_MGT / Tpcickc_MGT	MGTCLK	PIPEGTRESETDONEA
		PIPEGTRESETDONEB
		PIPEPHYSTATUSA
		PIPEPHYSTATUSB
		PIPERXCHARISKA[1:0]
		PIPERXCHARISKB[1:0]
		PIPERXDATAA[15:0]
		PIPERXDATAB[15:0]
		PIPERXENTERELECIDLEA
		PIPERXENTERELECIDLEB
		PIPERXSTATUSUSA[2:0]
		PIPERXSTATUSUSB[2:0]
Tpcicck_PWR / Tpcickc_PWR	USERCLK	CFGPMWAKEN
		CFGTURNOFFOKN
Tpcicck_SCAN / Tpcickc_SCAN	USERCLK	SCANEN
		SCANIN[4:0]
		SCANRESETMASK
Tpcidck_LOCKED / Tpcickd_LOCKED	USERCLK	CLOCKLOCKED
Tpcidck_RESET / Tpcickd_RESET	USERCLK	SYSRESETN
Tpcidck_RXRAM / Tpcickd_RXRAM	USERCLK	MIMRXRDATA[34:0]
Tpcidck_TRNFC / Tpcickd_TRNFC	USERCLK	TRNFCSEL[2:0]
Tpcidck_TRNRD / Tpcickd_TRNRD	USERCLK	TRNRDSTRDYN
Tpcidck_TRNRM / Tpcickd_TRNRM	USERCLK	TRNRMPOKN
Tpcidck_TRNTD / Tpcickd_TRNTD	USERCLK	TRNTD[31:0]
Tpcidck_TRNTE / Tpcickd_TRNTE	USERCLK	TRNTEOFN
		TRNTERRFWDN

Table I-1: PCIE_A1 Timing Parameters (Cont'd)

Name	Clock Domain	Signal Grouping
Tpcidck_TRNTS / Tpcickd_TRNTS	USERCLK	TRNTSOFN
		TRNTSRCDSNCN
		TRNTSRCRDYN
		TRNTSTRN
Tpcidck_TXRAM / Tpcickd_TXRAM	USERCLK	MIMTXRDATA[35:0]
Sequential Clock to Output Times for Integrated Endpoint Block Outputs		
Tpcicko_CFG	USERCLK	CFGCOMMANDINTERRUPTDISABLE
		CFGDEVCONTROLMAXPAYLOAD[2:0]
		CFGDEVCONTROLMAXREADREQ[2:0]
Tpcicko_CFGBUS	USERCLK	CFGBUSNUMBER[7:0]
Tpcicko_CFGCOMMAND	USERCLK	CFGCOMMANDSERREN
Tpcicko_CFGDEV	USERCLK	CFGDEVCONTROLCORRERRREPORTINGEN
		CFGDEVCONTROLEXTTAGEN
		CFGDEVCONTROLFATALERRREPORTINGEN
		CFGDEVCONTROLNONFATALREPORTINGEN
		CFGDEVCONTROLNOSNOOPEN
		CFGDEVCONTROLPHANTOMEN
		CFGDEVCONTROLURERRREPORTINGEN
		CFGDEVICENUMBER[4:0]
		CFGDEVSTATUSCORRERRDETECTED
		CFGDEVSTATUSFATALERRDETECTED
		CFGDEVSTATUSNONFATALERRDETECTED
		CFGDEVSTATUSURDETECTED
Tpcicko_CFGDO	USERCLK	CFGDO[31:0]
Tpcicko_CFGDONE	USERCLK	CFGRDWRDONEN
Tpcicko_CFGERR	USERCLK	CFGERRCPLRDYN
Tpcicko_CFGFCN	USERCLK	CFGFUNCTIONNUMBER[2:0]
Tpcicko_CFGINT	USERCLK	CFGINTERRUPTDO[7:0]
		CFGINTERRUPTRDYN

Table I-1: PCIE_A1 Timing Parameters (Cont'd)

Name	Clock Domain	Signal Grouping
Tpcicko_CFGLINK	USERCLK	CFGLINKCONTOLRCB
		CFGLINKCONTROLASPMCONTROL[1:0]
		CFGLINKCONTROLCOMMONCLOCK
		CFGLINKCONTROLEXTENDEDSYNC
Tpcicko_CFGOFF	USERCLK	CFGTOTURNOFFN
Tpcicko_CFGSTATE	MGTCLK	CFGLTSSMSTATE[4:0]
	USERCLK	CFGPCIELINKSTATEN[2:0]
Tpcicko_DBG	USERCLK	DBGBADTLPLCRC
		DBGBADTLPSEQNUM
		DBGMLFRMDLENGTH
		DBGMLFRMDMPS
		DBGMLFRMDTCVC
		DBGMLFRMDUNRECTYPE
		DBGREGDETECTEDCORRECTABLE
		DBGREGDETECTEDFATAL
		DBGREGDETECTEDNONFATAL
		DBGREGDETECTEDUNSUPPORTED
		DBGURNOBARHIT
		DBGURPOISCFGWR
		DBGURUNSUPMSG
Tpcicko_ENA	USERCLK	CFGCOMMANDBUSMASTERENABLE
		CFGCOMMANDIOENABLE
		CFGCOMMANDMEMENABLE
		CFGDEVCONTROLENABLERO
		CFGINTERRUPTMMENABLE[2:0]
Tpcicko_MSG	USERCLK	CFGINTERRUPTMSIENABLE

Table I-1: PCIE_A1 Timing Parameters (Cont'd)

Name	Clock Domain	Signal Grouping
Tpcicko_PIPE	MGTCLK	PIPEGTTXELECIDLEA
		PIPEGTTXELECIDLEB
		PIPERXPOLARITYA
		PIPERXPOLARITYB
		PIPERXRESETA
		PIPERXRESETB
		PIPETXCHARDISPMODEA[1:0]
		PIPETXCHARDISPMODEB[1:0]
		PIPETXCHARDISPVLA[1:0]
		PIPETXCHARDISPVLB[1:0]
		PIPETXCHARISKA[1:0]
		PIPETXCHARISKB[1:0]
		PIPETXDATAA[15:0]
		PIPETXDATAB[15:0]
		PIPETXRCVRDETA
		PIPETXRCVRDETB
Tpcicko_PWR	MGTCLK	PIPEGTPOWERDOWNA[1:0]
		PIPEGTPOWERDOWNB[1:0]
	USERCLK	CFGDEVCONTROLAUXPOWEREN
Tpcicko_RXRAM	USERCLK	MIMRXRADDR[11:0]
		MIMRXREN
		MIMRXWADDR[11:0]
		MIMRXWDATA[34:0]
		MIMRXWEN
Tpcicko_SCANOUT	USERCLK	SCANOUT[4:0]

Table I-1: PCIE_A1 Timing Parameters (Cont'd)

Name	Clock Domain	Signal Grouping
Tpcicko_STATUS	USERCLK	DBGBADDLLPSTATUS DBGBADTLPSTATUS DBGDLPROTTOCOLSTATUS DBGFCPROTOCOLERRSTATUS DBGMLFRMDTLPSTATUS DBGPOISTLPSTATUS DBGRCVROVERFLOWSTATUS DBGRPLYROLLOVERSTATUS DBGRPLYTIMEOUTSTATUS DBGURSTATUS
Tpcicko_TRN	USERCLK	TRNFCCPLD[11:0] TRNFCCPLH[7:0] TRNFCNPD[11:0] TRNFCNPH[7:0] TRNFCPD[11:0] TRNFCPH[7:0] TRNLNKUPN TRNRBARHITN[6:0] TRNRD[31:0] TRNREOFN TRNRERRFWDN TRNRSOFN TRNRSRCDSN TRNRSRCRDYN TRNTBUFAV[5:0] TRNTCFGREQN TRNTDSTRDYN TRNTERRDROPN
Tpcicko_TXRAM	USERCLK	MIMTXRADDR[11:0] MIMTXREN MIMTXWADDR[11:0] MIMTXWDATA[35:0] MIMTXWEN

TRN to AXI Interface Migration Considerations

This appendix describes the differences in signal naming and behavior for users migrating to the Spartan®-6 FPGA Integrated Block for PCI Express®, v2.x from the Spartan-6 FPGA Integrated Block for PCI Express, v1.x.

High-Level Summary

The v2.x versions of the Virtex-6 FPGA Integrated Block for PCI Express update the main user interface from TRN to the standard AXI4-Stream signal naming and behavior. In addition, all control signals that were active Low have been changed to active High. This list summarizes the main changes to the core:

- Signal name changes
- Datapath DWORD ordering
- All control signals are active High
- Start-of-frame (SOF) signaling is implied
- Remainder signals are replaced with Keep signals

Step-by-Step Migration Guide

This section describes the steps that a user should take to migrate an existing user application based on TRN to the AXI4-Stream interface.

1. For each signal in [Table J-1](#) labeled “Name change only”, connect the appropriate user application signal to the newly named core signal.
2. For each signal in [Table J-1](#) labeled “Name change; Polarity”, add an inverter and connect the appropriate user application signal to the newly named core signal.
3. Swap the DWORD ordering on the datapath signals as described in [Data Path DWORD Ordering](#).
4. Leave disconnected the user application signal originally connected to trn_tsof_n.
5. Recreate trn_rsof_n as described in the [Start-Of-Frame Signaling](#) section and connect to the user application as was originally connected.
6. Make the necessary changes as described in the [Remainder/Keep Signaling](#) section.
7. If using the trn_rsrc_dsc_n signal in the original design, make the changes as described in [Packet Transfer Discontinue on Receive](#) section, otherwise leave disconnected.
8. Make the changes as described in the [Packet Reordering on Receive](#) section.

Signal Changes

Table J-1 details the main differences in signaling between TRN Local-Link to AXI4-Stream interface.

Table J-1: Interface Changes

TRN Name	AXI Name	Difference
Common Interface		
sys_reset_n	sys_reset	Name change; Polarity
trn_clk	user_clk_out	Name change only
trn_reset_n	user_reset_out	Name change; Polarity
trn_lnk_up_n	user_lnk_up	Name change; Polarity
trn_fc_ph[7:0]	fc_ph[7:0]	Name change only
trn_fc_pd[11:0]	fc_pd[11:0]	Name change only
trn_fc_nph[7:0]	fc_nph[7:0]	Name change only
trn_fc_npdi[11:0]	fc_npdi[11:0]	Name change only
trn_fc_cplh[7:0]	fc_cplh[7:0]	Name change only
trn_fc_cpld[11:0]	fc_cpld[11:0]	Name change only
trn_fc_sel[2:0]	fc_sel[2:0]	Name change only
Transmit Interface		
trn_tsof_n		No equivalent for 32- and 64-bit version (see text)
trn_teof_n	s_axis_tx_tlast	Name change; Polarity
trn_td[W-1:0] (W = 32, 64, or 128)	s_axis_tx_tdata[W-1:0]	Name change; DWORD Ordering (see text)
trn_trem_n (64-bit interface)	s_axis_tx_tkeep[7:0]	Name change; Functional differences (see text)
trn_trem_n[1:0] (128-bit interface)	s_axis_tx_tkeep[15:0]	Name change; Functional differences (see text)
trn_tsrdy_n	s_axis_tx_tvalid	Name change; Polarity
trn_tdst_rdy_n	s_axis_tx_tready	Name change; Polarity
trn_tsrdy_dsc_n	s_axis_tx_tuser[3]	Name change; Polarity
trn_tbuff_av[5:0]	tx_buf_av[5:0]	Name Change
trn_terr_drop_n	tx_terr_drop	Name change; Polarity
trn_tstr_n	s_axis_tx_tuser[2]	Name change; Polarity
trn_tcrg_req_n (64-bit interface only)	tx_cfg_req	Name change; Polarity
trn_tcrg_gnt_n (64-bit interface only)	tx_cfg_gnt	Name change; Polarity

Table J-1: Interface Changes (Cont'd)

TRN Name	AXI Name	Difference
trn_terrfd_n	s_axis_tx_tuser[1]	Name change; Polarity
Receive Interface		
trn_rsof_n		No equivalent for 32 and 64-bit versions
trn_reof_n	m_axis_rx_tlast	Name change; Polarity
trn_rd[W-1:0] (W = 32, 64, or 128)	m_axis_rx_tdata[W-1:0]	Name change; DWORD Ordering
trn_rrem_n (64-bit interface)	m_axis_rx_tkeep	Name change; Functional differences (see text)
trn_rrem_n[1:0] (128-bit interface)	m_axis_rx_tuser[14:10], m_axis_rx_tuser[21:17]	Name change; Functional differences (see text)
trn_rerrfd_n	m_axis_rx_tuser[1]	Name change; Polarity
trn_rsrc_rdy_n	m_axis_rx_tvalid	Name change; Polarity
trn_rdst_rdy_n	m_axis_rx_tready	Name change; Polarity
trn_rsrc_dsc_n		No equivalent
trn_rnp_ok_n	rx_np_ok	Name change; Polarity; Extra delay (see text)
trn_rbar_hit_n[6:0]	m_axis_rx_tuser[8:2]	Name change; Polarity
Configuration Interface		
cfg_rd_wr_done_n	cfg_rd_wr_done	Name change; Polarity
cfg_byte_en_n[3:0]	cfg_byte_en[3:0]	Name change; Polarity
cfg_wr_en_n	cfg_wr_en	Name change; Polarity
cfg_rd_en_n	cfg_rd_en	Name change; Polarity
cfg_PCIE_link_state_n[2:0]	cfg_PCIE_link_state[2:0]	Name change only
cfg_trn_pending_n	cfg_trn_pending	Name change; Polarity
cfg_to_turnoff_n	cfg_to_turnoff	Name change; Polarity
cfg_turnoff_ok_n	cfg_turnoff_ok	Name change; Polarity
cfg_pm_wake_n	cfg_pm_wake	Name change; Polarity
cfg_wr_rw1c_as_rw_n	cfg_wr_rw1c_as_rw	Name change; Polarity
cfg_interrupt_n	cfg_interrupt	Name change; Polarity
cfg_interrupt_rdy_n	cfg_interrupt_rdy	Name change; Polarity
cfg_interrupt_assert_n	cfg_interrupt_assert	Name change; Polarity
cfg_err_ecrc_n	cfg_err_ecrc	Name change; Polarity
cfg_err_ur_n	cfg_err_ur	Name change; Polarity
cfg_err_cpl_timeout_n	cfg_err_cpl_timeout	Name change; Polarity
cfg_err_cpl_unexpect_n	cfg_err_cpl_unexpect	Name change; Polarity

Table J-1: Interface Changes (Cont'd)

TRN Name	AXI Name	Difference
cfg_err_cpl_abort_n	cfg_err_cpl_abort	Name change; Polarity
cfg_err_posted_n	cfg_err_posted	Name change; Polarity
cfg_err_cor_n	cfg_err_cor	Name change; Polarity
cfg_err_cpl_rdy_n	cfg_err_cpl_rdy	Name change; Polarity
cfg_err_locked_n	cfg_err_locked	Name change; Polarity

Data Path DWORD Ordering

The AXI4-Stream interface swaps the DWORD locations but preserves byte ordering within an individual DWORD as compared to the TRN interface. This change only affects the 64-bit and 128-bit versions of the core. [Figure J-1](#) and [Figure J-2](#) illustrate the DWORD swap ordering from TRN to AXI4-Stream for both 64-bit and 128-bit versions.

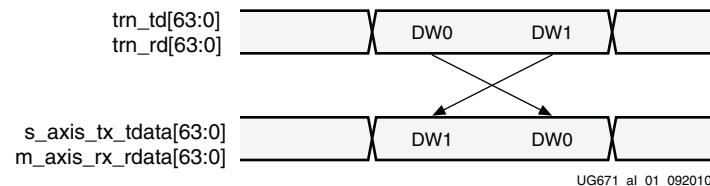


Figure J-1: TRN vs. AXI DWORD Ordering on Data Bus (64-bit)

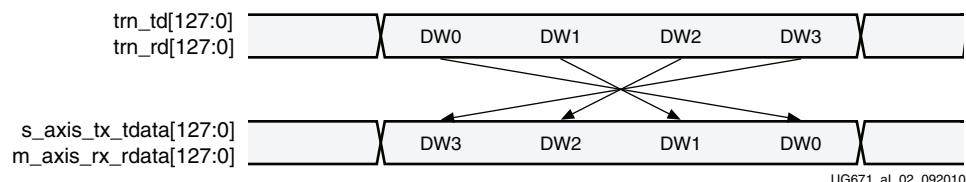


Figure J-2: TRN vs. AXI DWORD Ordering on Data Bus (128-bit)

Users migrating existing 64-bit and 128-bit TRN-based designs should swap DWORD locations for the `s_axis_tx_tdata[W-1:0]` and `s_axis_rx_rdata[W-1:0]` buses as they enter and exit the PCIe® core.

For example, existing user application pseudo-code:

```
usr_trn_rd[127:0] = trn_rd[127:0];
```

should be modified to:

```
usr_trn_rd[127:96] = s_axis_rx_rdata[31:0]
usr_trn_rd[95:64] = s_axis_rx_rdata[63:32]
usr_trn_rd[63:32] = s_axis_rx_rdata[95:64]
usr_trn_rd[31:0] = s_axis_rx_rdata[127:96]
```

Start-Of-Frame Signaling

The AXI4-Stream interface does not have equivalent signals for start-of-frame (trn_tsوف_n and trn_rsof_n) in the 32-bit and 64-bit versions. On the transmit side, existing TRN designs can simply leave the user trn_tsوف_n connection unconnected. On the receive side, existing TRN designs can recreate trn_rsof_n using simple logic, if necessary.

32- and 64-Bit Interfaces

First the user creates a sequential (clocked) signal called in_packet_reg. A combinatorial logic function using existing signals from the core can then be used to recreate trn_rsof_n as illustrated in this pseudo-code:

```
For every clock cycle (user_clk_out) do {
    if(reset)
        in_packet_reg = 0
    else if (m_axis_rx_tvalid and m_axis_rx_tready)
        in_packet_reg = !m_axis_rx_tlast
}

trn_rsof_n = !(m_axis_rx_tvalid & !in_packet_reg)
```

128-Bit Interface

The 128-bit interface provides an SOF signal. The user can invert is_sof[4] to recreate trn_rsof_n.

Remainder/Keep Signaling

This section covers the changes to the remainder signals trn_trem_n[1:0] and trn_rrem_n[1:0].

The AXI4-Stream interface uses “keep” signaling (byte enables) in place of remainder signaling. There are three key differences between the keep signals and the remainder signals as detailed in [Table J-2](#). There are also some differences between the 64-bit version and 128-bit version of the core. The 128-bit RX version replaces trn_rrem[1:0] with is_sof and is_eof, instead of keep signals. For simplicity, this section treats 64-bit and 128-bit transmit and receive operations separately.

Table J-2: Remainder Signal Differences

TRN Remainders 64-bit: trn_trem_n, trn_rrem_n 128-bit: trn_trem_n[1:0], trn_rrem_n[1:0]	AXI4-Stream Byte Enables 64-bit: s_axis_tx_tkeep[7:0], m_axis_rx_tkeep[7:0] 128-bit: s_axis_tx_tkeep[15:0], is_sof[4:0], is_eof[4:0]
Active Low	Active High
Acts on DWORDs	Acts on Bytes
Only valid on end-of-frame (EOF) cycles	Valid for every clock cycle that tvalid and tready are asserted

64-Bit Transmit

Existing TRN designs can do a simple conversion from the single trn_trem signal to s_axis_tx_tkeep[7:0]. Assuming the user currently has a signal named user_trn_trem that drives the trn_trem input, the listed pseudo-code illustrates the conversion to s_axis_tx_tkeep[7:0]. The user must drive s_axis_tx_tkeep[7:0] every clock cycle that tvalid is asserted.

```
if s_axis_tx_tlast == 1 //in a packet at EOF
    s_axis_tx_tkeep[7:0] = user_trn_trem_n ? 0Fh : FFh
else //in a packet but not EOF, or not in a packet
    s_axis_tx_tkeep = FFh
```

64-Bit Receive

Existing TRN designs can do a simple conversion on m_axis_rx_tkeep[7:0] to recreate the trn_rrem signal using combinatorial logic. The listed pseudo-code illustrates the conversion.

```
if m_axis_rx_tlast == 1
    trn_rrem_n = (m_axis_rx_tkeep[7:4] == Fh) ? 0b : 1b
else
    trn_rrem_n = 1b
```

128-Bit Transmit

Existing TRN designs can do a simple conversion from the single trn_trem[1:0] signal to s_axis_tx_tkeep[15:0]. Assuming the user currently has a signal named user_trn_trem[1:0] that drives the trn_trem[1:0] input, the listed pseudo-code illustrates the conversion to s_axis_tx_tkeep[15:0]. The user must drive s_axis_tx_tkeep[15:0] every clock cycle.

```
if s_axis_tx_tlast == 1 //in a packet at EOF
    if user_trn_trem_n[1:0]==00b
        s_axis_tx_tkeep[15:0] = FFFFh
    else if user_trn_trem_n[1:0] = 01b
        s_axis_tx_tkeep[15:0] = 0FFFh
    else if user_trn_trem_n[1:0] = 10b
        s_axis_tx_tkeep[15:0] = 00FFh
    else if user_trn_trem_n[1:0] = 11b
        s_axis_tx_tkeep[15:0] = 000Fh
    else //in a packet but not EOF, or not in a packet
        s_axis_tx_tkeep = FF FFh
```

128-Bit Receive

The 128-bit receive remainder signal trn_rrem[1:0] does not have an equivalent keep signal for AXI4-Stream. Instead, is_sof[4:0] (m_axis_rx_tuser[14:10]) and is_eof[4:0] (m_axis_rx_tuser[21:17]) are used. Existing TRN designs can do a conversion on the is_sof and is_eof signals to recreate the trn_rrem[1:0] signal using combinatorial logic. The listed pseudo-code illustrates the conversion. This pseudo-code assumes that the user has swapped the DWORD locations from the AXI4-Stream interface (see the usr_trn_rd[127:0] signal pseudo-code).

```

trn_rrem_n[1] = !is_sof[4] & !is_eof[4] | is_eof[4] & is_sof[3] |
is_eof[4] & !is_eof[3]]

trn_rrem_n[0] = !is_eof[2]

```

Note: is_eof[4] is equivalent to m_axis_rx_tlast.

Packet Transfer Discontinue on Receive

When the trn_rsrc_dsc_n signal in the TRN interface is asserted, it indicates to the user that a received packet has been discontinued. The AXI4-Stream interface has no equivalent signal. On both the TRN and AXI4-Stream cores, however, a packet is only discontinued on the receive interface if link connectivity is lost. Therefore, users can simply monitor the user_lnk_up signal to determine a receive packet discontinue condition.

On the TRN interface, the packet transmission on the data interface (trn_rd) stops immediately following assertion of trn_rsrc_dsc_n, and trn_reof_n might never be asserted. On the AXI4-Stream interface, the packet is padded out to the proper length of the TLP, and m_axis_rx_tlast is asserted even though the data is corrupted. [Figure J-3](#) and [Figure J-4](#) show the TRN and AXI signaling for packet discontinue. To recreate the trn_rsrc_dsc_n signal, the user can simply invert and add one clock cycle delay to user_lnk_up.

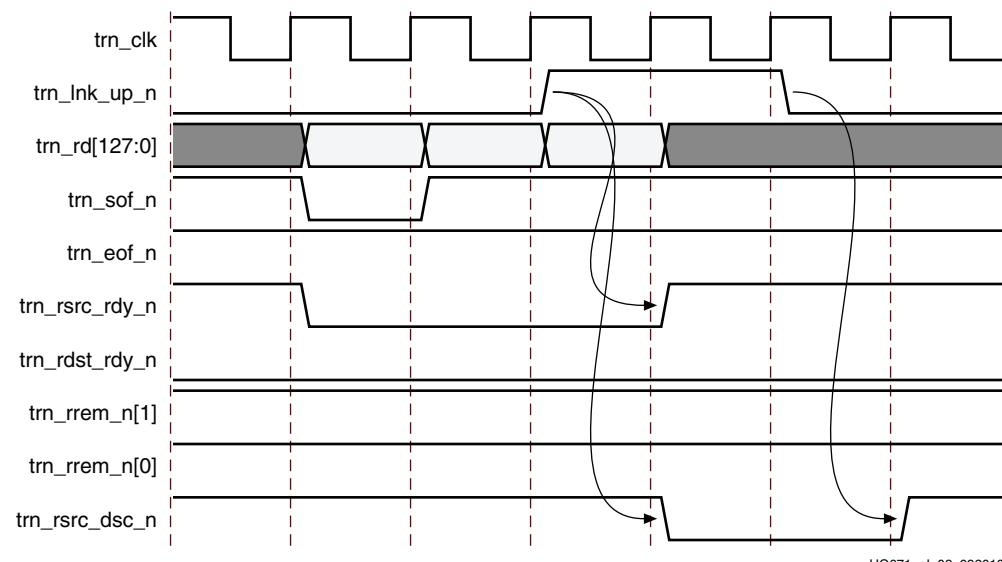


Figure J-3: Receive Discontinue on the TRN Interface

UG671_a1_03_092010

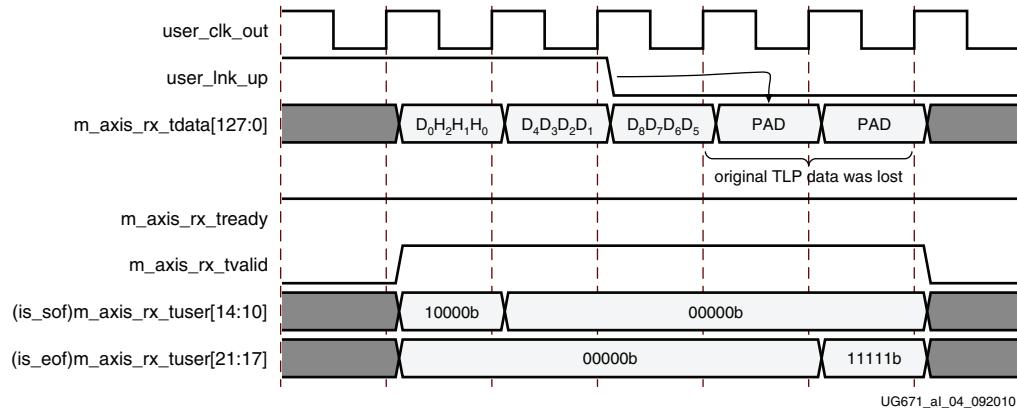


Figure J-4: Receive Discontinuity on the AXI4-Stream Interface

Packet Reordering on Receive

The TRN interface uses the `trn_rnp_ok_n` signal to reorder TLP traffic on the receive interface. The AXI4-Stream interface has an equivalent signal, `rx_np_ok`. Users need to account for two differences in the AXI4-Stream interface as shown in [Table J-3](#). Users have to account for these differences in their custom logic. If the user application does not use packet reordering, the user can simply tie `rx_np_ok` to 1b.

Table J-3: AXI4-Stream Interface Differences

TRN <code>trn_rnp_ok_n</code>	AXI4-Stream <code>rx_np_ok</code>
Active Low	Active High
Must be deasserted at least one clock cycle before <code>trn_reof_n</code> of the next-to-last Non-Posted TLP that the user can accept	Must be deasserted at least one clock cycle before <code>is_eof[4]</code> of the second-to-last Non-Posted TLP that the user can accept

System Reset

The system reset is usually provided by PERST#, which is an active Low signal. If the incoming reset signal is active Low, the user must invert this signal before connecting to the `sys_reset` signal on the core interface.